



OVP Processor Modeling Guide

Imperas Software Limited

Imperas Buildings, North Weston,
Thame, Oxfordshire, OX9 2HA, UK
docs@imperas.com



Author:	Imperas Software Limited
Version:	6.5.3.1
Filename:	OVP_Processor_Modeling_Guide.doc
Project:	OVP Processor Modeling Guide
Last Saved:	Monday, 20 July 2015
Keywords:	

Copyright Notice

Copyright © 2015 Imperas Software Limited All rights reserved. This software and documentation contain information that is the property of Imperas Software Limited. The software and documentation are furnished under a license agreement and may be used or copied only in accordance with the terms of the license agreement. No part of the software and documentation may be reproduced, transmitted, or translated, in any form or by any means, electronic, mechanical, manual, optical, or otherwise, without prior written permission of Imperas Software Limited, or as expressly provided by the license agreement.

Right to Copy Documentation

The license agreement with Imperas permits licensee to make copies of the documentation for its internal use only. Each copy shall include all copyrights, trademarks, service marks, and proprietary rights notices, if any.

Destination Control Statement

All technical data contained in this publication is subject to the export control laws of the United States of America. Disclosure to nationals of other countries contrary to United States law is prohibited. It is the reader's responsibility to determine the applicable regulations and to comply with them.

Disclaimer

IMPERAS SOFTWARE LIMITED., AND ITS LICENSORS MAKE NO WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, WITH REGARD TO THIS MATERIAL, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE.

Table of Contents

1	Preface.....	4
2	Introduction.....	5
3	Imperas Simulation Overview	6
4	Introduction to Processor Modeling.....	8
5	Implementing the Instruction Decoder	20
6	Implementing the Instruction Disassembler	29
7	Implementing Simple Behavior	34
8	Processor Flags and Register Dumping	43
9	Implementing Unconditional Jump Instructions.....	58
10	Implementing Conditional Jump Instructions.....	66
11	Implementing Memory Access Instructions	75
12	Modeling Exceptions	84
13	Modeling Mode-Dependent Behavior (Part 1)	103
14	Modeling Mode-Dependent Behavior (Part 2)	114
15	Implementing a Tick Timer	123
16	Modeling External Interrupts	139
17	Implementing the Debug Interface	147
18	Implementing Fixed-Mapped Virtual Memory	163
19	Implementing a Dynamic-Mapped TLB.....	172
20	Implementing a TLB LRU Replacement Policy.....	186
21	Implementing QuantumLeap-Compatible Models	192
22	Function Address Semihosting	198
23	Using Intercept Libraries for Instruction Set Enhancement	211
24	Adding an Extended Programmers View	221
25	Processor Configuration.....	227
26	Making High-Performance Processor Models.....	232

1 Preface

This document describes how to create processor models for use with OVPsim and Imperas simulation tools.

1.1 *Notation*

Code

Code extracts

1.2 *Recommended Reading*

Imperas simulation technology is based on just-in-time (JIT) compiler technology. The following book provides a good introduction to the concepts involved:

Virtual Machines, by James E. Smith, Ravi Nair
ISBN 1-55860-910-5
Publisher: Morgan Kaufmann/Elsevier

1.3 *Related Imperas & OVP Documents*

- VMI Morph Time Function Reference
- VMI Run Time Function Reference

2 Introduction

Imperas simulation technology enables very high performance simulation, debug and analysis of platforms containing multiple processors and peripheral models. The technology is designed to be extensible: you can create new models of processors and other platform components using interfaces and libraries supplied by Imperas. Processor models developed using this technology can be used both with Imperas simulation products and the freely-available OVPsim platform simulator.

This document describes how to use the OVP interfaces to create new processor models.

The documentation here is supported by C code samples in the `Examples` directory of your Imperas installation, and also to download from the OVPWorld website (www.ovpworld.org). The compilation make use of Makefiles, the instructions for which indicate the use of the command *make*, on Windows systems the MinGW *mingw32-make* command should be used in its place.

2.1 Prerequisites

Since models for use with Imperas and OVP tools are written in C, an important prerequisite is that you must be an expert in the C language.

In very rare circumstances it is beneficial to implement some highly-performance-critical routines directly in assembler. You should ensure you are familiar with the x86 instruction set and assembler usage if required.

GCC Compiler Versions

Linux32	4.5.2	i686-nptl-linux-gnu (Crosstool-ng)
Linux64	4.4.3	x86_64-unknown-linux-gnu (Crosstool-ng)
Windows32	4.4.7	mingw-w32-bin_i686-mingw
Windows64	4.4.7	mingw-w64-bin_i686-mingw

3 Imperas Simulation Overview

Before starting to create models for use with the Imperas simulation environment, you must understand how the components used in that environment interact. This section describes this in detail.

3.1 *Simulation Environments*

There are two simulation environments that can be used with models that you create:

- *OVPsim* allows processor models created using OVP modeling technology to be used in C harness or platform files to create executables that execute binaries compiled for those processor models. It can also simulate behavioral components (the subject of this guide). OVPsim can also be used in 3rd party simulation environments (for example, SystemC). It can also be used to create a test harness to help validate processor models under construction, or even to create custom simulation environments. OVPsim has less functionality than the Imperas Professional Simulator Products in some areas and has restricted commercial usage as stipulated in the OVP click-through license agreement.
- *Imperas Professional Simulator Products* enhance the basic capabilities provided by OVPsim, particularly in the areas of debugger integration, tool integration and multiprocessor simulation support (including QuantumLeap parallel simulation). Contact Imperas for more information.

3.2 *Processor Models*

The core simulation components are *processor models*. In order to create a new processor model, you must implement the following major components by writing C code using the Imperas *Virtual Machine Interface* (VMI) API:

- An *instruction decoder*, capable of decoding a single processor instruction. This is a required component for the *disassembler*, *morpher* and *debugger interface*, described in section 5.
- An *instruction disassembler*, capable of generating a text representation of an instruction, described in section 6.
- An *instruction morpher*, capable of describing the behavior of a single instruction, described in sections 7 - 16.
- A *debugger interface*, which provides functions required for the model to be debugged using gdb or the Imperas multiprocessor debugger, described in section 17.
- If a processor implements virtual memory, then the hardware structures that support that virtual memory (MMU and TLB, for example) should also form part of the processor model. This is described in sections 18 - 20.
- If processors are to be used with the QuantumLeap parallel simulation algorithm of the Imperas Professional Simulation products, some changes may be required (for example, to identify atomic instructions). This is described in section 21.

- A processor model must define its external hardware interface. This allows the model to associate its actions with external events and allows the simulator and other tools to interrogate the model and discover its interface. External inputs are covered in section 12.
- Most processor models will need to model members of a family of processors (family members are referred to as *variants*). To re-use your code as far as possible it is often convenient for one processor model to support multiple variants and options that can be configured from the platform. Model configuration is covered in section 25.

Processor models are compiled into a shared object (.so or .dll) which is then dynamically loaded by the simulation environment.

3.3 *Semihosting*

Semihosting allows behavior that would normally occur on a simulated system to be implemented using features of the host system instead. As a simple example, a real platform might contain a UART peripheral to receive output. When simulating this system, it is generally more convenient not to simulate the UART at all but instead to intercept any `write` call that a processor makes and redirect the output to the simulator log instead. Such behavior is specified in a *semihosting library* for a processor.

Implementation of semihosting libraries is described in section 22.

3.4 *Cache and Memory Subsystem Models*

Memory subsystem models such as caches can be modeled as loadable shared objects (or dynamic linked libraries on Windows) and separately instantiated. This makes it very easy to explore hardware options: what happens to the performance of this application if I double the size of the L2 cache?

Memory subsystem models can be either *full* or *transparent*. A full model implements memory contents: for example a full cache model would implement both cache tags and the cache line contents. A transparent model implements some state but not the memory contents: for example, a transparent cache model would implement the cache tags but *not* the line contents, which is useful for performance analysis models that simply count hits and misses.

Implementation of memory subsystem models is beyond the scope of this document.

4 Introduction to Processor Modeling

When creating a new processor model, it is important to follow a structured approach. We recommend the following:

1. Prepare the ground (section 4.1).
2. Create a processor outline model (section 4.2).
3. Create a test platform instantiating the processor model (section 4.3).
4. Create an application test case (section 4.4).
5. Run the application test case on the outline model (section 4.5).

Once the outline model has been prepared, further steps are required to complete the model. In detail, these are:

6. Implement the instruction decoder (chapter 5).
7. Implement the instruction disassembler (chapter 6).
8. Implement the behavior of each instruction (chapter 7 onwards).
9. Implement a debugger interface (chapter 17).
10. Provide supplementary simulation information (chapter 27).

Most processor models will also need one or more semihosting libraries to be implemented for them. This is described in section 22.

Items 1 to 5 are covered in this introductory section, which shows how to get up and running with a minimal processor model. Later chapters of the document will fill in details that are glossed over here in order to give a fast overview.

4.1 *Prerequisites*

Before starting implementation of a new processor model, we recommend that you do the following:

1. Identify the particular processor variant to be modeled (when variants exist).
2. Obtain a processor tool chain and understand how to use it for the variant you will be modeling (if a tool chain is available). Typically, you will find it useful to have an assembler, linker, object dump utility and C compiler. If you are writing a model for a completely new processor then it is possible that no supporting tools may yet exist: in this case, you will need to become familiar with the object code format of the processor and possibly implement a custom object file loader as part of the modeling project.
3. Obtain a golden reference model if possible. Validating a processor model is much easier if there is a golden reference against which comparisons can be made.

4.2 Creating a Processor Outline Model

A minimal processor outline model is available in the directory:

```
$IMPERAS_HOME/Examples/Models/Processor/1.or1kOutline
```

This model is for the freely-available OR1K processor (see http://opencores.org/or1k/Main_Page). This outline model in fact contains no code specific to the OR1K processor at all: it implements the bare minimum functionality to create a shared object usable by the Imperas simulation tools.

Take a copy of the outline model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/1.or1kOutline .
```

Compile the model using the `make` utility:

```
cd 1.or1kOutline
make
```

Running `make` compiles the model in the current directory (using `Makefile`) and links it with an Imperas stub library (`vmiStubs.static.a`) to create a shared object loadable by the Imperas tools (`model.so/model.dll`).

The model source code covered here and in following sections refers to Imperas header files in the directory:

```
$IMPERAS_HOME/ImpPublic/include/host/vmi
```

These header files comprise the Imperas *Virtual Modeling Interface* (VMI) API.

The outline model files are described in the following subsections.

4.2.1 File `or1kFunctions.h`

File `or1kFunctions.h` declares prototypes of functions used throughout the processor model. Functions that must have particular prototypes for use with the Imperas tools should be defined using macros from file `vmiAttrs.h` within the VMI header directory; for example, every processor must have constructor, which is declared as:

```
VMI_CONSTRUCTOR_FN(or1kConstructor);
```

The macro `VMI_CONSTRUCTOR_FN` is defined in `vmiAttrs.h` as:

```
#define VMI_CONSTRUCTOR_FN(_NAME) void _NAME( \
    vmiProcessorP   processor, \
    const char      *type, \
    Bool            simulateExceptions, \
    vmiSMPContextP  smpContext, \
    void            *parameterValues \
)
```

In other words, function `orlkConstructor` is a void function which is passed an argument `processor` which is of type `vmiProcessorP`, an argument `type` which is a constant string, a Boolean argument `simulateExceptions`, an argument `smpContext` of type `vmiSMPContextP` and an argument `parameterValues` of type `void*`. The `vmiProcessorP` type is an opaque type pointer representing the current processor state – we will see how this is used later in this section.

Always use the macros provided in the VMI header files to declare and define your functions: this protects any code you write from future changes to any of the Imperas function definitions.

4.2.2 File `orlkStructure.h`

File `orlkStructure.h` defines a structure that will be used to hold the state of a single OR1K processor. Because this is a generic model, at this point the structure is empty.

4.2.3 File `orlkMain.c`

File `orlkMain.c` implements two functions that must be present in every processor model: the *constructor* and *destructor*. The constructor function is called for each new instance of a processor. It should initialize the processor state (for example, by setting registers in the processor structure to a known state). The destructor is called at the end of simulation for each processor instance. It should perform any required processor-model-specific shutdown actions.

In this example, the constructor and destructor perform no action except to print that they have been called (using the message API defined in `vmiMessage.h`), and to allocate and free the model's *bus interface*:

```
VMI_CONSTRUCTOR_FN(orlkConstructor) {
    orlkP orlk = (orlkP)processor;

    vmiPrintf("%s called\n", FUNC_NAME);

    // create bus port specifications
    newBusPorts(orlk);
}

VMI_DESTRUCTOR_FN(orlkDestructor) {
    orlkP orlk = (orlkP)processor;
```

```
    vmiPrintf("%s called\n", FUNC_NAME);

    // free bus port specifications
    freeBusPorts(ork);
}
```

It is good practice to give each public model declaration a common, model specific prefix: this simplifies debugging the model in a simulation where several models of different types are in use. In this case, we have chosen the prefix `ork`.

File `orkMain.c` also implements a *bus port specification function* for the processor, which tells the simulator the number of bus ports that the OR1K has and their width. The OR1K, like many processors, has two bus ports. The first, called `INSTRUCTION` is a bus master port used to fetch instructions from memory. The second, called `DATA` is also a bus master port, used to read and write data to memory (on many systems, these two ports are connected to the same physical bus, so share the same address space). Objects representing these two ports are allocated and stored on the processor instance by function `newBusPorts`:

```
const static vmiBusPort busPorts[] = {
    {"INSTRUCTION", vmi_BP_MASTER, vmi_DOM_CODE, {32,32}, 1},
    {"DATA",        , vmi_BP_MASTER, vmi_DOM_DATA, {32,32}, 0},
};

static void newBusPorts(orkP ork) {

    Uns32 i;

    ork->busPorts = STYPE_CALLOC_N(vmiBusPort, NUM_MEMBERS(busPorts));

    for(i=0; i<NUM_MEMBERS(busPorts); i++) {
        ork->busPorts[i] = busPorts[i];
    }
}
```

The template structure `busPorts` describes the instruction and data bus ports. The `vmiBusPort` type is defined in `vmiPorts.h` as follows:

```
typedef enum vmiBusPortTypeE {
    vmi_BP_MASTER,
    vmi_BP_SLAVE,
    vmi_BP_MASTER_SLAVE
} vmiBusPortType;

typedef enum vmiDomainTypeE {
    vmi_DOM_CODE,      // std code domain port
    vmi_DOM_DATA,      // std data domain port
    vmi_DOM_OTHER      // other domain port
} vmiDomainType;

typedef struct vmiBusPorts {

    const char      *name;
    vmiBusPortType  type;
    vmiDomainType   domainType;
    struct          {Uns8 min; Uns8 max;} addrBits;
```

```
    Bool          mustBeConnected;
    memDomainP    domain;

    // space for documentation
    const char     *description;
    void           *descriptionDom;
} vmiBusPort;
```

The model fills the `name`, `type`, `domainType`, `addrBits.min`, `addrBits.max` and `mustBeConnected` fields; remaining fields are filled by the simulator as the model is instantiated. The fields have the following meanings:

1. `name`: the name of the port;
2. `type`: the port type (master or slave);
3. `domainType`: the port usage (code or data);
4. `addrBits.min`: the minimum width of a bus that can be connected;
5. `addrBits.max`: the maximum width of a bus that can be connected;
6. `mustBeConnected`: whether the port must be connected (if `False`, it may be left unconnected).

After the processor constructor has been called, the simulator obtains information about the model's bus ports by calling an iterator function here implemented by `orlkGetBusPortSpec` which returns a pointer to a `vmiBusPort` structure for each implemented port. Like most VMI iterators, it is called with zero to return the first object, with the previous object to return the next, and it returns zero when all objects have been returned:

```
VMI_BUS_PORT_SPECS_FN(orlkGetBusPortSpec) {
    orlkP orlk = (orlkP)processor;

    if(!prev) {
        // first port
        return orlk->busPorts;
    } else {
        // port other than the first
        Uns32 prevIndex = (prev->orlk->busPorts);
        Uns32 thisIndex = prevIndex+1;

        return (thisIndex<NUM_MEMBERS(busPorts)) ? &orlk->busPorts[thisIndex]:0;
    }
}
```

When simulation ends, the destructor frees the allocated bus port list by calling function `freeBusPorts`:

```
static void freeBusPorts(orlkP orlk) {
    if(orlk->busPorts) {
        STYPE_FREE(orlk->busPorts);
        orlk->busPorts = 0;
    }
}
```

```
}
```

4.2.4 File `or1kMorph.c`

File `or1kMorph.c` implements the OR1K morpher function. The morpher function is responsible for defining how each processor instruction should be evaluated. This is described in detail in section 7. The minimal example simply contains a call to function `vmimtExit`, which will cause the processor to terminate on the first instruction encountered.

4.2.5 File `or1kUtils.c`

File `or1kUtils.c` implements two simulation support functions required in every model: the *processor endianness* function and the *next instruction* function.

4.2.5.1 The Endianness Function

This function must return the endianness of the processor when fetching code. Currently supported options are `MEM_ENDIAN_BIG` and `MEM_ENDIAN_LITTLE`. This OR1K model is big endian, so the function is defined as:

```
VMI_ENDIAN_FN(or1kGetEndian) {  
    return MEM_ENDIAN_BIG;  
}
```

Some models have endianness dependent upon the current processor state. For this reason, the endianness callback is passed the current processor as an argument so that its state can be accessed if required.

This function can be called to request both the endianness of instruction fetches and the endianness of loads and stores: which is required is specified by the second argument to the `VMI_ENDIAN_FN`, a Boolean called `isFetch`. This is `True` for an instruction fetch and `False` for a data access.

4.2.5.2 The Next Instruction Function

Given an instruction address, this function must return the *next* instruction address. This function is used by the simulator to step through the simulated code when generating (morphing) equivalent native code.

For processors with variable-length instructions (for example, x86 variants) the next instruction address function will be required to perform a full or partial instruction decode in order to determine the next instruction address. On RISC processors, the instruction size may be constant, so no decode is required. See section 5 for details of implementing an instruction decoder.

The minimal processor model assumes a constant instruction size of four bytes and is therefore implemented like this:

```
VMI_NEXT_PC_FN(or1kNextInstruction) {  
    Uns32 nextAddress = (Uns32)(thisPC + 4);  
    return nextAddress;  
}
```

The next instruction function must correctly handle *instruction wraparound*. In the example above, it would be incorrect to implement the function as:

```
return thisPC + 4;
```

(this would not wrap round as required after 0xffffffffc, because the Addr type of thisPC is 64 bits, not 32).

4.2.6 File or1kInfo.c

File or1kInfo.c implements the processor information function, which returns information about the model in several categories:

```
#include "vmi/vmiAttrs.h"
#include "vmi/vmiModelInfo.h"

#include "or1kFunctions.h"

VMI_PROC_INFO_FN(or1kProcInfo) {
    static vmiProcessorInfo info;

    info.vlnv.vendor      = "ovpworld.org"      ;
    info.vlnv.library     = "processor"         ;
    info.vlnv.name        = "or1k"             ;
    info.vlnv.version     = "1.0"              ;

    info.semihost.vendor  = "ovpworld.org"      ;
    info.semihost.library = "semihosting"       ;
    info.semihost.name    = "or1kNewlib"       ;
    info.semihost.version = "1.0"              ;

    info.elfCode          = 33906;
    info.endianFixed      = True;
    info.endian           = MEM_ENDIAN_BIG;
    info.gdbPath          = "$IMPERAS_HOME/lib/$IMPERAS_ARCH/CrossCompiler/or32-elf/bin/or32-elf-gdb" VMI_EXE_SUFFIX;
    info.gdbFlags         = 0;
    info.gdbInitCommands  = 0;

    info.QLQualified      = True;
    info.debugSnapAddress = 0;
    return &info;
}
```

The example uses a static structure; in a more complex model the member values could depend on the current mode of the processor.

4.2.6.1 Model Location.

The first member specifies where the model will be stored, using the "Vendor, Library, Name, Version" (VLNV) notation.

4.2.6.2 Default Semihost Library

The second member specifies the location of the default semihost library (see Chapter 21) to be used with this model, using VLNv notation.

4.2.6.3 ELF Code and Endian

This field lets the simulator check the compatibility of an application program before it is loaded for execution by this model. Setting `endianFixed` to zero indicates that the processor can be either endian. This is not to be confused with the Endianness Function which returns the *current endianness* of the processor.

4.2.6.4 gdbPath

This field specifies the path to the debugger to be used with the model. The `VMI_EXI_SUFFIX` macro can be used to conditionally add the `.exe` file suffix required for an executable on a Windows host.

4.2.6.5 gdbFlags

This field specifies any flags to be supplied on the debugger command line when it is invoked.

4.2.6.6 gdbInitCommands

This field specifies any commands to be sent to the debugger after starting the executable, but before debugging begins. For example if the debugger supports several architectures, “set arch <specific_architecture>” can be used to choose one. If more than one command is required, separate then using the new-line “\n” character.

4.2.6.7 QuantumLeap qualified

This field specifies that the processor is able to run in parallel mode. See section 21.6.

4.2.6.8 Debugger snap address

Some processors do not use byte addressing; all instructions fall on 2-byte or 4-byte boundaries. A processor of this kind might use the least significant bits of its address to indicate special processor modes. If the gdb that is used with this processor requests breakpoints with the least significant bits set (also indicating the special processor modes), the simulator will not correctly detect the processor executing at a breakpoint address. The `debugSnapAddress` field can be used to work around this problem; set this field to force breakpoint addresses from the debugger to the appropriate boundary:

<code>debugSnapAddress</code>	meaning
0 (default) or 1	Breakpoint lies on 8-bit boundary
2	Breakpoint lies on 16-bit boundary
4	Breakpoint lies on 32-bit boundary
8	Breakpoint lies on 64-bit boundary
(Other values are illegal)	

4.2.7 File `or1kAttrs.c`

File `or1kAttrs.c` implements the *VMI instruction attributes* object for the OR1K processor. This is a C structure of type `vmiIASAttrs`, the type of which is defined in the VMI header file `vmiAttrs.h`. The structure encapsulates all required information about the processor in a form that is usable by the Imperas simulation products.

```

const vmiIASAttr modelAttrs = {

    //////////////////////////////////////
    // VERSION & SIZE ATTRIBUTES
    //////////////////////////////////////

    .versionString = VMI_VERSION,
    .modelType     = VMI_PROCESSOR_MODEL,
    .dictNames     = dictNames,
    .cpuSize       = sizeof(ork),

    //////////////////////////////////////
    // CREATE/DELETE ROUTINES
    //////////////////////////////////////

    .constructorCB = orkConstructor,
    .destructorCB  = orkDestructor,

    //////////////////////////////////////
    // MORPHER CORE ROUTINES
    //////////////////////////////////////

    .morphCB = orkMorphInstruction,

    //////////////////////////////////////
    // SIMULATION SUPPORT ROUTINES
    //////////////////////////////////////

    .getEndianCB = orkGetEndian,
    .nextPCCB    = orkNextInstruction,

    //////////////////////////////////////
    // PORT ROUTINES
    //////////////////////////////////////

    .busPortSpecsCB = orkGetBusPortSpec,

    //////////////////////////////////////
    // PROCESSOR INFO ROUTINE
    //////////////////////////////////////

    .procInfoCB = orkProcInfo,
};

```

Note that all fields in the structure are initialized *by name*. This is done so that source code changes are not required if new fields are added to the structure in future.

In the case of the minimal processor model, the structure contains:

1. A VMI version string, `VMI_VERSION` (defined in `vmiVersion.h`). This is used when the model is loaded by Imperas simulation products to ensure compatibility.
2. The type of model, `VMI_PROCESSOR_MODEL` (defined in `vmiTypes.h`). This is used by Imperas simulation products to ensure the correct kind of model is being loaded.
3. A list of *dictionary names* used by the model, `dictNames`. Dictionaries provide a mechanism to efficiently model *modal* processors and are discussed in chapter 14. Every processor must have at least one dictionary name, specified in a null-

terminated array of constant strings. In the case of the minimal processor model, the dictionary names are specified as:

```
static const char *dictNames[] = {"NORMAL", 0};
```

so there is a single dictionary called *NORMAL* in this model.

4. An indication of the size required for the processor structure defined in `orlkStructure.h`, `sizeof(orlk)`.
5. References to all the callbacks required to implement the processor model. For this minimal model, there are references to seven functions – `orlkConstructor`, `orlkDestructor`, `orlkMorphInstruction`, `orlkGetEndian`, `orlkNextInstruction`, `orlkGetBusPortSpec` and `orlkProcInfo`.

4.3 Implementing a Test Platform using OVPsim

Section 4.2 described how an outline processor model was compiled using the `make` command. In order to validate and debug the outline model, it is useful to have a test platform to drive it. The easiest way to create a test platform is using OVPsim.

Within the `1.orlkOutline` directory is a sample test platform source file, `platform/platform.c`. The test platform can be compiled to produce an executable, `platform.IMPERAS_ARCH.exe`, by using this command in the `orlkOutline` directory:

```
make -C platform
```

File `platform.c` has a main function as follows:

```
int main(int argc, char ** argv) {

    // check for the program argument
    if(argc!=2) {
        icmPrintf("%s: expected application name argument\n", argv[0]);
        return -1;
    }

    // initialize CpuManager
    icmInitPlatform(ICM_VERSION, 0, NULL, 0, "platform");

    const char *modelFile = "model." IMPERAS_SHRSUF;

    // create a processor
    icmProcessorP processor = icmNewProcessor(
        "cpu1",           // CPU name
        TYPE_NAME,       // CPU type
        0,               // CPU cpuId
        0,               // CPU model flags
        32,              // address bits
        modelFile,       // model file
        "modelAttrs",    // morpher attributes
        0,               // no CPU attributes
        0,               // user-defined attributes
        0,               // semi-hosting file
        0                // semi-hosting attributes
    );

    // load the processor object file
    icmLoadProcessorMemory(processor, argv[1], False, False, True);
}
```

```
// run processor, one instruction at a time
while(simulate(processor, 1)) {
    // keep going while processor is still running
}

icmTerminate();

return 0;
}
```

This main function does the following:

1. It initializes the CpuManager interface by calling `icmInitPlatform`.
2. It creates a single instance of the processor by calling `icmNewProcessor`. The object file name and the name of the `vmiIASAttrs` object within that object file are specified by `MORPHER_FILE` and `MORPHER_SYMBOL`, defined in the platform makefile.
3. It loads a simulated application into the processor memory space. The application name is passed as the first argument when the platform is run. The application is an ELF format file and is loaded using `icmLoadProcessorMemory`.
4. It calls a routine `simulate` to simulate the processor one instruction at a time;
5. Finally, it calls `icmTerminate` to end the simulation.

The function `simulate` calls the CpuManager routine `icmSimulate` to simulate for a number of instructions, as follows:

```
static Bool simulate(icmProcessorP processor, Uns64 instructions) {

    icmStopReason stopReason = icmSimulate(processor, instructions);

    switch(stopReason) {

        case ICM_SR_SCHED:
            // hit the scheduler limit
            return True;

        case ICM_SR_EXIT:
            // processor has exited
            return False;

        case ICM_SR_FINISH:
            // simulation must end
            return False;

        default:
            icmPrintf("unimplemented stopReason %u\n", stopReason);
            return False;
    }
}
```

OVPsim/CpuManager usage is explained in detail in the *OVPsim and Imperas CpuManager User Guide*.

4.4 Creating an Application Test Case

A test case must be created using the application tool chain. Because the OR1K processor is supported by Imperas tools and shipped as an example, there is already an encapsulated tool chain that you can use to compile test cases for it.

Within the `or1kOutline` directory is a sample test case, `application/application.c`, which simply prints a message and exits. The application can be compiled by using this command in the `or1kOutline` directory:

```
make -C application
```

The result is an ELF format file for the OR1K called `application.OR1K.elf`.

4.5 Running the Application Test Case with the Processor Model

Having compiled the outline processor model, test platform and application, you are now ready to run a simulation. Do this by running:

```
platform/platform.IMPERAS_ARCH.exe --program application/application.OR1K.elf
```

in the `1.or1kOutline` directory. You should see the following output:

```
or1kConstructor called  
Warning (PC_NRI) No register information callback given for processor 'cpul'  
or1kDestructor called
```

The output from the test case shows debug messages from the processor constructor and destructor. There is also a warning message about a missing callback in the model attributes structure (the *register information callback*, used to identify registers of particular interest to the simulator, such as the program counter); this can be ignored at this stage. As yet, the outline model has no functionality so it exits on execution of the first instruction. The steps you need to perform to make the model execute the application correctly are covered in the following chapters.

5 Implementing the Instruction Decoder

A key component of every processor model is the instruction decoder. The result of the decoder is used by several other model components, specifically:

1. The *morpher*, which generates equivalent native code for each simulated instruction.
2. The *disassembler*, which creates a text string representation of an instruction.
3. The *next address function*, which determines the address of the next instruction after a given address (for processors with variable-size instructions only).

5.1 The Template Decoder Model

A template model for the OR1K processor with a decoder can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/2.or1kDecoder
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/2.or1kDecoder .
```

Compile the model using the `make` command:

```
cd 2.or1kDecoder
make
```

The processor model is based on the previous outline model, with the changes listed below.

5.1.1 File `or1kInstructions.h`

File `or1kInstructions.h` defines C macros that extract fields from OR1K instructions, which have a fixed width of 32 bits. In this example, macros are defined only for some of the arithmetic and bitwise instructions of the OR1K: the file will be expanded in subsequent chapters.

You may follow the template in `or1kInstructions.h` or use any other technique for decoding instructions appropriate to your model.

5.1.2 File `or1kDecode.h`

File `or1kDecode.h` defines types used by the OR1K decoder and the decode function itself, as follows:

```
#define OR1K_DISPATCH_FN(_NAME) void _NAME( \
    or1kP or1k,          \
    Addr  thisPC,         \
    Uns32 instr,          \
    void *userData        \
)
typedef OR1K_DISPATCH_FN(( *or1kDispatchFn ));
```

or1kDispatchFn is a generic function type used by the decoder to handle decoded instructions.

```
typedef enum or1kInstructionTypeE {  
  
    // arithmetic instructions  
    OR1K_IT_ADDI,  
    OR1K_IT_ADDIC,  
    OR1K_IT_ANDI,  
    OR1K_IT_ORI,  
    OR1K_IT_XORI,  
    OR1K_IT_MULI,  
  
    // KEEP LAST: for sizing the array  
    OR1K_IT_LAST  
  
} or1kInstructionType;
```

or1kInstructionType enumerates the instruction types that the decoder will find. The enumeration will grow to cover many instructions as the model develops. Currently, some simple binary operations are decoded.

```
typedef const or1kDispatchFn or1kDispatchTableC[OR1K_IT_LAST];  
typedef or1kDispatchTableC *or1kDispatchTableCP;
```

or1kDispatchTableC defines an array type used by clients of the decoder to associate a callback function with each instruction type. The decoder identifies the instruction type and then calls the appropriate callback function from the table.

```
Bool or1kDecode(  
    or1kP          or1k,  
    Uns32          thisPC,  
    or1kDispatchTableCP table,  
    or1kDispatchFn defaultCB,  
    void           *userData  
);
```

or1kDecode is the decoder function entry point. It decodes the OR1K instruction at address `thisPC` and, if the decode is successful, dispatches a call to the appropriate function from the dispatch table passed as the `table` argument and returns `True`. If no matching instruction is found, the function calls `defaultCB` and returns `False`.

5.1.3 File `or1kDecode.c`

File `or1kDecode.c` implements the OR1K instruction decoder using the decoder utility API defined in the VMI header file `vmi/vmiDecode.h`. The decoder API works as follows:

1. A new *decode table* is created using `vmidNewDecodeTable`.
2. A set of *decode entries* is added to the table using `vmidNewEntryFmtBin`. Each entry added identifies a single instruction type.
3. Instructions are decoded using a call to `vmidDecode`, passing the decode table and an instruction pattern, which returns an identifier describing the instruction type found.

Refer to the *Imperas VMI Morph Time Function Reference* for more detailed information about the decoder API.

File `orlkDecode.c` contains the following code:

```
#define DECODE_ENTRY(_PRIORITY, _NAME, _FMT) \
    vmidNewEntryFmtBin( \
        table, \
        #_NAME, \
        ORLK_IT_##_NAME, \
        _FMT, \
        _PRIORITY \
    )
```

This macro wraps a call to the decoder utility API function `vmidNewEntryFmtBin`, defined in `vmi/vmiDecode.h` as:

```
Bool vmidNewEntryFmtBin(
    vmidDecodeTableP table,
    const char      *name,
    Uns32           matchValue,
    const char      *format,
    Int32           priority
);
```

The arguments are:

1. A *decode table* into which to add a new decode entry;
2. A *name* for the new entry;
3. A *value to return if the entry matches* (typically an enumeration member, in this example a member of the `orlkEnstructionType` enumeration);
4. A *format string*, which specifies the bit pattern for a matching entry. Characters in this string may have the following meanings:
 - a. **0**: the corresponding bit in the instruction must be 0;
 - b. **1**: the corresponding bit in the instruction must be 1;
 - c. **| , / <space> <tab>**: formatting character (ignored);
 - d. **Any other character**: the corresponding bit can be either 1 or 0.
5. A *priority* for the entry. This allows instructions to be defined that are subsets of others. For example, a processor might have an instruction

```
move r1, r2
```

which actually decodes as

```
ori r1, r2, 0
```

(in other words, the `move` instruction is just a special case of the `ori` instruction). The above situation can be handled by adding two entries to the decode table, one for `ori` (with lower priority) and one for `move` (with higher priority).

The macro is used to create the decode table in function `createDecodeTable`:

```
static vmidDecodeTableP createDecodeTable(void) {
    vmidDecodeTableP table = vmidNewDecodeTable(32, ORLK_IT_LAST);

    // handle arithmetic instructions (second argument constant)
```

```

    DECODE_ENTRY(0, ADDI, " |100111..... |");
    DECODE_ENTRY(0, ADDIC, " |101000..... |");
    DECODE_ENTRY(0, ANDI, " |101001..... |");
    DECODE_ENTRY(0, ORI, " |101010..... |");
    DECODE_ENTRY(0, XORI, " |101011..... |");
    DECODE_ENTRY(0, MULI, " |101100..... |");

    return table;
}

```

The function creates a new decode table, specifying that the value `OR1K_IT_LAST` should be returned if there is no match for a particular instruction pattern.

The macro expansions refer to the `or1kInstructionType` enumeration in `or1kDecode.h`. For example, the line:

```
DECODE_ENTRY(0, ADDI, " |100111..... |");
```

Expands to:

```

vmidNewEntryFmtBin(
    table, "ADDI", " |100111..... |", OP_ADDI, 0
);

```

The instruction bit pattern specifies that the six most significant bits of an `ADDI` instruction are `'b100111'`, and the remaining 26 bits can be any value (indicated by the `'.'` character in the format). The vertical bar characters are for formatting only and have no significance as part of the pattern.

```

static or1kInstructionType decode(Uns32 instruction) {

    // get the OR1K decode table
    static vmidDecodeTableP decodeTable;
    if(!decodeTable) {
        decodeTable = createDecodeTable();
    }

    // decode the instruction to get the type
    or1kInstructionType type = vmidDecode(decodeTable, instruction);

    return type;
}

```

Function `decode` creates the decode table (if it has not already been created) and decodes the passed instruction using it.

```

Bool or1kDecode(
    or1kP          or1k,
    Uns32          thisPC,
    or1kDispatchTableCP table,
    or1kDispatchFn defaultCB,
    void          *userData
) {
    // get the instruction at the passed address - always 4 bytes on OR1K
    vmiProcessorP processor = (vmiProcessorP)or1k;
    Uns32 instruction = vmicxtFetch4Byte(processor, thisPC);
    or1kInstructionType type = decode(instruction);
}

```

```
// apply the callback, or the default if no match
if(type!=OR1K_IT_LAST) {
    ((*table)[type])(orlk, thisPC, instruction, userData);
    return True;
} else {
    defaultCB(orlk, thisPC, instruction, userData);
    return False;
}
}
```

The routine `orlkDecode` implements the decoder. It does the following:

1. It calls `vmicxtFetch4Byte` to get the four-byte instruction for the passed processor at the given address.
2. It calls the utility function `decode` to get the instruction type;
3. It calls the dispatch function from the table associated with the matching entry, returning `True` if a match is found.
4. If no matching entry is found, it calls `defaultCB` and returns `False`.

5.1.4 File `orlkMorph.c`

This file implements the OR1K morpher function. The morpher function is responsible for defining how each processor instruction should be evaluated. This is described in detail in section 7; this example does the following:

```
static OR1K_DISPATCH_FN(morphNOP) {
    // no action for a NOP
}
```

`morphNOP` is a dispatcher function that implements a NOP. In this example, it is used for each of the arithmetic functions in `orlkDecode.h`. This implies that each arithmetic function is currently implemented as a NOP.

```
static orlkDispatchTableC dispatchTable = {

    // handle arithmetic instructions
    [OR1K_IT_ADDI] = morphNOP,
    [OR1K_IT_ADDIC] = morphNOP,
    [OR1K_IT_ANDI] = morphNOP,
    [OR1K_IT_ORI] = morphNOP,
    [OR1K_IT_XORI] = morphNOP,
    [OR1K_IT_MULI] = morphNOP
};
```

This table associates function `morphNOP` with each of the instructions decoded by the decoder.

```
static void undecoded(
    vmiProcessorP processor,
    Uns32          thisPC,
    Uns32          instruction
) {
    // report that undecoded instruction was encountered
    vmiPrintf(
```



```
    "CPU '%s' 0x%08x:0x%08x *** undecoded instruction: exiting ***\n",
    vmirtProcessorName(processor),
    thisPC,
    instruction
);

// exit the CPU
vmirtExit(processor);
}
```

Function `undecoded` prints a message when an undecoded instruction is encountered and halts the current processor by calling function `vmirtExit` (defined in `vmiRt.h`, the VMI *Run Time Function API*).

```
static OR1K_DISPATCH_FN(morphDefault) {

    // print message warning about the undecoded instruction
    vmimtArgProcessor();
    vmimtArgUns32((Uns32)thisPC);
    vmimtArgUns32(instr);
    vmimtCall((vmiCallFn)undecoded);
}
```

Function `morphDefault` is a dispatcher function that is called for unimplemented instructions in `or1kDecode.c`. It creates native code to call the `undecoded` function previously defined. Code morphing is explained in detail in chapter 7.

```
VMI_MORPH_FN(or1kMorphInstruction) {
    or1kDecode((or1kP)processor, thisPC, &dispatchTable, morphDefault, 0);
}
```

Function `or1kMorphInstruction` is the morpher entry point (referenced in the `vmiIASAttrs` structure for this processor model, in `or1kAttrs.h`). It calls the decoder passing the dispatch table and default dispatch function.

5.1.5 File `or1kDisassemble.c`

This file implements the OR1K disassembler function. The disassembler function is responsible for generating a string disassembly of a given instruction. This is described in detail in section 6; this example does the following:

```
static OR1K_DISPATCH_FN(disDefault) {

    // get output buffer for disassembly
    char *result = (char *)userData;

    // default disassembly just shows instruction pattern
    sprintf(result, "??? instruction:0x%08x", instr);
}
```

Function `disDefault` writes text representing disassembled output from the instruction to the string buffer passed as the `userData` argument to the dispatcher function. In this example, a simple string giving the instruction hexadecimal value is written.

```
static or1kDispatchTableC dispatchTable = {
```

```
// handle arithmetic instructions
[OR1K_IT_ADDI] = disDefault,
[OR1K_IT_ADDIC] = disDefault,
[OR1K_IT_ANDI] = disDefault,
[OR1K_IT_ORI] = disDefault,
[OR1K_IT_XORI] = disDefault,
[OR1K_IT_MULI] = disDefault
};
```

This table associates function `disDefault` with each of the instructions decoded by the decoder.

```
VMI_DISASSEMBLE_FN(orlkDisassemble) {

    // static buffer to hold disassembly result
    static char result[256];

    // disassemble, filling the result buffer
    orlkDecode((orlkP)processor, thisPC, &dispatchTable, disDefault, &result);

    // return the result
    return result;
}
```

Function `orlkDisassemble` disassembles one instruction. It calls the decoder passing the dispatch table and default callback. It also passes the address of a static string buffer as the `userData` parameter to the decoder; this is used to store the result in `orlkDisDefault` above.

The prototype for `orlkDisassemble` is in file `orlkFunctions.h`, and is referenced in the `vmiIASAttr` structure defined in `orlkAttrs.c`. The reason for this will be seen when an application example is run using the new model.

5.2 *Running the Application Test Case with the Processor Model*

The compilation invoked for the decoder processor model will also compile up the default application and test platform. To compile them individually use

```
make -C platform
make -C application
```

The platform is almost identical to the previous example; the one difference is in the call to `icmNewProcessor` in `platform/platform.c`:

```
// create a processor
icmProcessorP processor = icmNewProcessor(
    "cpu1",           // CPU name
    TYPE_NAME,        // CPU type
    0,                // CPU cpuId
    0,                // CPU model flags
    32,               // address bits
    modelFile,        // model file
    "modelAttrs",     // morpher attributes
    ICM_ATTR_TRACE,   // enable tracing
```

```

    0,          // user-defined attributes
    0,          // semi-hosting file
    0           // semi-hosting attributes
);

```

The value `ICM_ATTR_TRACE` in the seventh argument enables a trace mode where the model disassembly routine is called just before each instruction is executed, which is why the disassembler routine was added to the model `vmiIASAttr` structure.

Run:

```
platform/platform.IMPERAS_ARCH.exe --program application/application.OR1K.elf
```

in the `2.or1kDecoder` directory. You should see the following output:

```

orkConstructor called
Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 'cpul', 0x0000000000000100: ??? instruction:0x9c400000
Info 'cpul', 0x0000000000000104: ??? instruction:0x9c600000
Info 'cpul', 0x0000000000000108: ??? instruction:0x9c800000
Info 'cpul', 0x000000000000010c: ??? instruction:0x9ca00000
Info 'cpul', 0x0000000000000110: ??? instruction:0x9cc00000
Info 'cpul', 0x0000000000000114: ??? instruction:0x9ce00000
Info 'cpul', 0x0000000000000118: ??? instruction:0x9d000000
Info 'cpul', 0x000000000000011c: ??? instruction:0x9d200000
Info 'cpul', 0x0000000000000120: ??? instruction:0x9d400000
Info 'cpul', 0x0000000000000124: ??? instruction:0x9d600000
Info 'cpul', 0x0000000000000128: ??? instruction:0x9d800000
Info 'cpul', 0x000000000000012c: ??? instruction:0x9da00000
Info 'cpul', 0x0000000000000130: ??? instruction:0x9dc00000
Info 'cpul', 0x0000000000000134: ??? instruction:0x9de00000
Info 'cpul', 0x0000000000000138: ??? instruction:0x9e000000
Info 'cpul', 0x000000000000013c: ??? instruction:0x9e200000
Info 'cpul', 0x0000000000000140: ??? instruction:0x9e400000
Info 'cpul', 0x0000000000000144: ??? instruction:0x9e600000
Info 'cpul', 0x0000000000000148: ??? instruction:0x9e800000
Info 'cpul', 0x000000000000014c: ??? instruction:0x9ea00000
Info 'cpul', 0x0000000000000150: ??? instruction:0x9ec00000
Info 'cpul', 0x0000000000000154: ??? instruction:0x9ee00000
Info 'cpul', 0x0000000000000158: ??? instruction:0x9f000000
Info 'cpul', 0x000000000000015c: ??? instruction:0x9f200000
Info 'cpul', 0x0000000000000160: ??? instruction:0x9f400000
Info 'cpul', 0x0000000000000164: ??? instruction:0x9f600000
Info 'cpul', 0x0000000000000168: ??? instruction:0x9f800000
Info 'cpul', 0x000000000000016c: ??? instruction:0x9fa00000
Info 'cpul', 0x0000000000000170: ??? instruction:0x9fc00000
Info 'cpul', 0x0000000000000174: ??? instruction:0x9fe00000
Info 'cpul', 0x0000000000000178: ??? instruction:0x1820ffff
CPU 'cpul' 0x00000178:0x1820ffff *** undecoded instruction: exiting ***
orkDestructor called

```

After the constructor line, there is a line of trace output for every instruction that was successfully decoded by the decoder in this example. Each trace line gives the instruction address (starting with `0x100`, the start address specified in the ELF file) and the instruction disassembly, produced using the disassembler we defined. At address `0x178`, the processor encounters the first instruction not recognized by the decoder and is halted. In the next chapter, we will see how to terminate simulation more elegantly than this.

5.3 *More Complex Decoders*

The decoder in this OR1K example is quite simple because the OR1K instruction set has a small number of similar instructions of constant size (32 bits). CISC processors with variable-length instructions require a more sophisticated decoder. A good approach is to have multiple decode tables (a level 1 table is used to decode the first byte of the instruction, on the basis of this alternate level 2 tables are used, and so on). When decoders are complex, it is often useful to fill a data structure with information about the decoded instruction to use in later stages (instruction translation and disassembly); see the ARC 600/700, ARM or MIPS processor models on the ovpworld.com website for good examples.

All instruction fetches performed by a decoder should use `vmicxtFetch4Byte`, or related routines defined in `vmiCxt.h`. A single decode may perform several calls to fetch routines if required: in the example of the CISC processor, there may be an initial call to `vmicxtFetch1Byte` to fetch the first byte of an instruction, then a possible further call to `vmicxtFetch1Byte` to fetch the next byte, and so on.

6 Implementing the Instruction Disassembler

Having implemented an initial decoder framework, the next step is to start implementing the details of the instruction disassembler.

6.1 *The Template Disassembler Model*

A template model for the OR1K processor with a decoder and disassembler can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/3.or1kDisassembler
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/3.or1kDisassembler .
```

Compile the model using make:

```
cd 3.or1kDisassembler
make
```

The processor model is based on the previous decoder model, with the changes listed below.

6.1.1 File `or1kDisassemble.c`

`or1kDisassemble.c` now contains implementations of the arithmetic instruction disassembly routines, as follows:

```
static OR1K_DISPATCH_FN(disADDI) {doBinopSLit(userData, instr, "l.addi");}
static OR1K_DISPATCH_FN(disADDIC) {doBinopSLit(userData, instr, "l.addic");}
static OR1K_DISPATCH_FN(disANDI) {doBinopULit(userData, instr, "l.andi");}
static OR1K_DISPATCH_FN(disORI) {doBinopULit(userData, instr, "l.ori");}
static OR1K_DISPATCH_FN(disXORI) {doBinopSLit(userData, instr, "l.xori");}
static OR1K_DISPATCH_FN(disMULI) {doBinopSLit(userData, instr, "l.muli");}
```

Each of these disassembler callback functions handles a separate instruction type. They make use of two utility routines, `doBinopSLit` and `doBinopULit`, which produce disassembly for sign-extending and non-sign-extending 16-bit literal constants in instructions, respectively:

```
//
// Disassemble a binary/signed literal OR1K instruction
//
static void doBinopSLit(char *buffer, Uns32 instr, char *orlkop) {

    Uns32 rd = OP2_D(instr);
    Uns32 ra = OP2_A(instr);
    Int16 i = OP2_I(instr);

    sprintf(buffer, "%-8s r%u,r%u,0x%x", orlkop, rd, ra, i);
}
```

```
//
// Disassemble a binary/unsigned literal ORlK instruction
//
static void doBinopULit(char *buffer, Uns32 instr, char *orlkop) {

    Uns32 rd = OP2_D(instr);
    Uns32 ra = OP2_A(instr);
    Uns16 k  = OP2_I(instr);

    sprintf(buffer, "%-8s r%u,r%u,0x%x", orlkop, rd, ra, k);
}

```

Note that this example uses fixed-width types (Uns16, Uns32 etc). These types are defined in the include file:

```
$IMPERAS_HOME/ImpPublic/include/host/impTypes.h.
```

```
static orlkDispatchTableC dispatchTable = {

    // handle arithmetic instructions
    [ORlK_IT_ADDI] = disADDI,
    [ORlK_IT_ADDIC] = disADDIC,
    [ORlK_IT_ANDI] = disANDI,
    [ORlK_IT_ORI] = disORI,
    [ORlK_IT_XORI] = disXORI,
    [ORlK_IT_MULI] = disMULI
};

```

The dispatch table associates the new disassembly callback functions with the instructions decoded by the decoder.

6.2 *Running the Application Test Case with the Processor Model*

The compilation invoked for the decoder processor model will also compile up the default application and test platform. To compile them individually use

```
make -C platform
make -C application

```

The platform is identical to the previous example. Run:

```
platform/platform.IMPERAS_ARCH.exe --program application/application.ORlK.elf

```

in the 3.orlkDisassembler directory. You should see the following output:

```
orlkConstructor called
Warning (PC_NRI) No register information callback given for processor 'cpu1'
Info 'cpu1', 0x0000000000000100: l.addi    r2,r0,0x0
Info 'cpu1', 0x0000000000000104: l.addi    r3,r0,0x0
Info 'cpu1', 0x0000000000000108: l.addi    r4,r0,0x0
Info 'cpu1', 0x000000000000010c: l.addi    r5,r0,0x0
Info 'cpu1', 0x0000000000000110: l.addi    r6,r0,0x0
Info 'cpu1', 0x0000000000000114: l.addi    r7,r0,0x0
Info 'cpu1', 0x0000000000000118: l.addi    r8,r0,0x0

```

```

Info 'cpu1', 0x0000000000000011c: l.addi    r9,r0,0x0
Info 'cpu1', 0x00000000000000120: l.addi    r10,r0,0x0
Info 'cpu1', 0x00000000000000124: l.addi    r11,r0,0x0
Info 'cpu1', 0x00000000000000128: l.addi    r12,r0,0x0
Info 'cpu1', 0x0000000000000012c: l.addi    r13,r0,0x0
Info 'cpu1', 0x00000000000000130: l.addi    r14,r0,0x0
Info 'cpu1', 0x00000000000000134: l.addi    r15,r0,0x0
Info 'cpu1', 0x00000000000000138: l.addi    r16,r0,0x0
Info 'cpu1', 0x0000000000000013c: l.addi    r17,r0,0x0
Info 'cpu1', 0x00000000000000140: l.addi    r18,r0,0x0
Info 'cpu1', 0x00000000000000144: l.addi    r19,r0,0x0
Info 'cpu1', 0x00000000000000148: l.addi    r20,r0,0x0
Info 'cpu1', 0x0000000000000014c: l.addi    r21,r0,0x0
Info 'cpu1', 0x00000000000000150: l.addi    r22,r0,0x0
Info 'cpu1', 0x00000000000000154: l.addi    r23,r0,0x0
Info 'cpu1', 0x00000000000000158: l.addi    r24,r0,0x0
Info 'cpu1', 0x0000000000000015c: l.addi    r25,r0,0x0
Info 'cpu1', 0x00000000000000160: l.addi    r26,r0,0x0
Info 'cpu1', 0x00000000000000164: l.addi    r27,r0,0x0
Info 'cpu1', 0x00000000000000168: l.addi    r28,r0,0x0
Info 'cpu1', 0x0000000000000016c: l.addi    r29,r0,0x0
Info 'cpu1', 0x00000000000000170: l.addi    r30,r0,0x0
Info 'cpu1', 0x00000000000000174: l.addi    r31,r0,0x0
Info 'cpu1', 0x00000000000000178: ??? instruction:0x1820ffff
CPU 'cpu1' 0x000000178:0x1820ffff *** undecoded instruction: exiting ***

```

This reveals that the first instructions executed in the application are OR1K `addi` instructions, which clear the processor GPRs.

6.3 Creating Disassembler Point Tests

When a class of instructions has been added to a decoder (such as the arithmetic instructions above) it is good practice to fully test the disassembly behavior of the entire class before implementing any behavior for that instruction class. This is done most easily by assembler-level tests. File `asmtest.S` in directory

`3.or1kDisassembler/application` is an OR1K assembler file that is a good starting point for a disassembler test:

```

.global _start
_start:
    l.addi    r1,r2,0
    l.addi    r1,r2,1
    l.addi    r1,r2,-1
    l.addic   r1,r2,1
    l.addic   r1,r2,-1
    l.addic   r1,r2,0
    l.andi    r1,r2,1
    l.andi    r1,r2,-1
    l.andi    r1,r2,0
    l.ori     r1,r2,1
    l.ori     r1,r2,-1
    l.ori     r1,r2,0
    l.xori    r1,r2,1
    l.xori    r1,r2,-1
    l.xori    r1,r2,0
    l.muli    r1,r2,1
    l.muli    r1,r2,-1
    l.muli    r1,r2,0

.global exit
exit:

```

```
l.add      r1,r2,0
```

Assemble this file using:

```
cd application
make asmtest.OR1K.elf
cd ..
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```
orlkConstructor called
Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 'cpul', 0x0000000001000074: l.addi    r1,r2,0x0
Info 'cpul', 0x0000000001000078: l.addi    r1,r2,0x1
Info 'cpul', 0x000000000100007c: l.addi    r1,r2,0xffffffff
Info 'cpul', 0x0000000001000080: l.addic   r1,r2,0x1
Info 'cpul', 0x0000000001000084: l.addic   r1,r2,0xffffffff
Info 'cpul', 0x0000000001000088: l.addic   r1,r2,0x0
Info 'cpul', 0x000000000100008c: l.andi    r1,r2,0x1
Info 'cpul', 0x0000000001000090: l.andi    r1,r2,0xffff
Info 'cpul', 0x0000000001000094: l.andi    r1,r2,0x0
Info 'cpul', 0x0000000001000098: l.ori     r1,r2,0x1
Info 'cpul', 0x000000000100009c: l.ori     r1,r2,0xffff
Info 'cpul', 0x00000000010000a0: l.ori     r1,r2,0x0
Info 'cpul', 0x00000000010000a4: l.xori    r1,r2,0x1
Info 'cpul', 0x00000000010000a8: l.xori    r1,r2,0xffffffff
Info 'cpul', 0x00000000010000ac: l.xori    r1,r2,0x0
Info 'cpul', 0x00000000010000b0: l.muli    r1,r2,0x1
Info 'cpul', 0x00000000010000b4: l.muli    r1,r2,0xffffffff
Info 'cpul', 0x00000000010000b8: l.muli    r1,r2,0x0
Info 'cpul', 0x00000000010000bc: l.addi    r1,r2,0x0
Processor 'cpul' terminated at 'exit', address 0x10000bc
orlkDestructor called
```

It is good practice to make the output from the disassembler conform as closely to possible to the output generated by existing tools (for example, the OR1K objdump executable). This simplifies verification because output generated by the disassembler can be automatically compared against a golden log generated by the existing tool.

6.3.1 Elegant Test Termination using Semihosting

Note that the assembler test terminated more elegantly than the previous run: instead of:

```
CPU 'cpul' 0x00000178:0x0400037b *** undecoded instruction: exiting ***
```

We saw:

```
Processor 'cpul' terminated at 'exit', address 0x10000bc
```

This was possible because the test platform used with this example was modified to use *semihosting*, which will be briefly introduced here and covered in detail in chapter 22.

Imperas semihosting allows the default behavior of specified functions or instructions to be modified using a semihosting shared object library that is loaded by the simulator in

addition to the processor model. In this case, we defined a global label, `exit`, on the last instruction of the assembler test. This label can be used in conjunction with a standard Imperas semihosting shared object library, located at the following location under `$IMPERAS_HOME`:

```
$IMPERAS_LIB/ImperasLib/ovpworld.org/modelSupport/imperasExit/1.0/model.$SHRSUF
```

What this semihosting library does is terminate simulation immediately after any instruction labeled `exit`. To use the semihosting library, `platform/platform.c` has been modified as follows to select the `imperasExit` semihost library from the VLNV library and load it as part of the processor instantiation:

```
const char *modelFile      = "model." IMPERAS_SHRSUF;
const char *seminhostFile = icmGetVlnvString(NULL, "ovpworld.org",
                                              "modelSupport", "imperasExit", "1.0", "model");

// create a processor
icmProcessorP processor = icmNewProcessor(
    "cpu1",           // CPU name
    TYPE_NAME,       // CPU type
    0,               // CPU cpuId
    0,               // CPU model flags
    32,              // address bits
    modelFile,       // model file
    "modelAttrs",    // morpher attributes
    ICM_ATTR_TRACE,  // enable tracing or register values
    0,               // user-defined attributes
    seminhostFile,   // semi-hosting file
    "modelAttrs"     // semi-hosting attributes
);
```

You may use the `imperasExit` semihosting library with any processor model: it is not specific to the OR1K processor we are creating here.

7 Implementing Simple Behavior

When the processor decoder and disassembler are working correctly for a subset of processor instructions, you can start to implement behavior for those instructions. This chapter shows how this is done for simple instructions using the Imperas *code morphing* technology.

7.1 An Introduction to Code Morphing

Conventional processor models written in an HDL or similar modeling language might be implemented by a loop that is activated by a clock signal. On each activation of the clock, the model might fetch the next instruction, decode it, and call specific functions to perform the instruction (update model registers, read and write memory, and so on). If the model is cycle-accurate, there may be further complications of modeling pipelines, branch prediction and so on.

Although models written in this conventional style can be accurate and straightforward in structure, they are not fast: even a simple instruction accurate model written in C will probably run no faster than a few million instructions per second. Unfortunately, platform testing may require the execution of billions of instructions, which makes this style of model too slow.

Processor models designed for the Imperas tool set instead use just-in-time (JIT) *code morphing* technology. This works as follows:

1. As each new processor instruction is encountered during program execution, the instruction is translated (morphed) into equivalent native machine code. The exact translations to be made are specified by the processor modeler using the Imperas Virtual Machine Interface (VMI) API.
2. Contiguous sections of translated processor instructions are gathered into *code blocks*, which are held in a *dictionary* for the processor.
3. If a processor performs a jump to a simulated address that has already been translated to a code block held in the dictionary, there is no need to perform the translation again: the simulator simply re-executes the existing code block.

Imperas technology handles the generation of native machine code and the efficient management of code blocks and dictionaries to give extremely fast simulation. Depending on the complexity of the processor being simulated, speeds of many hundreds of processor MIPS can be achieved. This is possible because, as simulation proceeds, *run time* (execution of translated code blocks) dominates *morph time* (JIT compilation).

To support the JIT compiler, you must implement the *morpher*, which is responsible for defining how each processor instruction should be evaluated.

7.2 The Template Simple Behavioral Model

A template model for the OR1K processor with a decoder, disassembler and behavior can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/4.or1kBehaviorSimple
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/4.or1kBehaviorSimple .
```

Compile the model using make:

```
cd 4.or1kBehaviorSimple
make
```

The processor model is based on the previous disassembler model, with the changes listed in following sections.

7.2.1 File `or1kStructure.h`

The processor structure defined in file `or1kStructure.h` is where you define the registers and other state of the model. For this example, we need to model the 32 OR1K general-purpose registers. The structure is therefore declared like this:

```
#define OR1K_REGS 32          // basic OR1K registers
#define OR1K_BITS 32         // register size in bits

// processor structure
typedef struct or1kS {
    Uns32      regs[OR1K_REGS]; // basic registers
    vmiBusPortP busPorts;       // bus port descriptions
} or1k, *or1kP;
```

The OR1K general purpose registers are declared as a C array of `Uns32` values, `regs`.

As we will see in section 7.2.3, Imperas API routines for generating morphed code need to know about the *register byte offsets* of the register fields within the processor structure. In this case, the C structure offsets are as follows:

Register	Byte Offset
<code>regs[0]</code>	0
<code>regs[1]</code>	4
<code>regs[2]</code>	8
... etc ...	

To simplify calculation of these offsets, `or1kStructure.h` defines the following macros for use in *variable* C expressions:

```
// macros to specify target registers in VARIABLE expressions
#define OR1K_OFFSET(_F)      VMI_CPU_REG(or1kP, _F)
```

```
#define OR1K_REG(_R)          OR1K_OFFSET(regs[_R])
```

For example:

```
OR1K_REG(3)
```

Is used to identify OR1K general purpose register `r3` in morpher VMI API calls (see section 7.2.3 for examples).

7.2.2 File `or1kMain.c`

Now that the OR1K structure has real state implemented, the constructor in file `or1kMain.c` should be upgraded to initialize that state.

The constructor is called by the simulator whenever a new instance of the OR1K processor model is created. It is passed a pointer to the new processor model instance using a generic parameter called `processor` of type `vmiProcessorP`. In order to initialize the processor, the generic `processor` pointer should be cast to a specific `or1kP` pointer so that fields in the structure can be set.

By default, the new processor model instance is entirely zeroed out. In this example, we initialize every general purpose register in the OR1K from `r2` to `r31` with the pattern `0xdeadbeef`:

```
VMI_CONSTRUCTOR_FN(or1kConstructor) {  
  
    or1kP or1k = (or1kP)processor;  
    Uns32 i;  
  
    for(i=2; i<OR1K_REGS; i++) {  
        or1k->regs[i] = 0xdeadbeef;  
    }  
}
```

Register `r0` is left unmodified (zeroed out) because it is hardwired to zero in the OR1K. Register `r1` is also left zeroed out because this is the stack pointer register, implicitly initialized to zero.

(Note that the constructor and destructor no longer print that they have been called.)

7.2.3 File `or1kMorph.c`

`or1kMorph.c` now contains implementations of the arithmetic instruction code morphing routines, as follows:

```
static OR1K_DISPATCH_FN(morphADDI) {doBinopSLit(instr, vmi_ADD);}  
static OR1K_DISPATCH_FN(morphADDIC) {doBinopSLit(instr, vmi_ADC);}  
static OR1K_DISPATCH_FN(morphANDI) {doBinopULit(instr, vmi_AND);}  
static OR1K_DISPATCH_FN(morphORI) {doBinopULit(instr, vmi_OR);}  
static OR1K_DISPATCH_FN(morphXORI) {doBinopSLit(instr, vmi_XOR);}  
static OR1K_DISPATCH_FN(morphMULI) {doBinopSLit(instr, vmi_IMUL);}
```

Each of these callback functions handles a separate instruction type, and is associated with decoded instructions using `dispatchTable`:

```
static orlKDispatchTableC dispatchTable = {  
  
    // handle arithmetic instructions  
    [OR1K_IT_ADDI] = morphADDI,  
    [OR1K_IT_ADDIC] = morphADDIC,  
    [OR1K_IT_ANDI] = morphANDI,  
    [OR1K_IT_ORI] = morphORI,  
    [OR1K_IT_XORI] = morphXORI,  
    [OR1K_IT_MULI] = morphMULI  
};
```

In a similar fashion to the disassembler, the code morphing routines make use of two utility routines, `doBinopSLit` and `doBinopULit`, which implement behavior for sign-extending and non-sign-extending 16-bit literal constants in instructions, respectively:

```
//  
// Emit code to implement a binary/signed literal OR1K instruction  
//  
static void doBinopSLit(Uns32 instr, vmiBinop op) {  
  
    Uns32 rd      = OP2_D(instr);  
    Uns32 ra      = OP2_A(instr);  
    Int16 i       = OP2_I(instr);  
    vmiReg target = (rd==0) ? VMI_NOREG : OR1K_REG(rd);  
  
    if(ra==0) {  
        vmimtBinopRCC(OR1K_BITS, op, target, 0, i, 0);  
    } else {  
        vmimtBinopRRC(OR1K_BITS, op, target, OR1K_REG(ra), i, 0);  
    }  
}  
  
//  
// Emit code to implement a binary/unsigned literal OR1K instruction  
//  
static void doBinopULit(Uns32 instr, vmiBinop op) {  
  
    Uns32 rd      = OP2_D(instr);  
    Uns32 ra      = OP2_A(instr);  
    Uns16 k       = OP2_I(instr);  
    vmiReg target = (rd==0) ? VMI_NOREG : OR1K_REG(rd);  
  
    if(ra==0) {  
        vmimtBinopRCC(OR1K_BITS, op, target, 0, k, 0);  
    } else {  
        vmimtBinopRRC(OR1K_BITS, op, target, OR1K_REG(ra), k, 0);  
    }  
}
```

These functions use routines from the Imperas *morph time function* API (`vmiMt.h`) to *describe the behavior of the arithmetic instructions we are implementing in this example*. The destination register number (`rd`), argument register number (`ra`) and constant value (`k`) are encoded as fields in the instruction. It is very important to understand that `vmimt-`prefixed routines do not themselves perform any arithmetic operation on the processor registers: instead, *they describe the action to be performed*. The action descriptions are

used as input to the Imperas JIT compiler to generate native code that, when executed, performs the required arithmetic operation.

To further clarify this example, we will consider lines from `doBinopULit` in detail.

```
Uns32 rd    = OP2_D(instr);
Uns32 ra    = OP2_A(instr);
Uns16 k     = OP2_I(instr);
```

These three lines extract fields of interest from the instruction, using macros defined in `or1kInstruction.h`.

```
vmiReg target = (rd==0) ? VMI_NOREG : OR1K_REG(rd);
```

In the OR1K processor, register `r0` is hardwired to the constant value 0. Any attempt to write to this register should be discarded: this is indicated to the morph time API functions by using the special value `VMI_NOREG` as a target register. If the target register is writable, the macro `OR1K_REG(rd)` from `or1kStructure.h` is used to specify it, as described in section 7.2.1.

```
if(ra==0) {
    vmimtBinopRCC(OR1K_BITS, op, target, 0, k, 0);
```

If `r0` is used for argument register `ra`, the action of the instruction can be simplified to

```
rd := 0 <op> k
```

Operations of this form are described by the morph time API function `vmimtBinopRCC`. The operation `op` can be any of the operations specified in the `vmiBinop` enumeration, declared in file `vmiTypes.h`:

```
typedef enum {
    // ARITHMETIC OPERATIONS
    vmi_ADD,      // d <- a + b
    vmi_ADC,      // d <- a + b + C
    vmi_SUB,      // d <- a - b
    vmi_SBB,      // d <- a - b - C
    vmi_RSBB,     // d <- b - a - C
    vmi_RSUB,     // d <- b - a
    vmi_IMUL,     // d <- a * b (signed)
    vmi_MUL,      // d <- a * b (unsigned)
    vmi_IDIV,     // d <- a / b (signed)
    vmi_DIV,      // d <- a / b (unsigned)
    vmi_CMP,      // a - b

    // SATURATED ARITHMETIC OPERATIONS
    vmi_ADDSQ,    // d <- saturate_signed(a + b)
    vmi_SUBSQ,    // d <- saturate_signed(a - b)
    vmi_RSUBSQ,   // d <- saturate_signed(b - a)
    vmi_ADDUQ,    // d <- saturate_unsigned(a + b)
    vmi_SUBUQ,    // d <- saturate_unsigned(a - b)
    vmi_RSUBUQ,   // d <- saturate_unsigned(b - a)

    // HALVING ARITHMETIC OPERATIONS
```

```

vmi_ADDSH,      // d <- ((signed)(a + b)) / 2
vmi_SUBSH,      // d <- ((signed)(a - b)) / 2
vmi_RSUBSH,     // d <- ((signed)(b - a)) / 2
vmi_ADDUH,      // d <- ((unsigned)(a + b)) / 2
vmi_SUBUH,      // d <- ((unsigned)(a - b)) / 2
vmi_RSUBUH,     // d <- ((unsigned)(b - a)) / 2
vmi_ADDSHR,     // d <- round(((signed)(a + b)) / 2)
vmi_SUBSHR,     // d <- round(((signed)(a - b)) / 2)
vmi_RSUBSHR,    // d <- round(((signed)(b - a)) / 2)
vmi_ADDUHR,     // d <- round(((unsigned)(a + b)) / 2)
vmi_SUBUHR,     // d <- round(((unsigned)(a - b)) / 2)
vmi_RSUBUHR,    // d <- round(((unsigned)(b - a)) / 2)

// BITWISE OPERATIONS
vmi_OR,         // d <- a | b
vmi_AND,        // d <- a & b
vmi_XOR,        // d <- a ^ b
vmi_ORN,        // d <- a | ~b
vmi_ANDN,       // d <- a & ~b
vmi_XORN,       // d <- a ^ ~b
vmi_NOR,        // d <- ~(a | b)
vmi_NAND,       // d <- ~(a & b)
vmi_XNOR,       // d <- ~(a ^ b)

// SHIFT/ROTATE OPERATIONS
vmi_ROL,        // d <- a << b | a >> <bits>-b
vmi_ROR,        // d <- a >> b | a << <bits>-b
vmi_RCL,        // (d,c) <- (a,c) << b | (a,c) >> <bits>-b
vmi_RCR,        // (d,c) <- (a,c) >> b | (a,c) << <bits>-b
vmi_SHL,        // d <- a << b
vmi_SHR,        // d <- (unsigned)a >> b
vmi_SAR,        // d <- (signed)a >> b
} vmiBinop;

```

The final part of function `doBinopULit` handles the general case where `ra` does not specify OR1K register `r0`:

```

} else {
    vmimtBinopRRC(OR1K_BITS, op, target, OR1K_REG(ra), k, 0);
}

```

Operations of this form are described by the routine `vmimtBinopRRC`.

Refer to the *Imperas VMI Morph Time Reference* manual for more detailed information on all of the morph-time functions available in this API.

7.2.4 File `platform/platform.c`

The test platform for this example, `platform/platform.c`, has been changed as follows:

```

#define MODEL_FLAGS (ICM_ATTR_TRACE | ICM_ATTR_TRACE_REGS_AFTER)

// create a processor
icmProcessorP processor = icmNewProcessor(
    "cpu1",           // CPU name
    TYPE_NAME,       // CPU type
    0,               // CPU cpuId
    0,               // CPU model flags
    32,              // address bits

```

```
    modelFile,          // model file
    "modelAttrs",       // morpher attributes
    MODEL_FLAGS,        // enable tracing or register values
    0,                  // user-defined attributes
    semihostFile,       // semi-hosting file
    "modelAttrs"        // semi-hosting attributes
);
```

We have added the value `ICM_ATTR_TRACE_REGS_AFTER` to the flags passed to the processor instance. This enables dumping of the processor state after each instruction is executed.

7.3 *Running the Application Test Case with the Processor Model*

The compilation of the simple behavioral processor model will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run using:

```
platform/platform.IMPERAS_ARCH.exe --program application/application.OR1K.elf
```

in the `4.or1kBehaviorSimple` directory. You should see the following output:

```
Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 'cpul', 0x0000000000000100: l.addi    r2,r0,0x0
Info 'cpul' REGISTERS
CPU cpul (instruction 1):
    0: 00000000 00000000 00000000 deadbeef
   16: deadbeef deadbeef deadbeef deadbeef
   32: deadbeef deadbeef deadbeef deadbeef
   48: deadbeef deadbeef deadbeef deadbeef
   64: deadbeef deadbeef deadbeef deadbeef
   80: deadbeef deadbeef deadbeef deadbeef
   96: deadbeef deadbeef deadbeef deadbeef
  112: deadbeef deadbeef deadbeef deadbeef
  128: 08daa790
Info 'cpul', 0x0000000000000104: l.addi    r3,r0,0x0
Info 'cpul' REGISTERS
CPU cpul (instruction 2):
    0: 00000000 00000000 00000000 00000000
   16: deadbeef deadbeef deadbeef deadbeef
   32: deadbeef deadbeef deadbeef deadbeef
   48: deadbeef deadbeef deadbeef deadbeef
   64: deadbeef deadbeef deadbeef deadbeef
   80: deadbeef deadbeef deadbeef deadbeef
   96: deadbeef deadbeef deadbeef deadbeef
  112: deadbeef deadbeef deadbeef deadbeef
  128: 08daa790
Info 'cpul', 0x0000000000000108: l.addi    r4,r0,0x0
Info 'cpul' REGISTERS
CPU cpul (instruction 3):
    0: 00000000 00000000 00000000 00000000
   16: 00000000 deadbeef deadbeef deadbeef
```



```

32: deadbeef deadbeef deadbeef deadbeef
48: deadbeef deadbeef deadbeef deadbeef
64: deadbeef deadbeef deadbeef deadbeef
80: deadbeef deadbeef deadbeef deadbeef
96: deadbeef deadbeef deadbeef deadbeef
112: deadbeef deadbeef deadbeef deadbeef
128: 08daa790

... (many lines cut) ...

Info 'cpul', 0x00000000000000178: ??? instruction:0x1820ffff
CPU 'cpul' 0x000000178:0x1820ffff *** undecoded instruction: exiting ***
Info 'cpul' REGISTERS
CPU cpul (instruction 31):
  0: 00000000 00000000 00000000 00000000
 16: 00000000 00000000 00000000 00000000
 32: 00000000 00000000 00000000 00000000
 48: 00000000 00000000 00000000 00000000
 64: 00000000 00000000 00000000 00000000
 80: 00000000 00000000 00000000 00000000
 96: 00000000 00000000 00000000 00000000
112: 00000000 00000000 00000000 00000000
128: 08daa790

```

We now see the processor model executing instructions for the first time. In detail, the sequence when generating output is¹:

1. The instruction about to be executed is disassembled;
2. The instruction is executed;
3. The register state of the processor is dumped.

Because this ORK1 model has no register dump functionality specified at this point, the register values printed after each instruction are simply the raw contents of the OR1K structure. This includes the pointer value, `busPorts`, at offset 128. Because this is a pointer, its value will change from run to run. The next chapter describes how a model-specific register dump callback is written.

The initial instructions of the application zero out registers `r2-r31` of the OR1K processor using `l.addi` instructions. As this happens, we see each register value change from `0xdeadbeef` (set in the processor constructor) to `0x00000000`.

7.4 Instruction Temporaries

Some instructions cannot be implemented as a single VMI operation and instead require a sequence of operations and intermediate *temporaries* to generate the correct result. For

¹ In older versions of OVPsim, the processor flag `ICM_ATTR_TRACE_REGS` was used to enable register dumping. This flag caused register values to be dumped *before* the trace and instruction execution, i.e. the sequence for each instruction was:

1. dump register state;
2. disassemble instruction to be executed;
3. execute the instruction

This sequence can still be obtained using the new processor flag `ICM_ATTR_TRACE_REGS_BEFORE` if required. The old flag `ICM_ATTR_TRACE_REGS` is now a deprecated alias for `ICM_ATTR_TRACE_REGS_AFTER`.

example, suppose that there is a *signed halfword multiply* instruction, which works according to the following pseudo-code:

```
T132..0 = sign_extend(R115..0);
T232..0 = sign_extend(R215..0);
R332..0 = T132..0 * T232..0;
```

In other words, the instruction sign-extends the lower half of the two arguments and then multiplies those sign-extended values to produce the result. Implementing this instruction requires the use of two temporaries that are not true processor registers but instead represent intermediate values that are required only within an instruction.

The way to implement this is to introduce two new pseudo-registers into the processor structure as follows:

```
#define OR1K_REGS 32          // basic OR1K registers
#define OR1K_BITS 32         // register size in bits
#define OR1K_TNUM 2          // number of temporaries

// processor structure
typedef struct or1kS {
    Uns32 regs[OR1K_REGS]; // basic registers
    Uns32 temp[OR1K_TNUM]; // temporary pseudo-registers
    vmiBusPortP busPorts;  // bus port descriptions
} or1k, *or1kP;
```

The temporaries are specified to the morpher as follows:

```
// macros to specify target registers in VARIABLE expressions
#define OR1K_OFFSET(_F)      VMI_CPU_REG(or1kP, _F)
#define OR1K_REG(_R)         OR1K_OFFSET(regs[_R])
#define OR1K_OFFSET_TEMP(_F) VMI_CPU_TEMP(or1kP, _F)
#define OR1K_TEMP(_I)        OR1K_OFFSET_TEMP(temp[_I])
```

The macro VMI_CPU_TEMP identifies *temporaries* in exactly the same way that macro VMI_CPU_REG identifies *true registers*. Because the morpher knows that these values are temporaries and not true registers, it can generate more efficient code (the temporary values do not need to be written back to the processor structure at the end of the instruction).

These temporaries could then be used to implement the signed halfword multiply instruction as follows:

```
...
vmiReg target = (rd==0) ? VMI_NOREG : OR1K_REG(rd);

// generate intermediates
vmimtMoveExtendRR(OR1K_BITS, OR1K_TEMP(0), OR1K_BITS/2, OR1K_REG(ra), True);
vmimtMoveExtendRR(OR1K_BITS, OR1K_TEMP(1), OR1K_BITS/2, OR1K_REG(rb), True);

// generate result
vmimtBinopRRR(OR1K_BITS, vmi_IMUL, target, OR1K_TEMP(0), OR1K_TEMP(1), 0);
```

8 Processor Flags and Register Dumping

In general, arithmetic operations can both take as input and generate as output *flag values*. For example an add-with-carry operation has a carry flag input, and might generate a carry flag output. This chapter enhances the previous simple behavioral model to handle flag values for arithmetic instructions, and shows how to implement a model specific *register dump* routine to simplify model validation.

8.1 The Template Flags Model

A template model for the OR1K processor implementing instruction flags can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/5.or1kBehaviorFlags
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/5.or1kBehaviorFlags .
```

Compile the model using make:

```
cd 5.or1kBehaviorFlags
make
```

The processor model is based on the previous model, with the changes listed in following sections.

8.1.1 File `or1kStructure.h`

For this example, we need to model the 32 OR1K general-purpose registers and three boolean flags: *carry*, *overflow* and *branch*. The structure is therefore modified like this:

```
#define OR1K_REGS 32          // basic OR1K registers
#define OR1K_BITS 32         // register size in bits

// processor structure
typedef struct or1kS {

    Bool        carryFlag;    // carry flag
    Bool        overflowFlag; // overflow flag
    Bool        branchFlag;   // branch flag

    Uns32       regs[OR1K_REGS]; // basic registers

    vmiBusPortP busPorts;      // bus port descriptions
} or1k, *or1kP;
```

The C structure byte offsets of the various fields are now as follows:

Register	Byte Offset
carryFlag	0

```
overflowFlag 1
branchFlag   2
regs[0]      4
regs[1]      8
... etc ...
```

To simplify calculation of these offsets, `orlkStructure.h` now has the following macros for use in *variable* C expressions:

```
// macros to specify target registers in VARIABLE expressions
#define ORLK_OFFSET(_F)      VMI_CPU_REG(orlkP, _F)
#define ORLK_REG(_R)        ORLK_OFFSET(regs[_R])
#define ORLK_CARRY          ORLK_OFFSET(carryFlag)
#define ORLK_OVERFLOW       ORLK_OFFSET(overflowFlag)
```

In constant expression contexts (for example static structure initializers) these variants should be used instead:

```
// macros to specify target registers in CONSTANT expressions
#define ORLK_OFFSET_CONST(_F) VMI_CPU_REG_CONST(orlkP, _F)
#define ORLK_REG_CONST(_R)   ORLK_OFFSET_CONST(regs[_R])
#define ORLK_CARRY_CONST     ORLK_OFFSET_CONST(carryFlag)
#define ORLK_OVERFLOW_CONST  ORLK_OFFSET_CONST(overflowFlag)
```

8.1.2 File `orlkMorph.c`

How flags should be handled in an arithmetic operation is indicated by a *flags* argument to the VMI morph-time API call describing the operation. This argument is a pointer to a structure type defined in `vmiTypes.h`:

```
typedef enum {
    vmi_CF=0,           // carry flag
    vmi_PF=1,           // parity flag
    vmi_ZF=2,           // zero flag
    vmi_SF=3,           // sign flag
    vmi_OF=4,           // overflow flag
    vmi_LF=5            // KEEP LAST
} vmiFlag;

typedef enum {
    vmi_FN_NONE =0x00, // empty negate mask
    vmi_FN_CF_IN =0x01, // negate carry in flag
    vmi_FN_CF_OUT=0x02, // negate carry out flag
    vmi_FN_PF   =0x04, // negate parity flag
    vmi_FN_ZF   =0x08, // negate zero flag
    vmi_FN_SF   =0x10, // negate sign flag
    vmi_FN_OF   =0x20, // negate overflow flag
} vmiFlagNegate;

typedef struct vmiFlagsS {
    vmiReg    cin;           // register specifying carry in
    vmiReg    f[vmi_LF];    // registers to hold operation results
    vmiFlagNegate negate;    // bitmask of negated flags
} vmiFlags;
```

The `vmiFlag` enumeration lists all the flags that can be generated by an arithmetic operation: *carry*, *parity*, *zero*, *sign* and *overflow*. The `vmiFlagNegate` enumeration describes how flags are negated on input to and output from the operation.

The `vmiFlags` structure contains the following:

1. A field `cin` of type `vmiReg`. This field specifies *the register offset in a processor structure of a flag byte to use for the carry in value*.
2. An array of `vmiReg` values indexed by `vmiFlag` type. This field specifies *the register offsets in a processor structure of flag bytes into which generated flags should be written*.
3. A bitmask of type `vmiFlagNegate` specifying how flags should be negated on input to and output from the operation.

In other words, the `vmiFlags` structure allows you to specify boolean flag locations within your processor structure that can provide and accept flag values in arithmetic operations. These flags should always be declared in the processor structure as type `Bool`. The utility functions `doBinopSLit` and `doBinopULit` have been modified to accept a flags argument and pass it on to the VMI morph-time API functions:

```
//
// Emit code to implement a binary/signed literal ORlK instruction
//
static void doBinopSLit(Uns32 instr, vmiBinop op, vmiFlagsCP flags) {

    Uns32 rd      = OP2_D(instr);
    Uns32 ra      = OP2_A(instr);
    Int16 i       = OP2_I(instr);
    vmiReg target = (rd==0) ? VMI_NOREG : ORlK_REG(rd);

    if(ra==0) {
        vmimtBinopRCC(ORlK_BITS, op, target, 0, i, flags);
    } else {
        vmimtBinopRRC(ORlK_BITS, op, target, ORlK_REG(ra), i, flags);
    }
}

//
// Emit code to implement a binary/unsigned literal ORlK instruction
//
static void doBinopULit(Uns32 instr, vmiBinop op, vmiFlagsCP flags) {

    Uns32 rd      = OP2_D(instr);
    Uns32 ra      = OP2_A(instr);
    Uns16 k       = OP2_I(instr);
    vmiReg target = (rd==0) ? VMI_NOREG : ORlK_REG(rd);

    if(ra==0) {
        vmimtBinopRCC(ORlK_BITS, op, target, 0, k, flags);
    } else {
        vmimtBinopRRC(ORlK_BITS, op, target, ORlK_REG(ra), k, flags);
    }
}
```

In the specific example of the OR1K processor, the bitwise logical operations do not use or affect any processor flags. This is indicated by passing a null pointer for the `flags` argument:

```
static OR1K_DISPATCH_FN(morphANDI) {doBinopULit(instr, vmi_AND, 0);}
static OR1K_DISPATCH_FN(morphORI) {doBinopULit(instr, vmi_OR, 0);}
static OR1K_DISPATCH_FN(morphXORI) {doBinopSLit(instr, vmi_XOR, 0);}
```

The remaining arithmetic operations can generate *carry* and *overflow* flags, and (in the case of instruction `l.adc`) take a carry flag as input. Other possible output flags do not exist on the OR1K. This is indicated using a `vmiFlags` structure `flagsCO`:

```
const vmiFlags flagsCO = {
    OR1K_CARRY_CONST,          // offset to carry in flag
    {
        OR1K_CARRY_CONST,      // offset to carry out flag
        VMI_NOFLAG_CONST,      // parity flag not used
        VMI_NOFLAG_CONST,      // zero flag not used
        VMI_NOFLAG_CONST,      // sign flag not used
        OR1K_OVERFLOW_CONST    // offset to overflow flag
    }
};

static OR1K_DISPATCH_FN(morphADDI) {doBinopSLit(instr, vmi_ADD, &flagsCO);}
static OR1K_DISPATCH_FN(morphADDIC) {doBinopSLit(instr, vmi_ADC, &flagsCO);}
static OR1K_DISPATCH_FN(morphMULI) {doBinopSLit(instr, vmi_IMUL, &flagsCO);}
```

In detail, `flagsCO` specifies that:

1. Any input carry required by the arithmetic operation should be obtained from the processor structure at offset `OR1K_CARRY_CONST`, specified in `or1kStructure.h`. This corresponds to the `carry` Boolean field in the structure.
2. Any output carry generated by the arithmetic operation should be written to the processor structure at offset `OR1K_CARRY_CONST`.
3. Any output overflow generated by the arithmetic operation should be written to the processor structure at offset `OR1K_OVERFLOW_CONST`.
4. Any other output flags generated by the arithmetic operations should be discarded (indicated by using the special value `VMI_NOFLAG_CONST` in the appropriate `vmiFlags` structure field).
5. The carry flag should not be negated when used as an input and no flags should be negated on output. Therefore, the `negate` field of `flagsCO` is initialized to the default zero value (`vmi_FN_NONE`) by omitting it from the structure initializer.

Note what happens in `doBinopSLit` and `doBinopULit` when the output register `rd` is `r0`. Recall that `r0` is hardwired to zero on the OR1K processor. What should happen to the processor flags for an instruction where the output register is `r0`? The result should be *discarded* but changes to the flag values *preserved*. This can be indicated to the VMI morph-time API by specifying the special value `VMI_NOREG` as the destination register to `vmimtBinopRCC` or `vmimtBinopRRC`.

8.2 Validating Flag Behavior with Tests

For even apparently simple instructions like `l.addic`, it is clear that there are already a number of separate cases to be tested. An ideal test plan should cover the following options in various combinations:

1. target register `rd` of `r0-r31`
2. target register `rd` of `r0`
3. source register `ra` of `r1-r31`
4. source register `ra` of `r0`
5. validate carry output generated when required
6. validate overflow output generated when required
7. validate carry input used when required

File `asmtest.S` in directory `5.or1kBehaviorFlags/application` is an example of how this could be done.

```
.global _start
_start:
    // TEST PROLOGUE
    l.addi    r1,r0,0        // r1 = 0
    l.addi    r2,r0,1        // r2 = 1
    l.addi    r3,r0,-1       // r3 = -1
    l.addi    r4,r0,0x800    // r4 = 0x00000800
    l.muli    r4,r4,0x800    // r4 = 0x00400000
    l.muli    r4,r4,0x200    // r4 = 0x80000000
    l.addi    r5,r4,-1       // r5 = 0x7fffffff

    l.addic   r20,r0,-1
    l.addic   r20,r0,0
    l.addic   r20,r0,1
    l.addic   r20,r1,-1
    l.addic   r20,r1,0
    l.addic   r20,r1,1
    l.addic   r20,r2,-1
    l.addic   r20,r2,0
    l.addic   r20,r2,1
    l.addic   r20,r3,-1
    l.addic   r20,r3,0
    l.addic   r20,r3,1
    l.addic   r20,r4,-1
    l.addic   r20,r4,0
    l.addic   r20,r4,1
    l.addic   r20,r5,-1
    l.addic   r20,r5,0
    l.addic   r20,r5,1

    l.addic   r0,r0,-1
    l.addic   r0,r0,0
    l.addic   r0,r0,1
    l.addic   r0,r1,-1
    l.addic   r0,r1,0
    l.addic   r0,r1,1
    l.addic   r0,r2,-1
    l.addic   r0,r2,0
    l.addic   r0,r2,1
    l.addic   r0,r3,-1
    l.addic   r0,r3,0
```

```

        l.addic      r0,r3,1
        l.addic      r0,r4,-1
        l.addic      r0,r4,0
        l.addic      r0,r4,1
        l.addic      r0,r5,-1
        l.addic      r0,r5,0
        l.addic      r0,r5,1

.global exit
exit:
        l.addi       r1,r2,0

```

Assemble this file using:

```
make -C application
```

Make & run the platform using the assembler executable file:

```
make -C platform
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```

Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 'cpul', 0x0000000001000074: l.addi   r1,r0,0x0
Info 'cpul' REGISTERS
CPU cpul (instruction 1):
    0: 00000000 00000000 00000000 deadbeef
   16: deadbeef deadbeef deadbeef deadbeef
   32: deadbeef deadbeef deadbeef deadbeef
   48: deadbeef deadbeef deadbeef deadbeef
   64: deadbeef deadbeef deadbeef deadbeef
   80: deadbeef deadbeef deadbeef deadbeef
   96: deadbeef deadbeef deadbeef deadbeef
  112: deadbeef deadbeef deadbeef deadbeef
  128: deadbeef 08daa790

... (many lines cut) ...

Info 'cpul', 0x0000000001000118: l.addic  r0,r5,0x0
Info 'cpul' REGISTERS
CPU cpul (instruction 42):
    0: 00000100 00000000 00000000 00000001
   16: ffffffff 80000000 7fffffff deadbeef
   32: deadbeef deadbeef deadbeef deadbeef
   48: deadbeef deadbeef deadbeef deadbeef
   64: deadbeef deadbeef deadbeef deadbeef
   80: deadbeef 80000000 deadbeef deadbeef
   96: deadbeef deadbeef deadbeef deadbeef
  112: deadbeef deadbeef deadbeef deadbeef
  128: deadbeef 08daa790

Info 'cpul', 0x000000000100011c: l.addic  r0,r5,0x1
Info 'cpul' REGISTERS
CPU cpul (instruction 43):
    0: 00000100 00000000 00000000 00000001
   16: ffffffff 80000000 7fffffff deadbeef
   32: deadbeef deadbeef deadbeef deadbeef
   48: deadbeef deadbeef deadbeef deadbeef
   64: deadbeef deadbeef deadbeef deadbeef
   80: deadbeef 80000000 deadbeef deadbeef

```



```
    96: deadbeef deadbeef deadbeef deadbeef
   112: deadbeef deadbeef deadbeef deadbeef
   128: deadbeef 08daa790
Info 'cpul', 0x0000000001000120: l.addi    r1,r2,0x0
Processor 'cpul' terminated at 'exit', address 0x1000120
Info 'cpul' REGISTERS
CPU cpul (instruction 44):
    0: 00000000 00000000 00000001 00000001
   16: ffffffff 80000000 7fffffff deadbeef
   32: deadbeef deadbeef deadbeef deadbeef
   48: deadbeef deadbeef deadbeef deadbeef
   64: deadbeef deadbeef deadbeef deadbeef
   80: deadbeef 80000000 deadbeef deadbeef
   96: deadbeef deadbeef deadbeef deadbeef
  112: deadbeef deadbeef deadbeef deadbeef
  128: deadbeef 08daa790
```

8.3 *Model-Specific Dump Format*

Comparing the output from the above example with that from the simple behavioral model (section 7.3), there is a significant difference in format because each register dump now has 34 words (136 bytes) instead of 33 words (132 bytes). This is because adding the flags to the processor structure has increased its size. It also isn't clear what the dump is showing: which values represent general purpose registers, which represent flags, and which are supplemental values (for example the `busPorts` pointer) which do not represent true processor state at all? To address this problem, we need to add a model-specific register dump routine. A template model for the OR1K with this routine added can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/6.or1kBehaviorDump
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/6.or1kBehaviorDump .
```

Compile the model using make:

```
cd 6.or1kBehaviorDump
make
```

The processor model is based on the previous model, with the changes listed in following sections.

8.3.1 File `or1kStructure.h`

While implementing the OR1K register dump routine, we will update the processor model to partially implement the OR1K status register (`sr`). This is a 32-bit register which must be added to the processor definition in `or1kStructure.h` as follows:

```
#define OR1K_REGS 32          // basic OR1K registers
#define OR1K_BITS 32         // register size in bits

// processor structure
typedef struct or1kS {
```

```

    Bool      carryFlag;      // carry flag
    Bool      overflowFlag;   // overflow flag
    Bool      branchFlag;    // branch flag

    Uns32     regs[OR1K_REGS]; // basic registers

    Uns32      SR;              // status register

    vmiBusPortP busPorts;     // bus port descriptions
} orlk, *orlkP;

```

The status register bits are conveniently accessed using these macros:

```

// Bit definitions for the SR register
#define SPR_SR_CID      0xf0000000 // Context ID
#define SPR_SR_SUMRA    0x00010000 // Supervisor SPR read access
#define SPR_SR_FO       0x00008000 // Fixed one
#define SPR_SR_EPH      0x00004000 // Exception Prefix High
#define SPR_SR_DSX      0x00002000 // Delay Slot Exception
#define SPR_SR_OVE      0x00001000 // Overflow flag Exception
#define SPR_SR_OV       0x00000800 // Overflow flag
#define SPR_SR_CY       0x00000400 // Carry flag
#define SPR_SR_F        0x00000200 // Condition Flag
#define SPR_SR_CE       0x00000100 // CID Enable
#define SPR_SR_LEE      0x00000080 // Little Endian Enable
#define SPR_SR_IME      0x00000040 // Instruction MMU Enable
#define SPR_SR_DME      0x00000020 // Data MMU Enable
#define SPR_SR_ICE      0x00000010 // Instruction Cache Enable
#define SPR_SR_DCE      0x00000008 // Data Cache Enable
#define SPR_SR_IEE      0x00000004 // Interrupt Exception Enable
#define SPR_SR_TEE      0x00000002 // Tick timer Exception Enable
#define SPR_SR_SM       0x00000001 // Supervisor Mode

```

8.3.2 File orlkUtils.c

This file implements the OR1K register dump function using the VMI_DEBUG_FN macro, defined in vmiDbg.h. The function is as follows:

```

VMI_DEBUG_FN(orklDumpRegisters) {

    orlkP orlk = (orklP)processor;
    Uns32 i     = 0;

    vmiPrintf("----- \n");

    // print general-purpose registers
    while(i<OR1K_REGS) {
        vmiPrintf(" R%-2u: %08x", i, orlk->regs[i]);
        i++;
        if(!(i&3)) {
            vmiPrintf("\n");
        } else {
            vmiPrintf(" ");
        }
    }

    // newline if required before derived registers
    if(i&3) {

```

```
        vmiPrintf("\n");
    }

    // flags
    vmiPrintf(
        " BF:%u CF:%u OF:%u ",
        orlk->branchFlag,
        orlk->carryFlag,
        orlk->overflowFlag
    );

    // program counter
    vmiPrintf(" PC : %08x  ", (Uns32)vmirtGetPC(processor));

    // status register
    vmiPrintf(" SR : %08x  ", orlkGetSR(ork));

    vmiPrintf("\n----- \n\n");
}
```

The register dump function is passed a single `vmiProcessorP` argument, indicating the processor for which to dump registers. The first step is to cast this to an `orkP` type:

```
orkP orlk = (orkP)processor;
```

Next, the function prints out the values of the OR1K general purpose registers, naming them `r0`, `r1`, `r2` and so on. All output is generated using the VMI routine `vmiPrintf`, defined in `vmiMessage.h`:

```
while(i<OR1K_REGS) {
    vmiPrintf(" R%-2u: %08x", i, orlk->regs[i]);
    i++;
    if(!(i&3)) {
        vmiPrintf("\n");
    } else {
        vmiPrintf(" ");
    }
}
```

Next, the function prints the current settings of the *branch*, *carry* and *overflow* flags:

```
// flags
vmiPrintf(
    " BF:%u CF:%u OF:%u ",
    orlk->branchFlag,
    orlk->carryFlag,
    orlk->overflowFlag
);
```

8.3.2.1 Printing the Program Counter (PC)

The OR1K processor has a program counter register, PC, which we would like to print in the dump routine. Until this point, we have not modeled the processor program counter at all; how should it be done?

One solution would be to introduce an extra `pc` field into the processor structure, which we could update at the start of every instruction using a morph-time operation. For example:

```
// processor structure
typedef struct orlks {

    Bool        carryFlag;        // carry flag
    Bool        overflowFlag;     // overflow flag
    Bool        branchFlag;       // branch flag

    Uns32       regs[OR1K_REGS];  // basic registers

    Uns32       SR;               // status register
    Uns32       PC;              // program counter

    vmiBusPortP busPorts;         // bus port descriptions
} orlk, *orlkP;

#define OR1K_PC OR1K_REG(PC)

VMI_MORPH_FN(orkMorphInstruction) {
    vmimtMoveRC(OR1K_BITS, OR1K_PC, (Uns32)thisPC);
    orlkDecode((orlkP)processor, thisPC, OR1K_MORPH, 0);
}
```

However, this is unnecessarily inefficient: we have already seen from the instruction trace in previous examples that the simulator always knows the address of the current instruction. Instead of maintaining the program counter value in the model, it would be much better just to ask the simulator for the current program counter value when we need it. A routine to give exactly what is required is available in the VMI *run-time* interface (defined in file `vmiRt.h`):

```
//
// Return the current program counter for a processor
//
Addr vmirtGetPC(vmiProcessorP processor);
```

The OR1K register dump function uses this as follows:

```
vmiPrintf(" PC : %08x ", (Uns32)vmirtGetPC(processor));
```

This highlights a *very important point*: when writing a processor model, do not explicitly model register values that are *infrequently referenced* and can *easily be created on demand*. This is *always* the case for the program counter and very often the case for processor status registers. Failure to do this will result in processor models which are much slower than they need to be.

8.3.2.2 Printing the Status Register (sr)

As a second example of creating register values on demand, the OR1K also contains a status register, `sr`. This register encodes the values of the three OR1K flags (carry, overflow and branch) in addition to other status information (whether the processor is in

supervisor mode, for example). The OR1K register dump function prints the current value of the status register like this:

```
vmiPrintf(" SR : %08x ", orlkGetSR(ork));
```

The routine `orkGetSR` is implemented in `orkUtils.c` like this:

```
Uns32 orlkGetSR(orkP ork) {  
    fillSR(ork);  
    return ork->SR;  
}
```

The routine `fillSR` updates the current value of the `sr` register field in the processor structure so that it includes the three boolean flags:

```
#define SET_BIT(_R, _C, _M) \  
    if(_C) {                \  
        (_R) = (_R) | (_M); \  
    } else {                \  
        (_R) = (_R) & ~(_M); \  
    } \  
  
inline static void fillSR(orkP ork) {  
    SET_BIT(ork->SR, ork->branchFlag,  SPR_SR_F);  
    SET_BIT(ork->SR, ork->carryFlag,   SPR_SR_CY);  
    SET_BIT(ork->SR, ork->overflowFlag, SPR_SR_OV);  
}
```

In other words, when the model requires the current value of the OR1K status register `sr`, it should call the routine `orkGetSR`, which assembles the value by combining some bits stored in the processor structure `SR` field with the current values of the three flag registers. This is much more efficient than regenerating the full value of `sr` after each instruction that could possibly modify flag values.

For completeness, `orkUtils.c` also implements a public function to set the `sr` register, `orkSetSR`. This isn't used in this example, but will be required in the full model.

```
#define GET_BIT(_R, _M) \  
    (((_R) & (_M)) ? 1 : 0) \  
  
void orkSetSR(orkP ork, Uns32 value) {  
  
    // it is never possible to clear the fixed-one (FO) bit  
    value |= SPR_SR_FO;  
  
    // set the SR  
    ork->SR = value;  
  
    // set the current branch flag, carry flag and overflow flag from the SR  
    ork->branchFlag = GET_BIT(value, SPR_SR_F);  
    ork->carryFlag  = GET_BIT(value, SPR_SR_CY);  
    ork->overflowFlag = GET_BIT(value, SPR_SR_OV);  
}
```

The function `orlkSetSR` extracts the flag bits from the new value of the status register `sr` and copies them into the flag fields in the processor model structure so that consistency is maintained.

8.3.3 File `orlkMain.c`

The constructor has been changed to initialize the new status register `sr`:

```
VMI_CONSTRUCTOR_FN(orlkConstructor) {  
  
    orlkP orlk = (orlkP)processor;  
    Uns32 i;  
  
    // initialize general purpose registers  
    for(i=2; i<OR1K_REGS; i++) {  
        orlk->regs[i] = 0xdeadbeef;  
    }  
  
    // initialize status register SR  
    orlk->SR = SPR_SR_FO | SPR_SR_SM;  
  
    // create bus port specifications  
    newBusPorts(orlk);  
}
```

8.3.4 File `orlkAttrs.c`

The register dump routine has been added to the `vmiIASAttr` structure for the OR1K:

```
const vmiIASAttr modelAttrs = {  
  
    ///////////////////////////////////////  
    // VERSION & SIZE ATTRIBUTES  
    ///////////////////////////////////////  
  
    .versionString = VMI_VERSION,  
    .modelType     = VMI_PROCESSOR_MODEL,  
    .dictNames     = dictNames,  
    .cpuSize       = sizeof(orlk),  
  
    ///////////////////////////////////////  
    // CREATE/DELETE ROUTINES  
    ///////////////////////////////////////  
  
    .constructorCB = orlkConstructor,  
    .destructorCB  = orlkDestructor,  
  
    ///////////////////////////////////////  
    // MORPHER CORE ROUTINES  
    ///////////////////////////////////////  
  
    .morphCB = orlkMorphInstruction,  
  
    ///////////////////////////////////////  
    // SIMULATION SUPPORT ROUTINES  
    ///////////////////////////////////////  
  
    .getEndianCB = orlkGetEndian,  
    .nextPCCB    = orlkNextInstruction,  
    .disCB       = orlkDisassemble,  
}
```

```

////////////////////////////////////
// REGISTER ACCESS SUPPORT ROUTINES (DEBUGGER & SEMIHOSTING)
////////////////////////////////////

.debugCB = orlkdumpRegisters,
};

```

8.4 Validating Register Dumping with Point Tests

The generation the assembler test case in directory 6.orlkdumpBehaviorDump will also will also compile the default application and test platform, to compile individually, as before use:

```

make -C platform
make -C application

```

Run the platform using the assembler executable file:

```

platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf

```

The output from this should be as follows:

```

Warning (PC_NRI) No register information callback given for processor 'cpu1'
Info 'cpu1', 0x0000000001000074: l.addi    r1,r0,0x0
Info 'cpu1' REGISTERS
-----
R0 : 00000000    R1 : 00000000    R2 : deadbeef    R3 : deadbeef
R4 : deadbeef    R5 : deadbeef    R6 : deadbeef    R7 : deadbeef
R8 : deadbeef    R9 : deadbeef    R10: deadbeef    R11: deadbeef
R12: deadbeef    R13: deadbeef    R14: deadbeef    R15: deadbeef
R16: deadbeef    R17: deadbeef    R18: deadbeef    R19: deadbeef
R20: deadbeef    R21: deadbeef    R22: deadbeef    R23: deadbeef
R24: deadbeef    R25: deadbeef    R26: deadbeef    R27: deadbeef
R28: deadbeef    R29: deadbeef    R30: deadbeef    R31: deadbeef
BF:0 CF:0 OF:0  PC : 01000078  SR : 00008001
-----

Info 'cpu1', 0x0000000001000078: l.addi    r2,r0,0x1
Info 'cpu1' REGISTERS
-----
R0 : 00000000    R1 : 00000000    R2 : 00000001    R3 : deadbeef
R4 : deadbeef    R5 : deadbeef    R6 : deadbeef    R7 : deadbeef
R8 : deadbeef    R9 : deadbeef    R10: deadbeef    R11: deadbeef
R12: deadbeef    R13: deadbeef    R14: deadbeef    R15: deadbeef
R16: deadbeef    R17: deadbeef    R18: deadbeef    R19: deadbeef
R20: deadbeef    R21: deadbeef    R22: deadbeef    R23: deadbeef
R24: deadbeef    R25: deadbeef    R26: deadbeef    R27: deadbeef
R28: deadbeef    R29: deadbeef    R30: deadbeef    R31: deadbeef
BF:0 CF:0 OF:0  PC : 0100007c  SR : 00008001
-----

... etc ...

```

Now the trace output is much easier to understand because registers are printed with meaningful names.

8.5 Derived Flags

We have seen that the VMI API allows any of the *sign*, *carry*, *overflow*, *zero* or *parity* flags to be generated by an operation. It is often required to derive more complex flags from these: for example, it may be required to implement an unsigned below-or-equal condition flag, which is true if either the carry flag is set or the zero flag is set. The best approach is as follows:

1. generate the sign, carry, overflow, zero or parity flags as required as true processor registers;
2. use binary operations with width 8 to generate the derived flag using the basic flags as arguments as described below.

As an example, suppose that the OR1K model has been modified to implement sign and zero flags and a new temporary flag as follows:

```
#define OR1K_REGS 32          // basic OR1K registers
#define OR1K_BITS 32         // register size in bits

// processor structure
typedef struct or1kS {

    Bool        carryFlag;    // carry flag
    Bool        overflowFlag; // overflow flag
    Bool        zeroFlag;     // carry flag
    Bool        signFlag;     // overflow flag
    Bool        branchFlag;   // branch flag
    Bool        tempFlag;     // temporary flag

    Uns32       regs[OR1K_REGS]; // basic registers

    vmiBusPortP busPorts;      // bus port descriptions
} or1k, *or1kP;
```

and that new accessor macros for these flags have been added:

```
// macros to specify target registers in VARIABLE expressions
#define OR1K_OFFSET(_F)      VMI_CPU_REG(or1kP, _F)
#define OR1K_OFFSET_TEMP(_F) VMI_CPU_TEMP(or1kP, _F)
#define OR1K_REG(_R)        OR1K_OFFSET(regs[_R])
#define OR1K_CARRY          OR1K_OFFSET(carryFlag)
#define OR1K_OVERFLOW       OR1K_OFFSET(overflowFlag)
#define OR1K_ZERO           OR1K_OFFSET(zeroFlag)
#define OR1K_SIGN           OR1K_OFFSET(signFlag)
#define OR1K_TF             OR1K_OFFSET_TEMP(tempFlag)
```

Given these changes, use the following sequences to generate a derived flag in tempFlag:

Unsigned below-or-equal (CF==1) || (ZF==1):

```
vmimtBinopRRR(8, vmi_OR, OR1K_TF, OR1K_CARRY, OR1K_ZERO, 0);
```

Signed less-than (SF!=OF):

```
vmimtBinopRRR(8, vmi_XOR, OR1K_TF, OR1K_SIGN, OR1K_OVERFLOW, 0);
```


Signed less-than-or-equal ((ZF==1) || (SF!=OF)):

```
vmimtBinopRRR(8, vmi_XOR, OR1K_TF, OR1K_SIGN, OR1K_OVERFLOW, 0);  
vmimtBinopRR(8, vmi_OR, OR1K_TF, OR1K_ZERO, 0);
```

Complement of any flag:

```
vmimtBinopRRC(8, vmi_XOR, OR1K_TF, <flag_reg>, 1, 0);
```

Note that the recommended way to complement a flag is to exclusive-or it with 1.

9 Implementing Unconditional Jump Instructions

Up to this point, the OR1K examples have executed straight line code only. We will now implement unconditional jump instructions to allow simple non-linear programs to be run.

9.1 The Template Unconditional Jump Model

A template model for the OR1K processor implementing unconditional jump instructions can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/7.or1kBehaviorUncondJump
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/7.or1kBehaviorUncondJump .
```

Compile the model using make:

```
cd 7.or1kBehaviorUncondJump
make
```

The processor model is based on the previous model, with the changes listed in following sections.

9.1.1 File `or1kInstructions.h`

This file has been modified to include macros to decode the OR1K jump instructions.

9.1.2 File `or1kStructure.h`

The OR1K has a *link register*, `r9`, which is set to required return address in a *jump-and-link* (call) instruction:

```
#define OR1K_LINK 9           // link register index (R9)
#define OR1K_LINKREG OR1K_REG(OR1K_LINK)
```

9.1.3 File `or1kDecode.h`

The OR1K processor supports jump instructions with *delay slots* (instructions that are executed after a jump instruction before the jump is taken). As we will see shortly, the behavior of some instructions depends on whether they are encountered in a delay slot. To support this, the dispatch function template and the prototype for function `or1kDecode` have been changed:

```
#define OR1K_DISPATCH_FN(_NAME) void _NAME( \
    or1kP or1k,           \
    Uns32 thisPC,          \
    Uns32 instr,           \
    void *userData,        \
    Bool inDelaySlot       \
)
```

```
Bool orlkDecode(  
    orlkP          orlk,  
    Uns32          thisPC,  
    orlkDispatchTableCP table,  
    orlkDispatchFn defaultCB,  
    void           *userData,  
    Bool           inDelaySlot  
);
```

9.1.4 File `orlkDecode.c`

The OR1K jump instructions have been added to the dispatch table in a very similar manner as for previous instructions.

The implementation of `orlkDecode` has been upgraded to pass a flag indicating whether the current instruction is a delay slot instruction through to the dispatch function:

```
Bool orlkDecode(  
    orlkP          orlk,  
    Uns32          thisPC,  
    orlkDispatchTableCP table,  
    orlkDispatchFn defaultCB,  
    void           *userData,  
    Bool           inDelaySlot  
) {  
    // get the instruction at the passed address - always 4 bytes on OR1K  
    vmiProcessorP processor = (vmiProcessorP)orlk;  
    Uns32 instruction = vmicxtFetch4Byte(processor, thisPC);  
    orlkInstructionType type = decode(instruction);  
  
    // apply the callback, or the default if no match  
    if(type!=OR1K_IT_LAST) {  
        ((*table)[type])(orlk, thisPC, instruction, userData, inDelaySlot);  
        return True;  
    } else {  
        defaultCB(orlk, thisPC, instruction, userData, inDelaySlot);  
        return False;  
    }  
}
```

9.1.5 File `orlkDisassemble.c`

This file has been upgraded to implement disassembly callback functions for the jump instructions in a similar manner to previous instructions.

Note that the direct jump instructions (`l.j` and `l.jal`) make use of the current instruction address (`thisPC`) to calculate an absolute jump target address from the relative displacement that is encoded with the instruction:

```
static void doBranchJump(char *buffer, Uns32 thisPC, Uns32 instr, char *orlkop){  
  
    Int32 n          = OP7_N(instr);  
    Int32 offset      = ((n<<6) >> 4);  
    Uns32 toAddress    = thisPC + offset;  
  
    sprintf(buffer, "%-8s 0x%08x", orlkop, toAddress);  
}
```

The top-level disassembler function, `orlkDisassemble`, passes `False` for the new `inDelaySlot` argument to `orlkDecode` (instruction disassembly does not depend on whether the instruction is in a delay slot, so this value is ignored by all the disassembler callback functions).

9.1.6 File `orlkMorph.c`

This file has been upgraded to implement morph callback functions for the jump instructions, as described below.

The main morpher entry point function, `orlkMorphInstruction`, has been modified to pass the `inDelaySlot` argument down to the decoder (`orlkDecode`).

```
VMI_MORPH_FN(orlkMorphInstruction) {
    orlkDecode(
        (orlkP)processor, thisPC, &dispatchTable, morphDefault, 0, inDelaySlot
    );
}
```

When performing just-in-time compilation using the model morph callback, the simulator always knows whether the current instruction is a delay slot instruction. It provides this information to the model morpher entry point function as an argument, `inDelaySlot`, of the call to the `VMI_MORPH_FN` of the model.

9.1.6.1 Direct Unconditional Jump Instructions (`l.j` and `l.jal`)

The OR1K supports two direct unconditional jump instructions that we will implement now. Instruction `l.j` is a simple jump to a target address. Instruction `l.jal` is a jump-and-link instruction: there is a jump to a target address and a return address is saved in the link register (`r9`). Both these instructions are implemented with a single function:

`doJump`:

```
static void doJump(
    Uns32 instr,
    Uns32 thisPC,
    Bool link,
    Bool inDelaySlot
) {
    Int32 n = OP7_N(instr);
    Int32 offset = ((n<<6) >> 4);
    Uns32 toAddress = thisPC + offset;
    Uns32 nextAddress = thisPC + 8;
    vmiReg linkReg = link ? OR1K_LINKREG : VMI_NOREG;
    vmiJumpHint hint;

    // select an appropriate jump hint
    if(link) {
        hint = vmi_JH_CALL;
    } else {
        hint = vmi_JH_NONE;
    }

    if(inDelaySlot) {
        // jump in the delay slot does nothing
    }
}
```

```
    } else {  
        vmimtUncondJumpDelaySlot(  
            1,                // slotOps  
            nextAddress,      // linkPC  
            toAddress,        // toAddress  
            linkReg,          // linkReg  
            hint,             // hint  
            0                 // slotCB  
        );  
    }  
}
```

Whether the required instruction is a jump or a jump-and-link is specified by the `link` argument:

```
static OR1K_DISPATCH_FN(morphJ)    {doJump(instr, thisPC, False, inDelaySlot);}  
static OR1K_DISPATCH_FN(morphJAL) {doJump(instr, thisPC, True, inDelaySlot);}
```

For these direct jumps, the jump target is calculated from the current instruction address (`thisPC`) plus a signed offset encoded in a field in the instruction:

```
Int32    n          = OP7_N(instr);  
Int32    offset     = ((n<<6) >> 4);  
Uns32    toAddress  = thisPC + offset;
```

The main work of `doJump` is in these lines:

```
if(inDelaySlot) {  
    // jump in the delay slot does nothing  
} else {  
    vmimtUncondJumpDelaySlot(  
        1,                // slotOps  
        nextAddress,      // linkPC  
        toAddress,        // toAddress  
        linkReg,          // linkReg  
        hint,             // hint  
        0                 // slotCB  
    );  
}
```

If the current instruction is a delay slot instruction, *both `l.j` and `l.jal` have no effect*. It is therefore important that we know whether the current instruction is a delay slot so that appropriate action can be taken. This is why the decoder was modified to pass this information down to the dispatcher callback functions.

The morph-time function `vmimtUncondJumpDelaySlot` is used to describe the jump to the simulator. This function has six arguments:

1. `slotOps` is the number of instructions in the delay slot of this jump instruction. These OR1K instructions have one delay slot instruction. A value of 0 for `slotOps` specifies a jump with no delay slot instructions.
2. `linkPC` is used only if the jump is a jump-and-link, in which case it specifies the address that should be placed in the link register. For the OR1K, this is the address of the instruction after the delay slot instruction, i.e. `thisPC+8`.

3. `toAddress` is the jump target address.
4. `linkReg` is used to specify the link register for the jump, if this is a jump-and-link. If there is no link register (this is a simple jump), the value `VMI_NOREG` should be passed.
5. `hint` is used to help the simulator understand what kind of jump this is. In this chapter, we will see three values used:
 - a. `vmi_JH_CALL`: the jump is a *call* to a function;
 - b. `vmi_JH_RETURN`: the jump is a *return* from a function;
 - c. `vmi_JH_NONE`: the jump is neither a call nor a return.

Jump hints do not affect the behavior of a simulation but do improve performance (the example in later section 11.2.3 demonstrates this).

In this function, the instruction `l.jal` has a call hint of `vmi_JH_CALL`, and instruction `l.j` has a call hint of `vmi_JH_NONE`.

6. `slotCB`, if non-NULL, specifies a *post-delay-slot callback* function, taking the current processor as its only argument. The function is called just before the delayed branch is taken. If the branch is *not* taken for any reason (for example, if there is a simulated exception in the delay slot instruction), then the callback is *not* called.

The post-delay-slot callback is typically used to update processor state that should only be changed if the branch is taken. For example, if the instruction implements a switch to kernel mode then the state change reflecting this should typically be done in the post-delay-slot callback.

9.1.6.2 Indirect unconditional Jump Instructions (`l.jr` and `l.jalr`)

The OR1K also has two indirect conditional jump instructions. Instruction `l.jr` is a jump to a target address specified in a register. Instruction `l.jalr` is a jump-and-link instruction: there is a jump to a target address specified in a register, and a return address is saved in the link register (`r9`). Both these instructions are implemented with a single function: `doJumpReg`:

```
static void doJumpReg(
    Uns32 instr,
    Uns32 thisPC,
    Bool link,
    Bool inDelaySlot
) {
    Uns32 rb = OP8_B(instr);
    Uns32 nextAddress = thisPC + 8;
    vmiReg linkReg = link ? OR1K_LINKREG : VMI_NOREG;
    vmiJumpHint hint;

    // select an appropriate jump hint
    if(link) {
        hint = vmi_JH_CALL;
    } else if(rb==OR1K_LINK) {
        hint = vmi_JH_RETURN;
    } else {
        hint = vmi_JH_NONE;
    }

    if(inDelaySlot) {
        // jump in the delay slot does nothing
    }
}
```

```
    } else {  
        vmimtUncondJumpRegDelaySlot(  
            1,                // slotOps  
            nextAddress,      // linkPC  
            OR1K_REG(rb),     // toReg  
            linkReg,          // linkReg  
            hint,             // hint  
            0                 // slotCB  
        );  
    }  
}
```

Whether the required instruction is a jump or a jump-and-link is specified by the `link` argument:

```
static OR1K_DISPATCH_FN(morphJAL) {doJump(instr, thisPC, True, inDelaySlot);}   
static OR1K_DISPATCH_FN(morphJALR) {doJumpReg(instr, thisPC, True, inDelaySlot);}
```

For these indirect jumps, the jump target address is in a register encoded within the instruction:

```
Uns32      rb      = OP8_B(instr);
```

The main work of `doJumpReg` is in these lines:

```
if(inDelaySlot) {  
    // jump in the delay slot does nothing  
} else {  
    vmimtUncondJumpRegDelaySlot(  
        1,                // slotOps  
        nextAddress,      // linkPC  
        OR1K_REG(rb),     // toReg  
        linkReg,          // linkReg  
        hint,             // hint  
        0                 // slotCB  
    );  
}
```

Just as for direct jumps, indirect jumps have no effect in the delay slot of another jump.

The morph-time function `vmimtUncondJumpRegDelaySlot` is used to describe the jump to the simulator. This function has six arguments; all except the third argument are exactly the same as for `vmimtUncondJumpDelaySlot` (described in section 9.1.6.1). The third argument is used to specify the register containing the jump target address.

The jump hint to use with the indirect jump is determined as follows:

```
// select an appropriate jump hint  
if(link) {  
    hint = vmi_JH_CALL;  
} else if(rb==OR1K_LINK) {  
    hint = vmi_JH_RETURN;  
} else {  
    hint = vmi_JH_NONE;  
}
```

In other words, the jump hint indicates the type of the jump using the following rules:

1. if the this is a jump-and-link, then assume the jump is a *function call*;
2. otherwise, if this is an indirect jump using the OR1K link register (r9), then assume the jump is a *function return*;
3. otherwise, assume the jump is neither a call nor a return.

9.2 Validating Unconditional Jumps with Point Tests

Generate the assembler test case in directory 7.or1kBehaviorUncondJump which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```
Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 'cpul', 0x0000000001000074: l.addi    r1,r0,0x0
Info 'cpul', 0x0000000001000078: l.addi    r2,r0,0x0
Info 'cpul', 0x000000000100007c: l.jal     0x01000090
Info 'cpul', 0x0000000001000080: l.addi    r1,r1,0x1
Info 'cpul', 0x0000000001000090: l.j       0x0100009c
Info 'cpul', 0x0000000001000094: l.addi    r1,r1,0x1
Info 'cpul', 0x000000000100009c: l.addi    r8,r9,0x0
Info 'cpul', 0x00000000010000a0: l.jal     0x010000b8
Info 'cpul', 0x00000000010000a4: l.addi    r1,r1,0x1
Info 'cpul', 0x00000000010000b8: l.addi    r10,r9,0x4
Info 'cpul', 0x00000000010000bc: l.jr      r9
Info 'cpul', 0x00000000010000c0: l.addi    r1,r1,0x1
Info 'cpul', 0x00000000010000a8: l.addi    r9,r8,0x0
Info 'cpul', 0x00000000010000ac: l.jr      r9
Info 'cpul', 0x00000000010000b0: l.addi    r1,r1,0x1
Info 'cpul', 0x0000000001000084: l.jalr     r10
Info 'cpul', 0x0000000001000088: l.addi    r1,r1,0x1
Info 'cpul', 0x00000000010000ac: l.jr      r9
Info 'cpul', 0x00000000010000b0: l.addi    r1,r1,0x1
Info 'cpul', 0x000000000100008c: l.addi    r1,r1,0x0
Processor 'cpul' terminated at 'exit', address 0x100008c
-----
R0 : 00000000   R1 : 00000007   R2 : 00000000   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : 01000084   R9 : 0100008c   R10: 010000ac   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: deadbeef   R31: deadbeef
BF:0 CF:0 OF:0  PC : 01000090   SR : 00008001
-----
```


The file `platform/platform.c` has been slightly modified from the previous example in two ways:

1. `MODEL_FLAGS` has been specified as `ICM_ATTR_TRACE`. This enables instruction tracing but turns off register dumping on every instruction;
2. At the end of the simulation, there is a call to `icmDumpRegisters` to display the final processor register state.

The test case application/`asmtest.S` is as follows:

```

////////////////////////////////////
// MAIN ROUTINE
////////////////////////////////////
.global _start
_start:
    l.addi    r1,r0,0        // r1 = 0 (counts taken jumps)
    l.addi    r2,r0,0        // r2 = 0 (counts untaken jumps)
    l.jal     func1          // call func1
    l.addi    r1,r1,1        // increment r1 (delay slot instruction)
    l.jalr    r10            // call indirect using r10 (i.e. r10_addr)
    l.addi    r1,r1,1        // increment r1 (delay slot instruction)

    //////////////////////////////////////
    // EXIT FROM POINT TEST
    //////////////////////////////////////
.global exit
exit:
    l.addi    r1,r2,0

    //////////////////////////////////////
    // FUNCTION func1
    //////////////////////////////////////
func1:
    l.j       forward       // jump forward in this function
    l.addi    r1,r1,1        // increment r1 (delay slot instruction)
    l.addi    r2,r2,1        // ** not executed **
forward:
    l.addi    r8,r9,0        // save return address in r8
    l.jal     func2          // call func2
    l.addi    r1,r1,1        // increment r1 (delay slot instruction)
    l.addi    r9,r8,0        // restore return address from r8
r10_addr:
    l.jr      r9             // return from function
    l.addi    r1,r1,1        // increment r1 (delay slot instruction)
    l.addi    r2,r2,1        // ** not executed **

    //////////////////////////////////////
    // FUNCTION func2
    //////////////////////////////////////
func2:
    l.addi    r10,r9,4       // save return address+4 in r10
    l.jr      r9             // return from function
    l.addi    r1,r1,1        // increment r1 (delay slot instruction)
    l.addi    r2,r2,1        // ** not executed **

```

The test case has been designed to execute each of the jump instructions at least once. To exercise the `l.jalr` instruction, `func2` stores the address of label `r10_addr` in register `r10` to provide an appropriate target for the subsequent `l.jalr`.

10 Implementing Conditional Jump Instructions

In this chapter, we will implement comparison operations and conditional jumps for the OR1K. On this processor, conditional jumps are performed using two sets of instructions:

1. two registers (or a register and a constant) are compared using an instruction with the `l.sf` prefix (for example, `l.sfeq` compares two registers for equality). An internal *branch flag* is set based on the comparison result.
2. instructions `l.bf` and `l.bnf` then conditionally branch if the flag is true or false, respectively.

10.1 The Template Conditional Jump Model

A template model for the OR1K processor implementing conditional jump instructions can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/8.or1kBehaviorCondJump
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/8.or1kBehaviorCondJump .
```

Compile the model using make:

```
cd 8.or1kBehaviorCondJump
make
```

The processor model is based on the previous model, with the changes listed in following sections.

Decoder, disassembler and morpher files have been updated in this example to implement register-register arithmetic/bitwise and `l.nop` instructions (previously, only register-constant arithmetic/bitwise instructions were implemented). These instructions require no significant new knowledge or techniques and are not directly relevant to this chapter, so no further detail has been given here.

10.1.1 File `or1kInstructions.h`

This file has been modified to include macros to decode the OR1K instructions, `l.sf*` (opcode form 5) and `l.sfi*` (opcode form 6). Branch instructions use opcode form 7 macros (used previously for unconditional jumps).

10.1.2 File `or1kStructure.h`

The OR1K branch flag is set by the comparison operations and used by the conditional branch instructions. The flag is represented by a boolean field in the `or1kS` structure, accessed by the `OR1K_BRANCH` macro:

```
#define OR1K_BRANCH OR1K_OFFSET(branchFlag)
```

10.1.3 Files or1kDecode.h and or1kDecode.c

The OR1K branch and comparison instructions have been added to the dispatch table in a very similar manner as for previous instructions. The comparison operation decode differs slightly: the decode for OR1K_IT_SF and OR1K_IT_SFI instructions match 16 distinct instruction types with a different CMPOP field, but only 10 of these comparison operations are valid. Therefore, to decode a comparison operation, both the decode pattern *and* an opcode validity test are applied. This is implemented using a lookup table as follows:

```
typedef struct cmpInfoS {
    const char *name; // comparison operation name (for disassembly)
    vmiCondition cond; // equivalent VMI condition code
} cmpInfo;

const static cmpInfo cmpLut[16] = {
    {"eq", vmi_COND_Z }, // 0x0: a == b
    {"ne", vmi_COND_NZ }, // 0x1: a != b
    {"gtu", vmi_COND_NBE}, // 0x2: a > b (unsigned)
    {"geu", vmi_COND_NB }, // 0x3: a >= b (unsigned)
    {"ltu", vmi_COND_B }, // 0x4: a < b (unsigned)
    {"leu", vmi_COND_BE }, // 0x5: a <= b (unsigned)
    {0, vmi_COND_Z }, // 0x6: invalid
    {0, vmi_COND_Z }, // 0x7: invalid
    {0, vmi_COND_Z }, // 0x8: invalid
    {0, vmi_COND_Z }, // 0x9: invalid
    {"gts", vmi_COND_NLE}, // 0xa: a > b (signed)
    {"ges", vmi_COND_NL }, // 0xb: a >= b (signed)
    {"lts", vmi_COND_L }, // 0xc: a < b (signed)
    {"les", vmi_COND_LE }, // 0xd: a <= b (signed)
    {0, vmi_COND_Z }, // 0xe: invalid
    {0, vmi_COND_Z }, // 0xf: invalid
};

static const cmpInfo *getCmpInfo(Uns32 cmpOp) {
    return &cmpLut[cmpOp];
}

...

static or1kInstructionType decode(Uns32 instruction) {

    // get the OR1K decode table
    static vmiDecodeTableP decodeTable;
    if(!decodeTable) {
        decodeTable = createDecodeTable();
    }

    // decode the instruction to get the type
    or1kInstructionType type = vmiDecode(decodeTable, instruction);

    // some arguments to l.sf and l.sfi are invalid: filter them here
    if((type==OR1K_IT_SF) && !getCmpInfo(OP5_CMPPOP(instruction))->name) {
        type = OR1K_IT_LAST;
    } else if((type==OR1K_IT_SFI) && !getCmpInfo(OP6_CMPPOP(instruction))->name){
        type = OR1K_IT_LAST;
    }

    return type;
}
```

Once function `decode` has determined the apparent instruction type, it then performs a further check for `OR1K_IT_SF` and `OR1K_IT_SFI` instructions to validate the comparison operation. If the operation is invalid, it returns `OR1K_IT_LAST` (indicating that the instruction is invalid).

Two of the utility functions for decoding comparison instructions are made available in `orlkDecode.h`:

```
//
// Return the type string for a specific OR1K comparison operation, or NULL
// if the operation is invalid
//
const char *orlkDecodeCmpName(Uns32 instruction) {
    if(IS_OP_SF(instruction)) {
        return getCmpInfo(OP5_CMPOP(instruction))->name;
    } else if(IS_OP_SFI(instruction)) {
        return getCmpInfo(OP6_CMPOP(instruction))->name;
    } else {
        return 0;
    }
}

//
// Return the equivalent vmiCondition for a specific OR1K comparison operation,
// or vmi_COND_Z if the operation is invalid
//
vmiCondition orlkDecodeCmpCondition(Uns32 instruction) {
    if(IS_OP_SF(instruction)) {
        return getCmpInfo(OP5_CMPOP(instruction))->cond;
    } else if(IS_OP_SFI(instruction)) {
        return getCmpInfo(OP6_CMPOP(instruction))->cond;
    } else {
        return vmi_COND_Z;
    }
}
```

10.1.4 File `orlkDisassemble.c`

This file has been upgraded to implement disassembly callback functions for the comparison and branch instructions. The branch instructions `l.bf` and `l.bnf` use the `doBranchJump` utility function previously used for unconditional branches. The comparison instructions use the `orlkDecodeCmpName` utility function from `orlkDecode.h` to obtain the specific comparison instruction name:

```
static void doCompareReg(char *buffer, Uns32 instr, char *orlkopFmt) {

    Uns32 ra = OP5_A(instr);
    Uns32 rb = OP5_B(instr);
    char opBuf[16];

    sprintf(opBuf, orlkopFmt, orlkDecodeCmpName(instr));
    sprintf(buffer, "%-8s r%u,r%u", opBuf, ra, rb);
}

static void doCompareLit(char *buffer, Uns32 instr, char *orlkopFmt) {
```

```

Uns32 ra = OP6_A(instr);
Int16 i  = OP6_I(instr);
char  opBuf[16];

sprintf(opBuf, orlkopFmt, orlkDecodeCmpName(instr));
sprintf(buffer, "%-8si r%u,0x%x", opBuf, ra, i);
}

OR1K_DISPATCH_FN(orlkDisSF)    {doCompareReg(userData, instr, "l.sf%s");}
OR1K_DISPATCH_FN(orlkDisSFI)  {doCompareLit(userData, instr, "l.sf%si");}

```

10.1.5 File or1kMorph.c

This file has been upgraded to implement morph callback functions for the comparison and branch instructions, as described below.

10.1.5.1 Conditional Branch Instructions (l.bf and l.bnf)

The two branch instructions are implemented with a single function: doBranch:

```

static void doBranch(
    Uns32 instr,
    Uns32 thisPC,
    Bool  jumpIfTrue,
    Bool  inDelaySlot
) {
    Int32 n      = OP7_N(instr);
    Int32 offset = ((n<<6) >> 4);
    Uns32 toAddress = thisPC + offset;

    if(inDelaySlot) {
        // jump in the delay slot does nothing
    } else {
        vmimtCondJumpDelaySlot(
            1,                // slotOps
            OR1K_BRANCH,      // flagReg
            jumpIfTrue,       // jumpIfTrue
            0,                // linkPC
            toAddress,        // toAddress
            VMI_NOREG,        // linkReg
            vmi_JH_NONE,      // hint
            0                  // slotCB
        );
    }
}

```

Whether the branch should be taken when the flag is set or cleared is specified by the jumpIfTrue argument:

```

static OR1K_DISPATCH_FN(morphBF) {doBranch(instr, thisPC, True, inDelaySlot);}
static OR1K_DISPATCH_FN(morphBNF) {doBranch(instr, thisPC, False, inDelaySlot);}

```

The main work of doBranch is in these lines:

```

if(inDelaySlot) {
    // jump in the delay slot does nothing
} else {
    vmimtCondJumpDelaySlot(

```

```

        1,                // slotOps
        OR1K_BRANCH,      // flagReg
        jumpIfTrue,       // jumpIfTrue
        0,                // linkPC
        toAddress,        // toAddress
        VMI_NOREG,        // linkReg
        vmi_JH_NONE,      // hint
        0                 // slotCB
    );
}

```

If the current instruction is a delay slot instruction, both `l.bf` and `l.bnf` have no effect. Otherwise, the morph-time function `vmimtCondJumpDelaySlot` is used to describe the jump to the simulator. This function has eight arguments:

1. `slotOps` is the number of instructions in the delay slot of this jump instruction.
2. `flagReg` specifies a register in the processor model that is used to determine whether the branch is taken. Here, we use the OR1K branch flag.
3. `jumpIfTrue` indicates how the branch register is used. If `jumpIfTrue` is non-zero, the jump will be taken if the branch register is non zero. Otherwise, the jump will be taken if the branch register is zero.
4. `linkPC` is used only if the jump is a jump-and-link, in which case it specifies the address that should be placed in the link register. This does not apply for OR1K conditional branches.
5. `toAddress` is the jump target address.
6. `linkReg` is used to specify the link register for the jump, if this is a jump-and-link. If there is no link register (as in this case), the value `VMI_NOREG` should be passed.
7. `hint` is used to help the simulator understand what kind of jump this is – see chapter 9 for more details.
8. `slotCB`, if non-NULL, specifies a *post-delay-slot callback* function, taking the current processor as its only argument. The function is called just before the delayed branch is taken. If the branch is *not* taken, then the callback is *not* called. The post-delay-slot callback is typically used to update processor state that should only be changed if the branch is taken. For example, if the instruction implements a switch to kernel mode then the state change reflecting this should typically be done in the post-delay-slot callback.

10.1.5.2 Comparison Instructions

The comparison instructions are implemented with `doCompareReg` (for register-register comparisons) and `doCompareLit` (for register-constant comparisons). Both functions use the `or1kDecodeCmpCondition` utility function from `or1kDecode.h` to get the equivalent `vmiCondition` for use with the morph-time comparison functions `vmimtCompareRC` and `vmimtCompareRR`:

```

static void doCompareLit(Uns32 instr) {
    Uns32    ra    = OP6_A(instr);
    Int16    i     = OP6_I(instr);
    vmiCondition cond = or1kDecodeCmpCondition(instr);

    vmimtCompareRC(OR1K_BITS, cond, OR1K_REG(ra), i, OR1K_BRANCH);
}

```

```
static void doCompareReg(Uns32 instr) {  
  
    Uns32      ra  = OP5_A(instr);  
    Uns32      rb  = OP5_B(instr);  
    vmiCondition cond = or1kDecodeCmpCondition(instr);  
  
    if(rb==0) {  
        vmimtCompareRC(  
            OR1K_BITS, cond, OR1K_REG(ra), 0, OR1K_BRANCH  
        );  
    } else {  
        vmimtCompareRR(  
            OR1K_BITS, cond, OR1K_REG(ra), OR1K_REG(rb), OR1K_BRANCH  
        );  
    }  
}
```

Function `doCompareLit` uses the morph-time function `vmimtCompareRC` to describe the jump to the simulator. This function has five arguments:

1. `bits` is the bit width of the registers to be compared. All OR1K registers are `OR1K_BITS` bits wide (32 in this model).
2. `cond` describes the comparison that should be made. The members of the `vmiCondition` enumeration are specified in `vmiTypes.h`.
3. `ra` specifies the first register argument of the comparison
4. `c` specifies the second constant argument of the comparison
5. `flag` specifies the `Uns8` register that should be written with 1 if the condition is true and 0 if it is false. In this model, the `OR1K_BRANCH` register is written.

The main work of `doCompareReg` is in these lines:

```
if(rb==0) {  
    vmimtCompareRC(  
        OR1K_BITS, cond, OR1K_REG(ra), 0, OR1K_BRANCH  
    );  
} else {  
    vmimtCompareRR(  
        OR1K_BITS, cond, OR1K_REG(ra), OR1K_REG(rb), OR1K_BRANCH  
    );  
}
```

`doCompareReg` firsts tests whether the second register argument is register 0. If so, then the comparison is actually with a constant 0 (since `r0` is hardwired to 0), so `vmimtCompareRC` is used to perform the comparison. Otherwise `vmimtCompareRR` is used, which takes identical arguments to `vmimtCompareRC` except that argument 4 is a register instead of a constant.

10.2 Validating Conditional Jumps with Point Tests

Generate the assembler test case in directory `8.or1kBehaviorCondJump` which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform  
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```
Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 'cpul', 0x0000000001000074: l.addi    r1,r0,0x3
Info 'cpul', 0x0000000001000078: l.addi    r2,r0,0x2
Info 'cpul', 0x000000000100007c: l.jal     0x01000108
Info 'cpul', 0x0000000001000080: l.nop     0x0
Info 'cpul', 0x0000000001000108: l.addi    r29,r0,0x0
Info 'cpul', 0x000000000100010c: l.addi    r30,r0,0x1
Info 'cpul', 0x0000000001000110: l.addi    r31,r0,0x1
Info 'cpul', 0x0000000001000114: l.sfreq   r1,r2
Info 'cpul', 0x0000000001000118: l.bf      0x01000124
Info 'cpul', 0x000000000100011c: l.xor     r29,r29,r31
Info 'cpul', 0x0000000001000120: l.xor     r29,r29,r31
Info 'cpul', 0x0000000001000124: l.bnf      0x01000130
Info 'cpul', 0x0000000001000128: l.xor     r30,r30,r31
Info 'cpul', 0x0000000001000130: l.add     r31,r31,r31
Info 'cpul', 0x0000000001000134: l.sfne     r1,r2
Info 'cpul', 0x0000000001000138: l.bf      0x01000144
Info 'cpul', 0x000000000100013c: l.xor     r29,r29,r31
Info 'cpul', 0x0000000001000144: l.bnf      0x01000150
Info 'cpul', 0x0000000001000148: l.xor     r30,r30,r31
Info 'cpul', 0x000000000100014c: l.xor     r30,r30,r31
Info 'cpul', 0x0000000001000150: l.add     r31,r31,r31
. . . etc . . .
Info 'cpul', 0x0000000001000234: l.sfles   r1,r2
Info 'cpul', 0x0000000001000238: l.bf      0x01000244
Info 'cpul', 0x000000000100023c: l.xor     r29,r29,r31
Info 'cpul', 0x0000000001000240: l.xor     r29,r29,r31
Info 'cpul', 0x0000000001000244: l.bnf      0x01000250
Info 'cpul', 0x0000000001000248: l.xor     r30,r30,r31
Info 'cpul', 0x0000000001000250: l.add     r31,r31,r31
Info 'cpul', 0x0000000001000254: l.jr      r9
Info 'cpul', 0x0000000001000258: l.nop     0x0
Info 'cpul', 0x00000000010000fc: l.addi    r14,r29,0x0
Info 'cpul', 0x0000000001000100: l.addi    r15,r30,0x0
Info 'cpul', 0x0000000001000104: l.nop     0x0
Processor 'cpul' terminated at 'exit', address 0x1000104
-----
R0 : 00000000   R1 : ffffffff   R2 : ffffffff   R3 : 000000ce
R4 : 00000330   R5 : 000002a9   R6 : 00000157   R7 : 00000332
R8 : 000000cc   R9 : 010000fc   R10: 00000332   R11: 000000cc
R12: 000002a9  R13: 00000157   R14: 000000ce   R15: 00000330
R16: deadbeef  R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef  R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef  R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef  R29: 000000ce   R30: 00000330   R31: 00000400
BF:0 CF:0 OF:0 PC : 01000108   SR : 00008001
-----
```

The test case application/asmtest.S is as follows:

```
////////////////////////////////////
// MAIN ROUTINE
```



```

////////////////////////////////////
.global _start
_start:
    // test1: r1=3, r2=2
    l.addi    r1,r0,3          // r1=3
    l.addi    r2,r0,2          // r2=2
    l.jal     test             // call test
    l.nop     (delay slot instruction)
    l.addi    r3,r29,0          // move bf taken mask to r3
    l.addi    r4,r30,0          // move bnf taken mask to r4

    . . . etc . . .

    // test6: r1=-3, r2=-4
    l.addi    r1,r0,-3          // r1=-3
    l.addi    r2,r0,-4          // r2=-4
    l.jal     test             // call test
    l.nop     (delay slot instruction)
    l.addi    r14,r29,0          // move bf taken mask to r14
    l.addi    r15,r30,0          // move bnf taken mask to r15

    //////////////////////////////////////
    // EXIT FROM POINT TEST
    //////////////////////////////////////
.global exit
exit:
    l.nop

    //////////////////////////////////////
    // FUNCTION test
    //////////////////////////////////////
test:
    l.addi    r29,r0,0          // clear output mask r29 (bf taken)
    l.addi    r30,r0,1          // clear output mask r30 (bnf taken)
    l.addi    r31,r0,1          // initialize bitmask

    // test for sfreq
    l.sfreq   r1,r2             // r1==r2?
    l.bf      sfreqF             // go if true
    l.xor     r29,r29,r31        // add mask (delay slot instruction)
    l.xor     r29,r29,r31        // remove mask
sfreqF:      l.bnf      sfreqNF   // go if false
    l.xor     r30,r30,r31        // add mask (delay slot instruction)
    l.xor     r30,r30,r31        // remove mask
sfreqNF:     l.add     r31,r31,r31 // shift mask

    . . . etc . . .

    // test for sfles
    l.sfles   r1,r2             // r1<=r2? (signed)
    l.bf      sflesF             // go if true
    l.xor     r29,r29,r31        // add mask (delay slot instruction)
    l.xor     r29,r29,r31        // remove mask
sflesF:      l.bnf      sflesNF   // go if false
    l.xor     r30,r30,r31        // add mask (delay slot instruction)
    l.xor     r30,r30,r31        // remove mask
sflesNF:     l.add     r31,r31,r31 // shift mask

    l.jr      r9                // return, results in r29 and r30
    l.nop     (delay slot instruction)

```

The test case has been design to exercise all register-register comparison instructions, with a variety of input operands, and build up masks indicating how the comparison

results are treated by both the `l.bf` and `l.bnf` instructions. For example, this is an instruction sequence that is executed when function `test` is called for the first time, when `r1=3` and `r2=2`:

```
Info 'cpu1', 0x0000000001000114: l.sfeq    r1,r2        // test for equality
Info 'cpu1', 0x0000000001000118: l.bf      0x01000124    // branch if equal
Info 'cpu1', 0x000000000100011c: l.xor     r29,r29,r31    // (delay slot insn)
Info 'cpu1', 0x0000000001000120: l.xor     r29,r29,r31    // *** branch NOT taken
Info 'cpu1', 0x0000000001000124: l.bnf     0x01000130    // branch if not equal
Info 'cpu1', 0x0000000001000128: l.xor     r30,r30,r31    // (delay slot insn)
Info 'cpu1', 0x0000000001000130: l.add     r31,r31,r31    // *** branch TAKEN
```

In this example, the delay slot instruction is executed whether the branch is taken or not. It is also possible to describe branches that *annul* the delay slot instruction if the branch is not taken – refer to these routines in the *Imperas VMI Morph Time Reference* manual for more information:

```
vmimtCondJumpDelaySlotAnnul
vmimtCondJumpRegDelaySlotAnnul
vmimtSkipIfAnnul
```

11 Implementing Memory Access Instructions

In this chapter, we will implement memory load and store instructions for the OR1K. The processor supports six load instructions:

1. `l.lwz`: load 4 bytes; zero extend to 32 bits;
2. `l.lws`: load 4 bytes, sign extend to 32 bits (same as `l.lwz` on 32-bit core);
3. `l.lhz`: load 2 bytes, zero extend to 32 bits;
4. `l.lhs`: load 2 bytes, sign extend to 32 bits;
5. `l.lbz`: load 1 byte, zero extend to 32 bits;
6. `l.lbs`: load 1 byte, sign extend to 32 bits.

There are three store instructions:

1. `l.sw`: store 4 bytes;
2. `l.sh`: store 2 bytes;
3. `l.sb`: store 1 byte.

Accesses longer than one byte must be aligned with memory, otherwise the access generates an alignment exception – chapter 12 shows how this requirement can be modeled efficiently.

11.1 The Template Memory Access Model

A template model for the OR1K processor implementing memory access instructions can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/9.or1kBehaviorLoadStore
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/9.or1kBehaviorLoadStore.
```

Compile the model using make:

```
cd 9.or1kBehaviorLoadStore
make
```

The processor model is based on the previous model, with the changes listed in following sections.

Decoder, disassembler and morpher files have been updated in this example to implement shift/rotate and `l.movhi` instructions. The shift/rotate instructions are very similar to the arithmetic/bitwise instructions discussed previously. Instruction `l.movhi` implements a form of constant load to the high part of a register, common on RISC processors. These instructions require no significant new knowledge or techniques and are not directly relevant to this chapter, so no further detail has been given here.

11.1.1 File `or1kInstructions.h`

This file has been modified to include macros to decode the OR1K load and store instructions described at the start of this chapter.

11.1.2 File `or1kDecode.c`

The OR1K load and store instructions have been added to the dispatch table in a very similar manner as for previous instructions.

11.1.3 File `or1kDisassemble.c`

This file has been upgraded to implement disassembly callback functions for the load and store instructions. Once more, the code is very similar to previous instructions.

11.1.4 File `or1kMorph.c`

This file has been upgraded as described below.

11.1.4.1 Load Instructions

The six load instructions are implemented with a single function: `doLoad`:

```
static void doLoad(Uns32 instr, Uns32 bytes, Bool signExtend) {

    Uns32    rd    = OP2_D(instr);
    Uns32    ra    = OP2_A(instr);
    Int16    i     = OP2_I(instr);
    memEndian endian = getEndian();
    vmiReg    rdReg = rd ? OR1K_REG(rd) : VMI_NOREG;
    vmiReg    raReg = ra ? OR1K_REG(ra) : VMI_NOREG;

    vmimtLoadRRO(
        OR1K_BITS,          // destBits
        bytes*8,            // memBits
        i,                  // offset
        rdReg,              // rd
        raReg,              // ra
        endian,             // endian
        signExtend,         // signExtend
        True                // checkAlign
    );
}
```

The size of the load in bytes and whether sign extension is required are passed as arguments by the morph callback functions:

```
static OR1K_DISPATCH_FN(morphLWZ) {doLoad(instr, 4, False);}
static OR1K_DISPATCH_FN(morphLWS) {doLoad(instr, 4, True );}
static OR1K_DISPATCH_FN(morphLBZ) {doLoad(instr, 1, False);}
static OR1K_DISPATCH_FN(morphLBS) {doLoad(instr, 1, True );}
static OR1K_DISPATCH_FN(morphLHZ) {doLoad(instr, 2, False);}
static OR1K_DISPATCH_FN(morphLHS) {doLoad(instr, 2, True );}
```

Each load is specified by a call to the function `vmimtLoadRRO` from the Imperas Morph Time Function API. This takes eight arguments, as follows:

7. `destBits`: the size in bits of the destination register for the load;

8. `memBits`: the size in bits of the value in memory;
9. `offset`: a constant offset to be added to the address register `ra` to give the full memory address;
10. `rd`: the destination register for the load (if `rd` is `VMI_NOREG`, the load is performed but the fetched value discarded);
11. `ra`: a register holding the address from which to load (or `VMI_NOREG` if the load is from an address specified by `offset` only);
12. `endian`: the endianness of the load. This can be either `MEM_ENDIAN_BIG` or `MEM_ENDIAN_LITTLE`.
13. `signExtend`: whether the memory value should be assign extended if smaller than the register – if `False`, then the value is zero extended.
14. `checkAlign`: whether accesses to memory must be aligned. If this value is `True`, then any unaligned access will either cause simulation to terminate or a simulated exception to be taken: this is described in chapter 12.

For the OR1K processor, the address from which to load is calculated by adding address register `ra` to the constant value `c` from the instruction.

The endianness of the load is specified by function `getEndian`. This model supports big-endian only:

```
static memEndian getEndian(void) {  
    return MEM_ENDIAN_BIG;  
}
```

11.1.4.2 Store Instructions

The three store instructions are implemented with a single function: `doStore`:

```
static void doStore(Uns32 instr, Uns32 bytes) {  
  
    Uns32    ra    = OP10_A(instr);  
    Uns32    rb    = OP10_B(instr);  
    Int16    i     = OP10_I(instr);  
    memEndian endian = getEndian();  
    vmiReg    raReg = ra ? OR1K_REG(ra) : VMI_NOREG;  
  
    if(rb==0) {  
        vmimtStoreRCO(  
            bytes*8,      // bits  
            i,            // offset  
            raReg,        // ra  
            0,            // c  
            endian,       // endian  
            True          // checkAlign  
        );  
    } else {  
        vmimtStoreRRO(  
            bytes*8,      // bits  
            i,            // offset  
            raReg,        // ra  
            OR1K_REG(rb), // rb  
            endian,       // endian  
            True          // checkAlign  
        );  
    }  
}
```

```
}  
}
```

The size of the store in bytes is passed as an argument by the morph callback functions:

```
static OR1K_DISPATCH_FN(morphSW)    {doStore(instr, 4);}  
static OR1K_DISPATCH_FN(morphSB)    {doStore(instr, 1);}  
static OR1K_DISPATCH_FN(morphSH)    {doStore(instr, 2);}
```

A store of any register except `r0` is specified using `vmimtStoreRRO` from the Imperas Morph Time Function API. This takes six arguments, as follows:

1. `bits`: the size in bits of the destination register to be stored;
2. `offset`: a constant offset to be added to the address register `ra` to give the full memory address;
3. `ra`: a register holding the address to which to store (or `VMI_NOREG` if the store is to an address specified by `offset` only);
4. `rb`: the register to be stored;
5. `endian`: the endianness of the store. This can be either `MEM_ENDIAN_BIG` or `MEM_ENDIAN_LITTLE`.
6. `checkAlign`: whether accesses to memory must be aligned. If this value is `True`, then any unaligned access will either cause simulation to terminate or a simulated exception to be taken: this is described in chapter 12.

When register `r0` is being stored, the Imperas Morph Time Function API function `vmimtStoreRCO` is used instead. This takes identical arguments to `vmimtStoreRRO`, except that the fourth argument is a constant value instead of a register.

11.1.5 File `platform/platform.c`

The test platform for this example, `platform/platform.c`, has been changed as follows:

```
//  
// Main simulation routine  
//  
int main(int argc, char ** argv) {  
  
    // check for the application program name argument  
    if(argc!=2) {  
        icmPrintf("%s: expected application name argument\n", argv[0]);  
    }  
  
    // initialize CpuManager  
    icmInitPlatform(ICM_VERSION, 0, NULL, 0, "platform");  
  
    const char *modelFile = "model." IMPERAS_SHRSUF;  
    const char *semlhostFile = icmGetVlnvString(NULL, "ovpworld.org",  
                                                "modelSupport", "imperasExit", "1.0", "model");  
  
    // create a processor  
    icmProcessorP processor = icmNewProcessor(  
        "cpu1",           // CPU name  
        TYPE_NAME,       // CPU type  
        0,               // CPU cpuId  
        0,               // CPU model flags
```

```
    32,                // address bits
    modelFile,         // model file
    "modelAttrs",      // morpher attributes
    MODEL_FLAGS,       // use default attributes
    0,                // user-defined attributes
    semihostFile,      // semi-hosting file
    "modelAttrs"       // semi-hosting attributes
);

// load the processor object file
icmLoadProcessorMemory(processor, argv[1], False, False);

// run processor until done (no instruction limit)
while(simulate(processor, -1)) {
    // keep going while processor is still running
}

// dump the final register contents
icmDumpRegisters(processor);

// report the total number of instructions executed
icmPrintf(
    "processor has executed " FMT_64u " instructions\n",
    icmGetProcessorICount(processor)
);

icmTerminate();

return 0;
}
```

The significant change is in the call to function `simulate`:

```
while(simulate(processor, -1)) {
```

Previously, each call to `simulate` requested a single instruction to be executed. In this example, we use the value `-1` instead, which indicates that the simulator can execute an unlimited number of instructions before returning². In this case, this means that the call will only return when the program has completed.

11.2 Fibonacci Example

To demonstrate the load and store functions, we will use an assembler program that calculates Fibonacci numbers³ using a naive recursive algorithm (normally, instruction point tests should be created and tested first, or course). Once the basic example is working, we will use it to demonstrate simulator performance and the effect of *jump hints* (first encountered in chapter 9).

² To be precise, the second argument to `icmSimulate` (and `simulate`) is an `Uns64`, so a value of `-1` (sign-extended to 64 bits) in fact specifies that $2^{64}-1$ instructions should be executed – a very large number, but not quite unlimited.

³ Fibonacci numbers are defined as follows:

1. for $N \leq 1$: $\text{fib}(N) = N$;
2. for $N > 1$: $\text{fib}(N) = \text{fib}(N-1) + \text{fib}(N-2)$

11.2.1 Basic Example

Generate the assembler test case in directory `9.or1kBehaviorLoadStore` which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```
Warning (PC_NRI) No register information callback given for processor 'cpul'
Processor 'cpul' terminated at 'exit', address 0x1000080
```

```
-----
R0 : 00000000   R1 : 00000262   R2 : 00000179   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : 01000080  R10: deadbeef  R11: deadbeef
R12: deadbeef  R13: deadbeef  R14: deadbeef  R15: deadbeef
R16: deadbeef  R17: deadbeef  R18: deadbeef  R19: deadbeef
R20: deadbeef  R21: deadbeef  R22: deadbeef  R23: deadbeef
R24: deadbeef  R25: deadbeef  R26: deadbeef  R27: deadbeef
R28: deadbeef  R29: deadbeef  R30: deadbeef  R31: 00000000
BF:1 CF:1 OF:0 PC : 01000084  SR : 00008601
-----
```

```
processor has executed 22687 instructions
```

The test case application/asmtest.S is as follows:

```
////////////////////////////////////
// MAIN ROUTINE
////////////////////////////////////
.global _start
_start:
    l.addi    r31,r0,0        // initialize stack pointer to 0
    l.jal     fib             // calculate fib(15)
    l.addi    r1,r0,15        // r1 = 15 (delay slot)

    //////////////////////////////////////
    // EXIT FROM POINT TEST
    //////////////////////////////////////
.global exit
exit:
    l.nop

    //////////////////////////////////////
    // FUNCTION fib - calculate Fibonacci number of N, passed in r1.
    //                      result is returned in r1, r2 is destroyed
    //////////////////////////////////////
fib:
    l.sflsi   r1,1            // r1<=1? (signed)
    l.bf      done            // done if so, result is r1
    l.nop                      // (delay slot)

    l.addi    r31,r31,-12     // create stack frame
    l.sw      0(r31),r9       // save link register
    l.sw      4(r31),r1       // save input r1
```



```

        l.jal      fib          // calculate fib(N-1)
        l.addi     r1,r1,-1     // r1 = N-1 (delay slot)
        l.sw       8(r31),r1    // save fib(N-1)

        l.lwz      r1,4(r31)    // restore initial N
        l.jal      fib          // calculate fib(N-2)
        l.addi     r1,r1,-2     // r1 = N-2 (delay slot)

        l.lwz      r2,8(r31)    // restore fib(N-1)
        l.add      r1,r1,r2     // r1 = fib(N-2) + fib(N-1)

        l.lwz      r9,0(r31)    // restore link register
        l.addi     r31,r31,12   // destroy stack frame

done:    l.jr       r9          // return, result in r1
        l.nop                // (delay slot instruction)

```

The testcase calculates the value of `fib(15)`, returning the value in register `r1` (0x262, or 610 decimal). Register `r2` is used as an intermediate and is destroyed; register `r31` is used as a stack pointer. Because this is a naive recursive implementation, each call to `fib` creates up to two further recursive calls, and the current link register value (`r9`) and input value (`r1`) need to be preserved in a stack frame at each level using the load and store instructions we have just implemented.

11.2.2 Validating Simulation Performance

This Fibonacci implementation rapidly becomes computationally complex. Even when calculating a relatively small Fibonacci number, such as `fib(15)`, 22,687 instructions are performed. We can therefore use the example to test the basic simulation speed of the processor model.

Modify lines 25 and 26 of the test case `application/asmtest.S` as follows:

```

        l.jal      fib          // calculate fib(40)
        l.addi     r1,r0,40     // r1 = 40 (delay slot)

```

Regenerate the assembler test case and run it like this:

```

make -C application
time platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf

```

The output from this (after a few seconds) should be as follows:

```

Warning (PC_NRI) No register information callback given for processor 'cpul'
Processor 'cpul' terminated at 'exit', address 0x1000080
-----
R0 : 00000000  R1 : 06197ecb  R2 : 03c50ea2  R3 : deadbeef
R4 : deadbeef  R5 : deadbeef  R6 : deadbeef  R7 : deadbeef
R8 : deadbeef  R9 : 01000080  R10: deadbeef  R11: deadbeef
R12: deadbeef  R13: deadbeef  R14: deadbeef  R15: deadbeef
R16: deadbeef  R17: deadbeef  R18: deadbeef  R19: deadbeef
R20: deadbeef  R21: deadbeef  R22: deadbeef  R23: deadbeef
R24: deadbeef  R25: deadbeef  R26: deadbeef  R27: deadbeef
R28: deadbeef  R29: deadbeef  R30: deadbeef  R31: 00000000
BF:1 CF:1 OF:0 PC : 01000084  SR : 00008601

```

```
-----  
processor has executed 3808343229 instructions
```

The program has calculated `fib(40)` as `0x6197ecb` (102,334,155 decimal) using 3,808,343,229 simulated instructions. On a 2.8GHz Intel Core2 processor, `time` shows this takes about 7 seconds, giving a simulation speed for this example of about 544 simulated MIPS.

11.2.3 Demonstrating Jump Hint Effectiveness

Chapter 9 showed how *jump hints* should be used to tell the simulator what kind of jump is being performed (a call, a return or a simple jump that is neither a call nor a return). Now we have a test case that executes many calls and returns, we can demonstrate how effective these jump hints are when correctly applied. To do this, we will temporarily remove the jump hints from the processor model and then rerun `fib(40)` to see the effect.

Modify functions `doJump` and `doJumpReg` in file `or1kMorph.c` to remove the jump hints like this:

```
static void doJump(  
    Uns32 instr,  
    Uns32 thisPC,  
    Bool  link,  
    Bool  inDelaySlot  
) {  
    Int32      n          = OP7_N(instr);  
    Int32      offset     = ((n<<6) >> 4);  
    Uns32      toAddress   = thisPC + offset;  
    Uns32      nextAddress = thisPC + 8;  
    vmiReg      linkReg    = link ? OR1K_LINKREG : VMI_NOREG;  
    vmiJumpHint hint;  
  
    // provide no jump hint!  
    hint = vmi_JH_NONE;  
  
    if(inDelaySlot) {  
        // jump in the delay slot does nothing  
    } else {  
        vmimtUncondJumpDelaySlot(  
            1,                // slotOps  
            nextAddress,      // linkPC  
            toAddress,        // toAddress  
            linkReg,          // linkReg  
            hint,             // hint  
            0                 // slotCB  
        );  
    }  
}  
  
static void doJumpReg(  
    Uns32 instr,  
    Uns32 thisPC,  
    Bool  link,  
    Bool  inDelaySlot  
) {  
    Uns32      rb          = OP8_B(instr);
```

```
Uns32      nextAddress = thisPC + 8;
vmiReg      linkReg      = link ? OR1K_LINKREG : VMI_NOREG;
vmiJumpHint hint;

// provide no jump hint!
hint = vmi_JH_NONE;

if(inDelaySlot) {
    // jump in the delay slot does nothing
} else {
    vmimtUncondJumpRegDelaySlot(
        1,                // slotOps
        nextAddress,       // linkPC
        OR1K_REG(rb),      // toReg
        linkReg,           // linkReg
        hint,              // hint
        0                  // slotCB
    );
}
}
```

Rebuild the processor model and rerun `fib(40)` as follows:

```
make
time platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The program output is identical to before (jump hints do not affect behavior), but simulation speed is slower: on a 2.8 GHz Intel Core2 processor, `time` shows this takes about 9 seconds, giving a simulation speed for this example of about 423 simulated MIPS, *120 simulated MIPS slower than previously*.

When creating a new processor model, use a Fibonacci test case to validate that jump hints are working correctly. If performance is unchanged (or slower!) with jump hints present then they are not being used correctly.

12 Modeling Exceptions

In chapter 11, we implemented load and store instructions and noted that on the OR1K processor all loads and stores should be aligned to the load/store size. In this chapter, we will model the processor *exception behavior* that happens when an unaligned load or store is encountered.

This chapter also shows how to write exception handlers for arithmetic exceptions such as a divide by zero.

12.1 Basic Example

Firstly, we will examine the simulator behavior when no special action is taken to handle exceptions. Directory `10.or1kBehaviorExceptions/application` contains the following example in file `asmtest.S`:

```
.org 0x200
// Alignment Exception Handler (AT 0x200)
// increment count of alignment exceptions
l.addi    r30,r30,1
l.addi    r1,r1,1
l.rfe
// return from exception

.org 0x10000
// Application Code (AT 0x10000)

.global _start
_start:
    l.ori    r30,r0,0        // r30 = 0 (counts alignment exceptions)
    l.movhi  r1,0x8000       // r1 = 0x80000000
    l.movhi  r2,0x1234       // r2 = 0x12340000
    l.ori    r2,r2,0x5678    // r2 = 0x12345678
    l.ori    r3,r0,0        // r3 = 0 (loop count)

loop:
    l.sb     0(r1),r2        // do one-byte store
    l.sh     0(r1),r2        // do two-byte store
    l.sw     0(r1),r2        // do four-byte store

    l.addi    r3,r3,1        // increment loop count
    l.sfeqi   r3,10         // r3==10?
    l.bnf     loop          // go if not
    l.addi    r1,r1,1        // increment store address (delay slot)

    l.div     r30,r30,r0     // divide by zero

.global exit
exit:
    l.nop
```

This example uses `r1` to hold a write pointer, initially at address `0x80000000`. It then executes a loop ten times, trying to do one-byte, two-byte and four-byte writes to the pointer. Each time round the loop the pointer `r1` is incremented. Finally, it performs an arithmetic divide-by-zero.

Generate the assembler test case in directory 10.or1kBehaviorExceptions and create the test platform as follows:

```
make
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```
Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 'cpul', 0x00000000000010000: l.ori    r30,r0,0x0
Info 'cpul', 0x00000000000010004: l.movhi  r1,0x8000
Info 'cpul', 0x00000000000010008: l.movhi  r2,0x1234
Info 'cpul', 0x0000000000001000c: l.ori    r2,r2,0x5678
Info 'cpul', 0x00000000000010010: l.ori    r3,r0,0x0
Info 'cpul', 0x00000000000010014: l.sb     0x0(r1),r2
Info 'cpul', 0x00000000000010018: l.sh     0x0(r1),r2
Info 'cpul', 0x0000000000001001c: l.sw     0x0(r1),r2
Info 'cpul', 0x00000000000010020: l.addi   r3,r3,0x1
Info 'cpul', 0x00000000000010024: l.sfeqi  r3,0xa
Info 'cpul', 0x00000000000010028: l.bnf    0x00010014
Info 'cpul', 0x0000000000001002c: l.addi   r1,r1,0x1
Info 'cpul', 0x00000000000010014: l.sb     0x0(r1),r2
Info 'cpul', 0x00000000000010018: l.sh     0x0(r1),r2
Processor Exception (PC_PRX) Processor 'cpul' 0x10018: l.sh     0x0(r1),r2
Processor Exception (PC_WAX) Misaligned 2-byte write to 0x80000001
-----
R0 : 00000000   R1 : 80000001   R2 : 12345678   R3 : 00000001
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : deadbeef   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 00000000   R31: deadbeef
PC : 00010018   SR : 00008001   ESR: deadbeef   EPC: deadbeef
BF:0 CF:0 OF:0
-----
processor has executed 14 instructions
```

The test case runs successfully until an attempt is made to perform a 2-byte store at a 1-byte-aligned address. Then the processor terminates with a Processor Exception error, which is signaled to the platform by returning the value ICM_SR_WR_ALIGN from icmSimulate. File platform/platform.c in directory 10.or1kBehaviorExceptions has been updated like this to handle the possible exception case:

```
static Bool simulate(icmProcessorP processor, Uns64 instructions) {
    icmStopReason stopReason = icmSimulate(processor, instructions);
    switch(stopReason) {
        case ICM_SR_SCHED:
```

```

        // hit the scheduler limit
        return True;

    case ICM_SR_EXIT:
        // processor has exited
        return False;

    case ICM_SR_FINISH:
        // simulation must end
        return False;

    case ICM_SR_RD_PRIV:
    case ICM_SR_WR_PRIV:
    case ICM_SR_RD_ALIGN:
    case ICM_SR_WR_ALIGN:
    case ICM_SR_ARITH:
        // unhandled processor exception: simulation must end
        return False;

    default:
        icmPrintf("unimplemented stopReason %u\n", stopReason);
        return False;
}

```

Note that the only action taken in the model to enforce alignment checking was to pass `True` to the `checkAlign` parameter to each of the load/store morph-time functions: the simulator did the rest automatically.

It is possible instead to envisage using `vmimt` calls to explicitly construct an alignment check (for example, the address to which to store could be created in a temporary register using `vmimtBinopRRC`, the low-order bits of the temporary register value could be extracted using `vmimtBinopRRC` with a `vmiBinop` of `vmi_AND` and a constant value 3, this could be checked for zero and the simulation terminated by `vmimtFinish` if not zero). This is not however a good approach: the simulator's built in checks are much more efficient than any model code that tries to do the same thing. *Always use simulator alignment checking capabilities in preference to coding your own.*

A real processor does not of course “exit” when an exception is encountered: it typically enters a privileged mode and jumps to an exception vector instead. The model files in directory `10.or1kBehaviorExceptions` have been enhanced to support this as follows.

12.1.1 File `or1kStructure.h`

To support exceptions, the OR1K has three new registers we need to model now:

1. `esr` (exception status register): this saves the value of the `sr` register on exception entry;
2. `epc` (exception program counter register): this saves the current program counter on exception entry.
3. `eeaddr` (exception effective address register): for an exception caused by an invalid memory operation, this records the address that caused the exception.

These have been added to the processor structure:

```
typedef struct orlKs {  
  
    Bool        carryFlag;        // carry flag  
    Bool        overflowFlag;     // overflow flag  
    Bool        branchFlag;       // branch flag  
  
    Uns32       regs[OR1K_REGS];  // basic registers  
  
    Uns32       SR;               // status register  
    Uns32       ESR;              // exception status register  
    Uns32       EPC;              // exception program counter register  
    Uns32       EEAR;             // exception effective address register  
  
    vmiBusPortP busPorts;         // bus port descriptions  
  
} orlK, *orlKP;
```

There are also new macros to access the registers when morphing code:

```
#define OR1K_ESR    OR1K_OFFSET(ESR)  
#define OR1K_EPC    OR1K_OFFSET(EPC)  
#define OR1K_EEAR   OR1K_OFFSET(EEAR)
```

Initialization routines (in `orlKMain.c`) and register dump routines (in `orlKUtils.c`) have also been modified to handle the new registers.

12.1.2 File `orlKFunctions.h`

Prototypes for model routines that are called when the simulator detects a memory exception have been added:

```
VMI_RD_PRIV_EXCEPT_FN(orlKRdPrivExceptionCB);  
VMI_WR_PRIV_EXCEPT_FN(orlKWrPrivExceptionCB);  
VMI_RD_ALIGN_EXCEPT_FN(orlKRdAlignExceptionCB);  
VMI_WR_ALIGN_EXCEPT_FN(orlKWrAlignExceptionCB);
```

The memory access handlers are called for load/store *privilege* exceptions and load/store *alignment* exceptions. In this example, we are only interested in alignment exceptions but privilege exceptions will be implemented as well for completeness.

```
VMI_ARITH_EXCEPT_FN(orlKArithExceptionCB);
```

This handler is called when the simulator detects an arithmetic exception at run time (for example, a divide by zero).

12.1.3 File `orlKExceptionTypes.h`

This is a new file giving information about the exceptions on the OR1K. There is an enumeration of the possible exception types:

```
typedef enum orlKExceptionE {  
    OR1K_EXCPT_RST,        // reset  
    OR1K_EXCPT_BUS,        // alignment
```

```

    OR1K_EXCPT_DPF,          // data privilege
    OR1K_EXCPT_IPF,          // instruction privilege
    OR1K_EXCPT_TTI,          // tick timer
    OR1K_EXCPT_ILL,          // illegal instruction
    OR1K_EXCPT_EXI,          // external interrupt
    OR1K_EXCPT_SYS,          // system call
    OR1K_EXCPT_LAST          // KEEP LAST: for sizing
} orlkException;

```

Vector addresses for each exception are also defined:

```

#define RST_ADDRESS 0x100    // reset exception vector
#define BUS_ADDRESS 0x200    // alignment exception vector
#define DPF_ADDRESS 0x300    // data privilege exception vector
#define IPF_ADDRESS 0x400    // instruction privilege exception vector
#define TTI_ADDRESS 0x500    // tick timer exception vector
#define ILL_ADDRESS 0x700    // illegal instruction exception vector
#define EXI_ADDRESS 0x800    // external interrupt exception vector
#define SYS_ADDRESS 0xc00    // sys exception vector

```

12.1.4 File `orlkExceptions.c`

This file implements the exception handler callbacks, which are called at run time when a potential simulated exception occurs. The purpose of the callbacks is to put the processor into the state that it would enter if the same exception was encountered on the real hardware: typically, this means entering a privileged mode, saving some exception context state and jumping to an exception vector address. This is exactly what we will implement now for the OR1K.

It is possible for *multiple exception conditions to be encountered in a single simulated instruction*: for example, a store may be attempted to an address that is both misaligned and read-only. To handle this situation, the memory exception handlers work as follows:

1. The *alignment* handler is called first. This returns an unsigned result indicating whether the privilege exception handler should be called subsequently;
2. If the alignment handler returns *non-zero*, and there is also a privilege exception condition, then the *privilege* exception handler will be called. If the alignment handler returns *zero*, the privilege exception handler will *not* be called.

A non-zero result from the alignment handler may also indicate that the load/store address requires *snapping*, or that the value to load or store requires *rotation*. This is discussed in detail later in this section.

The memory exception handler callbacks for the OR1K are as follows:

```

VMI_RD_PRIV_EXCEPT_FN(orkRdPrivExceptionCB) {
    if(MEM_AA_IS_TRUE_ACCESS(attrs)) {
        orlkP orlk = (orkP)processor;
        orlk->EEAR = (Uns32)address;
        orlkTakeException(ork, OR1K_EXCPT_DPF, 0);
    }
}

```



```
}

VMI_WR_PRIV_EXCEPT_FN(orkWrPrivExceptionCB) {
    if(MEM_AA_IS_TRUE_ACCESS(attrs)) {
        orkP ork = (orkP)processor;
        ork->EEAR = (Uns32)address;
        orkTakeException(ork, ORLK_EXCPT_DPF, 0);
    }
}

VMI_RD_ALIGN_EXCEPT_FN(orkRdAlignExceptionCB) {
    orkP ork = (orkP)processor;
    ork->EEAR = (Uns32)address;
    orkTakeException(ork, ORLK_EXCPT_BUS, 0);
    return 0;
}

VMI_WR_ALIGN_EXCEPT_FN(orkWrAlignExceptionCB) {
    orkP ork = (orkP)processor;
    ork->EEAR = (Uns32)address;
    orkTakeException(ork, ORLK_EXCPT_BUS, 0);
    return 0;
}
```

In this case, both alignment handlers return zero, which means that *alignment exceptions have priority over privilege exceptions* (in other words, a store to an address that is both misaligned and read-only will cause an alignment exception only). Each exception is implemented by a call to `orkTakeException` (implemented in `orkUtils.c`), passing the appropriate exception type (`ORKL_EXCPT_DPF` for privilege exceptions, `ORKL_EXCPT_BUS` for alignment exceptions) and a zero offset (explained in the description of `orkTakeException`). The faulting address is saved in the `ear` register in the processor structure.

The read and write privilege handlers are both passed an argument, `attrs`, of type `memAccessAttrs`, defined in `vmiTypes.h` as follows:

```
typedef enum memAccessAttrsE {
    MEM_AA_FALSE = 0x0, // this is an artifact access
    MEM_AA_TRUE  = 0x1, // this is a true processor access
    MEM_AA_FETCH = 0x2, // this access is a fetch
} memAccessAttrs;
```

The `memAccessAttrs` type tells the processor model what kind of access is being performed. There are four possible values:

1. `MEM_AA_TRUE`: this indicates that the exception handler is being called because of a true processor read or write, and that the model should take any action needed to model the exception.
2. `MEM_AA_FALSE`: this indicates that this access is not a true processor read or write, but is instead some kind of *artifact* access. For example, it might be an access being made by the simulator itself, or by an attached debugger reading memory. In this case, the processor model should not update its state to reflect an exception, but might need to take some other action to make the memory readable or writable. As an example, the OVP ARM processor model implements a TLB

model that maps memory pages on demand based on the contents of a page table stored in memory, and these mappings need to be made even for an artifact access (so that a debugger can query virtual memory address locations even if that virtual address is not currently mapped, for example).

3. `MEM_AA_TRUE|MEM_AA_FETCH`: this indicates that the exception handler is being called because of a true processor fetch. Processor state should be updated to model the exception.
4. `MEM_AA_FALSE|MEM_AA_FETCH`: this indicates that the exception handler is being called because of an artifact fetch (usually caused by the JIT code generation engine). In this case, the processor model should not update its state to reflect an exception, but might need to take some other action to make the memory readable or writable.

For the OR1K, the read and write exception handlers both validate that the access is a non-artifact access before updating any processor state:

```
VMI_RD_PRIV_EXCEPT_FN(orkRdPrivExceptionCB) {  
    if(MEM_AA_IS_TRUE_ACCESS(attrs)) {  
        orkP ork = (orkP)processor;  
        ork->EEAR = (Uns32)address;  
        orkTakeException(ork, OR1K_EXCPT_DPF, 0);  
    }  
}  
  
VMI_WR_PRIV_EXCEPT_FN(orkWrPrivExceptionCB) {  
    if(MEM_AA_IS_TRUE_ACCESS(attrs)) {  
        orkP ork = (orkP)processor;  
        ork->EEAR = (Uns32)address;  
        orkTakeException(ork, OR1K_EXCPT_DPF, 0);  
    }  
}
```

This OR1K model does not implement any structure such as a demand-mapped TLB, so no action is taken for artifact accesses.

File `orkExceptions.c` also implements an *arithmetic* exception handler:

```
VMI_ARITH_EXCEPT_FN(orkArithExceptionCB) {  
  
    orkP ork = (orkP)processor;  
  
    switch(exceptionType) {  
  
        // integer divide-by-zero and overflow should not generate exceptions  
        // but instead set the carry flag  
        case VMI_INTEGER_DIVIDE_BY_ZERO:  
        case VMI_INTEGER_OVERFLOW:  
            ork->carryFlag = 1;  
            return VMI_INTEGER_ABORT;  
  
        // not expecting any other arithmetic exception types  
        default:  
            return VMI_INTEGER_UNHANDLED;  
    }  
}
```

When an integer divide or overflow is encountered, the OR1K does not jump to an exception vector: instead, it indicates the error by setting the processor carry flag. Other processor types that jump to exception vectors can be simulated in a similar manner to the memory exception handlers (i.e. save the current program counter and other state, and then jump to an exception vector – see the discussion of `or1kTakeException` in section 12.1.5).

The return value from the arithmetic exception callback is an enumerated value defined in `vmiTypes.h`:

```
typedef enum vmiIntegerExceptionResultE {  
    VMI_INTEGER_UNHANDLED,      // not handled  
    VMI_INTEGER_ABORT,         // handled, abort current instruction  
    VMI_INTEGER_CONTINUE,      // handled, continue current instruction  
} vmiIntegerExceptionResult;
```

A return value of `VMI_INTEGER_UNHANDLED` indicates that the numeric exception was not expected by this model and simulation should terminate.

A return value of `VMI_INTEGER_ABORT` indicates that the handler accepted the exception, and simulation should abort the remainder of this simulated instruction and resume execution with the *next* simulated instruction (or at an exception vector address, if `vmirtSetPC` or `vmirtSetPCException` are used in the handler: see section 12.1.5).

A return value of `VMI_INTEGER_CONTINUE` indicates that the handler accepted the exception, and simulation should resume at the next *native* instruction address after the offending instruction.

When writing code that could cause simulated exceptions, or which makes an embedded call that could update the current program counter using `vmirtSetPC` or `vmirtSetPCException`, always remember that the part of the instruction *after* the embedded call or simulated exception *will not be executed* if the program counter has been modified by `vmirtSetPC` or `vmirtSetPCException`, or if there is an arithmetic exception for which the handler returns `VMI_INTEGER_ABORT`. *Care must be taken to leave the processor model in a consistent state in this case.*

As a contrived example, suppose that a processor is being modeled that has a single instruction that implements a pair of loads into registers from different addresses:

```
0x0020000: r1=(ra1), r2=(ra2)
```

The obvious way to implement this would be with two `vmimtLoadRRO` calls, for example:

```
vmimtLoadRRO(32, 32, 0, CPU_REG(r1), CPU_REG(ra1), endian, False, True);  
vmimtLoadRRO(32, 32, 0, CPU_REG(r2), CPU_REG(ra2), endian, False, True);
```

However, suppose that there is a memory access violation on the access using `ra2` (but not `ra1`) that caused control to be transferred to a simulated exception handler. In this case, the processor would be left in a state with the instruction half-executed, because the load to `r1` would already have been done.

To get correct model behavior in this case, the first load should save its result in a temporary, which is written to the target register only if the second load succeeds:

```
vmimtLoadRRO(32, 32, 0, CPU_TEMP1, CPU_REG(ra1), endian, False, True);
vmimtLoadRRO(32, 32, 0, CPU_REG(r2), CPU_REG(ra2), endian, False, True);
vmimtMoveRR(32, CPU_REG(r1), CPU_TEMP1);
```

12.1.5 Files `or1kUtils.h` and `or1kUtils.c`

The new routine `or1kTakeException` is implemented as:

```
void or1kTakeException(or1kP or1k, or1kException exception, Uns32 pcOffset) {
    Uns8 simD;
    Uns32 simPC = (Uns32)vmirtGetPCDS((vmiProcessorP)or1k, &simD);

    or1kEnterSupervisorMode(or1k);
    or1k->EPC = simPC + pcOffset;

    // set sr[DSX] for exception in a delay slot
    if(simD) {
        or1k->SR |= SPR_SR_DSX;
    }

    // jump to the vector
    vmirtSetPCEXception((vmiProcessorP)or1k, exceptions[exception].code);
}
```

Because this routine is called at run time (as opposed to morph time) it uses functions from the *Imperas Run Time Function API* to update the processor model state. In detail, it works as follows:

```
Uns8 simD;
Uns32 simPC = (Uns32)vmirtGetPCDS((vmiProcessorP)or1k, &simD);
```

The function `vmirtGetPCDS` returns the currently-executing instruction address with any delay-slot byte offset. For non-delay-slot instructions, `simPC` will be set to the current instruction address and `simD` set to zero. For delay-slot instructions, `simPC` will be set to *the address of the preceding jump or branch instruction* and `simD` will be *the byte offset of the current instruction from the preceding jump or branch*. Since all OR1K instructions are four bytes long, `simD` will therefore be 4 for a delay-slot instruction.

```
or1kEnterSupervisorMode(or1k);
```

This is a routine that puts the simulated processor into supervisor mode, described below.

```
or1k->EPC = simPC + pcOffset;
```

This line saves the current program counter in register `epc` (or the jump/branch instruction address for delay slot instructions). A call-specific offset is added to the value saved (this value is zero for the memory exceptions).

```
if(simD) {
    orlk->SR |= SPR_SR_DSX;
}
```

These lines set a special bit in the status register `sr` if the exception occurred in a delay slot instruction. Recovery from delay slot instruction exceptions requires special processing in application exception handlers, so they need some way to find out whether the original exception was in a delay slot instruction or not.

```
vmirtSetPCException((vmiProcessorP)ork, exceptions[exception].code);
```

This line uses `vmirtSetPCException` to force the processor to jump to the exception vector address associated with the exception type. A table maps exception types to vector addresses:

```
#define ORLK_EXCEPTION_INFO(_D) [ORK_EXCPT_##_D] = {name:##_D,code:##_D##_ADDRESS}

static const vmiExceptionInfo exceptions[ORK_EXCPT_LAST] = {
    ORLK_EXCEPTION_INFO(RST),
    ORLK_EXCEPTION_INFO(BUS),
    ORLK_EXCEPTION_INFO(DPF),
    ORLK_EXCEPTION_INFO(IPF),
    ORLK_EXCEPTION_INFO(TTI),
    ORLK_EXCEPTION_INFO(ILL),
    ORLK_EXCEPTION_INFO(EXI),
    ORLK_EXCEPTION_INFO(SYS),
};
```

The exceptions are described using an array of `vmiExceptionInfo` type structures (defined in `vmiTypes.h`). This structure type will be required when adding debugger integration support routines to the model (see chapter 17).

The new routine `orkEnterSupervisorMode` saves `sr` in `esr`, updates `sr` to mask out various exceptions that must be disabled in supervisor mode, and indicates that we are in supervisor mode by setting the `SM` bit in `sr`, like this:

```
void orkEnterSupervisorMode(orkP ork) {

    const Uns32 clearBits = (
        SPR_SR_IEE |    // interrupt enable
        SPR_SR_TEE |    // tick timer enable
        SPR_SR_DME |    // data MMU enable
        SPR_SR_IME |    // instruction MMU enable
        SPR_SR_OVE |    // overflow exception enable
    );

    // save the current status register in esr
    ork->ESR = orkGetSR(ork);
```

```
// mask out the 'clear' bits and mask in supervisor mode
orlk->SR = (orlk->ESR & ~clearBits) | SPR_SR_SM;
}
```

Note that `orlkGetSR` is used to get the value of the `sr` register. This routine ensures that the flag bits are present in the returned value.

12.1.6 File `orlkAttrs.c`

File `orlkAttrs.c` has been updated to include references to the five new exception handler callbacks in the `modelAttrs` structure:

```
const vmiIASAttr modelAttrs = {

    . . . skipped lines . . .

    //////////////////////////////////////
    // EXCEPTION ROUTINES
    //////////////////////////////////////

    .rdPrivExceptCB = orlkRdPrivExceptionCB,
    .wrPrivExceptCB = orlkWrPrivExceptionCB,
    .rdAlignExceptCB = orlkRdAlignExceptionCB,
    .wrAlignExceptCB = orlkWrAlignExceptionCB,
    .arithExceptCB = orlkArithExceptionCB,

    . . . skipped lines . . .
};
```

12.1.7 `l.rfe` and `l.sys` Instructions

The model has been enhanced to implement the `l.rfe` instruction. This OR1K instruction performs a return from an exception handler: it copies register `esr` to register `sr` and performs an unconditional jump to the address stored in `epc`. Decode and disassembly for this instruction are very similar to previous instructions so no further details will be given here. The functionality of the `l.rfe` instruction is implemented in `orlkMorph.c` like this:

```
static void doRFE(void) {

    // set sr from esr (must call orlkSetSR to do this)
    vmimtArgProcessor();
    vmimtArgReg(32, OR1K_ESR);
    vmimtCall((vmiCallFn)orlkSetSR);

    // return to exception program counter
    vmimtUncondJumpReg(0, OR1K_EPC, VMI_NOREG, vmi_JH_RETURNINT);
}
```

The first part of `doRFE` constructs a *call* to function `orlkSetSR`, passing the current processor as the first argument and the value of register `esr` from the processor model structure as the second argument. The effect of this is to assign the current value of `esr` to register `sr`. Remember that whenever `sr` is set, we must use `orlkSetSR` to do it because assigning to this register implicitly sets the flag fields we maintain separately in the model.

The second part of `doRFE` uses `vmimtUncondJumpReg` to perform a jump to the address in the `epc` register. Note that the `l.rfe` instruction is not followed by a delay slot instruction. As this instruction implements a return from an exception handler, a new jump hint type is used - `vmi_JH_RETURNINT`.

In general, it is possible to emit code to call *any* function from morphed code by using a sequence of `vmimtArg`-prefixed functions followed by a call to `vmimtCall`. This means that for many instructions there is an important implementation choice to be made: is it best to implement the instruction directly using `vmimtBinopRRR`, `vmimtBinopRRC` etc, or is it best to use `vmimtCall` to call a C function to do the work instead?

In general, the rule is that *if the behavior of the instruction requires more than a few `vmimt`-prefixed calls to implement, or is difficult to encode using `vmimt` operations, then use `vmimtCall` and a C function to implement the instruction behavior.*

One important exception is that, wherever possible, jumps should be implemented using `vmimt` jump primitives instead of using `vmirtSetPC`, which is significantly slower. Note that a single instruction can be implemented using a mixture of `vmimt` primitive operations and `vmirtCall` calls, as in this example.

The `l.sys` instruction is also now implemented (though not required for this example). This instruction enters supervisor mode and jumps to an exception vector at address `0xc00`, saving the *next instruction* address in `epc`. It is implemented in `or1kMorph.c` as:

```
static void doSYS(Uns32 instr, Uns32 thisPC) {  
  
    vmimtArgProcessor();  
    vmimtArgUns32(OR1K_EXCPT_SYS);  
    vmimtArgUns32(4);  
    vmimtCall((vmiCallFn)or1kTakeException);  
  
    // terminate the current code block  
    vmimtEndBlock();  
}
```

Note that the value 4 is passed as the third argument to `or1kTakeException` to ensure that the saved address in `epc` is the *next instruction* address.

The call to `or1kTakeException` in `doSYS` is followed by `vmimtEndBlock` to indicate to the simulator that it should not attempt to create a single block of translated code coalescing the `l.sys` instruction with the following instruction, as these two instructions will not be executed sequentially – control is always transferred to the system call handler by an `l.sys` instruction. Refer to the *Imperas VMI Morph Time Function Reference* for more detailed information about when to use `vmimtEndBlock`.

If the test case already contained code to implement simulated exceptions, why did the original run at the start of this chapter exit from `icmSimulate` with an unhandled

processor exception? The reason is that whether or not simulated exceptions should be enabled is specified by a model flag in the platform (most application code should not generate exceptions in the normal case, and it is usually desired that any exception is an error that should stop simulation).

To enable simulated exception modeling, modify the `MODEL_FLAGS` definition in `10.orlkBehaviorExceptions/platform/platform.c` as follows:

```
#define MODEL_FLAGS (ICM_ATTR_TRACE | ICM_ATTR_SIMEX)
```

Then rebuild the platform and resimulate:

```
make -C platform
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should now be as follows:

```
Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 'cpul', 0x0000000000010000: l.ori    r30,r0,0x0
Info 'cpul', 0x0000000000010004: l.movhi  r1,0x8000
Info 'cpul', 0x0000000000010008: l.movhi  r2,0x1234
Info 'cpul', 0x000000000001000c: l.ori    r2,r2,0x5678
Info 'cpul', 0x0000000000010010: l.ori    r3,r0,0x0
Info 'cpul', 0x0000000000010014: l.sb     0x0(r1),r2
Info 'cpul', 0x0000000000010018: l.sh     0x0(r1),r2
Info 'cpul', 0x000000000001001c: l.sw     0x0(r1),r2
Info 'cpul', 0x0000000000010020: l.addi    r3,r3,0x1
Info 'cpul', 0x0000000000010024: l.sfeqi   r3,0xa
Info 'cpul', 0x0000000000010028: l.bnf     0x00010014
Info 'cpul', 0x000000000001002c: l.addi    r1,r1,0x1
Info 'cpul', 0x0000000000010014: l.sb     0x0(r1),r2
Info 'cpul', 0x0000000000010018: l.sh     0x0(r1),r2
Info 'cpul', 0x0000000000000200: l.addi    r30,r30,0x1
Info 'cpul', 0x0000000000000204: l.addi    r1,r1,0x1
Info 'cpul', 0x0000000000000208: l.rfe
Info 'cpul', 0x0000000000010018: l.sh     0x0(r1),r2
Info 'cpul', 0x000000000001001c: l.sw     0x0(r1),r2
Info 'cpul', 0x0000000000000200: l.addi    r30,r30,0x1
Info 'cpul', 0x0000000000000204: l.addi    r1,r1,0x1
Info 'cpul', 0x0000000000000208: l.rfe
Info 'cpul', 0x000000000001001c: l.sw     0x0(r1),r2
Info 'cpul', 0x0000000000000200: l.addi    r30,r30,0x1
Info 'cpul', 0x0000000000000204: l.addi    r1,r1,0x1
Info 'cpul', 0x0000000000000208: l.rfe
Info 'cpul', 0x000000000001001c: l.sw     0x0(r1),r2
Info 'cpul', 0x0000000000010020: l.addi    r3,r3,0x1
Info 'cpul', 0x0000000000010024: l.sfeqi   r3,0xa
. . . etc . . .
Info 'cpul', 0x0000000000010020: l.addi    r3,r3,0x1
Info 'cpul', 0x0000000000010024: l.sfeqi   r3,0xa
Info 'cpul', 0x0000000000010028: l.bnf     0x00010014
Info 'cpul', 0x000000000001002c: l.addi    r1,r1,0x1
Info 'cpul', 0x000000000001002c: l.div     r30,r30,r0
Info 'cpul', 0x0000000000010030: l.nop     0x0
Processor 'cpul' terminated at 'exit', address 0x10034
-----
R0 : 00000000   R1 : 80000025   R2 : 12345678   R3 : 0000000a
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
```



```

R8 : deadbeef    R9 : deadbeef    R10: deadbeef    R11: deadbeef
R12: deadbeef    R13: deadbeef    R14: deadbeef    R15: deadbeef
R16: deadbeef    R17: deadbeef    R18: deadbeef    R19: deadbeef
R20: deadbeef    R21: deadbeef    R22: deadbeef    R23: deadbeef
R24: deadbeef    R25: deadbeef    R26: deadbeef    R27: deadbeef
R28: deadbeef    R29: deadbeef    R30: 0000001b    R31: deadbeef
PC : 00010034    SR : 00008201    ESR: 00008001    EPC: 0001001c
BF:1 CF:1 OF:0
-----
processor has executed 185 instructions

```

Now, instead of terminating after 14 instructions, the processor jumps to the exception handler at address 0x200 instead (execution within the exception handler is highlighted in bold in the trace above for clarity). The exception handler increments the address in `r1` and returns using `l.rfe`, which re-executes the faulting instruction (obviously, in a real application the exception handler would do something more sensible than this, but the example is sufficient for demonstration purposes). As a side effect, the exception handler also increments `r30` so that on termination we have a count of the number of times it was called (0x1b, i.e. 27 times).

The divide-by-zero in the penultimate instruction now does not cause the simulation to exit, but instead sets the carry flag.

12.2 Misaligned Load/Store Address Snapping and Value Rotation

On some processors, loads and stores to misaligned addresses do not cause exceptions but are instead *snapped* to the correct alignment for the data size (so two-byte load addresses are rounded down to a two-byte boundary, four-byte load addresses to a four-byte boundary, and so on). In addition, some processors (e.g. old ARM processors) *rotate* values read from misaligned addresses, with the rotate amount based on the misaligned byte offset.

One way to implement address snapping would be to copy each load/store address to a temporary, mask it to the appropriate size using `vmimtBinopRC` with a `vmiBinop` of `vmi_AND`, and then use this resulting address as an argument to the load/store operation. Unfortunately, this approach adds significant run-time overhead to each memory access. Implementing value rotation is even more complex.

A better solution is possible using *read and write address snap handlers* defined in the processor model. These are defined using a macro in `vmiAttrs.h`:

```

#define VMI_RD_WR_SNAP_FN(_NAME) Uns32 _NAME( \
    vmiProcessorP processor, \
    Addr          address, \
    Uns32         bytes \
)
typedef VMI_RD_WR_SNAP_FN( (*vmiRdWrSnapFn) );

```

The read and write handlers are specified using the `rdSnapCB` and `wrSnapCB` fields in the processor attributes structure, respectively:

```
vmiRdWrSnapFn  rdSnapCB;           // read alignment snap function
vmiRdWrSnapFn  wrSnapCB;           // write alignment snap function
```

The return value from each handler is an integer which indicates what address snapping or value rotation is required. The return value is constructed using macro `MEM_SNAP` in `vmiTypes.h`:

```
#define MEM_SNAP(_SNAP, _ROTATE) (((Uns8)(_SNAP)) | ((_ROTATE)<<8))
```

In this macro, `_SNAP` specifies an *address rounding granularity* in bytes (typically 1, 2, 4 or 8), and `_ROTATE` specifies a value rotation in bits. As an example:

```
MEM_SNAP(4, 24)
```

indicates that a read/write address should be snapped to 4-byte alignment. In addition, a value being written should be rotated left by 24 bits before it is written, and a value being read should be rotated left by 24 bits before being assigned to a processor register.

In detail, the read and write snap handlers are used by the simulator as follows:

1. If an access is made to a misaligned address, any defined *address snap handler* is called first. If the handler is defined and returns *non-zero*, then the read or written value is modified using the granularity and rotation specified by the result.
2. Otherwise (if there is no address snap handler, or the address snap handler returns zero) any defined *align exception handler* is called for a misaligned address access. This should either return 0 (if the read or write should be terminated, possibly because an exception is taken) or 1 (if the read or write should proceed, possibly with a modified value)⁴.
3. If the read/write address has insufficient privileges, and either the address was aligned, or the snap handler or align exception handler returns non-zero, then the *privilege exception handler* is called.

12.2.1 ARM Model Load/Store Address Snap Callback

This is the read snap address callback from the OVP ARM processor model:

```
VMI_RD_WR_SNAP_FN(armRdSnapCB) {
    armP          arm = (armP)processor;
    armUnalignedAction ua;

    if(ALIGN_ENABLED(arm) || ((ua=getUnalignedAction(arm))==ARM_UA_DABORT)) {
        // take exception
    }
```

⁴ In fact, the align exception handler returns a granularity/rotate value in the same format as for the snap handler. A return value of 1 therefore indicates 1-byte alignment with zero rotation.

```
    return 0;

} else if((ua==ARM_UA_ROTATE) && arm->configInfo.rotateUnaligned) {

    // read snaps address and loads rotated value
    Uns32 rotate = address&(bytes-1);

    if(getEndian(arm)==MEM_ENDIAN_LITTLE) {
        rotate = bytes - rotate;
    }

    return MEM_SNAP(bytes, rotate*8);

} else {

    // read snaps address
    return MEM_SNAP(bytes, 0);

}
```

This callback does three things:

1. If the current instruction should cause an alignment exception, it returns 0 (so that the read alignment exception handler will be called);
2. Otherwise, if this is an ARM variant in which unaligned reads cause rotation of the read value, it calculates the required rotation based on the address and returns a result aligned to the item byte size with that rotation;
3. Otherwise, it returns a result with an aligned address but no rotation.

12.3 Memory Aborts

In addition to alignment and privilege exceptions, there is one other type of exception that can be handled in a processor model: a *memory abort*. Memory aborts are generated by the memory subsystem, typically when there is no implemented memory at a particular address. Read and write abort handlers are specified using the `rdAbortExceptCB` and `wrAbortExceptCB` fields in the processor attributes structure, respectively:

```
vmiRdAbortExceptFn rdSnapCB; // read abort exception
vmiRdAbortExceptFn wrSnapCB; // write abort exception
```

The read and write abort handlers are called in one of two circumstances:

1. When a read or write privilege exception handler indicates the access should be retried, but the simulator determines that there is no accessible memory at the faulting address.
2. When an externally-implemented memory model indicates that a memory access has not succeeded (for example, by calling `icmAbortRead` or `icmAbortWrite`).

In the first case, read and write privilege exception handlers can indicate that a read or write should be retried on completion using a by-ref argument, `action`, which should be set to the value `VMI_LOAD_STORE_CONTINUE`. As an example, here is the read privilege exception handler from the OVP ARM model:

```
VMI_RD_PRIV_EXCEPT_FN(armRdPrivExceptionCB) {
```

```
armP arm = (armP)processor;

if(!armVMMiss(arm, domain, MEM_PRIV_R, address, bytes, attrs)) {
    *action = VMI_LOAD_STORE_CONTINUE;
}
}
```

The function `armVMMiss` attempts to map the faulting address using either a TLB or MPU entry, returning `True` if the address could not be mapped (indicating a miss). If there is no miss, the function uses the `action` argument to indicate that the load should be retried.

The read abort handler in the OVP ARM model triggers an external memory abort:

```
VMI_RD_ABORT_EXCEPT_FN(armRdAbortExceptionCB) {
    armP arm = (armP)processor;
    armExternalMemoryAbort(arm, address, isFetch ? MEM_PRIV_X : MEM_PRIV_R);
}
```

The full load/store exception escalation process, including address snapping, alignment, privilege and abort handlers, is shown in the following figure.

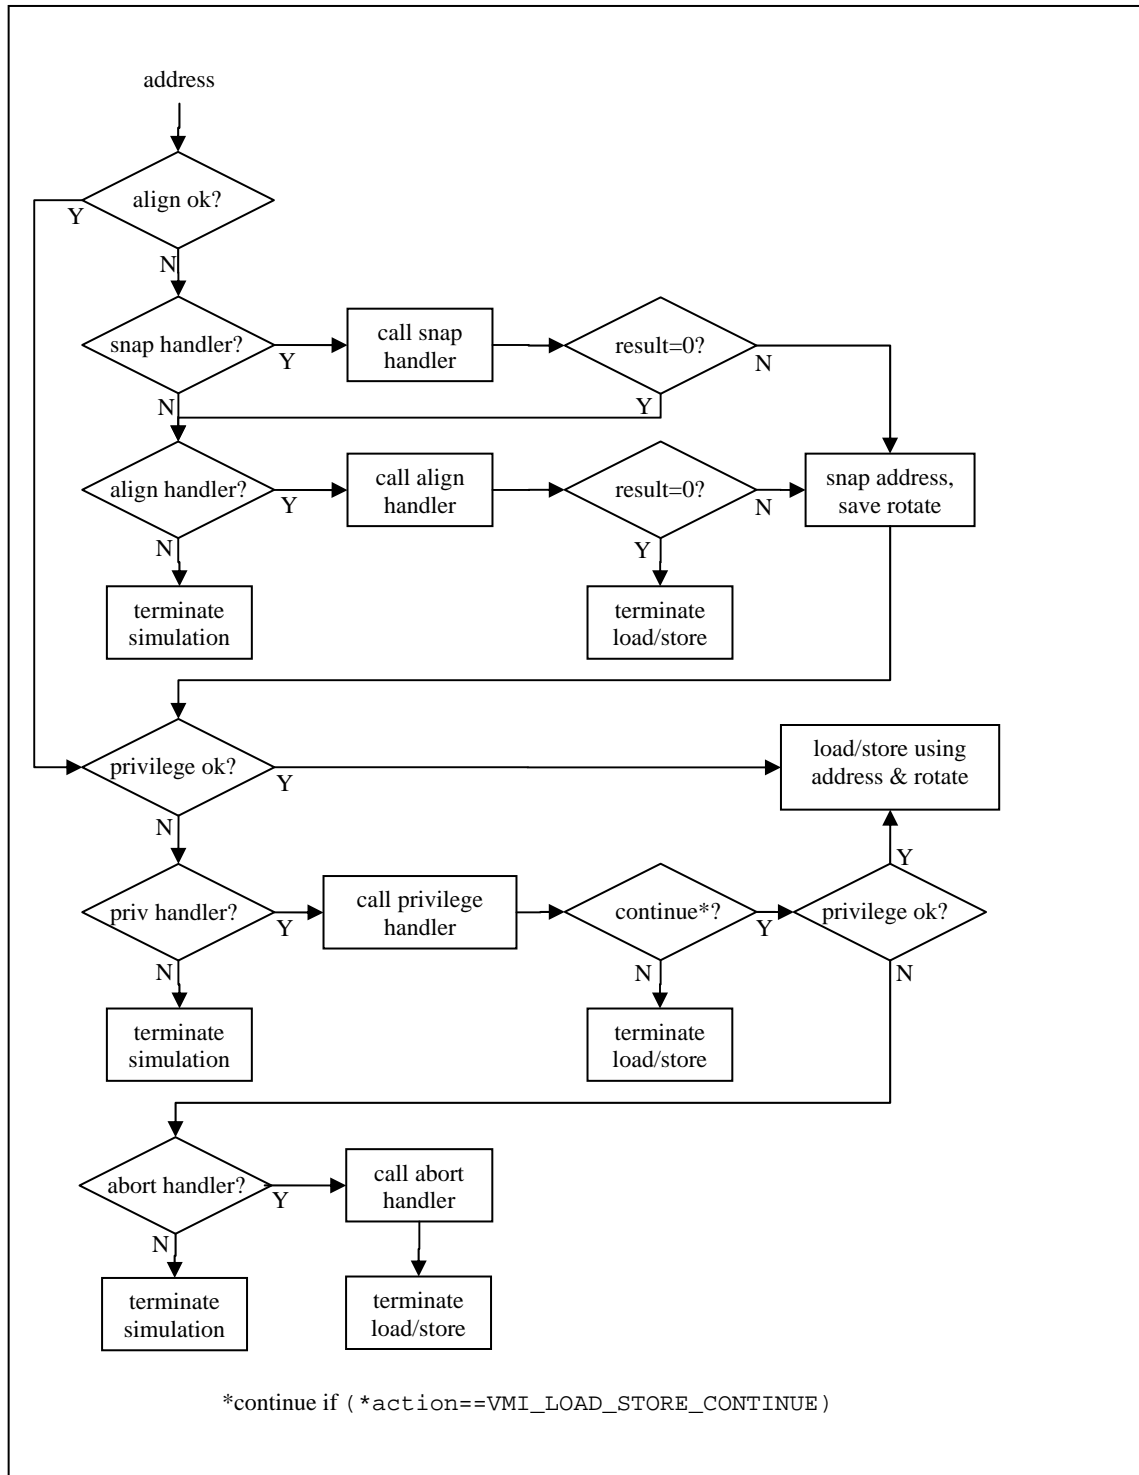


Figure 1: Load/Store Address Snapping and Exception Flow

12.4 Misaligned Fetch Address Snapping

Some processors snap misaligned fetch addresses to even boundaries. For example, the ARC 600/700 series processors snap all fetch addresses to a 2-byte aligned boundary.

There is a specific callback, `fetchSnapCB`, in the `vmiIASAttr` structure for fetch address snapping. The *address snapping callback* should be defined using the macro `VMI_FETCH_SNAP_FN`, defined in `vmiAttrs.h` as:

```
#define VMI_FETCH_SNAP_FN(_NAME) Addr _NAME( \
    vmiProcessorP processor, \
    Addr          thisPC    \
)
typedef VMI_FETCH_SNAP_FN(*vmiFetchSnapFn);
```

The address snapping callback takes a processor and an address argument and should return that address, appropriately snapped. For example, to snap addresses to a 2-byte boundary:

```
VMI_FETCH_SNAP_FN(orkFetchSnap) {
    return thisPC & ~1;
}
```

The `vmiIASAttr` structure should reference the fetch address snapping callback:

```
const vmiIASAttr modelAttrs = {
    . . . fields omitted . . .

    //////////////////////////////////////
    // MORPHER CORE ROUTINES
    //////////////////////////////////////

    .morphCB      = orkMorphInstruction,
    .fetchSnapCB = orkFetchSnap,

    . . . fields omitted . . .
};
```

13 Modeling Mode-Dependent Behavior (Part 1)

Up to now, all processor instructions have been modeled in a mode-independent way: the actions performed by each instruction have been independent of the current processor state. In real processors, there are usually instructions for which this is not the case. For example, some instructions may be intended for use only in a kernel or supervisor mode, and any attempt to use those instructions in user mode will generate a privileged instruction exception. For the OR1K, one such instruction has already been encountered: `l.rfe`, which should in fact only allow a return from exception in supervisor mode (so the implementation in chapter 12 was incorrect as it takes no account of this).

In this chapter, we will correct the functionality of `l.rfe` so that it takes account of the processor mode and also implement two mode modal instructions, `l.mfspr` and `l.mtspr`. Chapter 14 shows how modal instructions can be modeled differently to give higher performance.

13.1 The Template Modal Model

A template model for the OR1K processor implementing modal instructions can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/11.or1kBehaviorSPR
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/11.or1kBehaviorSPR .
```

Compile the model using make:

```
cd 11.or1kBehaviorSPR
make OPT=1
```

Note that the processor model has been built with compiler optimizations enabled (`OPT=1`) for this example, to get the fastest possible model. This is because we will use the model for performance testing at the end of this chapter.

The processor model is based on the previous model, with the changes listed in following sections.

13.2 Correcting `l.rfe` Behavior

The `l.rfe` instruction should return from an exception only in the case that the processor is in supervisor mode. In user mode, the processor should take an illegal instruction exception (at vector address `0x700`). To implement this behavior, function `doRFE` in `or1kMorph.c` has been updated as follows:

```
static void doRFE(void) {
    vmiLabelP inUserMode = vmimtNewLabel();
```

```

// test the SPR_SR_SM bit in OR1K_SR, setting OR1K_TEMPFLAG
vmimtBinopRRC(OR1K_BITS, vmi_AND, VMI_NOREG, OR1K_SR, SPR_SR_SM, &flagsTZ);

// go to label inUserMode if tempFlag set (SPR_SR_SM bit is zero)
vmimtCondJumpLabel(OR1K_TEMPFLAG, True, inUserMode);

////////////////////////////////////
// HERE IN SUPERVISOR MODE
////////////////////////////////////

// set sr from esr (must call or1kSetSR to do this)
vmimtArgProcessor();
vmimtArgReg(OR1K_BITS, OR1K_ESR);
vmimtCall((vmiCallFn)or1kSetSR);

// return to exception program counter
vmimtUncondJumpReg(0, OR1K_EPC, VMI_NOREG, vmi_JH_RETURNINT);

////////////////////////////////////
// HERE IN USER MODE
////////////////////////////////////

// insert the label targeted by vmimtCondJumpLabel above
vmimtInsertLabel(inUserMode);

// take illegal instruction exception
vmimtArgProcessor();
vmimtArgUns32(OR1K_EXCPT_ILL);
vmimtArgUns32(0);
vmimtCall((vmiCallFn)or1kTakeException);
}

```

This function uses an *intra-instruction jump* to execute one of two morphed-code subsequences depending on the setting of the SPR_SR_SM bit in the SR register. In detail, it works as follows:

```
vmiLabelP inUserMode = vmimtNewLabel();
```

This allocates a *label* that is used as a target of an intra-instruction jump.

```

// test the SPR_SR_SM bit in OR1K_SR, setting OR1K_TEMPFLAG
vmimtBinopRRC(OR1K_BITS, vmi_AND, VMI_NOREG, OR1K_SR, SPR_SR_SM, &flagsTZ);

```

This morphs code to perform a bitwise-and of the SR register (identified using OR1K_SR) and the constant SPR_SR_SM (defined in or1kStructure.h). Because argument 3 of vmimtBinopRRC is VMI_NOREG, the result is discarded. The last (flags) argument of vmimtBinopRRC is passed this vmiFlags structure:

```

const vmiFlags flagsTZ = {
    VMI_NOFLAG_CONST,          // carry in flag not used
    {
        VMI_NOFLAG_CONST,      // carry out flag not used
        VMI_NOFLAG_CONST,      // parity flag not used
        OR1K_TEMPFLAG_CONST,    // offset to zero flag
        VMI_NOFLAG_CONST,      // sign flag not used
        VMI_NOFLAG_CONST       // overflow flag not used
    }
}

```



```
};
```

The `vmiFlags` structure specifies that all flags generated by the bitwise-and should be discarded, except for the `zero` flag, which should be stored in a new temporary flag register in the OR1K processor structure (in `or1kStructure.h`):

```
typedef struct or1kS {

    Bool        carryFlag;        // carry flag
    Bool        overflowFlag;     // overflow flag
    Bool        branchFlag;       // branch flag
    Bool        tempFlag;          // temporary flag

    Uns32       regs[OR1K_REGS];  // basic registers

    Uns32       SR;               // status register
    Uns32       ESR;              // exception status register
    Uns32       EPC;              // exception program counter register
    Uns32       EEAR;             // exception effective address register

    vmiBusPortP busPorts;         // bus port descriptions
} or1k, *or1kP;

// macros to specify target registers in VARIABLE expressions
#define OR1K_OFFSET(_F)          VMI_CPU_REG(or1kP, _F)
#define OR1K_REG(_R)             OR1K_OFFSET(regs[_R])
#define OR1K_CARRY               OR1K_OFFSET(carryFlag)
#define OR1K_OVERFLOW            OR1K_OFFSET(overflowFlag)
#define OR1K_BRANCH              OR1K_OFFSET(branchFlag)
#define OR1K_TEMPFLAG          OR1K_OFFSET(tempFlag)
#define OR1K_LINKREG             OR1K_REG(OR1K_LINK)
#define OR1K_SR                  OR1K_OFFSET(SR)
#define OR1K_ESR                 OR1K_OFFSET(ESR)
#define OR1K_EPC                 OR1K_OFFSET(EPC)
#define OR1K_EEAR                OR1K_OFFSET(EEAR)

// macros to specify target registers in CONSTANT expressions
#define OR1K_OFFSET_CONST(_F)    VMI_CPU_REG_CONST(or1kP, _F)
#define OR1K_REG_CONST(_R)       OR1K_OFFSET_CONST(regs[_R])
#define OR1K_CARRY_CONST         OR1K_OFFSET_CONST(carryFlag)
#define OR1K_OVERFLOW_CONST      OR1K_OFFSET_CONST(overflowFlag)
#define OR1K_TEMPFLAG_CONST    OR1K_OFFSET_CONST(tempFlag)
```

Note that the new `tempFlag` field does not represent a true processor register: it is simply a temporary required for modeling purposes.

Having generated code that sets the value of the new `tempFlag` field if the processor is not in supervisor mode, `doRFE` then emits code to perform an intra-instruction jump if the flag is set:

```
// go to label inUserMode if tempFlag set (SPR_SR_SM bit is zero)
vmimtCondJumpLabel(OR1K_TEMPFLAG, True, inUserMode);
```

Next, code is generated to perform a return from exception in supervisor mode, just as in the previous example:

```
// set sr from esr (must call orlkSetSR to do this)
vmimtArgProcessor();
vmimtArgReg(OR1K_BITS, OR1K_ESR);
vmimtCall((vmiCallFn)orkSetSR);

// return to exception program counter
vmimtUncondJumpReg(0, OR1K_EPC, VMI_NOREG, vmi_JH_RETURNINT);
```

Now the label is inserted at the location where user mode code starts:

```
// insert the label targeted by vmimtCondJumpLabel above
vmimtInsertLabel(inUserMode);
```

And in user mode, `l.rfe` should cause an illegal instruction exception, implemented by a run-time call to `orkTakeException` (in `orkUtils.c`):

```
// take illegal instruction exception
vmimtArgProcessor();
vmimtArgUns32(OR1K_EXCPT_ILL);
vmimtArgUns32(0);
vmimtCall((vmiCallFn)orkTakeException);
```

The example at the end of this section shows the new code in action.

13.3 Implementing `l.mtspr`

The `l.mtspr` instruction implements a *move to special purpose register*. It allows a value in an OR1K general purpose register (`r0`, `r1` etc) to be written to a special purpose register (e.g. `sr`, `epc`, `esr`, and many other special purpose registers that are currently unimplemented). The target special purpose register is identified by a unique index number: for example, register `sr` has index `0x11`, register `epc` has index `0x20` and register `esr` has index `0x40`.

The index number of the special purpose register is calculated by adding an index register `ra` and a constant index `k`. This means that if the index register is anything other than `r0`, the special purpose register to update must be identified at *run time* (since there is no way to know the future value of `ra` when morphing code). However, if `ra` is `r0` (which is always zero) we know at *morph time* the SPR index (`k`) and therefore more efficient code can be created, as we will see below.

`l.mtspr` has been added to the decoder and disassembler in a similar way as for previous instructions. In `orkMorph.c`, the instruction is implemented by function `doMTSPR`:

```
static void doMTSPR(Uns32 instr, Uns32 thisPC) {

    Uns32 ra = OP10_A(instr);
    Uns32 rb = OP10_B(instr);
    Uns16 k  = OP10_I(instr);

    if(ra==0) {
        // faster variant when ra is r0
        doMTSPR_ra_0(thisPC, rb, k);
    } else {
```

```
// slower variant when ra is not r0
vmimtArgProcessor();
vmimtArgUns32(thisPC);
vmimtArgUns32(ra);
vmimtArgUns32(rb);
vmimtArgUns32(k);
vmimtCall((vmiCallFn)vmic_MTSPR);
}
}
```

In the case that `ra` (the index register) is zero, it calls `doMTSPR_ra_0` to emit code that targets a specific special purpose register. Otherwise a run-time call is created to function, `vmic_MTSPR`, which handles writing any special purpose register, like this:

```
static void vmic_MTSPR(orkP orlk, Uns32 thisPC, Uns32 ra, Uns32 rb, Uns32 k) {
    Uns32 sprNum = orlk->regs[ra] | k;

    switch(sprNum) {
        case SPR_OFF(SPR_SYS, SYS_SR):
            vmic_MTSPR_SYS_SR(ork, rb);
            break;

        case SPR_OFF(SPR_SYS, SYS_EPC):
            vmic_MTSPR_SYS_EPC(ork, rb);
            break;

        case SPR_OFF(SPR_SYS, SYS_EEAR):
            vmic_MTSPR_SYS_EEAR(ork, rb);
            break;

        case SPR_OFF(SPR_SYS, SYS_ESR):
            vmic_MTSPR_SYS_ESR(ork, rb);
            break;

        default:
            ignoreMTSPR(thisPC, ra, rb, k, sprNum);
            break;
    }
}
```

To avoid confusion between functions that should be called at *morph time* and those which are only applicable at *run time*, it can be helpful to use function prefixes: for example, use the prefix `vmic_` for a function that can only be called at run time.

The function constructs the special purpose register index and uses a case statement to determine the register to update. It then calls a specific update function for that register. As an example, the specific function to modify the `sr` register is:

```
static void vmic_MTSPR_SYS_SR(orkP orlk, Uns32 rb) {
    if(!IN_SUPERVISOR_MODE(ork)) {
        orlkTakeException(ork, ORLK_EXCPT_ILL, 0);
    } else {
        orlkSetSR(ork, orlk->regs[rb]);
    }
}
```

This function calls `orlkTakeException` if the processor is in user mode, which will cause the processor to take an illegal instruction exception. In supervisor mode, it calls `orlkSetSR` (from `orlkUtils.c`) to update the value of supervisor register `sr`.

In section 13.2, `l.rfe` was implemented using an intra-instruction conditional jump. It could just as well (and more clearly) have been implemented by a call to a run time function that performed the supervisor mode check, as above.

When the index register is `r0`, code to implement the assignment of the special purpose register is created by *morph time* function `doMTSPR_ra_0`:

```
static void doMTSPR_ra_0(Uns32 thisPC, Uns32 rb, Uns32 sprNum) {  
  
    switch(sprNum) {  
  
        case SPR_OFF(SPR_SYS, SYS_SR):  
            vmimtArgProcessor();  
            vmimtArgUns32(rb);  
            vmimtCall((vmiCallFn)vmic_MTSPR_SYS_SR);  
            break;  
  
        case SPR_OFF(SPR_SYS, SYS_EPC):  
            vmimtArgProcessor();  
            vmimtArgUns32(rb);  
            vmimtCall((vmiCallFn)vmic_MTSPR_SYS_EPC);  
            break;  
  
        case SPR_OFF(SPR_SYS, SYS_EEAR):  
            vmimtArgProcessor();  
            vmimtArgUns32(rb);  
            vmimtCall((vmiCallFn)vmic_MTSPR_SYS_EEAR);  
            break;  
  
        case SPR_OFF(SPR_SYS, SYS_ESR):  
            vmimtArgProcessor();  
            vmimtArgUns32(rb);  
            vmimtCall((vmiCallFn)vmic_MTSPR_SYS_ESR);  
            break;  
  
        default:  
            doMTSPR_ra_0_Default(thisPC, rb, sprNum);  
            break;  
    }  
}
```

Note that `doMTSPR_ra_0` determines the register to be written at *morph time*, and emits code that targets the specific register. `vmic_MTSPR` has to perform the equivalent check at *run time*, which will be slower. We will see this in the examples that follow.

13.4 Implementing *l.mfspr*

The `l.mfspr` instruction implements a *move from special purpose register*. It allows a value in an OR1K special purpose register to be assigned to a general purpose register. It is implemented in an analogous way to `l.mfspr` by function `doMFSPR` in `orlkMorph.c`.

13.5 Testing Illegal Instruction Exceptions

Directory 11.or1kBehaviorSPR/application contains the following example in file asmtest.S:

```
.org 0x700
// //////////////////////////////////////
// ILLEGAL INSTRUCTION EXCEPTION HANDLER (AT 0x700)
// //////////////////////////////////////
l.addi    r30,r30,1      // increment count of illegal instructions
l.sw      -4(r31),r1      // save value in r1
l.mfspr   r1,r0,0x20     // get epc in r1
l.addi    r1,r1,4        // move epc past faulting instruction
l.mtspr   r0,r1,0x20     // set epc from r1
l.lwz     r1,-4(r31)     // restore original r1
l.rfe                                // return from exception

.org 0x10000
// //////////////////////////////////////
// APPLICATION CODE (AT 0x10000)
// //////////////////////////////////////
.global _start
_start:
l.ori     r30,r0,0        // r30 = 0 (counts illegal instructions)
l.ori     r31,r0,0        // r31 = 0 (stack pointer)
l.mtspr   r0,r0,0x20     // clear epc

// //////////////////////////////////////
// SUPERVISOR MODE LOOP TEST
// //////////////////////////////////////
l.ori     r1,r0,2        // r1 = 2 (loop count)
loop1:
l.mfspr   r2,r0,0x20     // get epc in r2
l.addi    r2,r2,1        // increment r2
l.mtspr   r0,r2,0x20     // set epc from r2
l.addi    r1,r1,-1       // decrement r1
l.sfeqi   r1,0           // r1==0?
l.bnf     loop1          // go if not
l.nop                                // (delay slot)

// //////////////////////////////////////
// SUPERVISOR MODE FUNCTION CALL TEST
// //////////////////////////////////////
l.jal     incEPC         // incEPC (in supervisor mode)
l.nop                                // (delay slot)
l.mtspr   r0,r0,0x11     // clear supervisor mode

// //////////////////////////////////////
// USER MODE FUNCTION CALL TEST
// //////////////////////////////////////
l.jal     incEPC         // incEPC (in user mode)
l.nop                                // (delay slot)
l.rfe                                // *ILLEGAL* return from exception

.global exit
exit:
l.nop

// //////////////////////////////////////
// FUNCTION CALLED IN BOTH USER AND SUPERVISOR MODE
// //////////////////////////////////////
incEPC:
```

```

1.mfspr      r2,r0,0x20      // get epc in r2
1.addi       r2,r2,1         // increment r2
1.mtspr      r0,r2,0x20      // set epc from r2
1.jr         r9              // return
1.nop        // (delay slot)

```

This example begins execution at `_start` in supervisor mode. It then goes twice round `loop1`, incrementing the value of register `epc` (SPR index `0x20`) each time. These instructions are legal because the processor is in supervisor mode.

After the second loop iteration, the processor calls function `incEPC`, which also increments register `epc`.

It then clears supervisor mode with the instruction:

```

1.mtspr      r0,r0,0x11      // clear supervisor mode

```

Then, in user mode, the processor attempts to execute `l.rfe`. This fails, because it is now in user mode, and the handler at address `0x700` is executed. The handler updates the saved `epc` to skip the faulting instruction and returns (of course, a real handler would do something more useful than this). Finally, it calls `incEPC` again. The attempts to read and write `epc` in this function also fail, calling the handler.

Generate the assembler test case in directory `11.or1kBehaviorSPR` which will also compile the default application and test platform, to compile individually, as before use:

```

make -C platform
make -C application

```

Run the platform using the assembler executable file:

```

platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf

```

The output from this should be as follows (execution in the illegal instruction exception handler is highlighted in bold):

```

Warning (PC_NRI) No register information callback given for processor 'cpu1'
Info 'cpu1', 0x0000000000001000: l.ori    r30,r0,0x0
Info 'cpu1', 0x0000000000001004: l.ori    r31,r0,0x0
Info 'cpu1', 0x0000000000001008: l.mtspr  r0,r0,32
Info 'cpu1', 0x000000000000100c: l.ori    r1,r0,0x2
Info 'cpu1', 0x0000000000001010: l.mfspr  r2,r0,32
Info 'cpu1', 0x0000000000001014: l.addi   r2,r2,0x1
Info 'cpu1', 0x0000000000001018: l.mtspr  r0,r2,32
Info 'cpu1', 0x000000000000101c: l.addi   r1,r1,0xffffffff
Info 'cpu1', 0x0000000000001020: l.sfeqi  r1,0x0
Info 'cpu1', 0x0000000000001024: l.bnf    0x00010010
Info 'cpu1', 0x0000000000001028: l.nop    0x0
Info 'cpu1', 0x0000000000001010: l.mfspr  r2,r0,32
Info 'cpu1', 0x0000000000001014: l.addi   r2,r2,0x1
Info 'cpu1', 0x0000000000001018: l.mtspr  r0,r2,32
Info 'cpu1', 0x000000000000101c: l.addi   r1,r1,0xffffffff
Info 'cpu1', 0x0000000000001020: l.sfeqi  r1,0x0

```

```

Info 'cpu1', 0x0000000000010024: l.bnf      0x00010010
Info 'cpu1', 0x0000000000010028: l.nop      0x0
Info 'cpu1', 0x000000000001002c: l.jal      0x00010048
Info 'cpu1', 0x0000000000010030: l.nop      0x0
Info 'cpu1', 0x0000000000010048: l.mfspr   r2,r0,32
Info 'cpu1', 0x000000000001004c: l.addi     r2,r2,0x1
Info 'cpu1', 0x0000000000010050: l.mtspr   r0,r2,32
Info 'cpu1', 0x0000000000010054: l.jr      r9
Info 'cpu1', 0x0000000000010058: l.nop      0x0
Info 'cpu1', 0x0000000000010034: l.mtspr   r0,r0,17
Info 'cpu1', 0x0000000000010038: l.jal      0x00010048
Info 'cpu1', 0x000000000001003c: l.nop      0x0
Info 'cpu1', 0x0000000000010048: l.mfspr   r2,r0,32
Info 'cpu1', 0x0000000000000700: l.addi     r30,r30,0x1
Info 'cpu1', 0x0000000000000704: l.sw      0xffffffffc(r31),r1
Info 'cpu1', 0x0000000000000708: l.mfspr   r1,r0,32
Info 'cpu1', 0x000000000000070c: l.addi     r1,r1,0x4
Info 'cpu1', 0x0000000000000710: l.mtspr   r0,r1,32
Info 'cpu1', 0x0000000000000714: l.lwz     r1,0xffffffffc(r31)
Info 'cpu1', 0x0000000000000718: l.rfe
Info 'cpu1', 0x000000000001004c: l.addi     r2,r2,0x1
Info 'cpu1', 0x0000000000010050: l.mtspr   r0,r2,32
Info 'cpu1', 0x0000000000000700: l.addi     r30,r30,0x1
Info 'cpu1', 0x0000000000000704: l.sw      0xffffffffc(r31),r1
Info 'cpu1', 0x0000000000000708: l.mfspr   r1,r0,32
Info 'cpu1', 0x000000000000070c: l.addi     r1,r1,0x4
Info 'cpu1', 0x0000000000000710: l.mtspr   r0,r1,32
Info 'cpu1', 0x0000000000000714: l.lwz     r1,0xffffffffc(r31)
Info 'cpu1', 0x0000000000000718: l.rfe
Info 'cpu1', 0x0000000000010054: l.jr      r9
Info 'cpu1', 0x0000000000010058: l.nop      0x0
Info 'cpu1', 0x0000000000010040: l.rfe
Info 'cpu1', 0x0000000000000700: l.addi     r30,r30,0x1
Info 'cpu1', 0x0000000000000704: l.sw      0xffffffffc(r31),r1
Info 'cpu1', 0x0000000000000708: l.mfspr   r1,r0,32
Info 'cpu1', 0x000000000000070c: l.addi     r1,r1,0x4
Info 'cpu1', 0x0000000000000710: l.mtspr   r0,r1,32
Info 'cpu1', 0x0000000000000714: l.lwz     r1,0xffffffffc(r31)
Info 'cpu1', 0x0000000000000718: l.rfe
Info 'cpu1', 0x0000000000010044: l.nop      0x0
Processor 'cpu1' terminated at 'exit', address 0x10044

```

```

-----
R0 : 00000000   R1 : 00000000   R2 : 00000004   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : 00010040  R10: deadbeef  R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef  R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef  R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef  R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef  R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 00000003  R31: 00000000
PC : 00010048   SR : 00008000   ESR: 00008000   EPC: 00010044
BF:0 CF:0 OF:0
-----

```

processor has executed 56 instructions

13.6 Testing *l.mtspr* and *l.mfspr* Performance

When implementing *l.mtspr* and *l.mfspr*, we optimized the case where the index register is *r0* to improve performance. We can test the effect of this as follows.

13.6.1 Update platform/platform.c to Remove Tracing

Modify the MODEL_FLAGS in platform/platform.c to remove tracing as follows:

```
#define MODEL_FLAGS (ICM_ATTR_DEFAULT)
```

Then rebuild the platform:

```
make -C platform
```

13.6.2 Increase application/asmtest.s to Loop Count

Modify line 44 of application/asmtest.s to greatly increase the number of iterations of loop1 as follows:

```
l.movhi    r1,0x1000    // r1 = 0x10000000 (loop count)
```

This will cause the loop to be executed over 268 million times, which should take long enough to get meaningful performance numbers.

Then rebuild the test case and rerun:

```
make -C application asmtest.OR1K.elf
time platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be:

```
Warning (PC_NRI) No register information callback given for processor 'cpul'
Processor 'cpul' terminated at 'exit', address 0x10040
```

```
-----
R0 : 00000000   R1 : 00000000   R2 : 10000002   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : 00010040   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 00000003   R31: 00000000
PC : 00010048   SR : 00008000   ESR: 00008000   EPC: 00010044
BF:0 CF:0 OF:0
-----
```

```
processor has executed 1879048234 instructions
```

On a 2.8GHz Intel Core2 processor, time shows this takes about 3.5 seconds to execute 1,879,048,234 OR1K instructions, giving a simulation speed for this example of about 535 simulated MIPS.

13.6.3 Use Index Register r31

Now modify the inner loop in application/asmtest.s to use r31 as the index register instead of r0 (r31 happens to hold the value zero in this test case, but isn't *hard wired* to zero like r0):


```
l.mfspr      r2,r31,0x20    // get epc in r2
l.addi       r2,r2,1        // increment r2
l.mtspr      r31,r2,0x20    // set epc from r2
l.addi       r1,r1,-1       // decrement r1
l.sfeqi      r1,0           // r1==0?
l.bnf        loop1         // go if not
l.nop                // (delay slot)
```

This will cause the identification of the special purpose register to update to be deferred to run time, using calls to `vmic_MTSPR` and `vmic_MTSPR`, instead of calling the specific register SPR register access routines directly (e.g. `vmic_MTSPR_SYS_SR`).

Then rebuild the test case and rerun:

```
make -C application
time platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

Output is identical to before, but the run time is now 6.8 seconds, giving a simulation speed for this example of about 275 simulated MIPS. Although the change we made caused only a small amount of extra C code to be executed at run time (an extra function call and case statement) simulation performance is *260 MIPS slower than before*.

It is *very important* when creating a high-performance processor model to do as much work as possible at *morph time* and as little as possible at *run time*. The difference in simulation speed can be dramatic.

14 Modeling Mode-Dependent Behavior (Part 2)

In chapter 12.4, we saw how to model mode-dependent processor instructions with an example running at up to 535 simulated MIPS. This chapter shows how to get even faster performance on the same test case.

One significant problem with the implementation of the `l.rfe`, `l.mtspr` and `l.mfspr` instructions in chapter 12.4 is that they are coded to implement both kernel and user mode, and they select which behavior to perform at run time. The timing experiments at the end of the last chapter showed that it is possible to get dramatically faster performance if work can be moved from run time to morph time. Is it somehow possible to perform the supervisor-mode check at morph time to improve performance of these instructions?

First indications are promising. Recall that the instruction morpher callback function in `or1kMorph.c` has this implementation:

```
VMI_MORPH_FN(or1kMorphInstruction) {  
    or1kDecode((or1kP)processor, thisPC, OR1K_MORPH, 0, inDelaySlot);  
}
```

We are passing `processor` as the first argument to `or1kDecode`. This is in fact the current `or1kP` object for which code is being created (and is always supplied as an argument to a morpher callback declared with `VMI_MORPH_FN`). We know that we can determine whether an `or1k` is in supervisor mode using the `IN_SUPERVISOR_MODE` macro. Therefore, we can tell when morphing code whether user or supervisor mode code must be generated, which is exactly what we require. We will now see how this can be used to generate a faster model.

14.1 The Template Fast Modal Model

A template fast model for the OR1K processor implementing modal instructions can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/12.or1kBehaviorModeDict
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/12.or1kBehaviorModeDict .
```

Compile the model using `make`:

```
cd 12.or1kBehaviorModeDict  
make OPT=1
```

The processor model is based on the previous model, with the changes listed in following sections.

14.2 File *or1kMorph.c*

The dispatcher callback functions for `l.rfe`, `l.mtspr` and `l.mfspr` have been changed as follows:

```
static OR1K_DISPATCH_FN(morphRFE) {doRFE(or1k);}  
static OR1K_DISPATCH_FN(morphMTSPR) {doMTSPR(or1k, instr, thisPC);}  
static OR1K_DISPATCH_FN(morphMFSPR) {doMFSPR(or1k, instr, thisPC);}
```

The dispatcher callbacks have been modified to pass the current processor down to the specific morpher functions.

Function `doRFE` has been recoded to optimize for the current processor supervisor mode as follows:

```
static void doRFE(or1kP or1k) {  
  
    if(IN_SUPERVISOR_MODE(or1k)) {  
  
        // set sr from esr (must call or1kSetSR to do this)  
        vmimtArgProcessor();  
        vmimtArgReg(OR1K_BITS, OR1K_ESR);  
        vmimtCall((vmiCallFn)or1kSetSR);  
  
        // return to exception program counter  
        vmimtUncondJumpReg(0, OR1K_EPC, VMI_NOREG, vmi_JH_RETURNINT);  
  
    } else {  
  
        // take illegal instruction exception  
        vmimtArgProcessor();  
        vmimtArgUns32(OR1K_EXCPT_ILL);  
        vmimtArgUns32(0);  
        vmimtCall((vmiCallFn)or1kTakeException);  
  
    }  
}
```

What this does is as follows:

1. If the processor is currently in supervisor mode, it emits code to update register `sr` from register `esr` and return to the address in register `epc`, exactly as in the previous example.
2. Otherwise (the processor is in user mode), it emits code to jump to the illegal instruction exception vector.

Functions `doMTSPR` and `doMFSPR` have also been changed in an analogous way. Here is the new implementation of `doMTSPR`:

```
static void doMTSPR(or1kP or1k, Uns32 instr, Uns32 thisPC) {  
  
    Uns32 ra = OP10_A(instr);  
    Uns32 rb = OP10_B(instr);  
    Uns16 k  = OP10_I(instr);  
  
    if(ra==0) {
```

```

// faster variant when ra is r0 - select either supervisor mode or user
// mode function, based on current mode setting in sr
if(IN_SUPERVISOR_MODE(or1k)) {
    doMTSPR_ra_0_SM(thisPC, rb, k);
} else {
    doMTSPR_ra_0_UM(thisPC, rb, k);
}

} else {

    // slower variant when ra is not r0
    vmimtArgProcessor();
    vmimtArgUns32(thisPC);
    vmimtArgUns32(ra);
    vmimtArgUns32(rb);
    vmimtArgUns32(k);

    // select either supervisor mode or user mode callback, based on
    // current mode setting in sr
    if(IN_SUPERVISOR_MODE(or1k)) {
        vmimtCall((vmiCallFn)vmic_MTSPR_SM);
    } else {
        vmimtCall((vmiCallFn)vmic_MTSPR_UM);
    }
}
}

```

So user and supervisor mode behaviors are implemented separately. Here is the supervisor-mode function used when the SPR register to which to assign is known at morph time:

```

static void doMTSPR_ra_0_SM(Uns32 thisPC, Uns32 rb, Uns32 sprNum) {

    switch(sprNum) {

        case SPR_OFF(SPR_SYS, SYS_SR):
            vmimtArgProcessor();
            vmimtArgReg(OR1K_BITS, OR1K_REG(rb));
            vmimtCall((vmiCallFn)or1kSetSR);
            break;

        case SPR_OFF(SPR_SYS, SYS_EPC):
            vmimtMoveRR(OR1K_BITS, OR1K_EPC, OR1K_REG(rb));
            break;

        case SPR_OFF(SPR_SYS, SYS_EEAR):
            vmimtMoveRR(OR1K_BITS, OR1K_EEAR, OR1K_REG(rb));
            break;

        case SPR_OFF(SPR_SYS, SYS_ESR):
            vmimtMoveRR(OR1K_BITS, OR1K_ESR, OR1K_REG(rb));
            break;

        default:
            doMTSPR_ra_0_Default(thisPC, rb, sprNum);
            break;
    }
}

```

The morph-time function `vmimtMoveRR` is used here to specify a register-to-register assignment. The user-mode function simply calls the exception handler as follows:

```
static void doMTSPR_ra_0_UM(Uns32 thisPC, Uns32 rb, Uns32 sprNum) {
    switch(sprNum) {
        case SPR_OFF(SPR_SYS,SYS_SR):
        case SPR_OFF(SPR_SYS,SYS_EPC):
        case SPR_OFF(SPR_SYS,SYS_EEAR):
        case SPR_OFF(SPR_SYS,SYS_ESR):
            vmimtArgProcessor();
            vmimtArgUns32(OR1K_EXCPT_ILL);
            vmimtArgUns32(0);
            vmimtCall((vmiCallFn)or1kTakeException);
            break;
        default:
            doMTSPR_ra_0_Default(thisPC, rb, sprNum);
            break;
    }
}
```

14.3 Testing Optimized Illegal Instruction Exceptions

We will now test the optimized model using exactly the same test case as used in the previous example. Generate the assembler test case in directory

12.or1kBehaviorModeDict which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows (execution in the illegal instruction exception handler is highlighted in bold):

```
Warning (PC_NRI) No register information callback given for processor 'cpu1'
Info 'cpu1', 0x0000000000001000: l.ori    r30,r0,0x0
Info 'cpu1', 0x0000000000001004: l.ori    r31,r0,0x0
Info 'cpu1', 0x0000000000001008: l.mtspr  r0,r0,32
Info 'cpu1', 0x000000000000100c: l.ori    r1,r0,0x2
Info 'cpu1', 0x0000000000001010: l.mfspr  r2,r0,32
Info 'cpu1', 0x0000000000001014: l.addi   r2,r2,0x1
Info 'cpu1', 0x0000000000001018: l.mtspr  r0,r2,32
Info 'cpu1', 0x000000000000101c: l.addi   r1,r1,0xffffffff
Info 'cpu1', 0x0000000000001020: l.sfeqi  r1,0x0
Info 'cpu1', 0x0000000000001024: l.bnf    0x00010010
Info 'cpu1', 0x0000000000001028: l.nop    0x0
Info 'cpu1', 0x0000000000001010: l.mfspr  r2,r0,32
Info 'cpu1', 0x0000000000001014: l.addi   r2,r2,0x1
Info 'cpu1', 0x0000000000001018: l.mtspr  r0,r2,32
Info 'cpu1', 0x000000000000101c: l.addi   r1,r1,0xffffffff
Info 'cpu1', 0x0000000000001020: l.sfeqi  r1,0x0
```

```

Info 'cpul', 0x0000000000010024: l.bnf      0x00010010
Info 'cpul', 0x0000000000010028: l.nop      0x0
Info 'cpul', 0x000000000001002c: l.jal      0x00010048
Info 'cpul', 0x0000000000010030: l.nop      0x0
Info 'cpul', 0x0000000000010048: l.mfspr   r2,r0,32
Info 'cpul', 0x000000000001004c: l.addi     r2,r2,0x1
Info 'cpul', 0x0000000000010050: l.mtspr   r0,r2,32
Info 'cpul', 0x0000000000010054: l.jr      r9
Info 'cpul', 0x0000000000010058: l.nop      0x0
Info 'cpul', 0x0000000000010034: l.mtspr   r0,r0,17
Info 'cpul', 0x0000000000010038: l.jal      0x00010048
Info 'cpul', 0x000000000001003c: l.nop      0x0
Info 'cpul', 0x0000000000010048: l.mfspr   r2,r0,32
Info 'cpul', 0x000000000001004c: l.addi     r2,r2,0x1
Info 'cpul', 0x0000000000010050: l.mtspr   r0,r2,32
Info 'cpul', 0x0000000000010054: l.jr      r9
Info 'cpul', 0x0000000000010058: l.nop      0x0
Info 'cpul', 0x0000000000010040: l.rfe
Info 'cpul', 0x0000000000000700: l.addi     r30,r30,0x1
Info 'cpul', 0x0000000000000704: l.sw      0xffffffffc(r31),r1
Info 'cpul', 0x0000000000000708: l.mfspr   r1,r0,32
Info 'cpul', 0x000000000000070c: l.addi     r1,r1,0x4
Info 'cpul', 0x0000000000000710: l.mtspr   r0,r1,32
Info 'cpul', 0x0000000000000714: l.lwz     r1,0xffffffffc(r31)
Info 'cpul', 0x0000000000000718: l.rfe
Info 'cpul', 0x0000000000010044: l.nop      0x0
Processor 'cpul' terminated at 'exit', address 0x10044

```

```

-----
R0 : 00000000   R1 : 00000000   R2 : 00000004   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : 00010040   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 00000001   R31: 00000000
PC : 00010048   SR : 00008000   ESR: 00008000   EPC: 00010044
BF:0 CF:0 OF:0
-----

```

processor has executed 42 instructions

Comparing with the log output from chapter 12.4, it is obvious that something went wrong: only 42 instructions have been executed instead of 56, and the exception vector was executed only once instead of three times. The problem is that the application function `incEPC` at address `0x00010048` is called both in supervisor and user mode. When it is called first, the processor is in supervisor mode: at this point, a fragment of translated native code is created for the function body and saved in the processor code dictionary. On the second call (in user mode) the same fragment is re-executed. This is wrong because the code fragment was *specific to supervisor mode*.

14.4 Using Multiple Code Dictionaries

The problem we have seen in this example can be cured by using *two* code dictionaries: one for supervisor mode code fragments and one for user mode code fragments. To specify that we need multiple dictionaries, the list of dictionary names in `orlAttrs.c` has been changed to:

```
static const char *dictNames[] = {"SUPERVISOR", "USER", 0};
```

This definition says that we want this processor to have two dictionaries, the first (index 0) called `SUPERVISOR` and the second (index 1) called `USER`⁵.

When a processor starts executing it initially uses the *first* dictionary in this list (in this case `SUPERVISOR`). However, the current dictionary can be switched at any time using the VMI Run Time API function `vmirtSetMode`:

```
void vmirtSetMode(vmiProcessorP processor, Uns32 mode);
```

The `mode` argument in this prototype is the zero-based offset into the `dictNames` list of the new dictionary. So to switch to the supervisor mode dictionary, we would use:

```
vmirtSetMode((vmiProcessorP)orlk, 0);
```

And to switch to the user mode dictionary, we would use:

```
vmirtSetMode((vmiProcessorP)orlk, 1);
```

To use the new supervisor and user mode dictionaries, `orlkUtils.c` has been updated as follows:

```
static void setSRSwitchMode(orlkP orlk, Uns32 value) {
    Uns32 oldSM = orlk->SR & SPR_SR_SM;
    Uns32 newSM = value & SPR_SR_SM;

    // set the SR field
    orlk->SR = value;

    // switch mode if required
    if(MODAL && (oldSM != newSM)) {
        orlkMode newMode = newSM ? ORLK_MODE_SUPERVISOR : ORLK_MODE_USER;
        vmirtSetMode((vmiProcessorP)orlk, newMode);
    }
}

void orlkSetSR(orlkP orlk, Uns32 value) {
    // it is never possible to clear the fixed-one (FO) bit
    value |= SPR_SR_FO;

    // set the SR
    setSRSwitchMode(orlk, value);

    // set the current branch flag, carry flag and overflow flag from the SR
    orlk->branchFlag = GET_BIT(value, SPR_SR_F);
    orlk->carryFlag = GET_BIT(value, SPR_SR_CY);
    orlk->overflowFlag = GET_BIT(value, SPR_SR_OV);
}

void orlkEnterSupervisorMode(orlkP orlk) {
```

⁵ The processor model can have as many dictionaries as required (as long as there is at least one). The number of dictionaries is the number of strings in this zero-terminated array.

```
const Uns32 clearBits = (  
    SPR_SR_IEE |    // interrupt enable  
    SPR_SR_TEE |    // tick timer enable  
    SPR_SR_DME |    // data MMU enable  
    SPR_SR_IME |    // instruction MMU enable  
    SPR_SR_OVE      // overflow exception enable  
);  
  
// save the current status register in esr  
orlk->ESR = orlkGetSR(orlk);  
  
// mask out the 'clear' bits and mask in supervisor mode  
setSRSwitchMode(orlk, (orlk->ESR & ~clearBits) | SPR_SR_SM);  
}
```

The existing functions `orlkSetSR` and `orlkEnterSupervisorMode` have been modified so that they no longer update the SR field directly, but instead call a new function `setSRSwitchMode`. This new function determines whether the SM (supervisor mode) bit in the sr register has changed; if it has, it calls `vmirtSetMode` to switch code dictionaries. The index numbers for each mode are specified in `orlkStructure.h` as:

```
typedef enum orlkModeE {  
    ORLK_MODE_SUPERVISOR,  
    ORLK_MODE_USER,  
    ORLK_MODE_LAST  
} orlkMode;
```

To use the new modal code, we need to recompile the processor model with `MODAL` defined:

```
make clean  
make OPT=1 MODAL=1
```

And then rerun the example:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should now be correct: 56 instructions executed and three calls to the exception handler.

Now that the model is functionally correct, we should determine what effect the changes have had on its performance. To do this, redo exactly the steps described in section 13.6. On a 2.8 GHz Intel Core2 processor the results are as follows:

1. Using `r0` for the index register: 1,879,048,234 instructions in 0.9 seconds (2,090 simulated MIPS);
2. Using `r31` for the index register: 1,879,048,234 instructions in 7.0 seconds (268 simulated MIPS).

The first example (where the SPR is known at morph time) is vastly quicker – nearly two billion simulated instructions per second! Performance is similar for the second example - 268 simulated MIPS, as compared to 275 previously.

14.5 *Cautionary Notes about Code Dictionaries*

Although the use of multiple code dictionaries has significantly improved performance in these examples, this technique is not a panacea and should be used with discretion. This section describes the issues related to code dictionaries in more detail.

14.5.1 **vmirtSetMode is Slow**

Although we vastly accelerated `l.mtspr` and `l.mfspr`, that acceleration has a cost: we have added a call to `vmirtSetMode`, which slows down any instruction that causes a mode switch (i.e. which updates register `sr`). This can be seen using a small OR1K test case as follows:

```
.org 0xc00
////////////////////////////////////////////////////
// SYS VECTOR (AT 0xc00)
////////////////////////////////////////////////////
l.rfe                // return from exception (SWITCHES MODE)

.org 0x10000
////////////////////////////////////////////////////
// APPLICATION CODE (AT 0x10000)
////////////////////////////////////////////////////
.global _start
_start:
l.mtspr      r0,r0,0x11      // clear supervisor mode
l.movhi      r1,0x0100       // r1 = 0x01000000 (loop count)
loop1:
l.sys        0               // call sys (SWITCHES MODE)
l.addi       r1,r1,-1        // decrement r1
l.sfeqi      r1,0            // r1==0?
l.bnf        loop1          // go if not
l.nop                          // (delay slot)

.global exit
exit:
l.nop
```

This example performs a tight loop of six instructions of which two cause a mode switch.

On a 2.8 GHz Intel Core2 processor, this executes 100,663,299 simulated instructions in 5.9 seconds using the current model (a simulated speed of only 17 MIPS). On the processor model used in chapter 12.4, the run time is 4.4 seconds (a simulated speed of 23 MIPS).

Whether performance is better with multiple dictionaries therefore depends on the frequency of mode switching instructions compared with the frequency of instructions that can be optimized when there are multiple dictionaries: in some cases, application code may actually run slower when multiple dictionaries are used.

14.5.2 Model Code is More Complicated

The model needs to be carefully designed to ensure that dictionary code is always consistent with the simulated processor state. This is generally fairly easy as long as any code that could affect the mode is channeled through a single routine (`setSRSwitchMode` in this case). It is easy to get difficult-to-find bugs in poorly-structured models where calls to `vmirtSetMode` are not carefully controlled. As a general rule, there should only be one call to `vmirtSetMode` in a model, and this should be right next to code that updates the processor model register that affects the mode.

15 Implementing a Tick Timer

We have already seen in chapter 12 how synchronous exceptions (for example, alignment exceptions) can be efficiently modeled. We will now see how to model *asynchronous exceptions*, also known as interrupts. The VMI modeling API allows generic external exception behavior to be specified, as we will see in chapter 0. Additionally, it allows *tick timer* exceptions to be modeled very efficiently: the subject of this chapter.

15.1 OR1K Tick Timer Overview

The OR1K tick timer is controlled by two processor registers, the *tick timer mode register* (TTMR) and the *tick timer count register* (TTCR). These two SPR registers may be read and written using the `l.mfspr` and `l.mtspr` instructions we have seen previously (TTMR has SPR index 0x5000, TTCR has SPR index 0x5100).

The TTCR register is a 32-bit register that is incremented on each cycle when enabled by the TTMR register, as described below.

The TTMR register is subdivided into fields as follows:

Bit	31:30	29	28	27:0
Identifier	M	IE	IP	TP
R/W	R/W	R/W	R	R/W

The fields have the following meanings:

TP	Time Period 0x0000000: shortest comparison time period 0xffffffff: longest comparison time period
IP	Interrupt Pending 0: tick timer interrupt is not pending 1: tick timer interrupt is pending (IP can be <i>cleared</i> by writing 0 with <code>l.mtspr</code> , but may not be <i>set</i>)
IE	Interrupt Enable 0: tick timer does not generate interrupt 1: tick timer generates interrupt when <code>TTMR[TP]</code> matches <code>TTCR[27:0]</code>
M	Mode 00: timer is disabled 01: timer is restarted when <code>TTMR[TP]</code> matches <code>TTCR[27:0]</code> 10: timer stops when <code>TTMR[TP]</code> matches <code>TTCR[27:0]</code> 11: timer does not stop when <code>TTMR[TP]</code> matches <code>TTCR[27:0]</code> (if the timer is stopped in mode 10, writing to TTCR restarts it).

In our model, both TTCR and TTMR will be set to zero at reset, so the tick timer will initially be disabled.

15.2 Tick Timer Modeling Considerations

In true hardware, tick timers usually count processor *cycles*. In architectural models that are not cycle accurate, a common approximation is instead to count processor *instructions*. We will make this approximation in this OR1K model.

The tick timer could be modeled directly using the VMI API. At the start of every instruction, we could, for example, emit a call to a function that does the following:

1. Determine whether the counter is enabled by `TTMR[M]`.
2. If so, increment `TTCR` and compare `TTCR[27:0]` against `TTMR[TP]`.
3. If `TTCR[27:0]` and `TTMR[TP]` match, update state to stop the counter (if `TTMR[M]` is 10) and set `TTMR[IP]` (if `TTMR[IE]` is set).
4. If `TTMR[IP]` and `SR[TEE]` are set, make a call to the exception vector at `0x500`.

This would work perfectly well, but would be very slow. A much more efficient model can be made by using a combination of three routines from the VMI Run Time Function API:

```
Uns64 vmirtGetICount(vmiProcessorP processor);
```

`vmirtGetICount` returns a 64-bit count giving the total number of instructions that the processor has executed since simulation started.

```
void vmirtSetICountInterrupt(vmiProcessorP processor, Uns64 iCount);
```

`vmirtSetICountInterrupt` causes a model callback function to be executed after `iCount` more processor instructions have been simulated. The callback function is used to indicate whether the counter expiry alters the processor's behavior (whether an exception handler should be called, for example).

```
void vmirtClearICountInterrupt(vmiProcessorP processor);
```

`vmirtClearICountInterrupt` disables any instruction count interrupt previously enables using `vmirtSetICountInterrupt`.

A much more efficient model can be built using these functions as follows:

1. When `TTCR` or `TTMR` are written, determine the *implied* timer expiry count – in other words, after what count would the timer expire given the current `SPR` settings?
2. Use `vmirtSetICountInterrupt` to schedule a model callback after that count, or `vmirtClearICountInterrupt` to deschedule the callback if required.
3. When the callback is activated, schedule a call to the tick timer exception vector if the exception is enabled.
4. Do not model the `TTCR` register directly by incrementing it each instruction. Instead, derive the value if `TTCR` when requested using the processor instruction count returned by `vmirtGetICount` (in a similar manner as previously used for the status register `SR`).

Following sections describe the OR1K tick timer modeled using this approach.

15.3 The Template Tick Timer Model

A template model for the OR1K processor implementing a tick timer can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/13.or1kBehaviorTickTimer
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/13.or1kBehaviorTickTimer .
```

Compile the model using make:

```
cd 13.or1kBehaviorTickTimer
make
```

The processor model is based on the previous model, with the changes listed in following sections.

15.4 File or1kStructure.h

The OR1K processor structure has been updated as follows:

```
typedef struct or1kS {

    Bool        carryFlag;           // carry flag
    Bool        overflowFlag;        // overflow flag
    Bool        branchFlag;          // branch flag
    Bool        tempFlag;             // temporary flag

    Uns32        regs[OR1K_REGS];    // basic registers

    Uns32        SR;                  // status register
    Uns32        ESR;                 // exception status register
    Uns32        EPC;                 // exception program counter
    Uns32        EEAR;                // exception effective address register
    Uns32        TTCR;                // tick timer count register
    Uns32        TTCRSetCount;        // cycle count when TTCR set
    Bool        timerRunning;         // whether the timer is running

    union {                           // tick timer mode register
        Uns32 TTMR;
        struct {
            Uns32 TTMR_TP: 28;        // timeout count
            Uns32 TTMR_IP: 1;        // interrupt pending
            Uns32 TTMR_IE: 1;        // interrupt enable
            Uns32 TTMR_M : 2;        // timer mode
        };
    };

    vmiBusPortP busPorts;            // bus port descriptions
} or1k, *or1kP;
```

We have added fields `TTCR` and `TTMR` that will be used to model the tick timer SPR registers. There are also two fields `TTCRSetCount` and `timerRunning` which do not correspond to processor registers but which are modeling artifacts: `TTCRSetCount` records the processor instruction count when `TTCR` is written (required to derive the value of `TTCR` in later instructions); `timerRunning` is a boolean that indicates whether or not `TTCR` should be incremented each instruction.

15.5 File *or1kMorph.c*

This file has been modified to enhance `l.mtspr` and `l.mfspr` to allow reading and writing of `TTCR` and `TTMR` registers. `TTCR` is accessed by calling two new functions, `or1kGetTTCR` and `or1kSetTTCR`, implemented in `or1kExceptions.c` as described below. `TTMR` is written by `or1kSetTTMR`, also implemented in `or1kExceptions.c`.

15.6 File *or1kExceptions.c*

This file implements most of the new functionality to implement tick timer exceptions. The changes are as below.

```
typedef enum TTMRmodeE {
    TTMR_DISABLED = 0, // TTCR does not run
    TTMR_RESTART  = 1, // TTCR counts up until TTMR_TP, then restarts at 0
    TTMR_ONCE     = 2, // TTCR counts up until TTMR_TP, then stops
    TTMR_FREE     = 3  // TTCR counts up, overflowing at max
} TTMRmode;
```

This enumeration gives names for the four timer modes.

```
inline static Uns32 getThisICount(or1kP or1k) {
    return (Uns32)vmirtGetICount((vmiProcessorP)or1k);
}

inline static Uns32 getTTCR(or1kP or1k) {
    if(or1k->timerRunning) {
        return or1k->TTCR - or1k->TTCRSetCount + getThisICount(or1k);
    } else {
        return or1k->TTCR;
    }
}
```

`getTTCR` is an internal routine that returns the current effective value of the `TTCR` register. If the timer is running, `TTCR` is derived as follows:

1. Get the `TTCR` value recorded with the model.
2. Subtract the processor instruction count *when `TTCR` was written*.
3. Add the *current processor instruction count*.

If the timer is not running, the current `TTCR` value stored in the model is used.

```
static void setTTCR(or1kP or1k, Uns32 TTCR) {

    // update fields dependent on TTCR
    or1k->TTCR      = TTCR;
    or1k->TTCRSetCount = getThisICount(or1k);

    // if the timer is running, calculate the cycle delay to any interrupt
```

```
// (28 bits maximum) and schedule timer interrupt
if(ork->timerRunning) {
    Uns32 iCount = (ork->TTMR_TP-TTCR-1) & 0xffffffff;
    vmirtSetICountInterrupt((vmiProcessorP)ork, iCount);
} else {
    vmirtClearICountInterrupt((vmiProcessorP)ork);
}
}
```

setTTCR is an internal routine that is called when the TTCR register value is updated. It first saves the new TTCR value in the processor model and saves the current processor instruction count in TTCRSetCount (this is required so that the correct implied value of TTCR can be derived later). Next, if the timer is running, it calculates the implied timeout to counter expiry: the delta to the expiry instruction is the difference between TTMR_TP and TTCR[27:0] (masked to 28 bits), so it calls vmirtSetICountInterrupt to schedule a model callback after this number of instructions. If the timer is not running, it calls vmirtClearICountInterrupt to deschedule the callback.

```
Uns32 orkGetTTCR(orkP ork) {
    return getTTCR(ork);
}

void orkSetTTCR(orkP ork, Uns32 TTCR) {

    // restart the timer if mode is TTMR_ONCE
    if(ork->TTMR_M==TTMR_ONCE) {
        ork->timerRunning = True;
    }

    setTTCR(ork, TTCR);
}
```

These two routines implement the public interface to read and write the TTCR register. Note that writing TTCR when the timer mode is TTMR_ONCE causes the timer to be restarted if it is stopped.

```
void orkSetTTMR(orkP ork, Uns32 TTMR) {

    Uns32 TTCR = getTTCR(ork);

    // update TTMR, recording old and new values of TTMR_IP
    Bool oldIP = ork->TTMR_IP;
    ork->TTMR = TTMR;
    Bool newIP = ork->TTMR_IP;

    // TTMR_IP must not be set by l.mtspr!
    if(!oldIP && newIP) {
        ork->TTMR_IP = 0;
    }

    // start the timer if mode is TTMR_RESTART or TTMR_FREE
    // (for TTMR_ONCE, timer is restarted by write to TTCR)
    if((ork->TTMR_M==TTMR_RESTART) || (ork->TTMR_M==TTMR_FREE)) {
        ork->timerRunning = True;
    }

    setTTCR(ork, TTCR);
}
```

`orlkSetTTMR` implements the public interface to write `TTMR`. It first gets the current derived value of `TTCR`. It then sets the `TTMR` field in the processor structure, ensuring that the `TTMR_IP` bit does not change from 0 to 1 (`1.mtspr` cannot be used to set the interrupt pending bit, only to clear it). If the new mode is either `TTMR_RESTART` or `TTMR_FREE`, the timer is then restarted by setting `timerRunning`. Finally, `setTTCR` is called to reset the implied `TTCR` to the original value.

`TTCR` must be read using `getTTCR` and restored using `setTTCR` around the body of this routine for two reasons:

1. The way in which the derived value of `TTCR` is generated depends on the current setting of `timerRunning`. If `setTTCR` is not called, the next call to `getTTCR` will return a bogus value.
2. `setTTCR` is responsible for scheduling the instruction count callback, using `vmirtSetICountInterrupt`. If `setTTCR` isn't called, the instruction count callback will occur at the wrong time because changes to `TTMR` that affect the timeout (for example, and update of `TTMR[TP]`) won't be taken into account.

```
VMI_ICOUNT_FN(orlkICountPendingCB) {  
  
    orlkP orlk = (orlkP)processor;  
  
    switch(orlk->TTMR_M) {  
  
        case TTMR_RESTART:  
            // restart the timer from 0 on the NEXT instruction  
            setTTCR(orlk, -1);  
            break;  
  
        case TTMR_FREE:  
            // schedule the next interrupt event  
            setTTCR(orlk, getTTCR(orlk));  
            break;  
  
        case TTMR_ONCE:  
            // stop the timer on the NEXT instruction count  
            orlk->TTCR = (getTTCR(orlk)+1) & 0xffffffff;  
            orlk->timerRunning = False;  
            break;  
  
        case TTMR_DISABLED:  
            // how did we get here?  
            VMI_ABORT("timer interrupt, but timer was disabled");  
            break;  
    }  
  
    // if interrupt generation is enabled, set TTMR_IP  
    if(orlk->TTMR_IE) {  
        orlk->TTMR_IP = 1;  
    }  
  
    // handle exception if required  
    if(takeTEE(orlk)) {  
        vmirtDoSynchronousInterrupt(processor);  
    }  
}
```


Function `orlkICountPendingCB` is the callback that is called when the instruction count timeout specified by `vmirtSetICountInterrupt` has elapsed. The function prototype is specified in the VMI header file `vmiTypes.h` as follows:

```
#define VMI_ICOUNT_FN(_NAME) void _NAME( \  
    vmiProcessorP    processor,    \  
    vmiModelTimerP  timer,        \  
    Uns64            iCount,       \  
    void             *userData     \  
)
```

The arguments to this function are as follows:

1. The processor on which the timer has expired;
2. An argument `timer` of type `vmiModelTimerP`. This is an opaque type representing the implicit processor timer which is managed by the functions `vmirtSetICountInterrupt` and `vmirtClearICountInterrupt`;
3. An argument `iCount`, giving the current processor instruction count when the callback is activated;
4. A `userData` argument, which is always `NULL` for the implicit processor timer.

The function should update the processor state to reflect any changes caused by the timer expiry (for example, setting a pending-timer-interrupt bit). If necessary, it should signal that the processor needs to stop what it is doing and handle an exception by calling `vmirtDoSynchronousInterrupt`, as described below.

Based on the current timer mode setting when the timer expires, the processor state is updated in one of several ways:

```
case TTMR_RESTART:  
    // restart the timer from 0 on the NEXT instruction  
    setTTCR(orlk, -1);  
    break;
```

If the mode is `TTMR_RESTART`, the timer needs to restart from 0 at the *next* instruction. To do this, the callback sets `TTCR` to -1 now; when the timer is incremented before the next instruction is executed, it will have the value 0.

```
case TTMR_FREE:  
    // schedule the next interrupt event  
    setTTCR(orlk, getTTCR(orlk));  
    break;
```

With the timer free-running (mode is `TTMR_FREE`), `TTCR` is reset to its current value. This idiom ensures that another timeout is scheduled after `0x10000000` instructions.

```
case TTMR_ONCE:  
    // stop the timer on the NEXT instruction count  
    orlk->TTCR = (getTTCR(orlk)+1) & 0xffffffff;  
    orlk->timerRunning = False;  
    break;
```

With the timer in mode `TTMR_ONCE`, `TTCR` should be set to the value that it should hold from the next instruction onwards. Because the callback is invoked before execution of the faulting instruction, we need to increment the current value of `TTCR`.

Whether the timer expiry should cause a processor state change is determined by calling `takeTEE`, which is defined earlier in `orlkExceptions.c` as follows:

```
inline static Uns32 isTEEPending(orlkP orlk) {
    return (orlk->TTMR_IP && orlk->TTMR_IE);
}

inline static Bool isTEEEEnabled(orlkP orlk) {
    return (orlk->SR & SPR_SR_TEE);
}

inline static Uns32 takeTEE(orlkP orlk) {
    return isTEEPending(orlk) && isTEEEEnabled(orlk);
}
```

In other words, state change is required if `TTMR[IP]` and `SR[TEE]` are both set (the timer interrupt is both pending and enabled).

The instruction count timeout callback function `orlkICountPendingCB` **must not itself try to handle the interrupt** (for example, by calling `orlkTakeException`, which we first saw in chapter 12). Instead, *it must call `vmirtDoSynchronousInterrupt` to indicate that a timer exception is pending*. The timer interrupt must be handled by the *instruction fetch exception handler* function, specified by the `VMI_IFETCH_FN` macro in `vmiAttrs.h`:

```
#define VMI_IFETCH_FN(_NAME) vmifetchAction _NAME( \
    vmiProcessorP processor, \
    memDomainP domain, \
    Addr address, \
    Bool complete, \
    Bool annulled
)
```

Argument `domain` specifies the memory domain in which the fetch is being performed. The value of the domain can be used to control mode-specific fetch features (for example, how TLB mappings are performed). Argument `annulled` specifies whether the fetch is being made for an annulled delay slot instruction. Annulled instructions are sometimes treated differently (for example, they sometimes do not cause TLB misses). These two arguments are required to model some advanced feature, but are not discussed further here.

Type `vmifetchAction` is defined in `vmiTypes.h` as follows:

```
typedef enum vmifetchActionE {
    VMI_FETCH_NONE = 0,
    VMI_FETCH_EXCEPTION_COMPLETE = 1,
    VMI_FETCH_EXCEPTION_PENDING = 2
} vmifetchAction;
```

The instruction fetch exception handler is called in two phases. In the first phase (indicated by `complete` argument `False`), the function should determine *whether there is a pending exception on the processor that should prevent execution at the passed address and instead cause control to be transferred to an exception handler*. If there is such an exception pending, the function should return `VMI_FETCH_EXCEPTION_PENDING`; otherwise, it should return `VMI_FETCH_NONE`. In this phase, the instruction fetch handler *should not update the processor state*.

If the instruction fetch exception handler returns `VMI_FETCH_EXCEPTION_PENDING`, then it will subsequently be called again in a *second* phase (indicated by `complete` argument `True`). At this point, it should make any changes to the processor state required to handle the pending exception and return `VMI_FETCH_EXCEPTION_COMPLETE` to indicate that exception state has been updated.

Typically, the instruction fetch handler is required to handle a variety of exceptions: tick timer exceptions (as in this example), other external interrupts or synchronous exceptions such as invalid execute permission or alignment. In other words, the instruction count timeout callback is *specific to timer exceptions*, whereas the instruction fetch handler covers *all possible fetch exceptions*.

The initial implementation of the instruction fetch handler is as follows:

```
VMI_IFETCH_FN(orklIFetchExceptionCB) {  
    orklP orkl = (orklP)processor;  
  
    if(takeTEE(orkl)) {  
        // tick timer interrupt must be taken  
        if(complete) {  
            orklTakeException(orkl, OR1K_EXCPT_TTI, 0);  
            return VMI_FETCH_EXCEPTION_COMPLETE;  
        } else {  
            return VMI_FETCH_EXCEPTION_PENDING;  
        }  
    } else if(address & 3) {  
        // handle misaligned fetch exception  
        if(complete) {  
            orkl->EEAR = (Uns32)address;  
            orklTakeException(orkl, OR1K_EXCPT_BUS, 0);  
            return VMI_FETCH_EXCEPTION_COMPLETE;  
        } else {  
            return VMI_FETCH_EXCEPTION_PENDING;  
        }  
    }  
    } else if(!vmirtIsExecutable(processor, address)) {  
        // handle execute privilege exception  
        if(complete) {  
            orkl->EEAR = (Uns32)address;  
            orklTakeException(orkl, OR1K_EXCPT_IPF, 0);  
            return VMI_FETCH_EXCEPTION_COMPLETE;  
        } else {  

```

```
        return VMI_FETCH_EXCEPTION_PENDING;

    } else {

        // no fetch exception
        return VMI_FETCH_NONE;
    }
}
```

For a tick timer exception, the fetch exception handler causes control to be transferred immediately to the exception vector at `TTI_ADDRESS` without further execution of the instruction at the current address.

We have also implemented the instruction fetch alignment exception, which transfers control immediately to the exception vector at `BUS_ADDRESS` unless the fetch address is aligned to a 4-byte boundary, and the execute privilege exception, which transfers control immediately to the exception vector at `IPF_ADDRESS` if the fetch address does not have execute privileges⁶.

There is often also a requirement to transfer control to an exception handler vector *after the completion of the current instruction*. For example, the tick timer interrupt in the OR1K is enabled by a mask bit in the status register, `SR[TEE]`. What happens if `TTMR[IP]` is set and `SR[TEE]` is changed from 0 to 1 by execution of an `l.mtspr` instruction? In this case, the `l.mtspr` instruction should complete and the tick timer exception should occur before the *next* instruction is executed. To allow this behavior, there is one other useful public function defined in `or1kExceptions.c`:

```
void or1kInterruptNext(or1kP or1k) {
    if(takeTEE(or1k)) {
        vmirtDoSynchronousInterrupt((vmiProcessorP)or1k);
    }
}
```

`vmirtDoSynchronousInterrupt` causes the fetch exception handler to be invoked just before the next processor instruction is executed.

15.7 File `or1kUtils.c`

Function `or1kSetSR` has been modified as follows to handle the case described in the previous section where `TTMR[IP]` is set and `SR[TEE]` is changed from 0 to 1:

```
void or1kSetSR(or1kP or1k, Uns32 value) {

    // it is never possible to clear the fixed-one (FO) bit
    value |= SPR_SR_FO;

    // set the SR
    setSRSwitchMode(or1k, value);

    // set the current branch flag, carry flag and overflow flag from the SR
    or1k->branchFlag = GET_BIT(value, SPR_SR_F);
}
```

⁶ See section 18 for an example that exercises the execute privilege exception handler

```
orlk->carryFlag    = GET_BIT(value, SPR_SR_CY);
orlk->overflowFlag = GET_BIT(value, SPR_SR_OV);

// ensure any pending interrupt is taken before the next instruction
if(value & SPR_SR_TEE) {
    orlkInterruptNext(orlk);
}
}
```

Function `orlkDumpRegisters` has also been updated to write the `TTCR` and `TTMR` register values.

15.8 File *or1kAttrs.c*

The `vmiIASAttr` structure for the processor model has been modified to add both the instruction count timeout callback and the instruction fetch handler, as follows:

```
const vmiIASAttr modelAttrs = {

    . . . etc . . .

    //////////////////////////////////////
    // EXCEPTION ROUTINES
    //////////////////////////////////////

    .rdPrivExceptCB = orlkRdPrivExceptionCB,
    .wrPrivExceptCB = orlkWrPrivExceptionCB,
    .rdAlignExceptCB = orlkRdAlignExceptionCB,
    .wrAlignExceptCB = orlkWrAlignExceptionCB,
    .ifetchExceptCB = orlkIFetchExceptionCB,
    .arithExceptCB = orlkArithExceptionCB,
    .icountExceptCB = orlkICountPendingCB,

    . . . etc . . .
};
```

15.9 File *platform/platform.c*

The test platform for this example, `platform/platform.c`, has been changed as follows:

```
#define MODEL_FLAGS \
    (ICM_ATTR_TRACE | ICM_ATTR_TRACE_REGS_AFTER | ICM_ATTR_TRACE_ICOUNT)
```

These flags enable tracing of instructions and registers. `ICM_ATTR_TRACE_ICOUNT` prefixes each instruction in the trace with the instruction count for the processor.

15.10 Testing Tick Timer Exceptions

To test the tick timer behavior with mode `TTMR_ONCE`, generate the assembler test case in directory `13.orlkBehaviorTickTimer` which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows (much irrelevant output has been cut for conciseness):

```
Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 1: 'cpul', 0x00000000000010000: l.ori    r30,r0,0x0
Info 'cpul' REGISTERS
```

```
-----
R0 : 00000000   R1 : 00000000   R2 : deadbeef   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : deadbeef   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 00000000   R31: deadbeef
PC  : 00010004   SR  : 00008001   ESR: deadbeef   EPC: deadbeef
TCR: 00000000   TMR: 00000000   BF:0 CF:0 OF:0
-----
```

. . .

```
Info 8: 'cpul', 0x0000000000000c00: l.mtspr r0,r1,20480
  TCR: 00000000   TMR: a0000008   BF:0 CF:0 OF:0
Info 9: 'cpul', 0x0000000000000c04: l.rfe
  TCR: 00000000   TMR: a0000008   BF:0 CF:0 OF:0
Info 10: 'cpul', 0x000000000001001c: l.mtspr r0,r0,20736
  TCR: 00000000   TMR: a0000008   BF:0 CF:0 OF:0
Info 11: 'cpul', 0x0000000000010020: l.ori    r1,r0,0x8
  TCR: 00000001   TMR: a0000008   BF:0 CF:0 OF:0
Info 12: 'cpul', 0x0000000000010024: l.addi   r1,r1,0xffffffff
  TCR: 00000002   TMR: a0000008   BF:0 CF:1 OF:0
Info 13: 'cpul', 0x0000000000010028: l.sfeqi  r1,0x0
  TCR: 00000003   TMR: a0000008   BF:0 CF:1 OF:0
Info 14: 'cpul', 0x000000000001002c: l.bnf    0x00010024
  TCR: 00000004   TMR: a0000008   BF:0 CF:1 OF:0
Info 15: 'cpul', 0x0000000000010030: l.nop    0x0
  TCR: 00000005   TMR: a0000008   BF:0 CF:1 OF:0
Info 16: 'cpul', 0x0000000000010024: l.addi   r1,r1,0xffffffff
  TCR: 00000006   TMR: a0000008   BF:0 CF:1 OF:0
Info 17: 'cpul', 0x0000000000010028: l.sfeqi  r1,0x0
  TCR: 00000007   TMR: a0000008   BF:0 CF:1 OF:0
Info 18: 'cpul', 0x000000000001002c: *** FETCH EXCEPTION ***
  TCR: 00000008   TMR: b0000008   BF:0 CF:1 OF:0
Info 19: 'cpul', 0x0000000000000500: l.addi   r30,r30,0x1
Info 'cpul' REGISTERS
```

```
-----
R0 : 00000000   R1 : 00000006   R2 : deadbeef   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : deadbeef   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 00000001   R31: 00000000
PC  : 00000504   SR  : 00008001   ESR: 00008407   EPC: 0001002c
TCR: 00000008   TMR: b0000008   BF:0 CF:0 OF:0
-----
```

. . . etc . . .

The source code for this example is as follows:

```
.org 0x500
// Tick timer exception handler (at 0x500)
// increment count of timer exceptions
l.addi    r30,r30,1
l.sw      -4(r31),r1 // save value in r1
l.sw      -8(r31),r2 // save value in r2
l.mfspr   r1,r0,0x5000 // get ttmr in r1
l.movhi   r2,0xffff // r2 hi = ~TTMR_IP mask
l.ori     r2,r2,0xffff // r1 lo = ~TTMR_IP mask
l.and     r1,r1,r2 // clear TTMR_IP
l.mtspr   r0,r1,0x5000 // set ttmr from r1
l.lwz     r1,-4(r31) // restore original r1
l.lwz     r2,-8(r31) // restore original r2
l.sfeqi   r30,3 // r30==3?
l.bf      noReset // go if so
l.nop     // (delay slot)
l.mtspr   r0,r0,0x5100 // clear ttc (restarts counter)
noReset: l.rfe // return from exception

.org 0xc00
// SYSCALL VECTOR (AT 0xc00)
// set ttmr from r1
l.mtspr   r0,r1,0x5000
l.rfe // return from exception

.org 0x10000
// APPLICATION CODE (AT 0x10000)
.global _start
_start:
l.ori     r30,r0,0 // r30 = 0 (counts timer exceptions)
l.ori     r31,r0,0 // r31 = 0 (stack pointer)

l.ori     r1,r0,7 // r1 = 7 (SR_IEE | SR_TEE | SR_SM)
l.mtspr   r0,r1,0x11 // set sr from r1 (enables interrupts)
l.movhi   r1,0xa000 // M=ONCE,IE=True
l.ori     r1,r1,8 // TP=8 (count to match)
l.sys     0
l.mtspr   r0,r0,0x5100 // clear ttc (starts counter)

// TEST INTERRUPT WITH MODE ONCE
l.ori     r1,r0,8 // r1 = 8 (loop count)
loop1:
l.addi    r1,r1,-1 // decrement r1
l.sfeqi   r1,0 // r1==0?
l.bnf     loop1 // go if not
l.nop     // (delay slot)

.global exit
exit:
l.nop
```

Execution starts at label `_start`. At instruction 8, the `TTMR` register is set with `TTMR[M]=ONCE`, `TTMR[IE]=1` and `TTMR[TP]=8`. The counter does not start at this point because in mode `ONCE` it is activated only on a write to `TTCR`. At instruction 10, `TTCR` is written with 0, which starts the timer counting. At instruction 18, the timer expires generating the `FETCH EXCEPTION` message; instead of executing the instruction at address `0x1002c`, control is transferred to the exception handler.

The exception handler counts the number of exceptions. If fewer than 3 have occurred, it resets `TTCR`, which restarts the counter. In the full test case log, there are therefore three tick timer exceptions in total.

15.11 Explicit Processor Timers

This example uses the *implicit processor timer*, managed by functions `vmirtSetICountInterrupt` and `vmirtClearICountInterrupt`. It is also possible to create any number of additional *explicit* processor timers using functions described below. Each timer runs independently of the others.

Function `vmirtCreateModelTimer` creates a new timer for a processor. It is defined in `vmiRt.h` as follows:

```
vmiModelTimerP vmirtCreateModelTimer(  
    vmiProcessorP processor,  
    vmiICountFn   icountCB,  
    Uns32         scale,  
    void          *userData  
);
```

The argument `processor` is the processor to which the timer is to be attached. Argument `icountCB` is the timer expiry callback function. Argument `scale` is a scale factor by which the timer runs slower than the processor with which it is associated: for example, a scale value of 3 implies that the timer will appear to increment every three processor instructions. Argument `userData` is passed as the `userData` argument of the expiry function when it is called. The function returns an opaque type `vmiModelTimerP` which can be used to update the timer later; typically, this value will be saved in a field in the processor structure.

A previously-created timer can be cleared and deleted by `vmirtDeleteModelTimer`:

```
void vmirtDeleteModelTimer(vmiModelTimerP modelTimer);
```

The delay after which a timer will expire can be modified using `vmirtSetModelTimer`:

```
void vmirtSetModelTimer(vmiModelTimerP modelTimer, Uns64 iDelta);
```

This function sets the passed timer in exactly the same way that `vmirtSetICountInterrupt` sets the implicit timer. An explicit timer can be cleared using `vmirtClearModelTimer`:


```
void vmirtClearModelTimer(vmiModelTimerP modelTimer);
```

This function clears the passed timer in exactly the same way that `vmirtClearICountInterrupt` clears the implicit timer.

There are also three functions enabling the timer state to be queried. Function `vmirtIsModelTimerEnabled` returns a Boolean indicating if the timer is enabled (i.e. whether it has been activated using `vmirtSetModelTimer`):

```
Bool vmirtIsModelTimerEnabled(vmiModelTimerP modelTimer);
```

Function `vmirtGetModelTimerCurrentCount` returns the current timer value, either in terms of instructions or ticks (instructions scaled by the `scale` value when the timer was created):

```
Uns64 vmirtGetModelTimerCurrentCount(vmiModelTimerP modelTimer);
```

Function `vmirtGetModelTimerExpiryCount` returns the timer value at which the timer is scheduled to expire, either in terms of instructions or ticks (instructions scaled by the `scale` value when the timer was created):

```
Uns64 vmirtGetModelTimerExpiryCount(vmiModelTimerP modelTimer);
```

A basic model timer created using `vmirtCreateModelTimer` counts *the exact number of instructions executed by a processor*. In a quantized multiprocessor simulation, this might give the impression that time is running backwards, in the following scenario:

- CPU A reads a timer value *towards the end of a quantum* and uses it to calculate the current time.
- CPU A reaches the quantum end. CPU B starts simulating the same quantum.
- CPU B reads a timer value *towards the beginning of a quantum* and uses it to calculate the current time.

In this case, the time calculated by CPU A will appear to be later than that calculated by CPU B, even though it was obtained earlier. This is simply an artifact of multiprocessor quantized simulation, but this can cause problems for applications that rely upon a monotonically increasing view of time to work correctly. If the overall view of simulation time must increase monotonically, use function `vmirtCreateMonotonicModelTimer` to create the timer instead. It is defined in `vmiRt.h` as follows:

```
vmiModelTimerP vmirtCreateMonotonicModelTimer(  
    vmiProcessorP processor,  
    vmiICountFn    icountCB,  
    Uns32          scale,  
    void           *userData  
);
```

Monotonic timers are exactly like normal timers, except that the implied time seen when reading the timer is guaranteed to increase monotonically for all monotonic timers in the

platform. See the *VMI Run Time Function Reference* manual for more information about the algorithm used.

16 Modeling External Interrupts

The previous chapter showed how to model tick timer exceptions. The OR1K also supports generic external interrupts which we will model now.

16.1 OR1K PIC Overview

The OR1K programmable interrupt controller (PIC) is controlled by two processor registers, the *PIC mask register* (PICMR) and the *PIC status register* (PICSR). These two SPR registers may be read and written using the `l.mfspr` and `l.mtspr` instructions we have seen previously (PICMR has SPR index 0x4800, PICSR has SPR index 0x4802).

The PICMR register is used to mask or unmask up to 32 programmable interrupt sources.

The PICSR register is used by external interrupt sources to signal up to 32 interrupts. The OR1K is defined to allow either level-triggered or edge-triggered interrupts, or a mixture of both. In this implementation, we will support only level-triggered interrupts (so the external device will be responsible for all changes to bits in the PICSR register).

16.2 The Template External Interrupt Model

A template model for the OR1K processor implementing external interrupts can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/14.or1kBehaviorExternalInterrupt
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/14.or1kBehaviorExternalInterrupt .
```

Compile the model using make:

```
cd 14.or1kBehaviorExternalInterrupt
make
```

The processor model is based on the previous model, with the changes listed in following sections.

16.3 File *or1kStructure.h*

The OR1K processor structure has been updated to add fields for the PICMR and PICSR registers, and also macros to access the registers in code morphing routines.

16.4 File *or1kMorph.c*

This file has been modified to enhance `l.mtspr` and `l.mfspr` to allow reading and writing of PICSR and PICMR registers. PICSR can be read but not written (only external devices can modify this register). PICMR is written by calling `or1kSetPICMR`, implemented in `or1kExceptions.c`.

16.5 File *or1kExceptions.c*

This file implements most of the new functionality to implement external interrupts. The changes are as below.

The existing routines `or1kIFetchExceptionCB` and `or1kInterruptNext` have been modified to react to external interrupts in addition to timer interrupts.

`or1kIFetchExceptionCB` implements external interrupts as higher priority than timer interrupts:

```
VMI_IFETCH_FN(or1kIFetchExceptionCB) {  
  
    or1kP ork1 = (or1kP)processor;  
  
    if(takeIEE(ork1)) {  
  
        // external interrupt must be taken  
        if(complete) {  
            ork1TakeException(ork1, OR1K_EXCPT_EXI, 0);  
            return VMI_FETCH_EXCEPTION_COMPLETE;  
        } else {  
            return VMI_FETCH_EXCEPTION_PENDING;  
        }  
  
    } else if(takeTEE(ork1)) {  
  
        // tick timer interrupt must be taken  
        if(complete) {  
            ork1TakeException(ork1, OR1K_EXCPT_TTI, 0);  
            return VMI_FETCH_EXCEPTION_COMPLETE;  
        } else {  
            return VMI_FETCH_EXCEPTION_PENDING;  
        }  
  
    } else if(address & 3) {  
  
        // handle misaligned fetch exception  
        if(complete) {  
            ork1->EEAR = (Uns32)address;  
            ork1TakeException(ork1, OR1K_EXCPT_BUS, 0);  
            return VMI_FETCH_EXCEPTION_COMPLETE;  
        } else {  
            return VMI_FETCH_EXCEPTION_PENDING;  
        }  
  
    } else if(!vmirtIsExecutable(processor, address)) {  
  
        // handle execute privilege exception  
        if(complete) {  
            ork1->EEAR = (Uns32)address;  
            ork1TakeException(ork1, OR1K_EXCPT_IPF, 0);  
            return VMI_FETCH_EXCEPTION_COMPLETE;  
        } else {  
            return VMI_FETCH_EXCEPTION_PENDING;  
        }  
  
    } else {  
  
        return VMI_FETCH_NONE;  
    }  
}
```

```
}  
  
void orlkInterruptNext(orkP orlk) {  
  
    if(takeIEE(ork) || takeTEE(ork)) {  
        vmirtDoSynchronousInterrupt((vmiProcessorP)ork);  
    }  
}
```

Function takeIEE is implemented as:

```
inline static Uns32 isIEEPending(orkP orlk) {  
    return (ork->PICMR & ork->PICSRR);  
}  
  
inline static Bool isIEEEnabled(orkP orlk) {  
    return (ork->SR & SPR_SR_IEE);  
}  
  
inline static Uns32 takeIEE(orkP orlk) {  
    return isIEEPending(ork) && isIEEEnabled(ork);  
}
```

Function orlkSetPICMR is called whenever PICMR is written by instruction `l.mtspr`. Writing the programmable interrupt controller mask register could enable an interrupt that was previously disabled – `vmirtDoSynchronousInterrupt` is used to schedule an interrupt before the next instruction in this case:

```
void orlkSetPICMR(orkP orlk, Uns32 PICMR) {  
  
    orlk->PICMR = PICMR;  
  
    // take any pending external interrupt before the next instruction  
    if(takeIEE(ork)) {  
        vmirtDoSynchronousInterrupt((vmiProcessorP)ork);  
    }  
}
```

In order to allow interrupts to be raised externally to the model, it is necessary to register *net change functions* that are activated on external events. Each net change function should be defined using the `VMI_NET_CHANGE_FN` macro from `vmiTypes.h`:

```
#define VMI_NET_CHANGE_FN(_NAME) void _NAME( \  
    vmiProcessorP processor,    \  
    void          *userData,    \  
    Uns32         newValue     \  
)
```

The net change function is passed three arguments:

1. The processor that is being interrupted;
2. A processor-specific data pointer;
3. A new value for the net, the meaning of which is processor-specific.

For the OR1K model, a single net change function is currently used, defined as follows:

```
static VMI_NET_CHANGE_FN(externalInterrupt) {  
  
    orlkP orlk      = (orlkP)processor;  
    Uns32 deviceId = (Uns32)userData;  
  
    if(newValue) {  
        orlk->PICSR |= deviceId;  
    } else {  
        orlk->PICSR &= ~deviceId;  
    }  
  
    if(takeIEE(orlk)) {  
        vmirtDoSynchronousInterrupt(processor);  
    }  
}
```

For this model, the processor specific data pointer is used to hold a bit mask representing the interrupting device, and the value is a boolean indicating whether that mask has been enabled or disabled. The new mask value is calculated and applied to the processor PICSR register. If the external interrupt is active and enabled, the processor is notified that there is an exception that must be handled before the next instruction is executed by a call to `vmirtDoSynchronousInterrupt`⁷.

16.5.1 Connecting net ports

A net port on a model connects to a net in the platform. To specify nets to the simulator the model must provide an iterator function which returns the first or subsequent net port specifications, or 0 at the end of the list. The function must be registered in the model attributes table and should use this macro from `vmiPorts.h`:

```
#define VMI_NET_PORT_SPECS_FN(_NAME) vmiNetPortP _NAME ( \  
    vmiProcessorP processor,    \  
    vmiNetPortP prev           \  
)
```

Note that the iterator is also supplied with the processor pointer, so can adjust its behavior according to the configuration of the current module instance.

Each specification includes:

- Net port name.
- Net port type (*input*, *output* or *inout*).
- A callback function and user-data field (for an input).
- The address offset of a handle (used for output).
- Optional description.

The net port specification object is defined in `vmiPorts.h`.

⁷ Note the behavior described here applies to VMI interface version 2.0.16 and later. In older versions of the interface, the `VMI_NET_CHANGE_FN` callback returned a boolean to indicate whether the current processor should be interrupted, and the call to `vmirtDoSynchronousInterrupt` was not required. The interface has been changed so that an interrupt delivered to one member of an SMP cluster can in fact cause other members of that cluster to be interrupted.

```
typedef struct vmiNetPortS {
    char          *name;
    vmiNetPortType type;
    void          *userData;
    vmiNetChangeFn netChangeCB;          // callback for input or I/O net change
    ...
} vmiNetPort;
```

In this example (in `or1kExceptions.c`) the net port definitions are in a static structure. The accessor function iterates over this list:

```
////////////////////////////////////
// NET PORTS
////////////////////////////////////

static vmiNetPort netPorts[] = {
    {"intr0", vmi_NP_INPUT, (void*)1, externalInterrupt, 0, "External interrupt" },
    {"intr1", vmi_NP_INPUT, (void*)2, externalInterrupt, 0, "External interrupt" },
    {"intr2", vmi_NP_INPUT, (void*)4, externalInterrupt, 0, "External interrupt" },
    {"intr3", vmi_NP_INPUT, (void*)8, externalInterrupt, 0, "External interrupt" },
    VMI_END_NET_PORTS
};

//
// Get the next net port
//
VMI_NET_PORT_SPECS_FN(or1kGetNetPortSpec) {
    if (!prev) {
        return netPorts;
    }
    prev++;
    if (prev->name) {
        return prev;
    }
    return 0;
}
```

16.6 File `or1kUtils.c`

Function `or1kDumpRegisters` now also writes the new PICMR and PICSR registers.

16.7 File `platform/platform.c`

To stimulate the external interrupt signals, `platform/platform.c`, has been changed as follows. Net objects have been created in the platform and connected to the *intr0* and *intr1* interrupt inputs:

```
icmNetP intr0Net = icmNewNet("intr0Net");
icmNetP intr1Net = icmNewNet("intr1Net");
icmConnectProcessorNet(processor, intr0Net, "intr0", ICM_INPUT);
icmConnectProcessorNet(processor, intr1Net, "intr1", ICM_INPUT);
```

The simulation is controlled by this sequence:

```
// run processor for 9 instructions
simulate(processor, 9);
```

```

// raise intr0 for one instruction
icmWriteNet(intr0Net, 1);
simulate(processor, 1);
icmWriteNet(intr0Net, 0);

// run processor for 9 instructions
simulate(processor, 9);

// raise intr1 for one instruction
icmWriteNet(intr1Net, 1);
simulate(processor, 1);
icmWriteNet(intr1Net, 0);

// run processor until done (no instruction limit)
while(simulate(processor, -1)) {
    // keep going while processor is still running
}

```

The processor is first run for nine instructions. Then, *intr0* is raised (by *icmWriteNet*) and the processor stepped for one more instruction before *intr0* is lowered. We then run for nine more instructions before raising *intr1* for a single instruction. After that, the simulation is run to completion.

16.8 Testing External Exceptions

To test the external interrupt behavior, generate the assembler test case in directory `14.or1kBehaviorExternalInterrupt` which will also compile the default application and test platform, to compile individually, as before use:

```

make -C platform
make -C application

```

Run the platform using the assembler executable file:

```

platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf

```

The output from this should be as follows:

```

Warning (PC_NRI) No register information callback given for processor 'cpul'
Info 1: 'cpul', 0x0000000000010000: l.ori    r30,r0,0x0
Info 2: 'cpul', 0x0000000000010004: l.ori    r1,r0,0x7
Info 3: 'cpul', 0x0000000000010008: l.mtspr  r0,r1,17
Info 4: 'cpul', 0x000000000001000c: l.addi   r1,r0,0xffffffff
Info 5: 'cpul', 0x0000000000010010: l.mtspr  r0,r1,18432
Info 6: 'cpul', 0x0000000000010014: l.ori    r1,r0,0x8
Info 7: 'cpul', 0x0000000000010018: l.addi   r1,r1,0xffffffff
Info 8: 'cpul', 0x000000000001001c: l.sfeqi  r1,0x0
Info 9: 'cpul', 0x0000000000010020: l.bnf     0x00010018
Info 10: 'cpul', 0x0000000000010024: *** FETCH EXCEPTION ***
Info 11: 'cpul', 0x0000000000000800: l.addi   r30,r30,0x1
Info 12: 'cpul', 0x0000000000000804: l.rfe
Info 13: 'cpul', 0x0000000000010020: l.bnf     0x00010018
Info 14: 'cpul', 0x0000000000010024: l.nop     0x0
Info 15: 'cpul', 0x0000000000010018: l.addi   r1,r1,0xffffffff
Info 16: 'cpul', 0x000000000001001c: l.sfeqi  r1,0x0
Info 17: 'cpul', 0x0000000000010020: l.bnf     0x00010018

```



```

Info 18: 'cpul', 0x00000000000010024: l.nop      0x0
Info 19: 'cpul', 0x00000000000010018: l.addi    r1,r1,0xffffffff
Info 20: 'cpul', 0x0000000000001001c: *** FETCH EXCEPTION ***
Info 21: 'cpul', 0x0000000000000800: l.addi    r30,r30,0x1
Info 22: 'cpul', 0x0000000000000804: l.rfe
Info 23: 'cpul', 0x0000000000001001c: l.sfeqi   r1,0x0
Info 24: 'cpul', 0x00000000000010020: l.bnf     0x00010018
Info 25: 'cpul', 0x00000000000010024: l.nop      0x0
Info 26: 'cpul', 0x00000000000010018: l.addi    r1,r1,0xffffffff
Info 27: 'cpul', 0x0000000000001001c: l.sfeqi   r1,0x0
Info 28: 'cpul', 0x00000000000010020: l.bnf     0x00010018
Info 29: 'cpul', 0x00000000000010024: l.nop      0x0
Info 30: 'cpul', 0x00000000000010018: l.addi    r1,r1,0xffffffff
Info 31: 'cpul', 0x0000000000001001c: l.sfeqi   r1,0x0
Info 32: 'cpul', 0x00000000000010020: l.bnf     0x00010018
Info 33: 'cpul', 0x00000000000010024: l.nop      0x0
Info 34: 'cpul', 0x00000000000010018: l.addi    r1,r1,0xffffffff
Info 35: 'cpul', 0x0000000000001001c: l.sfeqi   r1,0x0
Info 36: 'cpul', 0x00000000000010020: l.bnf     0x00010018
Info 37: 'cpul', 0x00000000000010024: l.nop      0x0
Info 38: 'cpul', 0x00000000000010018: l.addi    r1,r1,0xffffffff
Info 39: 'cpul', 0x0000000000001001c: l.sfeqi   r1,0x0
Info 40: 'cpul', 0x00000000000010020: l.bnf     0x00010018
Info 41: 'cpul', 0x00000000000010024: l.nop      0x0
Info 42: 'cpul', 0x00000000000010018: l.addi    r1,r1,0xffffffff
Info 43: 'cpul', 0x0000000000001001c: l.sfeqi   r1,0x0
Info 44: 'cpul', 0x00000000000010020: l.bnf     0x00010018
Info 45: 'cpul', 0x00000000000010024: l.nop      0x0
Info 46: 'cpul', 0x00000000000010028: l.nop      0x0
Processor 'cpul' terminated at 'exit', address 0x10028
-----
R0 : 00000000   R1 : 00000000   R2 : deadbeef   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : deadbeef   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 00000002   R31: deadbeef
PC : 0001002c   SR : 00008607   ESR: 00008407   EPC: 0001001c
TCR: 00000000   TMR: 00000000   PSR: 00000000   PMR: ffffffff
BF:1 CF:1 OF:0
-----

processor has executed 46 instructions

```

The source code for this example is as follows:

```

.org 0x800
////////////////////////////////////
// EXTERNAL INTERRUPT HANDLER (AT 0x800)
////////////////////////////////////
l.addi    r30,r30,1      // increment count of external exceptions
l.rfe     // return from exception

.org 0x10000
////////////////////////////////////
// APPLICATION CODE (AT 0x10000)
////////////////////////////////////
.global _start
_start:
l.ori     r30,r0,0       // r30 = 0 (counts timer exceptions)

```

```
        l.ori      r1,r0,7          // r1 = 7 (SR_IEE | SR_TEE | SR_SM)
        l.mtspr    r0,r1,0x11       // set sr from r1 (enables interrupts)
        l.addi     r1,r0,-1         // r1 = -1
        l.mtspr    r0,r1,0x4800     // set picmr from r1 (enables interrupts)

loop1:   l.ori      r1,r0,8          // r1 = 8 (loop count)
        l.addi     r1,r1,-1         // decrement r1
        l.sfeqi    r1,0            // r1==0?
        l.bnf      loop1           // go if not
        l.nop                          // (delay slot)

.global exit
exit:    l.nop
```

Execution starts at label `_start`. The application enables external interrupts and starts executing a simple loop, `loop1`. The external interrupt exception handler at `0x800` counts the number of external interrupts in `r30`.

Note that external interrupts interrupt the flow of execution at instruction 10 and 20.

17 Implementing the Debug Interface

This section describes the implementation of a *debug interface* for the OR1K processor. This has two purposes:

1. It enables debuggers that support the *gdb remote serial protocol* (RSP) to be connected to the processor model;
2. It enables query functions in the Imperas ICM interface (for example `icmGetNextReg` and `icmGetRegByName`), which can in turn be used to implement custom debugger integrations or advanced test harnesses that are able to query and modify processor state.

17.1 The Template Debug Interface Model

A template model for the OR1K processor implementing a debugger interface can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/15.or1kDebugSupport
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/15.or1kDebugSupport .
```

Compile the model using make:

```
cd 15.or1kDebugSupport
make
```

The processor model is based on the previous model, with the changes listed in following sections.

17.2 File *or1kUtils.c*

File `or1kUtils.c` has been modified to implement processor *mode* and *exception* query callbacks. For each, there are two functions:

1. An iterator that lists each mode or exception type supported by the processor;
2. A function that returns the *currently active* mode or exception.

17.2.1 Processor Mode Iterator Function

The processor *mode iterator* is defined using the `VMI_MODE_INFO_FN` macro, defined in `vmiDbg.h` as follows:

```
#define VMI_MODE_INFO_FN(_NAME) vmiModeInfoCP _NAME( \
    vmiProcessorP processor, \
    vmiModeInfoCP prev \
)
typedef VMI_MODE_INFO_FN((*vmiModeInfoFn));
```

When called with a `NULL` value of `prev`, the function should return a description of the first mode supported by the processor. When called with a non-`NULL` value of `prev`, the

function should return a description of the *next* mode supported by the processor. When all modes have been returned, the function should return `NULL`. Each mode is described by returning a pointer to an object of type `vmiModeInfo`:

```
typedef struct vmiModeInfoS {  
    const char *name;           // exception name  
    Uns32      code;           // model-specific mode code  
} vmiModeInfo;
```

The `name` field gives a test name for the mode. The `code` field is model-specific, and typically will correspond to an enumeration in the model. For the OR1K, the mode iterator function is defined like this:

```
VMI_MODE_INFO_FN(or1kModeInfo) {  
  
    vmiModeInfoCP end = modes+OR1K_MODE_LAST;  
  
    // on the first call, start with the first member of the table  
    if(!prev) {  
        prev = modes-1;  
    }  
  
    // get the next member  
    vmiModeInfoCP this = prev+1;  
  
    // return the next member, or NULL if at the end of the list  
    return (this==end) ? 0 : this;  
}
```

A prototype of this function has been added to `or1kFunctions.h`.

This function refers to a constant static list of modes in `or1kUtils.c`:

```
#define OR1K_MODE_INFO(_D) [OR1K_MODE_##_D] = {name:##_D, code:OR1K_MODE_##_D}  
  
static const vmiModeInfo modes[OR1K_MODE_LAST] = {  
    OR1K_MODE_INFO(SUPERVISOR),  
    OR1K_MODE_INFO(USER)  
};
```

The `modes` array defines two modes, `SUPERVISOR` (with code `OR1K_MODE_SUPERVISOR`) and `USER` (with code `OR1K_MODE_USER`). Entry `OR1K_MODE_LAST` in the enumeration does not define a real mode but is instead used as a terminator for sizing of the `modes` array.

17.2.2 Processor Current Mode Query Function

The processor *current mode query function* is defined using the `VMI_GET_MODE_FN` macro, defined in `vmiDbg.h` as follows:

```
#define VMI_GET_MODE_FN(_NAME) vmiModeInfoCP _NAME(vmiProcessorP processor)  
typedef VMI_GET_MODE_FN((*vmiGetModeFn));
```

This function should return a `vmiModeInfoCP` description for the current mode. In the OR1K mode, the function is implemented like this:

```
VMI_GET_MODE_FN(or1kGetMode) {  
  
    or1kP    or1k    = (or1kP)processor;  
    Uns32    SM      = or1k->SR & SPR_SR_SM;  
    or1kMode newMode = SM ? OR1K_MODE_SUPERVISOR : OR1K_MODE_USER;  
  
    return modes+newMode;  
}
```

So depending on the current value of the SM bit in the status register, an appropriate entry from the `modes` table is selected.

A prototype of this function has been added to `or1kFunctions.h`.

17.2.3 Processor Exception Iterator Function

The processor *exception iterator* is defined using the `VMI_EXCEPTION_INFO_FN` macro, defined in `vmiDbg.h` as follows:

```
#define VMI_EXCEPTION_INFO_FN( _NAME ) vmiExceptionInfoCP _NAME( \  
    vmiProcessorP    processor,    \  
    vmiExceptionInfoCP prev        \  
)  
typedef VMI_EXCEPTION_INFO_FN((*vmiExceptionInfoFn));
```

The works in an analogous fashion to the mode iterator, described previously. Each exception is described by returning a pointer to an object of type `vmiExceptionInfo`:

```
typedef struct vmiExceptionInfoS {  
    const char *name;           // exception name  
    Uns32      code;           // model-specific exception code  
} vmiExceptionInfo;
```

Once again, the description contains a string description of the exception and a model-specific code. For the OR1K, the exception iterator function is defined like this:

```
VMI_EXCEPTION_INFO_FN(or1kExceptionInfo) {  
  
    vmiExceptionInfoCP end = exceptions+OR1K_EXCPT_LAST;  
  
    // on the first call, start with the first member of the table  
    if(!prev) {  
        prev = exceptions-1;  
    }  
  
    // get the next member  
    vmiExceptionInfoCP this = prev+1;  
  
    // return the next member, or NULL if at the end of the list  
    return (this==end) ? 0 : this;  
}
```

This function works in exactly the same way as the mode iterator, returning members of the exceptions array, previously described in section 12.1.5. A prototype of this function has been added to `orlkFunctions.h`.

17.2.4 Processor Current Exception Query Function

The processor *current exception query function* is defined using the `VMI_GET_EXCEPTION_FN` macro, defined in `vmiDbg.h` as follows:

```
#define VMI_GET_EXCEPTION_FN(_NAME) vmiExceptionInfoCP _NAME( \
    vmiProcessorP processor \
)
typedef VMI_GET_EXCEPTION_FN((*vmiGetExceptionFn));
```

This function should return a `vmiExceptionInfoCP` description for the current exception. In the OR1K mode, the function is implemented like this:

```
VMI_GET_EXCEPTION_FN(orlkGetException) {
    orlkP orlk = (orlkP)processor;
    return &exceptions[orlk->exception];
}
```

A prototype of this function has been added to `orlkFunctions.h`.

To implement this function, a new pseudo-register called `exception` has been added to the OR1K structure:

```
#include "orlkExceptionTypes.h"

// processor structure
typedef struct orlkS {

    . . . fields omitted . . .

    Uns32      SR;                // status register
    Uns32      ESR;               // exception status register
    Uns32      EPC;               // exception program counter register
    Uns32      EEAR;              // exception effective address register
    Uns32      PICMR;             // PIC mask register
    Uns32      PICSR;             // PIC status register
    Uns32      TTCR;              // tick timer count register
    Uns32      TTCRSetCount;      // cycle count when TTCR set
    Bool       timerRunning;      // whether the timer is running
    Bool       reset;             // whether the processor is being reset
    orlkException exception;      // current exception

    . . . fields omitted . . .

} orlk, *orlkP;
```

Finally, function `orlkTakeException` has been modified to update the new pseudo-register when an exception occurs:

```
void orlkTakeException(orlkP orlk, orlkException exception, Uns32 pcOffset) {

    Uns8  simD;
```

```
Uns32 simPC = (Uns32)vmirtGetPCDS((vmiProcessorP)orlk, &simD);

orlkEnterSupervisorMode(orlk);
orlk->EPC = simPC + pcOffset;

// set sr[DSX] for exception in a delay slot
if(simD) {
    orlk->SR |= SPR_SR_DSX;
}

// jump to the vector
orlk->exception = exception;
vmirtSetPCException((vmiProcessorP)orlk, exceptions[exception].code);
}
```

17.3 File *or1kRegisters.c*

In order to implement a debugger interface, functions need to be added to allow the debugger to *read*, *write* and *query* registers in the processor model. This is done by specifying an array of `vmiRegInfo` structures, one for each register in the processor that should be accessible to the debugger. The definition of this structure is as follows (in file `vmiDbg.h`):

```
typedef struct vmiRegInfoS {
    const char    *name;                // register identification name
    const char    *description;         // description string
    vmiRegGroupCP group;                // group for this register
    Uns32         gdbIndex;             // gdb ordinal index number
    vmiRegUsage   usage                 : 2; // any special usage for this register
    vmiRegAccess  access                : 2; // allowed access
    Bool         noSaveRestore         : 1; // does not participate in save/restore
    Uns32         unused               : 11; // unused fields
    Uns32         bits                 : 16; // size of register in bits
    vmiRegReadFn  readCB;               // read callback function
    vmiRegWriteFn writeCB;              // write callback function
    vmiReg        raw;                  // raw register value (if no callback)
    void          *userData;            // model-specific data pointer
} vmiRegInfo;
```

The fields in the structure are as follows:

1. **name**: the name to use to refer to the register.
2. **description**: this is a constant description string.
3. **group**: this is a pointer to a register group description (see below).
4. **gdbIndex**: each processor register has a unique index number in *gdb*, which should be provided here. You will need to examine the *gdb* processor-specific source code for your processor to find the value to enter here.
5. **usage**: any special usage for the register should be given using this field, which is a member of the `vmiRegUsage` enumeration in `vmiDbg.h`. Special usages are the program counter (`vmi_REG_PC`), stack pointer (`vmi_REG_SP`) and frame pointer (`vmi_REG_FP`).
6. **access**: this defines the register accessibility, a member of the `vmiRegAccess` enumeration in `vmiDbg.h`. Valid values are `vmi_RA_NONE` (no access), `vmi_RA_R` (read-only), `vmi_RA_W` (write-only) and `vmi_RA_RW` (read/write).

7. **noSaveRestore**: this Boolean value should be set to `True` if the register should not be automatically saved and restored (save/restore is a feature of Imperas Professional Tools and is not covered further in this manual).
8. **bits**: this is the size of the register in bits.
9. **readCB**: this is a read callback function that returns the current value of the register.
10. **writeCB**: this is a write callback function that is used to set the current value of the register.
11. **raw**: for a plain register that requires no special behavior for read or write, set `readCB` and/or `writeCB` to `NULL` and use this `vmiReg` field to specify the register location in the processor structure instead.
12. **userData**: this is a pointer to model-specific data that can be used if required in read and write callback functions.

The core data structure in `or1kRegisters.c` is an null-terminated array of these structures with one entry for each OR1K register that is made visible via the debug interface:

```
static const vmiRegInfo registers[] = {
    // registers visible in gdb
    OR1K_REG_INFO("R0",      0,  vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_R (regs[0],
OR1K_GROUP(GPR))      , "constant zero"),
    OR1K_REG_INFO("R1",      1,  vmi_REG_SP,   OR1K_BITS, OR1K_RAW_REG_RW(regs[1],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R2",      2,  vmi_REG_FP,   OR1K_BITS, OR1K_RAW_REG_RW(regs[2],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R3",      3,  vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[3],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R4",      4,  vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[4],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R5",      5,  vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[5],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R6",      6,  vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[6],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R7",      7,  vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[7],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R8",      8,  vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[8],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R9",      9,  vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[9],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R10",     10, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[10],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R11",     11, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[11],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R12",     12, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[12],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R13",     13, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[13],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R14",     14, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[14],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R15",     15, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[15],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R16",     16, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[16],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R17",     17, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[17],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R18",     18, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[18],
OR1K_GROUP(GPR))      , 0),
    OR1K_REG_INFO("R19",     19, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[19],
OR1K_GROUP(GPR))      , 0),

```



```

    OR1K_REG_INFO("R20", 20, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[20],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R21", 21, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[21],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R22", 22, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[22],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R23", 23, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[23],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R24", 24, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[24],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R25", 25, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[25],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R26", 26, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[26],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R27", 27, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[27],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R28", 28, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[28],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R29", 29, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[29],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R30", 30, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[30],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("R31", 31, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_RW(regs[31],
OR1K_GROUP(GPR)), 0),
    OR1K_REG_INFO("PC", 64, vmi_REG_PC, OR1K_BITS, OR1K_CB_REG_RW(readPC, writePC,
OR1K_GROUP(SYSTEM)), 0),
    OR1K_REG_INFO("SR", 65, vmi_REG_NONE, OR1K_BITS, OR1K_CB_REG_R(readSR,
OR1K_GROUP(SYSTEM)), "status register"),
    OR1K_REG_INFO("EPCR", 66, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_R(EPC,
OR1K_GROUP(SYSTEM)), "Exception PC"),
    OR1K_REG_INFO("EEAR", 67, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_R(EEAR,
OR1K_GROUP(SYSTEM)), "exception effective address "),

    // registers not visible in gdb
    OR1K_REG_INFO("EXCPT", 100, vmi_REG_NONE, OR1K_BITS, OR1K_RAW_REG_R(exception,
OR1K_GROUP(SYSTEM)), 0),

    {0},
};

```

The table contains two kinds of registers: those visible to gdb and those that are not. In the latter category, there is a single register, EXCPT, which is the new exception pseudo-register that records any active exception. Which group a particular register is in is determined by macro IS_GDB_HIDDEN_REG, which is True if the register is invisible to gdb:

```

#define OR1K_GDB_HIDDEN_INDEX 100
#define IS_GDB_HIDDEN_REG(_I) ((_I)>=OR1K_GDB_HIDDEN_INDEX)

```

Each register is assigned to a register group using the group field in the vmiRegInfo structure. Groups have no simulation purpose, but can be used by debuggers to gather registers into sets for display purposes. Each group is defined as follows:

```

typedef struct vmiRegGroups {
    const char *name;           // register group name
} vmiRegGroup;

```

In the OR1K, there are two groups, GPR and SYSTEM, defined in or1kRegisters.c like this:

```

typedef enum or1kRegGroupIdE {
    OR1K_RG_GPR,           // GPR group

```

```
    OR1K_RG_SYSTEM, // System register group
    OR1K_RG_LAST    // KEEP LAST: for sizing
} orlkRegGroupId;

static const vmiRegGroup groups[OR1K_RG_LAST+1] = {
    [OR1K_RG_GPR]    = {name: "GPR"    },
    [OR1K_RG_SYSTEM] = {name: "System"}
};
```

A *register group iterator* is defined which returns all the register groups supported by the model:

```
VMI_REG_GROUP_FN(orkRegGroup) {
    if(!prev) {
        return groups;
    } else if((prev+1)->name) {
        return prev+1;
    } else {
        return 0;
    }
}
```

Given an argument of type `vmiRegGroupCP` (i.e. a pointer to a member of the `groups` array) this function should return the *next* register group description in the array, or `NULL` if there are no more register group descriptions. If called with a `NULL` argument, it should return the first register group description in the array.

There is a macro used to select the appropriate group for a register in the table of register descriptions:

```
#define OR1K_GROUP(_G) &groups[OR1K_RG_##_G]
```

The contents of the register description array are made available to the simulator by implementing a *register structure iterator* function using the `VMI_REG_INFO_FN` macro, defined in `vmiDbg.h` as follows:

```
#define VMI_REG_INFO_FN(_NAME) vmiRegInfoCP _NAME( \
    vmiProcessorP processor, \
    vmiRegInfoCP prev, \
    Bool gdbFrame \
)
typedef VMI_REG_INFO_FN((*vmiRegInfoFn));
```

Given an argument of type `vmiRegInfoCP` (i.e. a pointer to a member of the `registers` array) this function should return the *next* register description in the array, or `NULL` if there are no more register descriptions. If called with a `NULL` argument, it should return the first register description in the array. The iterator is used both to access registers in a gdb frame⁸ and to provide more general register access. When used for a gdb frame, the registers to be returned *must comply with the fixed frame format expected by the gdb*. The

⁸ A *gdb frame* is the fixed-format packet associated with the *g* and *G* commands in RSP.

function is told whether or not it should return registers for the gdb frame by the `gdbFrame` argument.

The OR1K register structure iterator has this definition:

```
VMI_REG_INFO_FN(or1kRegInfo) {  
    return getNextRegister((or1kP)processor, prev, gdbFrame);  
}
```

A prototype of this function has been added to `or1kFunctions.h`. Function `getNextRegister` is defined as follows:

```
static vmiRegInfoCP getNextRegister(or1kP or1k, vmiRegInfoCP reg, Bool gdbFrame)  
{  
    do {  
        if(!reg) {  
            reg = registers;  
        } else if((reg+1)->name) {  
            reg = reg+1;  
        } else {  
            reg = 0;  
        }  
    } while(reg && !isRegSupported(or1k, reg, gdbFrame));  
    return reg;  
}
```

The function returns each member of registers in turn, skipping those that are defined to be unsupported by `isRegSupported`:

```
static Bool isRegSupported(or1kP or1k, vmiRegInfoCP reg, Bool gdbFrame) {  
    if(gdbFrame && IS_GDB_HIDDEN_REG(reg->gdbIndex)) {  
        // if this is a GDB frame request then registers that should be hidden  
        // from GDB should be ignored  
        return False;  
    } else {  
        // other registers are always supported  
        return True;  
    }  
}
```

So `isRegSupported` is responsible for hiding registers that should not be visible when this is a gdb frame request.

17.3.1 Register Read Callback Functions

Register read callback functions are defined using the `VMI_REG_READ_FN` macro, defined in `vmiDbg.h`:

```
#define VMI_REG_READ_FN(_NAME) Bool _NAME( \  
    vmiProcessorP processor, \  
    vmiRegInfoCP reg, Bool gdbFrame)
```

```
    vmiRegInfoCP reg,      \  
    void          *buffer  \  
)
```

Given a processor and a `vmiRegInfoCP` structure representing a processor register, the function should fill the passed buffer with the current value of the register in the passed processor. As an example, there is a special register read callback function for the program counter register that uses `vmirtGetPC` as follows:

```
static VMI_REG_READ_FN(readPC) {  
    *(Uns32*)buffer = (Uns32)vmirtGetPC(processor);  
    return True;  
}
```

Additionally, registers `SR` and `EPCR` are handled specially:

```
static VMI_REG_READ_FN(readSR) {  
    orlKP orlk = (orlKP)processor;  
    *(Uns32*)buffer = orlkGetSR(ork);  
    return True;  
}  
  
static VMI_REG_READ_FN(readEPC) {  
    orlKP orlk = (orlKP)processor;  
    *(Uns32*)buffer = orlk->EPC;  
    return True;  
}
```

17.3.2 Register Write Callback Functions

Register write callback functions are defined using the `VMI_REG_WRITE_FN` macro, defined in `vmiDbg.h`:

```
#define VMI_REG_WRITE_FN(_NAME) Bool _NAME( \  
    vmiProcessorP processor,  \  
    vmiRegInfoCP reg,        \  
    const void    *buffer     \  
)
```

Given a processor and a `vmiRegInfoCP` structure representing a processor register, the function should set the current value of the register in the passed processor from the buffer. As an example, there is a special register write callback function for the program counter register that uses `vmirtSetPC` as follows:

```
static VMI_REG_WRITE_FN(writePC) {  
    vmirtSetPC(processor, *(Uns32*)buffer);  
    return True;  
}
```

17.4 Raw Register Access

For registers whose values are held directly in the processor structure and which have no special behavior on access, there is no need to define read and write callback functions: instead, it is possible to define that these registers are accessible in their *raw* state. In this case, all this is required is to specify the `vmiReg` for the register in the `raw` field.

File `orlkRegisters.c` contains a helper macro used in the register table to specify a read/write register that can be accessed raw:

```
#define ORLK_RAW_REG_RW(_R, _G) \
    access    : vmi_RA_RW,      \
    raw       : ORLK_OFFSET(_R), \
    group     : _G
```

This macro is used to define access to registers R1 to R31.

Another macro defines read-only registers that can be accessed raw:

```
#define ORLK_RAW_REG_R(_R, _G) \
    access    : vmi_RA_R,      \
    raw       : ORLK_OFFSET(_R), \
    group     : _G
```

This macro is used to define access to registers R0, EPC and EEAR.

17.5 File `or1kAttrs.c`

The `modelAttrs` structure in `or1kAttrs.c` has been changed to include references to all the debug interface functions, as follows:

```
const vmiIASAttr modelAttrs = {
    . . . lines omitted . . .

    //////////////////////////////////////
    // DEBUGGER INTEGRATION SUPPORT ROUTINES
    //////////////////////////////////////

    .regGroupCB      = orlkRegGroup,
    .regInfoCB       = orlkRegInfo,
    .exceptionInfoCB = orlkExceptionInfo,
    .modeInfoCB      = orlkModeInfo,
    .getExceptionCB  = orlkGetException,
    .getModeCB       = orlkGetMode,
    .debugCB         = orlkDumpRegisters,
};
```

17.5.1 File `platform/platform.c`

The test platform for this example, `platform/platform.c`, has been changed as follows:

```
int main(int argc, char **argv) {
    // initialize CpuManager, specify debugging is required
    icmInitPlatform(ICM_VERSION, ICM_GDB_CONSOLE, "rsp", 0, "platform");
```

The third argument to `icmInitPlatform` is now `"rsp"`, to specify that a local host debugger should be attached to the platform before simulation is started using the RSP protocol. The first argument `"ICM_GDB_CONSOLE"` causes a GDB debugger to automatically be started. Alternatively, without the use of the automatic startup of the

GDB setting the fourth argument (the port number) to zero would cause the simulator to choose the next available port number on which to await a debugger connection (and print it in a GDBT_PORT message: see below).

There are also some lines added that test the register, exception and mode query functions:

```
queryRegisters(processor);
queryExceptions(processor);
queryModes(processor);
```

Function `queryRegisters` iterates over each register in each register group, printing their names:

```
static void queryRegisters(icmProcessorP processor) {
    icmPrintf("%s REGISTERS\n", icmGetProcessorName(processor, "/"));
    icmRegGroupP group = NULL;
    while((group=icmGetNextRegGroup(processor, group))) {
        icmPrintf("  GROUP %s\n", icmGetRegGroupName(group));
        icmRegInfoP reg = NULL;
        while((reg=icmGetNextRegInGroup(processor, group, reg))) {
            icmPrintf("    REGISTER %s\n", icmGetRegInfoName(reg));
        }
    }
    icmPrintf("\n");
}
```

Function `queryExceptions` uses the exception iterator and current exception query functions:

```
static void queryExceptions(icmProcessorP processor) {
    const char *name = icmGetProcessorName(processor, "/");
    if(!icmGetNextException(processor, 0)) {
        icmPrintf("%s HAS NO EXCEPTION INFORMATION\n", name);
    } else {
        icmPrintf("%s EXCEPTIONS\n", name);
        icmExceptionInfoP info = NULL;
        while((info=icmGetNextException(processor, info))) {
            icmPrintf(
                "  %s (code %u)\n",
                icmGetExceptionInfoName(info),
                icmGetExceptionInfoCode(info)
            );
        }
    }
}
```

```
        if((info=icmGetException(processor))) {
            icmPrintf(
                "current: %s (code %u)\n",
                icmGetExceptionInfoName(info),
                icmGetExceptionInfoCode(info)
            );
        }
    }

    icmPrintf("\n");
}
```

Function `queryModes` uses the mode iterator and current mode query functions:

```
static void queryModes(icmProcessorP processor) {

    const char *name = icmGetProcessorName(processor, "/");

    if(!icmGetNextMode(processor, 0)) {

        icmPrintf("%s HAS NO MODE INFORMATION\n", name);

    } else {

        icmPrintf("%s MODES\n", name);

        icmModeInfoP info = NULL;

        while((info=icmGetNextMode(processor, info))) {
            icmPrintf(
                "  %s (code %u)\n",
                icmGetModeInfoName(info),
                icmGetModeInfoCode(info)
            );
        }

        if((info=icmGetMode(processor))) {
            icmPrintf(
                "current: %s (code %u)\n",
                icmGetModeInfoName(info),
                icmGetModeInfoCode(info)
            );
        }

        icmPrintf("\n");
    }
}
```

17.6 Testing the Debugger Interface

To test the debugger interface, generate the assembler test case in directory `15.or1kDebugSupport` which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output will start with

```
Info (GDBT_PORT) Host: <hostname>, Port: <portnum>
Info (GDBT_WAIT) Waiting for remote debugger to connect...
```

17.6.1 Using Automatic GDB CONSOLE

If the ICM_GDB_CONSOLE argument is used, the simulator starts a GDB session in a new window. It will display the following output after connection of the debugger

```
Info (GDBT_CONNECTED) Client connected
```

We now have the debugger connected to the simulation, cause the simulation to start by stepping one instruction

```
(gdb) stepi
```

The remainder of the platform executes, the output should be as follows:

```
/cpul REGISTERS
GROUP GPR
REGISTER R0
REGISTER R1
REGISTER R2
REGISTER R3
REGISTER R4
REGISTER R5
REGISTER R6
REGISTER R7
REGISTER R8
REGISTER R9
REGISTER R10
REGISTER R11
REGISTER R12
REGISTER R13
REGISTER R14
REGISTER R15
REGISTER R16
REGISTER R17
REGISTER R18
REGISTER R19
REGISTER R20
REGISTER R21
REGISTER R22
REGISTER R23
REGISTER R24
REGISTER R25
REGISTER R26
REGISTER R27
REGISTER R28
REGISTER R29
REGISTER R30
REGISTER R31
GROUP System
REGISTER PC
REGISTER SR
REGISTER EPCR
```



```
REGISTER FEAR
REGISTER EXCPT

/cpul EXCEPTIONS
  RST (code 256)
  BUS (code 512)
  DPF (code 768)
  IPF (code 1024)
  TTI (code 1280)
  ILL (code 1792)
  EXI (code 2048)
  SYS (code 3072)
current: RST (code 256)

/cpul MODES
  SUPERVISOR (code 0)
  USER (code 1)
current: SUPERVISOR (code 0)
```

The output first lists all of the *registers* in the processor model (by group), then the *exceptions*, then the *modes*.

We can carry out normal debugging commands supported by *gdb* – for example, try setting a breakpoint at `fib`, continuing and disassembling. Also experiment with the *gdb* `info registers` command to show the current contents of all registers.

17.6.2 Using Manual Remote GDB Connection

If the “ICM_GDB_CONSOLE” argument is not used, the simulator waits for a debugger connection to be established. The GDB can be started and connected using the following instructions.

It is simpler if the following is carried out in a separate terminal, so that standard output from the *gdb* debug session and the simulation are separated. Run up the *gdb* debugger for the OR1K processor using the following command in the `15.or1kDebugSupport` directory:

```
${IMPERAS_HOME}/lib/${IMPERAS_ARCH}/CrossCompiler/or32-elf/bin/or32-elf-gdb
```

You should see the following output:

```
GNU gdb 5.3
Copyright 2002 Free Software Foundation, Inc.
GDB is free software, covered by the GNU General Public License, and you are
welcome to change it and/or distribute copies of it under certain conditions.
Type "show copying" to see the conditions.
There is absolutely no warranty for GDB. Type "show warranty" for details.
This GDB was configured as "--host=i686-pc-*****-gnu --target=or32-elf".
(gdb)
```

To connect to the simulation and start debugging the application, use the following commands, replacing `<portnum>` with the number of the port echoed by the simulator as it started (see the `GDBT_PORT` message above):

```
(gdb) file application/asmtest.OR1K.elf
(gdb) target remote localhost:<portnum>
```

The simulator will display the following output after connection of the debugger

```
Info (GDBT_CONNECTED) Client connected
```

We now have the debugger connected to the simulation, cause the simulation to start by stepping one instruction

```
(gdb) stepi
```

We can also carry out normal debugging commands supported by *gdb* – for example, try setting a breakpoint at `fib`, continuing and disassembling. Also experiment with the *gdb* `info registers` command to show the current contents of all registers.

18 Implementing Fixed-Mapped Virtual Memory

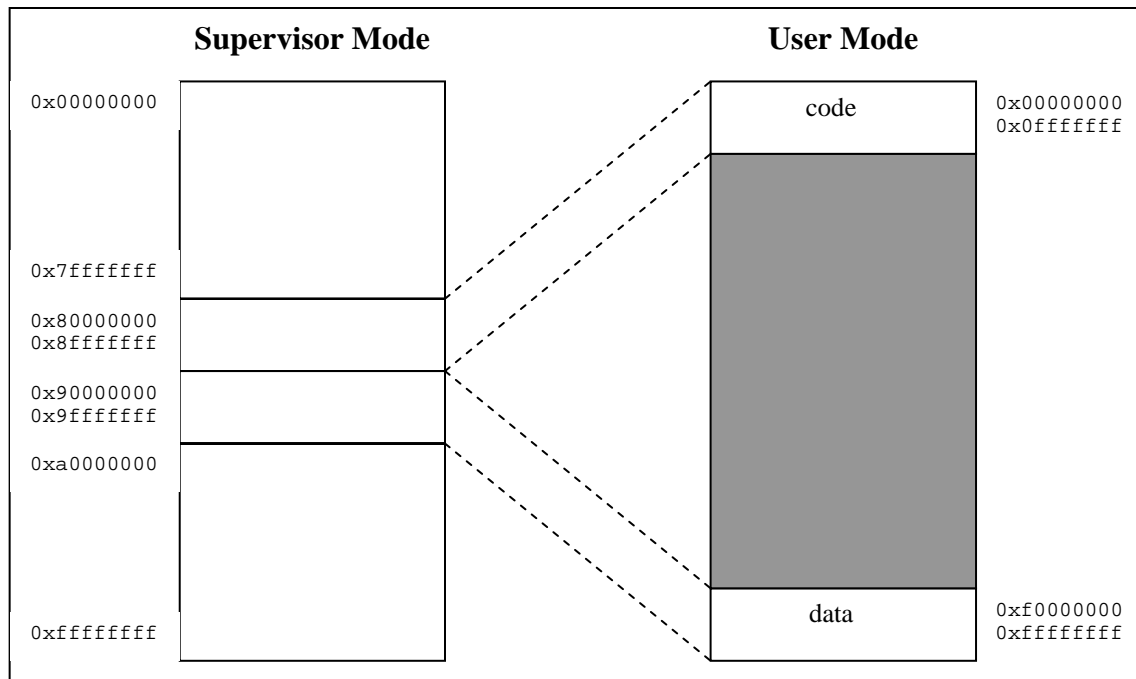
The OR1K example described so far in this document has only physical memory. It is also possible to implement *virtual memory* in a highly-efficient manner.

This chapter shows how to implement *fixed-mapped* virtual memory, where the virtual-to-physical address mappings are constant. Chapter 19 extends this to show how to implement *dynamic-mapped* virtual memory, where the virtual-to-physical address mappings can be changed at run time.

18.1 Example Memory Maps

The OR1K processor can be configured with a full TLB-based virtual memory system. Although it is perfectly possible to implement the full OR1K virtual memory algorithm using the VMI interface, the code to do this is too complex for an introductory example. Therefore, this chapter will implement a much simpler memory mapping scheme that demonstrates the required concepts without the complexity of the true virtual memory algorithm.

The fixed mapping scheme we will implement will look like this:



In other words, the example memory mapping scheme will be as follows:

1. In *supervisor* mode, the entire address space is mapped with full access permissions (read, write and execute)
2. In *user* mode, there are three distinct memory areas:
 - a. The address range 0x00000000:0x0fffffff is *code* memory which has *execute and read* permission, but is *not writable in user mode* (any attempt

to write to this address range should generate an exception). It is mapped to the physical address range 0x80000000:0x8fffffff (so, for example, address 0x80000000 in supervisor mode addresses the same location as address 0x00000000 in user mode).

- b. The address range 0xf0000000:0xffffffff is *data* memory which has *read and write* permission, but is *not executable in user mode* (any attempt to execute code in this address range should generate an exception). It is mapped to the physical address range 0x90000000:0x9fffffff (so, for example, address 0x90000000 in supervisor mode addresses the same location as address 0xf0000000 in user mode).
- c. The address range 0x10000000:0xefffffff is unmapped in user mode; any attempt to access this memory in any way should generate an exception.

18.2 The Template Fixed-Mapped Model

A template model for the OR1K processor implementing a fixed-mapped virtual memory scheme can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/16.or1kBehaviorVM
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/16.or1kBehaviorVM .
```

Compile the model using make:

```
cd 16.or1kBehaviorVM
make
```

The processor model is based on the previous model, with the changes listed in following sections.

18.3 or1kVM.c

The fixed-mapped virtual memory scheme described in the previous section is created by defining a *virtual memory constructor* function for the OR1K processor. This is implemented as follows in file `or1kVM.c`:

```
VMI_VMINIT_FN(or1kVMInit) {
    // for this example, the SUPERVISOR memory domain will be the physical
    // domain; the USER memory domain will be a derived virtual domain
    memDomainP physicalDomain = codeDomains[0];
    memDomainP supervisorDomain = physicalDomain;
    memDomainP userDomain = vmirtNewDomain("user", 32);

    // create an initial mapping that makes physical addresses 0x80000000:
    // 0x8fffffff visible at 0x00000000:0x0fffffff in the USER memory space
    vmirtAliasMemory(
        physicalDomain,
        userDomain,
```

```

        0x80000000,
        0x8fffffff,
        0x00000000,
        0
    );

    // remove write permissions from addresses 0x00000000:0xffffffff in the
    // user address space
    vmirtProtectMemory(
        userDomain,
        0x00000000,
        0xffffffff,
        MEM_PRIV_RX,
        MEM_PRIV_SET
    );

    // create an initial mapping that makes physical addresses 0x90000000:
    // 0x9fffffff visible at 0xf0000000:0xffffffff in the USER memory space
    vmirtAliasMemory(
        physicalDomain,
        userDomain,
        0x90000000,
        0x9fffffff,
        0xf0000000,
        0
    );

    // remove execute permissions from addresses 0xf0000000:0xffffffff in the
    // USER memory space
    vmirtProtectMemory(
        userDomain,
        0xf0000000,
        0xffffffff,
        MEM_PRIV_RW,
        MEM_PRIV_SET
    );

    // supervisorDomain should be used for both instruction and data in
    // SUPERVISOR mode
    codeDomains[OR1K_MODE_SUPERVISOR] = supervisorDomain;
    dataDomains[OR1K_MODE_SUPERVISOR] = supervisorDomain;

    // userDomain should be used for both instruction and data in USER mode
    codeDomains[OR1K_MODE_USER] = userDomain;
    dataDomains[OR1K_MODE_USER] = userDomain;
}

```

The virtual memory constructor is defined using the `VMI_VMINIT_FN` macro, defined as follows in `vmiAttrs.h`:

```

#define VMI_VMINIT_FN(_NAME) void _NAME( \
    vmiProcessorP processor,          \
    memDomainPP codeDomains,         \
    memDomainPP dataDomains          \
)

```

The virtual memory constructor is called by the simulator after the processor constructor has been called. It is passed three arguments:

1. the newly-constructed processor;
2. an array of memory domain objects for instruction fetch (the *code* domains);

3. an array of memory domain objects for data access (the *data* domains).

The number of entries in each of the code and data domain arrays is the same as the number of *code dictionaries* that the processor has. The number of code dictionaries is determined by the length of the dictionary names array in file `or1kAttrs.c`:

```
static const char *dictNames[] = {"SUPERVISOR", "USER", 0};
```

For the OR1K, each of the code and data domain arrays will therefore contain two entries. Each of the entries is used as follows:

1. entry 0 of `codeDomains` is used when fetching instructions in *supervisor* mode;
2. entry 1 of `codeDomains` is used when fetching instructions in *user* mode;
3. entry 0 of `dataDomains` is used for loads and stores in *supervisor* mode;
4. entry 1 of `dataDomains` is used for loads and stores in *user* mode;

By default, each entry in `codeDomains` and `dataDomains` is seeded with an identical *physical* domain⁹. This means that, when a virtual memory constructor is not specified, all processor modes access the same memory space. However, the default entries in the `codeDomains` and `dataDomains` arrays can be overridden with new memory domain objects in the virtual memory constructor to specify different mappings for each processor mode, as described in detail below.

```
memDomainP physicalDomain = codeDomains[0];
```

This line gets the physical domain automatically associated with the OR1K processor when it was created.

```
memDomainP supervisorDomain = physicalDomain;
```

In supervisor mode, we want to use the physical domain for all accesses.

```
memDomainP userDomain = vmirtNewDomain("user", 32);
```

In user mode, we do not want to use the physical domain, but instead we will use a new domain, with a 32-bit address width, created by calling `vmirtNewDomain`.

```
vmirtAliasMemory(  
    physicalDomain,  
    userDomain,  
    0x80000000,  
    0x8fffffff,  
    0x00000000,  
    0  
);
```

This call to `vmirtAliasMemory` maps addresses `0x80000000:0x8fffffff` in `physicalDomain` to addresses `0x80000000:0x8fffffff` in `userDomain`. See the *VMI*

⁹ In fact, it is possible in the ICM interface to specify separate physical code and physical data domains – see the *OVPsim and CpuManager User Guide*.

Run Time Function Reference documentation for more information about `vmirtAliasMemory`.

```
vmirtProtectMemory(  
    userDomain,  
    0x00000000,  
    0x0fffffff,  
    MEM_PRIV_RX,  
    MEM_PRIV_SET  
);
```

This call to `vmirtProtectMemory` sets the access permissions on the address range `0x80000000:0x8fffffff` in `userDomain` to execute and read (but not write).

```
vmirtAliasMemory(  
    physicalDomain,  
    userDomain,  
    0x90000000,  
    0x9fffffff,  
    0xf0000000,  
    0  
);  
  
vmirtProtectMemory(  
    userDomain,  
    0xf0000000,  
    0xffffffff,  
    MEM_PRIV_RW,  
    MEM_PRIV_SET  
);
```

These two calls map addresses `0x90000000:0x9fffffff` in `physicalDomain` to addresses `0xf0000000:0xffffffff` in `userDomain`, and give that address range read and write (but not execute) permission.

```
codeDomains[OR1K_MODE_SUPERVISOR] = supervisorDomain;  
dataDomains[OR1K_MODE_SUPERVISOR] = supervisorDomain;  
  
codeDomains[OR1K_MODE_USER] = userDomain;  
dataDomains[OR1K_MODE_USER] = userDomain;
```

These lines override the default domains to use for the various access types and processor modes, so that `supervisorDomain` is used for all supervisor mode accesses and `userDomain` is used for all user mode accesses.

18.4 *File or1kAttrs.c*

A prototype for `or1kVMInit` has been added to `or1kFunctions.h`, and is referenced in the attribute structure in `or1kAttrs.c`:

```
const vmiIASAttr modelAttrs = {  
  
    //////////////////////////////////////  
    // VERSION & SIZE ATTRIBUTES  
    //////////////////////////////////////  
}
```

```
.versionString = VMI_VERSION,
.modelType     = VMI_PROCESSOR_MODEL,
.dictNames     = dictNames,
.cpuSize       = sizeof(ork),

////////////////////////////////////
// CREATE/DELETE ROUTINES
////////////////////////////////////

.constructorCB = orkConstructor,
.vmInitCB      = orkVMInit,
.destructorCB  = orkDestructor,

... etc ...
```

18.4.1 File platform/platform.c

The test platform for this example, `platform/platform.c`, has been changed to enable both tracing and simulated exceptions in the OR1K processor:

```
#define MODEL_FLAGS (ICM_ATTR_TRACE | ICM_ATTR_TRACE_ICOUNT | ICM_ATTR_SIMEX)
```

Simulated exceptions are required because we want the processor to run exception handlers on illegal read, write or instruction fetch – if this is not enabled, then simulation would terminate on an illegal access.

18.5 Testing Fixed-Mapped Virtual Memory

To test the virtual memory behavior, generate the assembler test case in directory `16.orkBehaviorVM` which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```
Info 1: 'cpu1', 0x0000000000000000: l.ori    r31,r0,0x0
Info 2: 'cpu1', 0x0000000000000004: l.movhi  r1,0x8000
Info 3: 'cpu1', 0x0000000000000008: l.movhi  r2,0x0
Info 4: 'cpu1', 0x000000000000000c: l.ori    r2,r2,0xc0c
Info 5: 'cpu1', 0x0000000000000010: l.movhi  r3,0x0
Info 6: 'cpu1', 0x0000000000000014: l.ori    r3,r3,0xc2c
Info 7: 'cpu1', 0x0000000000000018: l.lwz    r4,0x0(r2)
Info 8: 'cpu1', 0x000000000000001c: l.sw     0x0(r1),r4
Info 9: 'cpu1', 0x0000000000000020: l.addi   r2,r2,0x4
Info 10: 'cpu1', 0x0000000000000024: l.addi   r1,r1,0x4
Info 11: 'cpu1', 0x0000000000000028: l.sfne   r2,r3
Info 12: 'cpu1', 0x000000000000002c: l.bf     0x00000018
... etc ...
Info 61: 'cpu1', 0x000000000000002c: l.bf     0x00000018
Info 62: 'cpu1', 0x0000000000000030: l.nop    0x0
Info 63: 'cpu1', 0x0000000000000034: l.mtspr  r0,r0,64
```



```

Info 64: 'cpul', 0x0000000000000038: l.mtspr r0,r0,32
Info 65: 'cpul', 0x000000000000003c: l.rfe
Info 66: 'cpul', 0x0000000000000000: l.movhi r1,0xf000
Info 67: 'cpul', 0x0000000000000004: l.ori r2,r0,0x1234
Info 68: 'cpul', 0x0000000000000008: l.sw 0x0(r1),r2
Info 69: 'cpul', 0x000000000000000c: l.sw 0x0(r0),r2
Info 70: 'cpul', 0x0000000000000300: l.mfspr r1,r0,32
Info 71: 'cpul', 0x0000000000000304: l.addi r1,r1,0x4
Info 72: 'cpul', 0x0000000000000308: l.mtspr r0,r1,32
Info 73: 'cpul', 0x000000000000030c: l.rfe
Info 74: 'cpul', 0x0000000000000010: l.movhi r1,0xf000
Info 75: 'cpul', 0x0000000000000014: l.jalr r1
Info 76: 'cpul', 0x0000000000000018: l.nop 0x0
Info 77: 'cpul', 0x00000000f0000000: *** FETCH EXCEPTION ***
Info 78: 'cpul', 0x0000000000000400: l.mtspr r0,r9,32
Info 79: 'cpul', 0x0000000000000404: l.rfe
Info 80: 'cpul', 0x000000000000001c: l.sys
Info 81: 'cpul', 0x00000000000000c0: l.movhi r1,0x9000
Info 82: 'cpul', 0x00000000000000c04: l.lwz r3,0x0(r1)
Info 83: 'cpul', 0x00000000000000c08: l.nop 0x0
Processor 'cpul' terminated at 'exit', address 0xc08

```

```

-----
R0 : 00000000 R1 : 90000000 R2 : 00001234 R3 : 00001234
R4 : 20000000 R5 : deadbeef R6 : deadbeef R7 : deadbeef
R8 : deadbeef R9 : 0000001c R10: deadbeef R11: deadbeef
R12: deadbeef R13: deadbeef R14: deadbeef R15: deadbeef
R16: deadbeef R17: deadbeef R18: deadbeef R19: deadbeef
R20: deadbeef R21: deadbeef R22: deadbeef R23: deadbeef
R24: deadbeef R25: deadbeef R26: deadbeef R27: deadbeef
R28: deadbeef R29: deadbeef R30: deadbeef R31: 00000000
PC : 00000c0c SR : 00008001 ESR: 00008000 EPC: 00000020
TCR: 00000000 TMR: 00000000 PSR: 00000000 PMR: 00000000
BF:0 CF:0 OF:0
-----

```

processor has executed 83 instructions

The source code for this example is as follows:

```

////////////////////////////////////////////////////
// KERNEL START CODE (AT 0x0)
////////////////////////////////////////////////////
.global _start
_start:
    l.ori        r31,r0,0           // r31 = 0 (stack pointer)

    // prepare to copy application to user space
    l.movhi      r1,0x8000          // r1 = 0x80000000
    l.movhi      r2,hi(appStart)    // r2 = appStart
    l.ori        r2,r2,lo(appStart)
    l.movhi      r3,hi(appEnd)      // r3 = appEnd
    l.ori        r3,r3,lo(appEnd)

    // copy application to user space
loop:   l.lwz      r4,0(r2)           // r4 = word of application code
        l.sw      0(r1),r4           // copy to 0x80000000
        l.addi    r2,r2,4            // increment src pointer
        l.addi    r1,r1,4            // increment dst pointer
        l.sfne    r2,r3             // r2!=r3?
        l.bf      loop              // go if true
        l.nop                                // (delay slot instruction)

```

```

        // run user code
        l.mtspr      r0,r0,0x40      // clear esr
        l.mtspr      r0,r0,0x20      // clear epc (resume address in user space)
        l.rfe        // return from exception (runs user code)

.org 0x300
        ///////////////////////////////////////////////////
        // DATA PRIVILEGE EXCEPTION VECTOR (AT 0x300) - SKIP INSTRUCTION
        ///////////////////////////////////////////////////
        l.mfspr      r1,r0,0x20      // get epc in r1
        l.addi       r1,r1,4          // increment address to skip faulting insn
        l.mtspr      r0,r1,0x20      // set epc
        l.rfe        // return from exception

.org 0x400
        ///////////////////////////////////////////////////
        // CODE PRIVILEGE EXCEPTION VECTOR (AT 0x400) - RESUME AT LINK ADDRESS
        ///////////////////////////////////////////////////
        l.mtspr      r0,r9,0x20      // set epc from link register (r9)
        l.rfe        // return from exception

.org 0xc00
        ///////////////////////////////////////////////////
        // SYSCALL VECTOR (AT 0xc00) - TERMINATE PROGRAM
        ///////////////////////////////////////////////////
        l.movhi      r1,0x9000      // r1 = 0x90000000
        l.lwz        r3,0(r1)        // load r3 from address 0x90000000
.global exit
exit:    l.nop

        ///////////////////////////////////////////////////
        // USER APPLICATION (IN KERNEL MEMORY)
        ///////////////////////////////////////////////////
appStart:
        l.movhi      r1,0xf000      // r1 = 0xf0000000
        l.ori        r2,r0,0x1234   // r2 = 0x00001234
        l.sw         0(r1),r2       // legal store to 0xf0000000
        l.sw         0(r0),r2       // attempt ILLEGAL store to 0x00000000
        l.movhi      r1,0xf000      // r1 = 0xf0000000
        l.jalr       r1             // attempt ILLEGAL jump to 0xf0000000
        l.nop        // (delay slot instruction)
        l.sys        0              // exit program
appEnd:

```

Execution starts at label `_start`. The application (running in supervisor mode) first executes a loop to copy the code between labels `appStart` and `appEnd` to address `0x80000000` (so this code will become visible in the user address space at address `0x00000000`). When this is done (at instruction 63), the processor clears `esr` and `epc` and executes an `l.rfe` instruction (return from exception) to start executing at address `0x00000000` in *user* mode:

```

        l.mtspr      r0,r0,0x40      // clear esr
        l.mtspr      r0,r0,0x20      // clear epc (resume address in user space)
        l.rfe        // return from exception (runs user code)

```

Note that the trace output now shows *user mode instruction addresses*:

```

Info 66: 'cpu1', 0x0000000000000000: l.movhi  r1,0xf000
Info 67: 'cpu1', 0x0000000000000004: l.ori    r2,r0,0x1234

```

```
Info 68: 'cpu1', 0x0000000000000008: l.sw      0x0(r1),r2
```

The application stores the value 0x1234 to address 0xf0000000 (legal in user mode) and then attempts the same store to address 0x00000000 (illegal in user mode, as this address has read and execute permissions only). This causes a data privilege exception at instruction 70. The data privilege exception handler simply increments `epcr` and returns from the exception (to skip the faulting instruction):

```
Info 69: 'cpu1', 0x000000000000000c: l.sw      0x0(r0),r2
Info 70: 'cpu1', 0x0000000000000300: l.mfspr  r1,r0,32
Info 71: 'cpu1', 0x0000000000000304: l.addi   r1,r1,0x4
Info 72: 'cpu1', 0x0000000000000308: l.mtspr  r0,r1,32
Info 73: 'cpu1', 0x000000000000030c: l.rfe
```

Next, the application attempts a call to address 0xf0000000 (also illegal in user mode, as this address has read and write permissions only). This causes an instruction privilege exception at instruction 78. The instruction privilege exception handler resumes execution at the link address (0x1c in the user address space):

```
Info 75: 'cpu1', 0x0000000000000014: l.jalr   r1
Info 76: 'cpu1', 0x0000000000000018: l.nop    0x0
Info 77: 'cpu1', 0x00000000f0000000: *** FETCH EXCEPTION ***
Info 78: 'cpu1', 0x0000000000000400: l.mtspr  r0,r9,32
Info 79: 'cpu1', 0x0000000000000404: l.rfe
```

Finally, the user mode program executes an `l.sys` instruction. The system call vector loads the contents of address 0x90000000 into register `r3` and then exits. On completion, register `r3` contains 0x1234, proving that address 0x90000000 in supervisor mode is correctly mapped to address 0xf0000000 in user mode:

```
-----
R0 : 00000000   R1 : 90000000   R2 : 00001234   R3 : 00001234
R4 : 20000000   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : 0000001c  R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: deadbeef   R31: 00000000
PC  : 00000c0c  SR  : 00008001  ESR: 00008000  EPC: 00000020
TCR: 00000000  TMR: 00000000  PSR: 00000000  PMR: 00000000
BF:0 CF:0 OF:0
-----
```

19 Implementing a Dynamic-Mapped TLB

Chapter 18 showed how to implement *fixed-mapped* virtual memory. This chapter extends that example to show how to implement *dynamic-mapped* memory, where the virtual-to-physical address mappings can be changed at run time. Specifically, we implement a simple TLB structure.

19.1 General TLB Concepts

Most processors that implement virtual memory have hardware support for virtual-to-physical address translation, implemented using a *translation lookaside buffer* (TLB). The TLB is typically an associative cache of a number of valid virtual-to-physical page mappings. If a memory access (fetch, read or write) uses a virtual address that is mapped in the TLB, the corresponding physical address location is obtained immediately and execution continues uninterrupted. Otherwise, if there is a *TLB miss*, a supervisor or kernel mode routine is typically called to handle the miss. The miss handler typically has access to a much larger table of virtual-to-physical address mappings in kernel memory: if the virtual address was present in this table but not in the TLB, the miss handler will eject an existing valid entry and replace it with the required entry from the page table; if the virtual address was not valid, a privilege exception will typically be generated. When an entry is being replaced, a key concept is the *replacement policy* that determines which existing valid entry is to be ejected¹⁰.

Note that a TLB miss is often not handled the same way as a cache miss: typically, a cache miss simply stalls the processor while the cache line is filled, whereas a TLB miss usually requires an exception handler to be called to process it and modify the TLB contents

TLB entries typically hold more than just virtual-to-physical mappings: usually, each has some access permissions associated with it, and often an address space id (ASID) that restricts the validity of the entry to a subset of running processes (this prevents the TLB having to be completely flushed and refilled on a process switch). There is often also a choice of page sizes supported.

19.2 The Simple Example TLB

As mentioned in chapter 18, the OR1K processor can be configured with a full virtual memory system implementing a TLB. However, although this can be modeled using the VMI interface, it is too complex for an introductory example; instead, we will model a very simple TLB structure that demonstrates the key concepts but is easier to understand.

The simple TLB will extend the fixed-mapped virtual memory example of chapter 18 as follows:

1. Memory at addresses in the range `0x10000000:0xffffffff` in user memory space will be divided into 4096-byte pages that will be mapped on demand.

¹⁰ In some processor architectures (e.g. ARM) the process of TLB maintenance is handled by special hardware.

2. The OR1K processor will implement a small TLB with four entries, so up to four different 4096-byte pages may be mapped concurrently.
3. At address 0x10000 in the supervisor address space, we will locate a page table structure. The data in this will be laid out as follows:
 - a. 0x10000: the number of entries in the table
 - b. 0x10004: virtual address for entry 1
 - c. 0x10008: virtual address for entry 2
 - d. 0x1000c: virtual address for entry 3
 - e. (and so on up to the number of entries in the table)

The page table will dynamically grow as more pages in the user address space need to be mapped.

4. The physical pages corresponding to the virtual addresses in the page table will be allocated in order starting at 0x10000000 in supervisor space. In other words, entry 1 in the table will be mapped to physical address 0x10000000, entry 2 will be mapped to physical address 0x10001000, entry 3 will be mapped to physical address 0x10002000, and so on.
5. In order to update the processor TLB, we will modify the behavior of the `l.nop 99` instruction. When `l.nop 99` is executed, we will assume that:
 - a. register `r1` holds a TLB index number (0-3) indicating the TLB entry to replace;
 - b. register `r2` holds a user space virtual address;
 - c. register `r3` holds a physical address.

Executing the `l.nop 99` should discard any current mapping for the numbered TLB entry in the processor, and create a new mapping for that entry, mapping one page of data at the virtual address in the user memory domain to the physical address in the physical memory domain.

6. The replacement policy in this example will be simple round-robin.
7. In this simple example, all dynamically mapped pages will have read/write privilege, and there will be no way to invalidate an entry in the TLB (other than to replace it with a new mapping). We will also not implement any means of freeing an allocated physical page. All these would of course be supported in a real operating system, but would over complicate this example.

19.3 The Template Simple TLB Model

A template model for the OR1K processor implementing a simple TLB can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/17.or1kBehaviorTLB
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/17.or1kBehaviorTLB .
```

Compile the model using make:

```
cd 17.or1kBehaviorTLB
make
```

The processor model is based on the previous model, with the changes listed in following sections.

19.4 *File or1kStructure.h*

Some new #defines describe some aspects of the TLB:

```
#define OR1K_TLB_SIZE 4           // size of TLB
#define OR1K_PAGE_SIZE 4096      // size of TLB-mapped page
```

There is also a new type, `tlbEntry`, implementing a single entry in the TLB:

```
// simple TLB entry
typedef struct tlbEntryS {
    Uns32 va;           // virtual address
    Uns32 pa;           // physical address
    Bool  valid;        // is entry valid?
} tlbEntry, *tlbEntryP;
```

The processor structure has an array of TLB entries, and also a new memory domain field giving the physical domain for use during dynamic mapping:

```
typedef struct or1kS {
    . . . lines omitted . . .

    memDomainP physicalDomain; // physical domain target for TLB mappings
    tlbEntry tlb[OR1K_TLB_SIZE]; // simulated TLB
} or1k, *or1kP;
```

19.5 *Files or1kVM.c and or1kVM.h*

The virtual memory constructor has been modified to save the physical domain for use in the TLB update function:

```
VMI_VMINIT_FN(or1kVMInit) {
    . . . lines omitted . . .

    // save physicalDomain on the OR1K structure for TLB usage
    or1kP or1k = (or1kP)processor;
    or1k->physicalDomain = physicalDomain;
}
```

A new function, `or1kSetTLBEntry`, is used to update the TLB mapping in the OR1K structure. The function takes as arguments the processor, a virtual address (`va`) and a physical address (`pa`):

```
void or1kSetTLBEntry(or1kP or1k, Uns32 tlbIndex, Uns32 va, Uns32 pa) {

    // get the current instruction count (for messages)
    Uns64 iCount = vmirtGetICount((vmiProcessorP)or1k);

    // clip tlbIndex, va and pa to valid values
    tlbIndex &= (OR1K_TLB_SIZE-1);
```

```
va      &= ~(OR1K_PAGE_SIZE-1);
pa      &= ~(OR1K_PAGE_SIZE-1);

// get the TLB entry to update and the memory domain affected
tlbEntryP entry      = &or1k->tlb[tlbIndex];
memDomainP physicalDomain = or1k->physicalDomain;
memDomainP tlbDomain   = or1k->tlbDomain;

// if TLB entry is already mapped, unmap it
if(entry->valid) {

    vmiPrintf(
        FMT_64u": DELETE entry %u mapping (va:0x%08x pa:0x%08x)\n",
        iCount, tlbIndex, entry->va, entry->pa
    );

    vmirtUnaliasMemoryVM(
        tlbDomain, entry->va, entry->va+OR1K_PAGE_SIZE-1, True, 0
    );
}

// update the TLB entry
entry->va      = va;
entry->pa      = pa;
entry->valid   = True;

vmiPrintf(
    FMT_64u": CREATE entry %u mapping (va:0x%08x pa:0x%08x)\n",
    iCount, tlbIndex, va, pa
);

// establish the new page mapping with read/write permissions
vmirtAliasMemoryVM(
    physicalDomain, tlbDomain, pa, pa+OR1K_PAGE_SIZE-1, va, 0,
    MEM_PRIV_RW, True, 0
);
};
```

The function first gets the TLB entry to update, and the physical and TLB domains:

```
tlbEntryP entry      = &or1k->tlb[tlbIndex];
memDomainP physicalDomain = or1k->physicalDomain;
memDomainP tlbDomain   = or1k->tlbDomain;
```

Next, it invalidates any current mapping for the TLB entry using `vmirtUnaliasMemoryVM`:

```
Bool vmirtUnaliasMemoryVM(
    memDomainP virtualDomain,
    Addr      virtualLowAddr,
    Addr      virtualHighAddr,
    Bool      isGlobal,
    Uns32     ASID
);
```

The arguments to this function are:

1. The domain in which to remove a virtual memory mapping.
2. The address range for which the mapping should be removed.

3. A boolean indicating whether the entry is a global entry, if `True`, or an ASID-managed entry, if `False` (see *TLB Modeling with ASID-Labeled Entries* later in this section). In this example, the TLB entries are assumed to be globally-mapped.
4. If the mapping to remove is ASID-managed, the corresponding ASID.

See the *VMI Run Time Function Reference* manual for more information about this function.

To monitor what is happening, the details of the deleted mapping are also printed out:

```
if(entry->valid) {  
    vmiPrintf(  
        FMT_64u": DELETE entry %u mapping (va:0x%08x pa:0x%08x)\n",  
        iCount, tlbIndex, entry->va, entry->pa  
    );  
    vmirtUnaliasMemoryVM(  
        tlbDomain, entry->va, entry->va+OR1K_PAGE_SIZE-1, True, 0  
    );  
}
```

Next, the TLB entry is updated to describe the new mapping, and the details of the new mapping are printed out:

```
entry->va    = va;  
entry->pa    = pa;  
entry->valid = True;  
  
vmiPrintf(  
    FMT_64u": CREATE entry %u mapping (va:0x%08x pa:0x%08x)\n",  
    iCount, tlbIndex, va, pa  
);
```

Finally, the new mapping is established with read/write access permissions, using function `vmirtAliasMemoryVM`:

```
Bool vmirtAliasMemoryVM(  
    memDomainP physicalDomain,  
    memDomainP virtualDomain,  
    Addr        physicalLowAddr,  
    Addr        physicalHighAddr,  
    Addr        virtualLowAddr,  
    memMRUSetP  mruSet,  
    memPriv     privilege,  
    Bool        isGlobal,  
    Uns32       ASID  
);
```

The arguments to this function are:

1. The *physical domain* to which to map.
2. The *virtual domain* in which the mapping should be created.
3. The address range in the physical domain of the region to map.
4. The base address in the virtual domain of the region to map.

5. A structure of type `memMRUSetP`, which automates the maintenance of region usage so that the *least-recently-used* TLB entry can be identified if required. In this example we are not interested in maintaining least-recently-used state (the replacement policy is round-robin) so this argument is `NULL`. See chapter 20 which discusses this in more detail.
6. An argument specifying the access privileges that the page should have.
7. A boolean indicating whether the entry is a global entry, if `True`, or an ASID-managed entry, if `False` (see *TLB Modeling with ASID-Labeled Entries* later in this section). In this example, the TLB entries are assumed to be globally-mapped.
8. If the mapping to create is ASID-managed, the corresponding ASID.

```
vmirtAliasMemoryVM(  
    physicalDomain, tlbDomain, pa, pa+OR1K_PAGE_SIZE-1, va, 0,  
    MEM_PRIV_RW, True, 0  
);
```

19.6 File *or1kMorph.c*

A new morpher callback function has been created to handle an `l.nop 99` instruction, calling the new TLB update function `or1kSetTLBEntry`. To enable an example application to be debugged, the `l.nop 98` instruction has also been subverted to print the current contents of register `r1`:

```
static void vmic_printVal(Uns32 val) {  
    vmiPrintf("    fib returns %u\n", val);  
}  
  
static void doNop(Uns32 instr) {  
  
    Uns32 code = OP3_I(instr);  
  
    // subvert nop 98 to print the contents of r1 and nop 99 to set a TLB entry  
    if(code==98) {  
        vmimtArgReg(OR1K_BITS, OR1K_REG(1));  
        vmimtCall((vmiCallFn)vmic_printVal);  
    } else if(code==99) {  
        vmimtArgProcessor();  
        vmimtArgReg(OR1K_BITS, OR1K_REG(1));    // r1: TLB entry index  
        vmimtArgReg(OR1K_BITS, OR1K_REG(2));    // r2: virtual address (va)  
        vmimtArgReg(OR1K_BITS, OR1K_REG(3));    // r3: physical address (pa)  
        vmimtCall((vmiCallFn)or1kSetTLBEntry);  
    }  
}  
  
static OR1K_DISPATCH_FN(morphNOP)    {doNop(instr);}
```

It can often be a very useful technique to subvert certain opcodes to assist application debugging. Many processor instruction sets contain unassigned instructions for just this purpose. Make sure not to leave debugging hacks in the final model!

19.7 Testing the Simple TLB Model

To test the TLB behavior, generate the assembler test case in directory 17.or1kBehaviorTLB which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```

    fib returns 1
369: CREATE entry 0 mapping (va:0x80001000 pa:0x10000000)
    fib returns 2
439: CREATE entry 1 mapping (va:0x80002000 pa:0x10001000)
    fib returns 1
    fib returns 3
549: CREATE entry 2 mapping (va:0x80003000 pa:0x10002000)
    fib returns 1
    fib returns 2
    fib returns 5
699: CREATE entry 3 mapping (va:0x80005000 pa:0x10003000)
    fib returns 1
    fib returns 2
    fib returns 1
    fib returns 3
    fib returns 8
919: DELETE entry 0 mapping (va:0x80001000 pa:0x10000000)
919: CREATE entry 0 mapping (va:0x80008000 pa:0x10004000)
    fib returns 1
1021: DELETE entry 1 mapping (va:0x80002000 pa:0x10001000)
1021: CREATE entry 1 mapping (va:0x80001000 pa:0x10000000)
    fib returns 2
1092: DELETE entry 2 mapping (va:0x80003000 pa:0x10002000)
1092: CREATE entry 2 mapping (va:0x80002000 pa:0x10001000)
    fib returns 1
    fib returns 3
. . . lines omitted . . .
    fib returns 233
62313: DELETE entry 1 mapping (va:0x80005000 pa:0x10003000)
62313: CREATE entry 1 mapping (va:0x800e9000 pa:0x1000b000)
    fib returns 610
62494: DELETE entry 2 mapping (va:0x8000d000 pa:0x10005000)
62494: CREATE entry 2 mapping (va:0x80262000 pa:0x1000d000)
Processor 'cpul' terminated at 'exit', address 0xc00
-----
R0 : 00000000   R1 : 00000262   R2 : 80262000   R3 : 00000001
R4 : 15000000   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : 00000008   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 000001f3   R31: 00000000
PC  : 00000c04   SR  : 00008201   ESR: 00008200   EPC: 0000000c
```

```
TCR: 00000000   TMR: 00000000   PSR: 00000000   PMR: 00000000
BF:1 CF:0 OF:0
-----
```

```
processor has executed 62512 instructions
```

The source code for this example, in file `application/asmtest.S`, breaks down into three main sections. Firstly, there is a section that initializes the processor and copies the user program to `0x80000000` before entering user mode and starting the program. This is exactly as in the example in chapter 18:

```

////////////////////////////////////
// KERNEL START CODE (AT 0x0)
////////////////////////////////////
.global _start
_start:
    l.ori      r31,r0,0          // r31 = 0 (stack pointer)
    l.ori      r30,r0,0          // r30 = 0 (TLB miss count)

    // prepare to copy application to user space
    l.movhi    r1,0x8000        // r1 = 0x80000000
    l.movhi    r2,hi(appStart)  // r2 = appStart
    l.ori      r2,r2,lo(appStart)
    l.movhi    r3,hi(appEnd)    // r3 = appEnd
    l.ori      r3,r3,lo(appEnd)

    // copy application to user space
loop:   l.lwz      r4,0(r2)       // r4 = word of application code
        l.sw       0(r1),r4      // copy to 0x80000000
        l.addi     r2,r2,4       // increment src pointer
        l.addi     r1,r1,4       // increment dst pointer
        l.sfne     r2,r3        // r2!=r3?
        l.bf       loop         // go if true
        l.nop                     // (delay slot instruction)

    // run user code
    l.mtspr     r0,r0,0x40       // clear esr
    l.mtspr     r0,r0,0x20       // clear epc (resume address in user space)
    l.rfe                     // return from exception (runs user code)

```

The user mode application calculates `fib(15)` using the naive recursive algorithm previously described in this document. On every return from the `fib` function, it reads and writes to an address in user memory, calculated as `0x80000000 + result*4096`. In other words, if the `fib` function is about to return 1, it reads and writes to address `0x80001000`, if the `fib` function is about to return 2, it reads and writes to address `0x80002000`, and so on:

```

appStart:
    l.jal      fib              // calculate fib(15)
    l.addi     r1,r0,15         // r1 = 15 (delay slot)
    l.sys      0                // exit the application

fib:   l.sflesi  r1,1           // r1<=1? (signed)
        l.bf     done          // done if so, result is r1
        l.nop                     // (delay slot)

        l.addi     r31,r31,-12   // create stack frame
        l.sw       0(r31),r9     // save link register

```

```

        l.sw      4(r31),r1      // save input r1

        l.jal     fib           // calculate fib(N-1)
        l.addi    r1,r1,-1      // r1 = N-1 (delay slot)
        l.sw      8(r31),r1      // save fib(N-1)

        l.lwz     r1,4(r31)     // restore initial N
        l.jal     fib           // calculate fib(N-2)
        l.addi    r1,r1,-2      // r1 = N-2 (delay slot)

        l.lwz     r2,8(r31)     // restore fib(N-1)
        l.add     r1,r1,r2      // r1 = fib(N-2) + fib(N-1)

        l.lwz     r9,0(r31)     // restore link register
        l.addi    r31,r31,12    // destroy stack frame

        l.nop     98            // whats in r1?
        l.muli    r2,r1,PAGE_SIZE // r2 = fib * page size
        l.movhi   r3,0x8000    // r3 = user heap base (0x80000000)
        l.add     r2,r2,r3      // r2 = fib page
        l.lwz     r3,0(r2)     // r3 = old fib count
        l.addi    r3,r3,1      // increment count
        l.sw      0(r2),r3      // save new count

done:   l.jr      r9            // return, result in r1
        l.nop                      // (delay slot instruction)
appEnd:

```

The choice of application is completely arbitrary, and has been selected merely to cause writes to a somewhat random sequence of pages in the TLB mapped region starting at address 0x80000000 in the user address space. The subverted instruction `l.nop 98` has been used to print the return value from the `fib` function to make it easier to follow the flow of execution.

The third section of the example is the data privilege exception handler:

```

.org 0x300
////////////////////////////////////////////////////////////////////
// DATA PRIVILEGE EXCEPTION VECTOR (AT 0x300) - UPDATE TLB
////////////////////////////////////////////////////////////////////
        l.sw      -4(r31),r1      // save r1
        l.sw      -8(r31),r2      // save r2
        l.sw      -12(r31),r3     // save r3
        l.sw      -16(r31),r4     // save r4
        l.sw      -20(r31),r5     // save r5
        l.sw      -24(r31),r6     // save r6
        l.sw      -28(r31),r7     // save r7

        l.addi    r30,r30,1      // increment TLB miss count

        l.mfspr   r2,r0,0x30     // r2 = eear
        l.addi    r1,r0,-PAGE_SIZE // r1 = page mask
        l.and     r2,r2,r1      // mask faulting va to page size
        l.movhi   r4,0x0001      // r4 = 0x10000 (page table address)
        l.lwz     r5,0(r4)       // r5 = table size
        l.addi    r7,r4,8        // r7 = current page table entry
        l.addi    r6,r0,0        // r6 = 0 (current entry index)

try:   l.sfeq     r6,r5          // last entry?
        l.bf     miss          // go if so (a miss)

```

```

        l.nop                // (delay slot instruction)
        l.lwz                r1,0(r7)        // r1 = current entry va
        l.sfeq               r1,r2          // does va match entry?
        l.bf                 hit            // go if so (a hit)
        l.nop                // (delay slot instruction)
        l.addi               r7,r7,4        // r7 = next page table entry
        l.j                  try           // try next entry
        l.addi               r6,r6,1        // r6 = next entry index (delay slot)

miss:   l.sw                 0(r7),r2        // save va into page table
        l.addi               r5,r5,1        // increment table size
        l.sw                 0(r4),r5       // save new table size

hit:    l.muli               r3,r6,PAGE_SIZE // r3 = current entry index * page size
        l.movhi              r5,0x1000     // r5 = heap base (0x10000000)
        l.add                r3,r3,r5      // r3 = pa
        l.lwz                r1,4(r4)      // r1 = tlbIndex
        l.nop                99            // set TLB entry
        l.addi               r1,r1,1        // increment tlbIndex
        l.andi               r1,r1,TLB_SIZE-1 // clip to TLB size
        l.sw                 4(r4),r1      // save tlbIndex for next call

        l.lwz                r1,-4(r31)    // restore r1
        l.lwz                r2,-8(r31)    // restore r2
        l.lwz                r3,-12(r31)   // restore r3
        l.lwz                r4,-16(r31)   // restore r4
        l.lwz                r5,-20(r31)   // restore r5
        l.lwz                r6,-24(r31)   // restore r6
        l.lwz                r7,-28(r31)   // restore r7
        l.rfe                // return from exception

```

The handler does the following:

1. it gets the fault address for the read or write;
2. it determines whether there is already a mapping for that fault address in the page table at address 0x10000 in supervisor memory - if not, it allocates a new page from the heap starting at 0x10000000 in supervisor memory, and creates a new page table entry mapping the virtual page address to the new page;
3. finally, it executes `l.nop 99` to update the processor simulated TLB entry.

The entry index to update is generated in a round-robin fashion. The fault handler keeps count of the number of times it has been called in `r30`. Examining the test output, we see that the first four page accesses in the TLB allocated region fill entries 0, 1, 2 and 3 of the TLB in sequence:

```

        fib returns 1
369: CREATE entry 0 mapping (va:0x80001000 pa:0x10000000)
        fib returns 2
439: CREATE entry 1 mapping (va:0x80002000 pa:0x10001000)
        fib returns 1
        fib returns 3
549: CREATE entry 2 mapping (va:0x80003000 pa:0x10002000)
        fib returns 1
        fib returns 2
        fib returns 5
699: CREATE entry 3 mapping (va:0x80005000 pa:0x10003000)

```

After that, every TLB miss must first cause an existing mapping to be deleted before the new mapping is established:

```

    fib returns 1
    fib returns 2
    fib returns 1
    fib returns 3
    fib returns 8
919: DELETE entry 0 mapping (va:0x80001000 pa:0x10000000)
919: CREATE entry 0 mapping (va:0x80008000 pa:0x10004000)
    fib returns 1
1021: DELETE entry 1 mapping (va:0x80002000 pa:0x10001000)
1021: CREATE entry 1 mapping (va:0x80001000 pa:0x10000000)
    fib returns 2
1092: DELETE entry 2 mapping (va:0x80003000 pa:0x10002000)
1092: CREATE entry 2 mapping (va:0x80002000 pa:0x10001000)
. . . etc . . .

```

At the end of simulation, register `r30` holds the number of TLB misses – 0x1f3 (499):

```

-----
R0 : 00000000  R1 : 00000262  R2 : 80262000  R3 : 00000001
R4 : 15000000  R5 : deadbeef  R6 : deadbeef  R7 : deadbeef
R8 : deadbeef  R9 : 00000008  R10: deadbeef  R11: deadbeef
R12: deadbeef  R13: deadbeef  R14: deadbeef  R15: deadbeef
R16: deadbeef  R17: deadbeef  R18: deadbeef  R19: deadbeef
R20: deadbeef  R21: deadbeef  R22: deadbeef  R23: deadbeef
R24: deadbeef  R25: deadbeef  R26: deadbeef  R27: deadbeef
R28: deadbeef  R29: deadbeef  R30: 000001f3  R31: 00000000
PC : 00000c04  SR : 00008201  ESR: 00008200  EPC: 0000000c
TCR: 00000000  TMR: 00000000  PSR: 00000000  PMR: 00000000
BF:1 CF:0 OF:0
-----

```

19.8 TLB Modeling with Multiple Processor Modes

In this example, only the user mode address space was TLB mapped. In the general case, more than one address space may be TLB mapped. This introduces an extra level of complexity because memory mappings need to be maintained in all TLB mapped address spaces.

There are two ways to address this problem, described in the following subsections.

19.8.1 Apply Changes in All TLB-Mapped Domains

One solution is to *apply all TLB changes to every TLB-mapped domain*. For example, if both kernel and user domains are TLB mapped, each call to `or1kSetTLBEntry` would require *two* calls to `vmirtAliasMemoryVM`, one to establish the mapping and protections in the kernel mode domain and one to establish them in the user mode domain. This solution is the simplest to implement initially.

19.8.2 Maintain Multiple Copies of the TLB

An alternative solution requires *maintaining multiple copies of the TLB, one for each TLB-mapped domain*, as follows:

1. All TLB changes are made only in the copy of the TLB for the current processor mode, and all calls to `vmirtAliasMemoryVM` and `vmirtUnaliasMemoryVM` are applied only to the current domain (as in the example above).
2. When the processor switches mode, all mappings from the *previously-current* copy of the TLB are replicated in the *new current* TLB, and corresponding memory mappings in the new current memory domain are established using `vmirtAliasMemoryVM` and `vmirtUnaliasMemoryVM`.

Maintaining multiple copies of the TLB is often faster than the simpler approach of applying changes to all TLB-mapped domains, but requires careful coding to avoid obscure bugs.

19.9 TLB Modeling with ASID-Labeled Entries

It is common hardware practice to label TLB entries with *address-space identifiers* (ASIDs). ASIDs allow a TLB to be partitioned efficiently between several processes: the processor has a current ASID register, and only entries that match that register are considered candidates for matching.

From VMI version 2.0.20 on, ASID-mapped virtual memory pages are supported directly by the simulator. You can set the current ASID for the processor using:

```
void vmirtSetProcessorASID(vmiProcessorP processor, Uns32 ASID);
```

The ASID can be any 32-bit value. When pages are mapped using `vmirtAliasMemoryVM`, they can either be specified to be *global* mappings (in which case the mapping is valid irrespective of ASID) or *ASID-managed* mappings (in which case the mapping is valid only whether the processor ASID matches the ASID specified when the virtual page was mapped). When the ASID is modified by `vmirtSetProcessorASID`, the simulator automatically invalidates any existing ASID-managed mappings for the old ASID: no special action needs to be taken in the model.

To specify an ASID-mapped mapping, use the `isGlobal` and `ASID` arguments to `vmirtAliasMemoryVM`; for example, the following line creates a mapping valid only when the processor ASID is 34:

```
vmirtAliasMemoryVM(  
    physicalDomain, tlbDomain, pa, pa+OR1K_PAGE_SIZE-1, va, 0,  
    MEM_PRIV_RW, False, 34  
);
```

19.9.1 Managing Virtual Address Aliases with Different ASID

When processors support ASID-based mapping, it is common for their TLBs to be populated with entries that map the *same virtual address* with *different ASIDs*. For example, an operating system may always place the first executable address of a program at the same virtual address, e.g. `0x80000000`. If the processor is running an OS with four currently-running user processes, there may therefore be four distinct mappings for

virtual address `0x80000000` in the TLB, one for each user process, each with a different ASID.

As described above, the simulator automatically invalidates mappings for an *old* ASID on a processor ASID switch. It does not, however, automatically re-establish any mapping for the *new* ASID that may have been specified in the past: it is up to the model to do this. In general, it is most efficient to re-establish mappings using a *lazy* scheme, as described in the next subsection.

19.10 Lazy Mapping of TLB Entries

In the example described in this chapter, memory mappings are managed using `vmirtAliasMemoryVM` and `vmirtUnaliasMemoryVM` whenever a TLB entry is updated. However, these memory management functions are quite compute-intensive and should be used sparingly. If TLB updates establish mappings that are not used by the running program before being replaced, then the processor model will run slower than necessary because time will be wasted setting up memory mappings that are never actually used. This is very often the case for modal processors, since mappings are typically set up in *kernel* mode and used only in *user* mode.

For best performance, it is therefore best to establish virtual memory mappings on demand in a *lazy* fashion. This is done as follows:

1. Each TLB entry has an additional boolean field, `mapped`. This boolean indicates whether the TLB entry has already had a mapping established for it by `vmirtAliasMemoryVM`. Any entry with `mapped` of `False` definitely has no mapping established. Any entry with `mapped` of `True` may or may not have a current valid mapping (if it is an ASID-managed entry, a mapping may have been established but later invalidated by a processor ASID switch, as described in the previous section).
2. Each TLB entry can be in one of three states:
 - a. `valid=0, mapped=0`: the TLB entry is not valid.
 - b. `valid=1, mapped=1`: the TLB entry is valid and possibly mapped.
 - c. `valid=1, mapped=0`: the TLB entry is valid but *not* currently mapped.
3. When a TLB entry is written, if `entry->mapped` is 1 then any existing mapping for that entry must be discarded using `vmirtUnaliasMemoryVM`. If the entry is an ASID-managed entry, then the `isGlobal` and `ASID` parameters must match those specified when the mapping was created. However, a new mapping is *not* established with `vmirtAliasMemoryVM` at this point: instead, `entry->mapped` is set to 0.
4. When the lazy scheme is in use, processor model exception handler callbacks will be called for processor addresses that are *valid but not mapped*. Therefore, each exception handler needs to allow for this by establishing a mapping for a valid but unmapped address if required. Template code for the OR1K could be as follows:


```
VMI_RD_PRIV_EXCEPT_FN(orkRdPrivExceptionCB) {  
  
    orkP ork = (orkP)processor;  
  
    // if the address is present in the TLB but not currently mapped, establish  
    // the mappings  
    if(orkTLBMapRead(ork, address, bytes)) {  
  
        // here if the read address range is mapped and readable - redo the read  
        *action = VMI_LOAD_STORE_CONTINUE;  
  
    } else if(MEM_AA_IS_TRUE_ACCESS(attrs)) {  
  
        // here if a true exception  
        ork->EEAR = (Uns32)address;  
        orkTakeException(ork, ORK_EXCPT_DPF, 0);  
    }  
}
```

Code for `orkTLBMapRead` (not shown) would do the following:

1. Establish whether the address range `address:address+bytes-1` lies in a TLB-mapped page that allows read access and is valid. The matching entry may or may not be already marked as mapped (it may be an ASID-managed entry that was automatically unmapped by an ASID switch).
2. If so, create the mapping using `vmirtAliasMemoryVM`, set `entry->mapped` to 1, and return `True`.
3. If not, return `False`.

Note the behavior of function `orkRdPrivExceptionCB` when `orkTLBMapRead` returns `True`. In this case, the function sets the by-ref parameter `action` to the value `VMI_LOAD_STORE_CONTINUE`. What this does is cause the failing read to be *retried* on return from the exception callback – since a TLB mapping has been established, this read should now succeed (but see the detailed description in section 12 for cases in which the access may still fail).

Also note that the TLB mapping and subsequent read are done for both *artifact* and *non-artifact* accesses, but any exception is taken only for *non-artifact* accesses. Refer to section 12 for a detailed description of how these access types are indicated by the `attrs` parameter to the exception callback.

20 Implementing a TLB LRU Replacement Policy

Chapter 19 showed how to implement a TLB with a round-robin replacement policy. Often, other replacement policies are used. One of the most common is the LRU replacement policy, where the entry to replace is the *least-recently-used* entry. The simulator has special support to enable an LRU replacement policy to be efficiently modeled, as demonstrated in this example.

20.1 Introduction to LRU Replacement Implementation

Before looking at the details of the LRU replacement policy implementation, it is useful to understand some of the associated concepts.

To model an LRU replacement policy, it is first required to have a *state variable* that represents the current order of entries in the table. For example, suppose that the TLB contains four entries, numbered 0, 1, 2 and 3. At a particular point in time, the four entries may have been accessed (in most-to-least-recent order) in any of 4! (i.e. 24) different ways. Suppose that the current state implies the following ordering:

0 3 1 2

If there is now a read of an address mapped by TLB entry 1, that entry should be promoted to the most-recently-used (MRU) position, yielding a new current state, implying the following ordering:

1 0 3 2

In general, it is possible to construct a *transition table* for each TLB entry that, when indexed by the current state, will return the new state:

```
newState = transitionTable[currentState];
```

When an entry has to be ejected from the table, the last entry implied by the state is the least recently used and should be selected (entry 2 in this example).

To implement an LRU replacement policy on a set of TLB entries, two things are therefore required:

1. A state variable that encodes the entry ordering for the set of LRU-managed entries.
2. A set of transition tables, one for each entry, used to get the next state when the current entry is promoted to the most-recently-used position.

Given this information, the simulator is able to manage the state variable automatically at each read, write or fetch access.

20.2 The Template LRU Replacement Policy Model

A template model for the OR1K processor implementing a TLB with LRU replacement policy can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/18.or1kBehaviorTLBMRU
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/18.or1kBehaviorTLBMRU .
```

Compile the model using make:

```
cd 18.or1kBehaviorTLBMRU
make
```

The processor model is based on the previous model, with the changes listed in following sections.

20.3 File or1kStructure.h

The `tlbEntry` structure has been modified to contain a new field of type `memMRUSet` (defined in `vmiTypes.h`):

```
typedef struct tlbEntryS {
    Uns32    va;                // virtual address
    Uns32    pa;                // physical address
    Bool     valid;             // is entry valid?
    memMRUSet set;              // MRU management
} tlbEntry, *tlbEntryP;
```

The `memMRUSet` type has this definition:

```
typedef struct memMRUSetS {
    const Uns32 *nextState;     // LUT giving next state from current state
    Uns32      *currentState;  // reference to current state
} memMRUSet;
```

The `memMRUSet` structure has two entries:

1. A transition table, `nextState`. Given a current state, this is indexed to find the next state.
2. A pointer to a current state variable, `currentState`.

The `or1k` structure has a new `mruState` field, used to represent the current state of the TLB entry set:

```
typedef struct or1kS {
    . . . lines omitted . . .

    memDomainP physicalDomain; // physical domain target for TLB mappings
    memDomainP tlbDomain;      // memory domain to which TLB mappings apply
    tlbEntry   tlb[OR1K_TLB_SIZE]; // simulated TLB
```

```

    Uns32      mruState;           // TLB MRU state variable

} orlk, *orlkP;

```

20.4 File or1kVM.c

Function `or1kVMInit` has some new code to initialize the `memMRUSet` structures inside the TLB entries:

```

VMI_VMINIT_FN(or1kVMInit) {
    . . . lines omitted . . .

    Uns32 i;
    for(i=0; i<OR1K_TLB_SIZE; i++) {
        orlk->tlb[i].set.nextState      = vmirtGetMRUStateTable(OR1K_TLB_SIZE, i);
        orlk->tlb[i].set.currentState = &orlk->mruState;
    }
}

```

For each entry, two things are done:

1. The `nextState` field is initialized with a transition table returned by a call to function `vmirtGetMRUStateTable`. Given the number of entries in the TLB (`OR1K_TLB_SIZE`) and the index number of this entry (`i`), this function returns a transition table encoding state transitions when entry `i` is promoted to the most-recently-used slot.
2. The `currentState` field is initialized to point to the `mruState` field in the `OR1K` processor structure.

Function `vmirtGetMRUStateTable` can be used to obtain transition tables for any number of entries up to and including 8. The transition table it returns is of type `const Uns32 *`.

You do not have to use `vmirtGetMRUStateTable` to obtain the transition table – it is provided for convenience only, and any other transition table can be provided if desired.

In this example, we have a single set of MRU-managed entries, and therefore there is a single state variable, `mruState`, in the processor structure. Multiple independent MRU-managed sets can be modeled: simply ensure that there is a separate state variable for each set in the processor model.

Function `or1kSetTLBEntry` has been modified so that, if it is passed a `tlbIndex` equal to the number of TLB entries (`OR1K_TLB_SIZE`), it selects the least-recently-used entry to discard:

```

void or1kSetTLBEntry(orlkP orlk, Uns32 tlbIndex, Uns32 va, Uns32 pa) {
    // get the current instruction count (for messages)
    Uns64 iCount = vmirtGetICount((vmiProcessorP)orlk);

    // either use the specified TLB index or, if tlbIndex is OR1K_TLB_SIZE,
    // replace the LRU entry

```

```
if(tlbIndex==OR1K_TLB_SIZE) {
    tlbIndex = vmirtGetNthStateIndex(
        OR1K_TLB_SIZE, orlk->mruState, OR1K_TLB_SIZE-1
    );
} else {
    tlbIndex &= (OR1K_TLB_SIZE-1);
}
```

Function `vmirtGetNthStateIndex` has the following prototype (in file `vmiRt.h`):

```
Uns8 vmirtGetNthStateIndex(Uns32 numEntries, Uns32 state, Uns32 position);
```

Given the number of entries in a transition table, a current state and a *position*, this function returns the entry at the passed position for that state. A position of 0 implies the most-recently-used entry and a position of `numEntries-1` implies the least-recently-used entry (the function can also return the entry at any intermediate position between the most and least recently used, though this is seldom useful).

In this example, we provide the current MRU state from the processor structure and request the least-recently-used entry (position `OR1K_TLB_SIZE-1`).

`vmirtGetNthStateIndex` works only with states managed by transition tables returned by `vmirtGetMRUStateTable`. If you use custom state tables, you will need to derive the least-recently-used state yourself.

In order to tell the simulator that memory accesses to a particular dynamic-mapped memory region should update an MRU state variable, there is one further change in function `orlkSetTLBEntry`:

```
vmirtAliasMemoryVM(
    physicalDomain, tlbDomain, pa, pa+OR1K_PAGE_SIZE-1, va,
    &orlk->tlb[tlbIndex].set, MEM_PRIV_RW, True, 0
);
```

The `mruSet` argument to `vmirtAliasMemoryVM` is a pointer to a `memMRUSet` structure – in this case, we select the structure in the current TLB entry.

20.5 Testing the LRU Replacement Policy Model

To test the TLB behavior, generate the assembler test case in directory `18.orlkBehaviorTLBMRU` which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```

    fib returns 1
369: CREATE entry 0 mapping (va:0x80001000 pa:0x10000000)
    fib returns 2
436: CREATE entry 1 mapping (va:0x80002000 pa:0x10001000)
    fib returns 1
    fib returns 3
543: CREATE entry 2 mapping (va:0x80003000 pa:0x10002000)
    fib returns 1
    fib returns 2
    fib returns 5
690: CREATE entry 3 mapping (va:0x80005000 pa:0x10003000)
    fib returns 1
    fib returns 2
    fib returns 1
    fib returns 3
    fib returns 8
907: DELETE entry 3 mapping (va:0x80005000 pa:0x10003000)
907: CREATE entry 3 mapping (va:0x80008000 pa:0x10004000)
    fib returns 1
    fib returns 2
    fib returns 1
    fib returns 3
    fib returns 1
    fib returns 2
    fib returns 5
1192: DELETE entry 3 mapping (va:0x80008000 pa:0x10004000)
1192: CREATE entry 3 mapping (va:0x80005000 pa:0x10003000)
. . . lines omitted . . .
    fib returns 233
53940: DELETE entry 1 mapping (va:0x80005000 pa:0x10003000)
53940: CREATE entry 1 mapping (va:0x8000e9000 pa:0x1000b000)
    fib returns 610
54118: DELETE entry 3 mapping (va:0x8000d000 pa:0x10005000)
54118: CREATE entry 3 mapping (va:0x80262000 pa:0x1000d000)
Processor 'cpul' terminated at 'exit', address 0xc00

```

```

-----
R0 : 00000000   R1 : 00000262   R2 : 80262000   R3 : 00000001
R4 : 15000000   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : 00000008   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: 00000156   R31: 00000000
PC  : 00000c04   SR  : 00008201   ESR: 00008200   EPC: 0000000c
TCR: 00000000   TMR: 00000000   PSR: 00000000   PMR: 00000000
BF:1 CF:0 OF:0
-----

```

processor has executed 54133 instructions

The source code for this example, in file application/asmtest.S, is identical to the previous example, except that TLB_SIZE is always used as the entry argument to the l.nop 99 instruction to indicate that the LRU entry should be replaced:

```

hit:    l.muli      r3,r6,PAGE_SIZE // r3 = current entry index * page size
        l.movhi     r5,0x1000      // r5 = heap base (0x10000000)
        l.add       r3,r3,r5       // r3 = pa
        l.addi      r1,r0,TLB_SIZE // r1 = TLB_SIZE (i.e. replace LRU entry)
        l.nop       99             // set TLB entry

```

When the application runs, it proceeds in a similar way to the run in chapter 19 until the point just after the `fib(8)` result; at this point, entry 3 is discarded (the `fib(5)` result entry) since this is the least-recently-used entry.

At the end of simulation, register `r30` holds the number of TLB misses – 0x156 (342):

R0 : 00000000	R1 : 00000262	R2 : 80262000	R3 : 00000001
R4 : 15000000	R5 : deadbeef	R6 : deadbeef	R7 : deadbeef
R8 : deadbeef	R9 : 00000008	R10: deadbeef	R11: deadbeef
R12: deadbeef	R13: deadbeef	R14: deadbeef	R15: deadbeef
R16: deadbeef	R17: deadbeef	R18: deadbeef	R19: deadbeef
R20: deadbeef	R21: deadbeef	R22: deadbeef	R23: deadbeef
R24: deadbeef	R25: deadbeef	R26: deadbeef	R27: deadbeef
R28: deadbeef	R29: deadbeef	R30: 00000156	R31: 00000000
PC : 00000c04	SR : 00008201	ESR: 00008200	EPC: 0000000c
TCR: 00000000	TMR: 00000000	PSR: 00000000	PMR: 00000000
BF:1 CF:0 OF:0			

Note that the number of misses in this example (342) is less than the number in the previous example (499), implying that the LRU replacement policy is giving some benefit over a simple round-robin policy.

21 Implementing QuantumLeap-Compatible Models

As of VMI version 6.0.0, Imperas Professional Simulation products implement a parallel simulation algorithm called *QuantumLeap*, which enables multicore platform simulation to be distributed over separate threads on multiple cores of the host machine for improved performance.

In order for processor models to run correctly under QuantumLeap, some care must be taken to indicate to the simulator instructions that need to be executed *atomically* or which access *shared state*. This chapter explains the changes required.

21.1 Introduction to Multiprocessor Simulation

The *OVPsim and Imperas CpuManager User Guide* describes in detail how multiprocessor platforms may be constructed and simulated using the ICM API function `icmSimulatePlatform`, and how QuantumLeap parallel simulation can be enabled. Refer to the chapter titled *Multiprocessor Support* in that document before reading further.

21.2 QuantumLeap Requirements

The QuantumLeap parallel simulation algorithm accelerates multiprocessor simulation by distributing execution of cores in the platform over threads on multiple cores of the host machine. This means that simulation of multiple cores is performed in parallel, increasing performance. Unless informed otherwise, the simulator kernel assumes that instructions can execute in parallel, except in the following cases:

1. The instruction makes an embedded call to a *synchronizing* function from the VMI interface (for example, see the *VMI Run Time Function Reference* manual for a description of which functions in that API are synchronizing).
2. The instruction activates a memory read or write callback function.
3. The instruction is intercepted (for example, by a Semihost library).

In any of the cases described above, the simulator will assume that synchronous execution is required. It will ensure that all other parallel processor threads are stopped before allowing the current processor thread to proceed, thereby guaranteeing that the current processor sees a consistent system state that cannot be asynchronously modified by execution of another processor. The simulator also enforces synchronous execution during code morphing and intercepted function calls.

The algorithm automatically ensures safe, deterministic parallel simulation in the vast majority of cases. However, there are three scenarios that need to be handled explicitly in the processor model:

1. Identification of test-and-set or swap atomic instructions.
2. Identification of load/store exclusive constructs.
3. Identification of instructions that access shared register state.

These three cases are covered in the next sections.

21.3 Test-and-Set or Swap Atomic Instructions

Traditionally, instruction sets have implemented *test-and-set* or *swap* instructions that enable a memory location to be read and updated in a single atomic instruction. Using the basic VMI morph-time API, a swap instruction could be described like this:

```
vmimtLoadRRO(32, 32, c, CPUX_TMP, CPUX_REG(ra), MEM_ENDIAN_BIG, True, True);  
vmimtStoreRRO(32, c, CPUX_REG(ra), CPUX_REG(rd), MEM_ENDIAN_BIG, True);  
vmimtMoveRR(CPUX_GBITS, CPUX_REG(rd), CPUX_TMP);
```

This sequence emits code that first loads the contents of a memory location into a processor temporary, then stores to the same location from a register, and finally moves the temporary to the same register. The overall effect is to swap the register and memory location.

The above emitted code sequence works correctly using the standard, single-threaded multiprocessor simulation algorithm. However, when QuantumLeap is enabled, there is a chance that another asynchronously-executing processor could modify the memory location between execution of the load and the store, leading to incorrect behavior (assuming the instruction is atomic in the real hardware). To prevent this occurring, the current instruction should be identified as atomic, using the `vmimtAtomic` function, as follows:

```
vmimtAtomic();  
vmimtLoadRRO(32, 32, c, CPUX_TMP, CPUX_REG(ra), MEM_ENDIAN_BIG, True, True);  
vmimtStoreRRO(32, c, CPUX_REG(ra), CPUX_REG(rd), MEM_ENDIAN_BIG, True);  
vmimtMoveRR(CPUX_GBITS, CPUX_REG(rd), CPUX_TMP);
```

Function `vmimtAtomic` indicates to the simulator that *all other processors in a multiprocessor simulation must be stopped while this instruction executes*. This ensures that the memory location content cannot be updated by an asynchronously-executing processor between the load and store.

Function `vmimtAtomic` can be called at any point in the emission of the current instruction. We could, for example, have inserted the call after the final `vmimtMoveRR` and achieved the same effect.

21.4 Load/Store Exclusive Constructs

Traditional *test-and-set* or *swap* instructions and now being replaced in more modern instruction sets with *load/store exclusive* blocks, because these scale better to highly-parallel systems. Load/store exclusive synchronization is typically implemented using a pair of instructions: an initial *load exclusive*, which loads from an address and sets up a monitor that detects any writes by other processors to that address, and a subsequent *store exclusive*, which commits a write to the address only if the monitor has not detected a write to the same address by another processor in the interim. Many features of load/store exclusive instructions are typically implementation-dependent (for example, the granularity of the exclusive region and the conditions that can cause an exclusive store to fail) but this is not significant here.

In following subsections, we will first describe how the basic load/store exclusive construct should be implemented assuming QuantumLeap is not in use, and then describe the changes required to support QuantumLeap.

21.4.1 Describing the Load Exclusive Instruction

The load exclusive instruction is typically implemented using a load which additionally sets a load-exclusive-active flag and records the load-exclusive address, like this:

```
// load from address in ra
vmimtLoadRRO(
    32, 32, 0, CPUX_REG(rd), CPUX_REG(ra), MEM_ENDIAN_BIG, True, True
);

// initiate load/store exclusive is active
vmimtMoveRC(8, CPUX_LDREX_FLAG, 1);

// save load/store exclusive address
vmimtMoveRR(32, CPUX_LDREX_ADDRESS, CPUX_REG(ra));
```

(In this example, CPUX_LDREX_FLAG and CPUX_LDREX_ADDRESS are assumed to be architectural registers that indicate whether a load/store exclusive is active and the load/store exclusive address, respectively.)

21.4.2 Describing the Store Exclusive Instruction

The store exclusive instruction typically does the following:

1. Validates that the load-exclusive-active flag is set (otherwise the store is skipped);
2. Validates that the load-exclusive address matches (otherwise the store is skipped);
3. Performs the store;
4. Clears the load-exclusive-active flag.

A simple implementation could therefore be:

```
vmiLabelP done = vmimtNewLabel();

// skip store if load/store exclusive is not active
vmimtCompareRCJumpLabel(8, vmi_COND_NE, CPUX_LDREX_FLAG, 1, done);

// skip store if load/store address does not match
vmimtCompareRR(32, vmi_COND_NE, CPUX_LDREX_ADDRESS, CPUX_REG(ra), CPUX_TMP);
vmimtCondJumpLabel(CPUX_TMP, True, done);

// store to address in ra
vmimtStoreRRO(32, 0, CPUX_REG(ra), CPUX_REG(rd), MEM_ENDIAN_BIG, True);

// jump to here if store should be skipped
vmimtInsertLabel(done);

// terminate load/store exclusive
vmimtMoveRC(8, CPUX_LDREX_FLAG, 0);
```

21.4.3 Handling the Address Monitor

The basic load/store exclusive instructions have now been implemented to the level required for single-processor simulation. The implementation is not yet sufficient for a

multiprocessor simulation, however, because there is no monitor installed on the load/store exclusive address to detect writes to that address *by other processors*.

When QuantumLeap is not active, there is no chance that another processor could write to the load/store exclusive address while the current processor is running because the simulation is single-threaded: if this processor is running, then all others must be suspended. However, if a load/store exclusive is in force when a processor reaches the end of its time-slice then it *is* possible for another processor to store to the monitored address *while this processor is suspended awaiting its next time slice*. This can be efficiently handled as follows:

1. If a load/store exclusive is in force when the processor reaches the end of its time slice, *install a memory callback on the load/store exclusive address to detect writes to that address by other processors*.
2. When a processor starts a new time slice, *remove any previously installed memory callback before resuming simulation*.

Here is a typical implementation of this algorithm:

```
//  
// Callback to abort load/store exclusive on a conflicting write by another  
// processor  
//  
static VMI_MEM_WATCH_FN(abortEA) {  
    if(processor) {  
        cpuxP cpux = (cpuxP)userData;  
        cpux->ldrexActive = False;  
        updateExclusiveAccessCallback(cpux, False);  
    }  
}  
  
//  
// Install or remove the exclusive access monitor callback  
//  
static void updateExclusiveAccessCallback(cpuxP cpux, Bool install) {  
  
    memDomainP domain = vmirtGetProcessorDataDomain((vmiProcessorP)cpux);  
    Uns32      simLow  = cpux->ldrexAddress;  
    Uns32      simHigh = simLow+3;  
  
    // install or remove a watchpoint on the current exclusive access address  
    if(install) {  
        vmirtAddWriteCallback(domain, 0, simLow, simHigh, abortEA, cpux);  
    } else {  
        vmirtRemoveWriteCallback(domain, 0, simLow, simHigh, abortEA, cpux);  
    }  
}  
  
//  
// This is called on simulator context switch (when this processor is either  
// about to start or about to stop simulation)  
//  
VMI_IASSWITCH_FN(cpuxContextSwitchCB) {  
    cpuxP cpux = (cpuxP)processor;  
    if(cpux->ldrexActive) {  
        updateExclusiveAccessCallback(cpux, (state==RS_SUSPEND));  
    }  
}
```

In the algorithm, field `ldrexActive` corresponds to VMI register `CPUX_LDREX_FLAG`, and field `ldrexAddress` corresponds to VMI register `CPUX_LDREX_ADDRESS`.

The context-switch function `cpuxContextSwitchCB` is defined using the macro `VMI_IASSWITCH_FN`, define in file `vmiAttrs.h`:

```
#define VMI_IASSWITCH_FN(_NAME) void _NAME( \
    vmiProcessorP processor, \
    vmiIASRunState state \
)
```

Parameter `state` indicates the new state of the processor:

```
typedef enum vmiIASRunStateE {
    RS_RUN,           // processor about to be run
    RS_SUSPEND        // processor has just been suspended
} vmiIASRunState;
```

When a processor is about to start its time slice, the context switch function is called with a state of `RS_RUN`; when it has completed its time slice, the context switch function is called with a state of `RS_SUSPEND`. The context-switch callback needs to be specified in the processor model attributes structure using the `switchCB` field:

```
const vmiIASAttr modelAttrs = {
    ... fields omitted ...

    .switchCB = cpuxContextSwitchCB,

    ... fields omitted ...
};
```

21.4.4 Load/Store Exclusive with QuantumLeap

The preceding subsections describe how to implement load/store exclusive instructions in the absence of QuantumLeap simulation. When QuantumLeap simulation is enabled, the algorithm as described is no longer sufficient because it is no longer true that other processors are stopped while the load/store exclusive block is active.

To make the algorithm compatible with QuantumLeap, the only change required is *to indicate that the initial load exclusive instruction is atomic*: the simulator is then able to correctly detect writes to an exclusive block by other processors. The load-exclusive instruction therefore should be changed like this:

```
// indicate load-exclusive is atomic
vmimtAtomic();

// load from address in ra
vmimtLoadRRO(
    32, 32, 0, CPUX_REG(rd), CPUX_REG(ra), MEM_ENDIAN_BIG, True, True
);

// initiate load/store exclusive is active
```

```
vmimtMoveRC(8, CPUX_LDREX_FLAG, 1);  
  
// save load/store exclusive address  
vmimtMoveRR(32, CPUX_LDREX_ADDRESS, CPUX_REG(ra));
```

21.5 Accessing Shared Register State

Apart from the explicit synchronization instructions described above, the only other area in which care needs to be taken when making models compatible with QuantumLeap is in *accesses to shared register state*. It is sometimes the case (particularly in multicore models) that a particular register is accessible to more than one core. If such a shared register is accessed in a read-modify-write fashion, or is updated non-atomically by several VMI morph-time calls in a single simulated instruction, then use `vmimtAtomic` to ensure that all other processors are stopped while the updates occur to prevent them from seeing invalid interim register state. Note that there is generally no need to use `vmimtAtomic` if the shared register is written or read in its entirety: such accesses will be atomic.

Take special care to ensure that embedded calls are not accessing dynamically-changing shared state in an uncontrolled way. If an embedded call needs to access such state, use `vmimtAtomic` together with `vmimtCall/vmimtCallResult` to ensure other processors are stopped while the call takes place. As mentioned above, many calls in the VMI API are *synchronizing*: that is, in a QuantumLeap simulation they will cause the current thread to suspend until all other asynchronously-executing threads have been safely stopped. This often means that embedded calls that access shared state are in fact implicitly synchronizing and do not need to be explicitly identified as such with `vmimtAtomic`.

21.6 Enabling QuantumLeap in a processor model

When the requirements described above have been met, the model must notify the simulator that it can support parallel simulation. To do this, set to `True` the `QLQualified` field in the `vmiProcessorInfo` structure for the model (see section 4.2.6.7). Note that unless this field is set, the simulator will not run in parallel mode.

22 Function Address Semihosting

Semihosting allows behavior that would normally occur on a simulated system to be implemented using features of the host system instead. As a simple example, a real platform might contain a UART peripheral to receive output. When simulating this system, it is generally more convenient not to simulate the UART at all but instead to intercept any `write` call that a processor makes and redirect the output to the simulator log instead.

This section will describe a semihosting support library for the OR1K processor when used with the popular Newlib library.

22.1 Interception

Semihosting is based on a more fundamental concept: *interception*. Using Imperas technology, it is possible to define *intercept libraries*, which are loadable shared objects (on Linux¹¹) or dynamic linked libraries (on Windows).

The intercept library can specify alternate behavior for a particular instruction type (for example a `TRAP` or `SYS` instruction), or when execution reaches a particular address (for example, an interrupt vector address), or when a particular function is executed (for example, a call to `write`).

Interception requires no application image file modification or special application compilation modes (except that, for *function address interception*, the application must be compiled with function symbols present). There can be several intercept libraries available for use with a processor (for example, there might be a Newlib semihosting intercept library and a uClibc semihosting intercept library). It is even possible to cascade multiple intercept libraries for a single processor¹².

Intercept libraries exist entirely separately from the processor model. This means that you do not need access to the processor model source to create a new intercept library. To access the processor model registers, the intercept library uses the Debug interface, also used to support the gdb RSP (see chapter 17). This of course implies that the Debug interface must be implemented as a prerequisite before an intercept library can be created.

22.2 The Template Semihosting Library

A template model for the OR1K Newlib semihosting intercept library can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/19.or1kSemiHosting
```

Take a copy of the template model:

¹¹ Imperas Professional tools are available on both Linux and Windows operating systems and fully support the OVP APIs.

¹² This feature is available only in the Imperas Professional tools.

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/19.or1kSemiHosting .
```

Compile the model using make:

```
cd 19.or1kSemiHosting
make
```

The processor model is the same as that described in example 15.or1kDebugSupport, with one small change, described in section 22.4.

There is a new directory, `semihosting`, which contains the source file for the Newlib `semihosting` intercept library (`or1kNewlib.c`). This is compiled to a Linux shared library, `or1kNewlib.so`, or Windows dll, `or1kNewlib.dll`. File `or1kNewlib.c` is described in the next section.

22.3 File *semihosting/or1kNewlib.c*

Every intercept library has an object of type `vmiosAttr` describing the interceptions it performs. The structure type is defined in `vmiosAttrs.h`:

```
typedef struct vmiosAttrS {
    ///////////////////////////////////////////////////////////////////
    // VERSION
    ///////////////////////////////////////////////////////////////////

    const char      *versionString;    // version string (THIS MUST BE FIRST)
    vmioModelType    modelType;         // Type of model (enum)
    const char      *packageName;      // package name
    Uns32           objectSize;        // size in bytes of VMIO object

    ///////////////////////////////////////////////////////////////////
    // CONSTRUCTOR/DESTRUCTOR ROUTINES
    ///////////////////////////////////////////////////////////////////

    vmiosConstructorFn constructorCB;   // constructor
    vmiosDestructorFn  destructorCB;    // destructor

    ///////////////////////////////////////////////////////////////////
    // INSTRUCTION INTERCEPT ROUTINES
    ///////////////////////////////////////////////////////////////////

    vmiosMorphFn       morphCB;         // morph override callback
    vmiosNextPCFn       nextPCCB;       // get next instruction address
    vmiosDisassFn       disCB;          // disassemble instruction

    ///////////////////////////////////////////////////////////////////
    // FORMAL PARAMETERS iterators to find parameters accepted by this model
    ///////////////////////////////////////////////////////////////////

    vmiosParamSpecFn    paramSpecsCB;   // callback for next formal param

    ///////////////////////////////////////////////////////////////////
    // ADDRESS INTERCEPT DEFINITIONS
    ///////////////////////////////////////////////////////////////////

    vmiosInterceptDesc intercepts[];    // null-terminated intercept list
} vmiosAttr;
```

22.3.1 OR1K Newlib Semihosting vmiosAttr Definition

For the OR1K Newlib semihosting intercept library, the vmiosAttr structure instance is as follows:

```
vmiosAttr modelAttrs = {

    //////////////////////////////////////
    // VERSION
    //////////////////////////////////////

    .versionString = VMI_VERSION,          // version string
    .modelType     = VMI_INTERCEPT_LIBRARY, // type
    .packageName  = "Newlib",              // description
    .objectSize    = sizeof(vmiosObject),  // size in bytes of OSS object

    //////////////////////////////////////
    // CONSTRUCTOR/DESTRUCTOR ROUTINES
    //////////////////////////////////////

    .constructorCB = constructor,          // object constructor
    .destructorCB  = destructor,           // object destructor

    //////////////////////////////////////
    // ADDRESS INTERCEPT DEFINITIONS
    //////////////////////////////////////

    .intercepts    =
    {
        // -----
        // Name      Address  Attributes  Callback
        // -----
        { "_close",  0,       OSIA_OPAQUE, closeInt    },
        { "_exit",   0,       OSIA_OPAQUE, exitInt     },
        { "_fstat",  0,       OSIA_OPAQUE, fstatInt    },
        { "_gettimeofday", 0,   OSIA_OPAQUE, gettimeofdayInt },
        { "_ioctl",  0,       OSIA_OPAQUE, ioctlInt    },
        { "_lseek",  0,       OSIA_OPAQUE, lseekInt    },
        { "_lstat",  0,       OSIA_OPAQUE, lstatInt    },
        { "_open",   0,       OSIA_OPAQUE, openInt     },
        { "_read",   0,       OSIA_OPAQUE, readInt     },
        { "_stat",   0,       OSIA_OPAQUE, statInt     },
        { "_time",   0,       OSIA_OPAQUE, timeInt     },
        { "_times",  0,       OSIA_OPAQUE, timesInt    },
        { "_unlink", 0,       OSIA_OPAQUE, unlinkInt   },
        { "_write",  0,       OSIA_OPAQUE, writeInt    },
        { 0,         0,       0,         0             },
    }
};
```

In detail, the sections of this file are described below.

```
.versionString = VMI_VERSION,          // version string
.modelType     = VMI_INTERCEPT_LIBRARY, // type
```

Each intercept library contains a reference to the current VMI version (from `vmiVersion.h`) and the type of model (in `vmiTypes.h`) so that the simulator can verify interface compatibility.


```
.packageName    = "Newlib",           // description
```

This is a descriptive name (used for message reporting only).

```
.objectSize     = sizeof(vmiosObject), // size in bytes of OSS object
```

The intercept library defines a custom structure, used to hold all data required for an instance of that library; this defines the size of that structure so that it can be automatically allocated by the simulator when the library is instantiated. For the Newlib semihosting intercept library, the structure is defined like this:

```
#define FILE_DES_NUM 128
#define REG_ARG_NUM  3

typedef struct vmiosObjectS {

    // first few argument registers (standard ABI)
    vmiRegInfoCP args[REG_ARG_NUM];

    // return register (standard ABI)
    vmiRegInfoCP result;

    // stack pointer (standard ABI)
    vmiRegInfoCP sp;

    // __impure_ptr address and domain
    Addr         impurePtrAddr;
    memDomainP   impurePtrDomain;

    // file descriptor table
    Int32 fileDescriptors[FILE_DES_NUM];

} vmiosObject;
```

(The fields of this structure will be covered in more detail on following sections.)

Next, the `vmiosAttr` structure contains references to *constructor* and *destructor* functions, called when an instance of the intercept library is created and destroyed, respectively:

```
.constructorCB = constructor,           // object constructor
.destructorCB  = destructor,           // object destructor
```

These functions will be covered in detail in a later subsection.

Finally, the `vmiosAttr` structure contains a list of *address intercept* definitions:

```
.intercepts    =
{
    // -----
    // Name      Address    Attributes    Callback
    // -----
    { "_close",  0,        OSIA_OPAQUE, closeInt    },
    { "__exit",  0,        OSIA_OPAQUE, exitInt     },
}
```

```

    { "_fstat",          0,          OSIA_OPAQUE, fstatInt      },
    { "_gettimeofday",  0,          OSIA_OPAQUE, gettimeofdayInt },
    { "_ioctl",         0,          OSIA_OPAQUE, ioctlInt       },
    { "_lseek",         0,          OSIA_OPAQUE, lseekInt        },
    { "_lstat",         0,          OSIA_OPAQUE, lstatInt         },
    { "_open",          0,          OSIA_OPAQUE, openInt          },
    { "_read",          0,          OSIA_OPAQUE, readInt           },
    { "_stat",          0,          OSIA_OPAQUE, statInt           },
    { "_time",          0,          OSIA_OPAQUE, timeInt            },
    { "_times",         0,          OSIA_OPAQUE, timesInt           },
    { "_unlink",        0,          OSIA_OPAQUE, unlinkInt          },
    { "_write",         0,          OSIA_OPAQUE, writeInt           },
    { 0,                0,          0,                0              },
}

```

Each entry in this null-terminated table of intercept definitions is of type `vmiosInterceptDesc`:

```

typedef enum vmiosInterceptAttrE {
    OSIA_NONE      = 0x0,          // no special attributes
    OSIA_OPAQUE    = 0x1,          // opaque intercept (otherwise transparent)
    OSIA_THREAD    = 0x2,          // run in thread (otherwise synchronous)
} vmiosInterceptAttr;

typedef struct vmiosInterceptDescS {
    const char      *name;          // for interception by name
    Addr            simAddress;     // for interception by address
    vmiosInterceptAttr attrs;       // intercepted function attributes
    vmiosInterceptFn interceptCB;   // interception callback
    void            *userData;      // client-specific data pointer
    Bool            skipPrologue;   // Use gdb to find the prologue
} vmiosInterceptDesc;

```

In detail, each entry has the following fields:

1. A function name: if non-NULL, this field specifies the *name of a function* in the application executable to which this row applies. For example, the first row in the Newlib semihosting intercept library applies to function `_close`.
2. A function address: if non-zero, this field specifies an *address* within the application executable to which this row applies. This field isn't used in this example, but is useful when interception of a known address, such as an exception handler address, is required.
3. A bitfield enumeration, `attrs`, comprised of bitwise-or of the following members:
 - `OSIA_OPAQUE`: this indicates whether the action performed by this intercept should *replace* any default behavior specified by the processor model (if present) or be performed *in addition* to the default behavior specified by the processor model (if absent). Semihosting libraries in general specify replacement behaviors, so this field is usually present.
 - `OSIA_THREAD`: this option is available only with Imperas Professional products. It indicates that the intercepted function should be run in a separate hardware thread, improving simulator performance.
4. An *intercept function*, which specifies the new behavior for the intercepted address. For example, the first row in the Newlib semihosting intercept library

associates intercept function `closeInt` with the application function `_close`. Intercept functions are covered in more detail below.

5. A client-specific data pointer, passed as an argument to the intercept function. This is unused in the current example.
6. A Boolean, `skipPrologue`: if False, then the exact symbol address is intercepted; if True, then a gdb for the processor is invoked by the simulator to calculate the intercept address *after any function prologue*. This is not required for this example.

The template Newlib semihosting intercept library supplies opaque function address intercepts for each of the following functions in the application: `_close`, `__exit`, `_fstat`, `_gettimeofday`, `_ioctl`, `_lseek`, `_lstat`, `_open`, `_read`, `_stat`, `_time`, `_times`, `_unlink` and `_write`. Because these functions are all *opaquely* intercepted, whenever the simulator executes at any of these function addresses, it will perform actions specified by the corresponding intercept functions instead of the normal processor model behavior.

22.3.2 OR1K Newlib Semihosting Constructor Definition

The constructor for the OR1K Newlib semihosting intercept library has this definition:

```
static VMIO_CONSTRUCTOR_FN(constructor) {  
  
    Uns32 i;  
  
    // first few argument registers (standard ABI)  
    object->args[0] = vmiosGetRegDesc(processor, "R3");  
    object->args[1] = vmiosGetRegDesc(processor, "R4");  
    object->args[2] = vmiosGetRegDesc(processor, "R5");  
  
    // return register (standard ABI)  
    object->result = vmiosGetRegDesc(processor, "R11");  
  
    // stack pointer (standard ABI)  
    object->sp = vmiosGetRegDesc(processor, "R1");  
  
    // __impure_ptr address  
    object->impurePtrDomain = vmirtAddressLookup(  
        processor, ERRNO_REF, &object->impurePtrAddr  
    );  
  
    // initialize stdin, stderr and stdout  
    object->fileDescriptors[0] = vmiosGetStdin(processor);  
    object->fileDescriptors[1] = vmiosGetStdout(processor);  
    object->fileDescriptors[2] = vmiosGetStderr(processor);  
  
    // initialize remaining file descriptors  
    for(i=3; i<FILE_DES_NUM; i++) {  
        object->fileDescriptors[i] = -1;  
    }  
}
```

The constructor first obtains *register description* objects that enable certain named registers with the processor model to be read and written. Because we are writing function address intercepts in this library, we need to be able to access several registers used in the standard processor ABI: the first few function argument registers (R3, R4 and

R5 for the OR1K), the function result register (R11 for the OR1K) and the stack pointer (R1 for the OR1K):

```
// first few argument registers (standard ABI)
object->args[0] = vmiosGetRegDesc(processor, "R3");
object->args[1] = vmiosGetRegDesc(processor, "R4");
object->args[2] = vmiosGetRegDesc(processor, "R5");

// return register (standard ABI)
object->result = vmiosGetRegDesc(processor, "R11");

// stack pointer (standard ABI)
object->sp = vmiosGetRegDesc(processor, "R1");
```

The `vmiRegInfoCP` object returned by `vmiosGetRegDesc` is in fact a register descriptor supplied by the debug interface, created in chapter 17.

Next, to support correct error return from functions implemented in the Newlib library, the constructor obtains the address of a special symbol, `__impure_ptr`, that is always defined in any Newlib application:

```
object->impurePtrDomain = vmirtAddressLookup(
    processor, ERRNO_REF, &object->impurePtrAddr
);
```

Newlib allows reentrant calls, so that instead of having a single `errno` variable to signal library errors there is instead a pointer, `__impure_ptr`, that points to the current `errno` value to update. Hence, if we are to semihost a call that could set `errno`, we need to obtain the address in `__impure_ptr` to determine what `errno` address to write.

Finally, the constructor initializes a file descriptor map for the semihost intercept library. This maps from file numbers expected by the application to native file pointers. Files 0, 1 and 2 are the standard input, standard output and standard error, respectively:

```
object->fileDescriptors[0] = vmiosGetStdin(processor);
object->fileDescriptors[1] = vmiosGetStdout(processor);
object->fileDescriptors[2] = vmiosGetStderr(processor);
```

Other files are initially closed:

```
for(i=3; i<FILE_DES_NUM; i++) {
    object->fileDescriptors[i] = -1;
}
```

See the *VMI OS Support Function Reference* manual for more information about all functions with the `vmios` prefix.

22.3.3 OR1K Newlib Semihosting Destructor Definition

The destructor for the OR1K Newlib semihosting intercept library is currently a void function:

```
static VMIO_DESTROY_FN(destructor) {  
}
```

Typically, the destructor would be used to print out statistics gathered by the intercept library and free any temporary structures that were allocated.

22.3.4 Function Address Intercept Example: `closeInt`

To understand how to write a function address intercept callback, see function `closeInt` in `or1kNewlib.c`, which supplies alternate behavior to be performed when the application `_close` function is executed. This function should close a file descriptor passed as the only argument. It is implemented like this:

```
static VMIO_INTERCEPT_FN(closeInt) {  
  
    Int32 fd;  
  
    // obtain function arguments  
    getArg(processor, object, 0, &fd);  
  
    // implement close  
    Int32 fdMap = mapFileDescriptor(processor, object, fd);  
    Int32 result = vmiosClose(processor, fdMap);  
  
    // null out the semihosted file descriptor if success  
    if(!result) {  
        object->fileDescriptors[fd] = -1;  
    }  
  
    // return result  
    setErrnoAndResult(processor, object, result, context);  
}
```

All function address intercept callbacks should be defined using the `VMIO_INTERCEPT_FN` macro, defined in file `vmiosAttrs.h`:

```
#define VMIO_INTERCEPT_FN(_NAME) void _NAME( \  
    vmiProcessorP processor,    \  
    vmiosObjectP object,       \  
    Addr          thisPC,      \  
    const char    *context,    \  
    void          *userData,    \  
    Bool          atOpaqueIntercept \  
)  
typedef VMIO_INTERCEPT_FN(( *vmiosInterceptFn ));
```

The function address intercept callback is called at *morph time* and should use the VMI Morph Time Function API to generate code to implement required behavior. The callback is passed six arguments:

1. The processor that is about to execute code at the intercepted address;
2. The current function intercept object;
3. The intercepted address;
4. A context string, which gives the function name being intercepted (e.g. `"_close"`).

5. The client-specific data pointer associated with the row of the function address intercept table defining this interception.
6. An indication of whether the current address is already opaquely intercepted. This may be required when intercept libraries are cascaded – for example, an intercept of a function call may expect there to be a corresponding return later, but this won't be the case if the call has been opaquely intercepted already.

This function first gets the first argument (argument 0) of `_close`, using a utility function `getArg`:

```
Int32 fd;
getArg(processor, object, 0, &fd);
```

`getArg` itself simply accesses a register value with the standard OR1K ABI, using function `vmiosRegRead`, which obtains the value of the register using the Debug interface:

```
static void getArg(
    vmiProcessorP processor,
    vmiosObjectP object,
    Uns32 index,
    void *result
) {
    if(index >= REG_ARG_NUM) {

        vmiMessage("P", "OR1K_ANS_NEWLIB",
            "No more than %u function arguments supported",
            REG_ARG_NUM
        );

        vmirtFinish(-1);

    } else {

        vmiosRegRead(processor, object->args[index], result);

    }
}
```

Next, `closeInt` calls another utility function, `mapFileDescriptor`, which translates from an application file number to the equivalent native one, and closes that file using the `vmios`-prefixed function `vmiosClose`:

```
Int32 fdMap = mapFileDescriptor(processor, object, fd);
Int32 result = vmiosClose(processor, fdMap);
```

`mapFileDescriptor` is very simple, and uses the table of file descriptors in the intercept library instance object to perform the mapping:

```
static Int32 mapFileDescriptor(
    vmiProcessorP processor,
    vmiosObjectP object,
    Uns32 i
) {
    if(i >= FILE_DES_NUM) {
```

```
        return -1;
    } else {
        return object->fileDescriptors[i];
    }
}
```

A mapping table used to translate from application to native file descriptors instead of using native descriptors directly for two reasons.

Firstly, it ensures that the application sees the full range of expected file numbers 0, 1, 2, ... FILE_DES_NUM-1, irrespective of what native files are in use.

Secondly, it prevents badly-behaved simulated applications closing or otherwise modifying unrelated native files of which it should have no knowledge.

Always use a mapping table to translate from application to native file descriptors in semihosting intercept libraries.

If the file close succeeded, `closeInt` then nulls out the corresponding entry in the file descriptor table:

```
if(!result) {
    object->fileDescriptors[fd] = -1;
}
```

Finally, a utility function `setErrnoAndResult` is called. This does two things:

1. It updates the current `errno` so that it is consistent with the command just executed;
2. It assigns the result of `vmiosClose` to the appropriate result register in the OR1K ABI:

```
setErrnoAndResult(processor, object, result, context);
```

`setErrnoAndResult` is defined as:

```
static void setErrnoAndResult(
    vmiProcessorP processor,
    vmiosObjectP object,
    Int32 result,
    const char *context
) {
    if(!object->impurePtrDomain) {
        vmiMessage("P", "OR1K_ICF_NEWLIB",
            "Interception of '%s' failed - %s not found "
            "(application does not appear to be compiled with newlib "
            "or has no symbols)",
            context, ERRNO_REF
        );

        vmirtFinish(-1);
    } else if(result<0) {

        memDomainP domain = object->impurePtrDomain;
        memEndian endian = vmirtGetProcessorDataEndian(processor);
        Int32 errnoValue = -result;
```

```

Uns32      impurePtrAddr;

result = -1;

// swap errno endianness if required
if(endian != ENDIAN_NATIVE) {
    errnoValue = swap4(errnoValue);
}

// read __impure_ptr value
vmirtReadNByteDomain(
    domain, object->impurePtrAddr, &impurePtrAddr,
    sizeof(impurePtrAddr), 0, False
);

// swap errno address endianness if required
if(endian != ENDIAN_NATIVE) {
    impurePtrAddr = swap4(impurePtrAddr);
}

// write back errno
vmirtWriteNByteDomain(
    domain, impurePtrAddr+OR1K_ERRNO_OFFSET, &errnoValue,
    sizeof(errnoValue), 0, True
);
}

vmiosRegWrite(processor, object->result, &result);
}

```

`setErrnoAndResult` reads the address in the `__impure_ptr` variable using `vmirtReadNByteDomain` and stores the required `errno` value to this address using `vmirtWriteNByteDomain`. It also performs any endianness conversions required if (as for the OR1K) the simulated processor endianness differs from the host.

The template Newlib semihosting intercept library has been designed to be relatively easy to port to any new processor: changes should be limited to the constructor (where details of the processor ABI will need to be encoded) and perhaps `getArg` (if parameter values are passed on the stack instead of in registers).

22.4 File *or1kSemiHost.c*

The OR1K model contains one new file, `or1kSemiHost.c`, which implements a single function that is required to support opaque function interception: `or1kIntReturnCB`.

```

VMI_INT_RETURN_FN(or1kIntReturnCB) {
    vmimtUncondJumpReg(0, OR1K_REG(OR1K_LINK), VMI_NOREG, vmi_JH_RETURN);
}

```

This intercept return function generates code that forces a return from an opaquely-intercepted function address after one instruction. For the OR1K, this is done simply by jumping to the link address. *Note that this function does not implement any true processor behavior: it is required merely to allow opaque function intercepts to work – therefore, there should be no attempt to model true hardware features such as delay slots.*

The new function is prototyped in `orlkFunctions.h` and referenced in `orlkAttrs.h`:

```
const vmiIASAttr modelAttrs = {  
    . . . lines omitted . . .  
  
    //////////////////////////////////////  
    // IMPERAS INTERCEPTED FUNCTION SUPPORT ROUTINES  
    //////////////////////////////////////  
  
    .intReturnCB = orlkIntReturnCB,  
  
    . . . lines omitted . . .  
};
```

22.5 Platform Makefile

The Makefile entry from the platform has been modified for this example so that the OR1K Newlib intercept library is passed as the `SEMIHOST_FILE` #define value when `platform/platform.c` is compiled.

22.6 Flow of Control for Opaque Address Intercepts

When an opaquely-intercepted function is encountered by the simulator, the flow of control will appear as follows:

1. The call to the intercepted address will be as normal.
2. The simulator will appear to execute one instruction at the intercepted function address. However, this single instruction will perform the entire behavior specified by the interception library for that address.
3. After the instruction completes, the processor will immediately return to the instruction after the call to the intercepted address – in other words, the entire functionality at the intercepted address will appear to have been replaced by a single instruction.

22.7 Testing the Semihosting Intercept Library

To test the semihosting intercept library, generate the C test case in directory `19.orlkSemiHosting` which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform  
make -C application
```

Run the platform using the C program executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/application.OR1K.elf
```

The output from this should be as follows:

```
Info 1: 'cpu1', 0x0000000000000100: 1.addi    r2,r0,0x0  
Info 2: 'cpu1', 0x0000000000000104: 1.addi    r3,r0,0x0  
Info 3: 'cpu1', 0x0000000000000108: 1.addi    r4,r0,0x0  
Info 4: 'cpu1', 0x000000000000010c: 1.addi    r5,r0,0x0
```

```
Info 5: 'cpul', 0x0000000000000110: l.addi    r6,r0,0x0
. . . lines omitted . . .
Info 340: 'cpul', 0x00000000000009854: l.ori     r3,r4,0x0
Info 341: 'cpul', 0x00000000000009858: l.addi    r4,r0,0x0
Info 342: 'cpul', 0x0000000000000985c: l.sw      0x0(r10),r4
Info 343: 'cpul', 0x00000000000009860: l.jal     0x00009bc4
Info 344: 'cpul', 0x00000000000009864: l.ori     r4,r5,0x0
Info 345: 'cpul', 0x00000000000009bc4: *** INTERCEPT *** (_fstat)
Info 346: 'cpul', 0x00000000000009868: l.sfnei   r11,0xffffffff
. . . lines omitted . . .
Info 1387: 'cpul', 0x000000000000099b4: l.sw      0x0(r10),r4
Info 1388: 'cpul', 0x000000000000099b8: l.ori     r4,r5,0x0
Info 1389: 'cpul', 0x000000000000099bc: l.jal     0x00009c48
Info 1390: 'cpul', 0x000000000000099c0: l.ori     r5,r6,0x0
Info 1391: 'cpul', 0x00000000000009c48: *** INTERCEPT *** (_write)
main called
Info 1392: 'cpul', 0x000000000000099c4: l.sfnei   r11,0xffffffff
. . . lines omitted . . .
Info 2224: 'cpul', 0x00000000000001718: l.jal     0x00009b3c
Info 2225: 'cpul', 0x0000000000000171c: l.ori     r3,r10,0x0
Info 2226: 'cpul', 0x00000000000009b3c: *** INTERCEPT *** (__exit)
processor has executed 2226 instructions
```

The application starts running normally. Then, after 345 instructions, there is a call to function `_fstat`, which is intercepted. The entire behavior of the intercepted `_fstat` appears to occur in a single instruction, and at instruction 346 execution has returned from `_fstat`. At instruction 1391, there is an intercepted call to `_write`, which actually writes the string `main called` to the standard output, before returning after one instruction. There are various other intercepted calls in this run, finishing with an intercepted call to `_exit` at instruction 2226, which terminates the current processor, ending simulation.

Note that intercepted calls are automatically reported in the trace output to make it clear where behavior is deviating from the standard processor model.

23 Using Intercept Libraries for Instruction Set Enhancement

The previous chapter described how to use Imperas intercept library technology to implement semihosting libraries. When using Imperas Professional products, it is also possible to use this technology to implement enhancements to processor model instruction sets, *even when the source of the processor model is unavailable*.

Note that this feature is not available in OVPSim.

23.1 The Template Instruction Set Enhancement Library

A template model for the OR1K processor with its instruction set enhanced by an external intercept library can be found at:

```
$IMPERAS_HOME/Examples/Models/Processor/20.or1kExchange
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/20.or1kExchange .
```

Compile the model using make:

```
cd 20.or1kExchange
make
```

The processor model is the same as that described in example 19.or1kSemiHosting.

There is a new directory, `exchange`, which contains the source file for the enhanced instruction set intercept library. This is compiled to a Linux shared library, `model.so`, or Windows dll, `model.dll`. File `or1kExchange.c` is described in the next section.

23.2 File `exchange/or1kExchange.c`

The standard OR1K processor contains no single instruction that enables a register to be exchanged with a memory location. It might be reasonable to want to enhance the instruction set to provide this capability, if the OR1K is to be used in a multiprocessor platform. This example adds an exchange instruction using an intercept library, so that the core functionality of the basic model is unaffected.

The `vmiosAttr` structure implementing the additional exchange instruction is defined as follows:

```
vmiosAttr modelAttrs = {
    ///////////////////////////////////////////////////////////////////
    //  VERSION
    ///////////////////////////////////////////////////////////////////
}
```

```

.versionString = VMI_VERSION,          // version string
.modelType     = VMI_INTERCEPT_LIBRARY, // type
.packageName   = "Exchange",           // description
.objectSize    = sizeof(vmiosObject),  // size in bytes of OSS object

////////////////////////////////////
// CONSTRUCTOR/DESTRUCTOR ROUTINES
////////////////////////////////////

.constructorCB = constructor,          // object constructor

////////////////////////////////////
// INSTRUCTION INTERCEPT ROUTINES
////////////////////////////////////

.morphCB       = exchangeMorph,        // morph callback
.nextPCCB      = exchangeNextPC,       // get next instruction address
.disCB         = exchangeDisass,       // disassemble instruction

////////////////////////////////////
// ADDRESS INTERCEPT DEFINITIONS
////////////////////////////////////

.intercepts    = {{0}}
};

```

When implementing supplementary instructions in an intercept library, both a *morpher callback* and a *disassembly callback* are required. Address intercepts are generally not required, so the address intercept table for this intercept library is empty.

23.2.1 OR1K Newlib Semihosting vmiosAttr Definition

The `vmiosObject` for this intercept library is defined as follows:

```

#define OR1K_GPR_NUM 32

typedef struct vmiosObjectS {

    // handles for the OR1K GPRs
    vmiRegInfoCP regs[OR1K_GPR_NUM];

    // enhanced instruction decode table
    vmidDecodeTableP table;

} vmiosObject;

```

Field `regs` is initialized in the constructor to hold a `vmiRegInfo` handle to each of the OR1K GPRs. Field `table` is also initialized in the constructor, and holds a decode table to decode supplementary instructions implemented by this intercept library.

23.2.2 Constructor Definition

The constructor has this definition:

```

static VMIOS_CONSTRUCTOR_FN(constructor) {

    Uns32 i;

```

```
for(i=0; i<OR1K_GPR_NUM; i++) {
    char regName[8];
    sprintf(regName, "R%u", i);
    object->regs[i] = vmiosGetRegDesc(processor, regName);
}

object->table = createDecodeTable();
}
```

The constructor first obtains *register description* objects for each of the OR1K general-purpose registers:

```
for(i=0; i<OR1K_GPR_NUM; i++) {
    char regName[8];
    sprintf(regName, "R%u", i);
    object->regs[i] = vmiosGetRegDesc(processor, regName);
}
```

Then it creates a decoder table to decode extra instructions implemented by this library:

```
object->table = createDecodeTable();
```

Function `createDecodeTable` uses the standard VMI *decoder* API functions to create a new decode table for a single exchange instruction:

```
static vmidDecodeTableP createDecodeTable(void) {

    vmidDecodeTableP table = vmidNewDecodeTable(OR1K_BITS, OR1K_EIT_LAST);

    // handle exchange instruction
    DECODE_ENTRY(0, EXW, "|111101.....|");

    return table;
}
```

We have selected an arbitrary unused instruction prefix for the new instruction. Refer to chapter 5 for detailed information about the VMI decoder function API.

23.2.3 The Morpher Callback: `exchangeMorph`

This example implements a *morpher callback function* that specifies behavior for the new exchange instruction. It is defined as follows:

```
static VMIOS_MORPH_FN(exchangeMorph) {

    Uns32 instruction = vmicxtFetch4Byte(processor, thisPC);
    or1kEnhancedInstrType type = vmidDecode(object->table, instruction);

    if(type==OR1K_EIT_EXW) {
        *opaque = True;
        *userData = (void *)instruction;
        return doExchange;
    } else {
        return 0;
    }
}
```

The intercept library morpher callback is called *before* the standard processor model callback. It first decodes the instruction at the current program counter address:

```
Uns32          instruction = vmicxtFetch4Byte(processor, thisPC);
orlkenhancedInstrType type  = vmidDecode(object->table, instruction);
```

If the instruction is the new exchange instruction, the function first indicates that this is an *opaque* intercept:

```
if(type==OR1K_EIT_EXW) {
    *opaque = True;
```

What this means is that the behavior in the intercept library will *replace* any default behavior in the processor model - if `opaque` was instead set to `False`, then this would be a *transparent* intercept, and behavior specified in the intercept library would be performed *in addition* to the standard behavior in the processor model.

The function then supplies the `userData` for the callback function that will implement the exchange instruction. In this case, the `userData` is simply the instruction pattern itself:

```
*userData = (void *)instruction;
```

Finally, it returns a run-time function that implements the new instruction:

```
return doExchange;
```

If the decoded instruction is *not* the new exchange instruction, the function returns a null pointer to indicate that the standard processor model behavior should be performed.

23.2.4 The Next Instruction Callback: `exchangeNextPC`

This example also implements a *next PC callback function* that specifies the next instruction address after for the new exchange instruction. It is defined as follows:

```
//
// Return instruction address after passed program counter
//
static VMIO_NEXT_PC_FN(exchangeNextPC) {

    // decode the instruction to get the type
    Uns32          instruction = vmicxtFetch4Byte(processor, thisPC);
    orlkenhancedInstrType type  = vmidDecode(object->table, instruction);

    if(type==OR1K_EIT_EXW) {
        *nextPC = thisPC+4;
        return True;
    } else {
        return False;
    }
}
```

The callback decodes the instruction at the passed address. If it is the exchange instruction, it sets the `nextPC` byref argument to the address of the instruction following the instruction and returns `True`; otherwise, it returns `False`.

The next instruction address callback is in fact only required if *the next instruction address differs from the address that would be calculated by the base model*, which is not the case for the OR1K model (since all instructions are four bytes long). The callback could therefore have been omitted for this example (and specified as 0 in the attribute structure). It has been specified in this case for example purposes only.

23.2.5 The Disassembler Callback: `exchangeDisass`

This example also implements a *disassembler callback function* that specifies disassembly for the new exchange instruction. It is defined as follows:

```
static VMIO$DISASSEMBLE_FN(exchangeDisass) {

    Uns32          instruction = vmicxtFetch4Byte(processor, thisPC);
    orlkEnhancedInstrType type   = vmidDecode(object->table, instruction);

    if(type==OR1K_EIT_EXW) {

        static char buffer[256];

        // extract instruction fields
        Uns32 ra = OPEX_A(instruction);
        Uns32 rb = OPEX_B(instruction);
        Int16 i  = OPEX_I(instruction);

        sprintf(buffer, "%-8s 0x%x(r%u),r%u", "l.exw", i, ra, rb);

        return buffer;

    } else {

        return 0;

    }
}
```

Once again, the callback first decodes the instruction at the passed address:

```
Uns32          instruction = vmicxtFetch4Byte(processor, thisPC);
orlkEnhancedInstrType type   = vmidDecode(object->table, instruction);
```

If the instruction is the new exchange instruction, a disassembly string is created in a static buffer and returned. The general approach is just the same as for standard instruction disassembly callbacks; refer to chapter 6 for more details. We have assumed that the new instruction takes arguments in the same format as the existing `l.sw` (store word) instruction:

```
if(type==OR1K_EIT_EXW) {

    static char buffer[256];

    // extract instruction fields
```

```
Uns32 ra = OPEX_A(instruction);
Uns32 rb = OPEX_B(instruction);
Int16 i  = OPEX_I(instruction);

sprintf(buffer, "%-8s 0x%x(r%u),r%u", "l.exw", i, ra, rb);

return buffer;
```

If the new instruction is *not* the new exchange instruction, the function returns a null pointer to indicate that standard processor model disassembly should be performed.

23.2.6 The Exchange Instruction Callback: doExchange

The exchange instruction is implemented by a run-time callback function, doExchange, defined as follows:

```
static VMIO_INTERCEPT_FN(doExchange) {

    memDomainP domain = vmirtGetProcessorDataDomain(processor);
    memEndian  endian = vmirtGetProcessorDataEndian(processor);

    Uns32 instruction = (Uns32)userData;

    Uns32 ra = OPEX_A(instruction);
    Uns32 rb = OPEX_B(instruction);
    Int16 i  = OPEX_I(instruction);

    Uns32 rbOld;
    vmiosRegRead(processor, object->regs[rb], &rbOld);

    Uns32 raValue;
    vmiosRegRead(processor, object->regs[ra], &raValue);

    Uns32 rbNew;
    vmirtReadNByteDomain(domain, raValue+i, &rbNew, sizeof(rbNew), 0, True);

    if(endian != ENDIAN_NATIVE) {
        rbOld = swap4(rbOld);
        rbNew = swap4(rbNew);
    }

    vmirtWriteNByteDomain(domain, raValue+i, &rbOld, sizeof(rbOld), 0, True);

    vmiosRegWrite(processor, object->regs[rb], &rbNew);
}
```

Because this instruction is going to access processor memory, the callback first gets the processor data domain and endianness:

```
memDomainP domain = vmirtGetProcessorDataDomain(processor);
memEndian  endian = vmirtGetProcessorDataEndian(processor);
```

It then recovers the instruction pattern (passed as the `userData` argument) and extracts the value of the offset, `ra` and `rb` fields from the instruction (assuming it follows the same format as the `l.sw` instruction):

```
Uns32 instruction = (Uns32)userData;
```



```
Uns32 ra = OPEX_A(instruction);
Uns32 rb = OPEX_B(instruction);
Int16 i  = OPEX_I(instruction);
```

Next, the current value of registers `rb` and `ra` are extracted from the processor model using the standard `vmiosRegRead` accessor function, introduced in the previous chapter:

```
Uns32 rbOld;
vmiosRegRead(processor, object->regs[rb], &rbOld);

Uns32 raValue;
vmiosRegRead(processor, object->regs[ra], &raValue);
```

Then the value of the four-byte word at address `ra+i` is read:

```
Uns32 rbNew;
vmirtReadNByteDomain(domain, raValue+i, &rbNew, sizeof(rbNew), 0, True);
```

Because the memory access routines such as `vmirtReadNByteDomain` access memory in byte order, the endianness of the old and new values of the memory location are then swapped if the OR1K processor endianness differs from the native host (which is little endian):

```
if(endian != ENDIAN_NATIVE) {
    rbOld = swap4(rbOld);
    rbNew = swap4(rbNew);
}
```

Next, address `ra+i` is written with the current value of register `rb`:

```
vmirtWriteNByteDomain(domain, raValue+i, &rbOld, sizeof(rbOld), 0, True);
```

Finally, register `rb` in the processor model is updated with the value read from memory using the standard `vmiosRegWrite` accessor function:

```
vmiosRegWrite(processor, object->regs[rb], &rbNew);
```

To keep this example simple, the callback function has ignored the possibility of exceptions on the read or write accesses. A production-quality version should validate alignment and protection before making any changes to the processor state.

23.3 The Platform File, *platform/platform.c*

The platform is similar to previous examples, except that a new line is included to install the new intercept library on the processor model using function `icmAddInterceptObject`:

```
int main(int argc, char ** argv) {

    if(argc!=2) {
        icmPrintf("%s: expected application name argument\n", argv[0]);
        return -1;
    }
}
```

```

    }

    icmInitPlatform(ICM_VERSION, 0, NULL, 0, "platform");

    const char *modelFile = "model." IMPERAS_SHRSUF;
    const char *semlhostFile = icmGetVlnvString(NULL, "ovpworld.org",
                                                "modelSupport", "imperasExit", "1.0", "model");

    icmProcessorP processor = icmNewProcessor(
        "cpul",           // CPU name
        TYPE_NAME,       // CPU type
        0,               // CPU cpuId
        0,               // CPU model flags
        32,              // address bits
        modelFile,       // model file
        "modelAttrs",    // morpher attributes
        MODEL_FLAGS,     // enable debug
        0,               // user-defined attributes
        semlhostFile,    // semi-hosting file
        "modelAttrs"     // semi-hosting attributes
    );

    // add intercept library for exchange instruction
    icmAddInterceptObject(
        processor, "exchange", "exchange/model." IMPERAS_SHRSUF, "modelAttrs", 0);
    );

    . . . etc . . .
}

```

Note that the new intercept library is installed *in addition* to the standard semlhost library on the processor model itself. This ability to install multiple intercept libraries is only available with the Imperas Professional products.

23.4 Testing the Intercept Library

To test the intercept library, generate the assembler test case in directory 20.or1kExchange which will compile the semlhost library and will also compile the default application and test platform, to compile individually, as before use:

```

make -C exchange
make -C platform
make -C application

```

Run the platform using the assembler program executable file:

```

platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf

```

The example is a very simple one: it saves a value on the stack and then uses the new exchange instruction to exchange that with a value in a register. Because the standard assembler does not know about the new exchange instruction, it is specified as a raw bit pattern using the `.word` assembler directive:

```

////////////////////////////////////
// MAIN ROUTINE
////////////////////////////////////
.global _start

```

```

_start:
    l.addi    r31,r0,0        // initialize stack pointer to 0
    l.addi    r1,r0,0x100     // r1 = 0x100
    l.sw      -4(r31),r1      // save r1 value on stack
    l.addi    r1,r0,0x200     // r1 = 0x200
    .word     0xf7ff0ffc      // exch r1,-4(r31)
    l.lwz     r2,-4(r31)      // get current value of -4(r31)

    //////////////////////////////////////
    // EXIT FROM POINT TEST
    //////////////////////////////////////
.global exit
exit:
    l.nop

```

The output from this should be as follows:

. . . lines omitted . . .

Info 4: 'cpu1', 0x0000000001000080: l.addi r1,r0,0x200

Info 'cpu1' REGISTERS

R0 : 00000000	R1 : 00000200	R2 : deadbeef	R3 : deadbeef
R4 : deadbeef	R5 : deadbeef	R6 : deadbeef	R7 : deadbeef
R8 : deadbeef	R9 : deadbeef	R10: deadbeef	R11: deadbeef
R12: deadbeef	R13: deadbeef	R14: deadbeef	R15: deadbeef
R16: deadbeef	R17: deadbeef	R18: deadbeef	R19: deadbeef
R20: deadbeef	R21: deadbeef	R22: deadbeef	R23: deadbeef
R24: deadbeef	R25: deadbeef	R26: deadbeef	R27: deadbeef
R28: deadbeef	R29: deadbeef	R30: deadbeef	R31: 00000000
PC : 01000084	SR : 00008001	ESR: deadbeef	EPC: deadbeef
TCR: 00000000	TMR: 00000000	PSR: 00000000	PMR: 00000000
BF:0 CF:0 OF:0			

Info 5: 'cpu1', 0x0000000001000084: l.exw 0xffffffffc(r31),r1

Info 'cpu1' REGISTERS

R0 : 00000000	R1 : 00000100	R2 : deadbeef	R3 : deadbeef
R4 : deadbeef	R5 : deadbeef	R6 : deadbeef	R7 : deadbeef
R8 : deadbeef	R9 : deadbeef	R10: deadbeef	R11: deadbeef
R12: deadbeef	R13: deadbeef	R14: deadbeef	R15: deadbeef
R16: deadbeef	R17: deadbeef	R18: deadbeef	R19: deadbeef
R20: deadbeef	R21: deadbeef	R22: deadbeef	R23: deadbeef
R24: deadbeef	R25: deadbeef	R26: deadbeef	R27: deadbeef
R28: deadbeef	R29: deadbeef	R30: deadbeef	R31: 00000000
PC : 01000088	SR : 00008001	ESR: deadbeef	EPC: deadbeef
TCR: 00000000	TMR: 00000000	PSR: 00000000	PMR: 00000000
BF:0 CF:0 OF:0			

Info 6: 'cpu1', 0x0000000001000088: l.lwz r2,0xffffffffc(r31)

Info 'cpu1' REGISTERS

R0 : 00000000	R1 : 00000100	R2 : 00000200	R3 : deadbeef
R4 : deadbeef	R5 : deadbeef	R6 : deadbeef	R7 : deadbeef
R8 : deadbeef	R9 : deadbeef	R10: deadbeef	R11: deadbeef
R12: deadbeef	R13: deadbeef	R14: deadbeef	R15: deadbeef
R16: deadbeef	R17: deadbeef	R18: deadbeef	R19: deadbeef
R20: deadbeef	R21: deadbeef	R22: deadbeef	R23: deadbeef
R24: deadbeef	R25: deadbeef	R26: deadbeef	R27: deadbeef
R28: deadbeef	R29: deadbeef	R30: deadbeef	R31: 00000000

```
PC : 0100008c   SR : 00008001   ESR: deadbeef   EPC: deadbeef
TCR: 00000000   TMR: 00000000   PSR: 00000000   PMR: 00000000
BF:0 CF:0 OF:0
```

```
-----
Info 7: 'cpul', 0x000000000100008c: 1.nop      0x0
Processor 'cpul' terminated at 'exit', address 0x100008c
Info 'cpul' REGISTERS
```

```
-----
R0 : 00000000   R1 : 00000100   R2 : 00000200   R3 : deadbeef
R4 : deadbeef   R5 : deadbeef   R6 : deadbeef   R7 : deadbeef
R8 : deadbeef   R9 : deadbeef   R10: deadbeef   R11: deadbeef
R12: deadbeef   R13: deadbeef   R14: deadbeef   R15: deadbeef
R16: deadbeef   R17: deadbeef   R18: deadbeef   R19: deadbeef
R20: deadbeef   R21: deadbeef   R22: deadbeef   R23: deadbeef
R24: deadbeef   R25: deadbeef   R26: deadbeef   R27: deadbeef
R28: deadbeef   R29: deadbeef   R30: deadbeef   R31: 00000000
PC : 01000090   SR : 00008001   ESR: deadbeef   EPC: deadbeef
TCR: 00000000   TMR: 00000000   PSR: 00000000   PMR: 00000000
BF:0 CF:0 OF:0
-----
```

```
processor has executed 7 instructions
```

24 Adding an Extended Programmers View

This section provides details of extending a programmers view on the processor model from the standard view provided by the debug interface. The extended programmer's view is supported by the Imperas Professional Tools such as the MP Debugger and interception plugins; it is not supported by OVPSim.

The extended programmers view is implemented using functions from the VMI Run Time API.

24.1 An Example Programmers View

A model for the OR1K processor with additional programmer's view can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/21.or1kProgrammersView
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/21.or1kProgrammersView .
```

Compile the model using make:

```
cd 21.or1kProgrammersView
make
```

The processor model is based on the model of example 15.or1kDebugSupport, with the changes listed in following sections.

24.2 File *or1kStructure.h*

The or1k structure has new fields, `viewObject` and `addrExEvent`, used to hold the created object parent and an address exception event respectively:

```
typedef struct or1kS {
    . . . lines omitted . . .

    vmiViewObjectP exObject;           // view object for address exception info
    vmiViewEventP  addrExEvent;        // event generated on an address exception
} or1k, *or1kP;
```

24.3 File *or1kView.c*

This is a new file in which the extended programmers view is defined for the processor. In this example we have added a new event and some associated objects.

The event `addrExEvent` is triggered if an address exception occurs.

The view of the address exception has been created as an individual object with other objects associated with it. An object allows events and other objects to be grouped.

```

void orlkCreateView(orkP orlk) {

    // get the base processor view object
    vmiProcessorP processor = (vmiProcessorP)ork;
    vmiViewObjectP processorObject = vmirtGetProcessorViewObject(processor);

    // add new view object
    orlk->exObject = vmirtAddViewObject(
        processorObject, "addressException", "Address exception"
    );

    // Create an event to be generated on an address exception
    orlk->addrExEvent = vmirtAddViewEvent(
        orlk->exObject, "address", "Address exception event trigger"
    );

    // Create an object to access the EEAR
    vmiViewObjectP eearObject = vmirtAddViewObject(ork->exObject, "eear", "");
    vmirtSetViewObjectRefValue(eearObject, VMI_VVT_UN32, &ork->EEAR);

    // Create an object to access the EPC
    vmiViewObjectP epcObject = vmirtAddViewObject(ork->exObject, "epc", "");
    vmirtSetViewObjectRefValue(epcObject, VMI_VVT_UN32, &ork->EPC);

    // Create an object to access the ESR
    vmiViewObjectP esrObject = vmirtAddViewObject(ork->exObject, "esr", "");
    vmirtSetViewObjectRefValue(esrObject, VMI_VVT_UN32, &ork->ESR);
}

```

24.4 File orlkExceptions.c

The address exception is generated in the exception functions.

```

//
// Read privilege exception handler callback function
//
VMI_RD_PRIV_EXCEPT_FN(orkRdPrivExceptionCB) {

    if(MEM_AA_IS_TRUE_ACCESS(attrs)) {

        orkP orlk = (orkP)processor;
        orlk->EEAR = (Uns32)address;
        orlkTakeException(ork, DPF_ADDRESS);

        if(ork->addrExEvent){
            vmirtTriggerViewEvent(ork->addrExEvent);
        }
    }
}

//
// Write privilege exception handler callback function
//
VMI_WR_PRIV_EXCEPT_FN(orkWrPrivExceptionCB) {

    if(MEM_AA_IS_TRUE_ACCESS(attrs)) {

        orkP orlk = (orkP)processor;
        orlk->EEAR = (Uns32)address;
        orlkTakeException(ork, DPF_ADDRESS);

        if(ork->addrExEvent) {

```

```
        vmirtTriggerViewEvent(ork->addrExEvent);
    }
}

//
// Read alignment exception handler callback function
//
VMI_RD_ALIGN_EXCEPT_FN(orkRdAlignExceptionCB) {

    orkP ork = (orkP)processor;
    ork->EEAR = (Uns32)address;
    orkTakeException(ork, BUS_ADDRESS);

    if(ork->addrExEvent) {
        vmirtTriggerViewEvent(ork->addrExEvent);
    }

    return 0;
}

//
// Write alignment exception handler callback function
//
VMI_WR_ALIGN_EXCEPT_FN(orkWrAlignExceptionCB) {

    orkP ork = (orkP)processor;
    ork->EEAR = (Uns32)address;
    orkTakeException(ork, BUS_ADDRESS);

    if(ork->addrExEvent) {
        vmirtTriggerViewEvent(ork->addrExEvent);
    }

    return 0;
}
```

24.5 Testing the Extended Programmers View

The extended programmers view is only accessible from the Imperas Professional Tools.

24.5.1 Running in OVP

The model can be executed in OVP. Generate the assembler test case in directory 21.orkProgrammersView which will also compile the default application and test platform, to compile individually, as before use:

```
make -C platform
make -C application
```

Run the platform using the assembler executable file:

```
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf
```

The output from this should be as follows:

```
. . . lines omitted . . .

Info 'cpu1', 0x00000000000010000: l.ori    r30,r0,0x0
```

```

Info 'cpul', 0x00000000000010004: l.movhi r1,0x8000
Info 'cpul', 0x00000000000010008: l.movhi r2,0x1234
Info 'cpul', 0x0000000000001000c: l.ori r2,r2,0x5678
Info 'cpul', 0x00000000000010010: l.ori r3,r0,0x0
Info 'cpul', 0x00000000000010014: l.sb 0x0(r1),r2
Info 'cpul', 0x00000000000010018: l.sh 0x0(r1),r2
Info 'cpul', 0x0000000000001001c: l.sw 0x0(r1),r2
Info 'cpul', 0x00000000000010020: l.addi r3,r3,0x1
Info 'cpul', 0x00000000000010024: l.sfeqi r3,0xa
Info 'cpul', 0x00000000000010028: l.bnf 0x00010014
Info 'cpul', 0x0000000000001002c: l.addi r1,r1,0x1
Info 'cpul', 0x00000000000010014: l.sb 0x0(r1),r2
Info 'cpul', 0x00000000000010018: l.sh 0x0(r1),r2
Info 'cpul', 0x0000000000000200: l.addi r30,r30,0x1
Info 'cpul', 0x0000000000000204: l.addi r1,r1,0x1
Info 'cpul', 0x0000000000000208: l.rfe

```

. . . lines omitted . . .

```

Info 'cpul', 0x0000000000001001c: l.sw 0x0(r1),r2
Info 'cpul', 0x00000000000010020: l.addi r3,r3,0x1
Info 'cpul', 0x00000000000010024: l.sfeqi r3,0xa
Info 'cpul', 0x00000000000010028: l.bnf 0x00010014
Info 'cpul', 0x0000000000001002c: l.addi r1,r1,0x1
Info 'cpul', 0x00000000000010030: l.div r30,r30,r0
Info 'cpul', 0x00000000000010034: l.nop 0x0
Processor 'cpul' terminated at 'exit', address 0x10034

```

R0 : 00000000	R1 : 80000025	R2 : 12345678	R3 : 0000000a
R4 : deadbeef	R5 : deadbeef	R6 : deadbeef	R7 : deadbeef
R8 : deadbeef	R9 : deadbeef	R10: deadbeef	R11: deadbeef
R12: deadbeef	R13: deadbeef	R14: deadbeef	R15: deadbeef
R16: deadbeef	R17: deadbeef	R18: deadbeef	R19: deadbeef
R20: deadbeef	R21: deadbeef	R22: deadbeef	R23: deadbeef
R24: deadbeef	R25: deadbeef	R26: deadbeef	R27: deadbeef
R28: deadbeef	R29: deadbeef	R30: 0000001b	R31: deadbeef
PC : 00010038	SR : 00008601	ESR: 00008001	EPC: 0001001c
TCR: 00000000	TMR: 00000000	PSR: 00000000	PMR: 00000000
BF:1	CF:1	OF:0	

processor has executed 185 instructions

24.5.2 Operation in Imperas MP Debugger

The same platform can be executed with the Imperas professional simulator for debugging with the Imperas MP Debugger

```

IMPERAS_RUNTIME=CpuManager \
platform/platform.IMPERAS_ARCH.exe --program application/asmtest.OR1K.elf \
--idebug

```

or, using the platform shared object loaded into the Imperas professional tools for debugging.

```

imperas.exe --idebug -icmobject model.so --icmargv --program asmtest.OR1K.elf

```

The debugger allows access to the extended programmers view using the *view* command.

The platform is created as a shared object to load into the Imperas MP Debugger as part of the normal platform build process

```
make -C platform
```

If the same platform is executed with the address exception event point enabled the debugger will stop execution and allow the values of the defined objects to be accessed. The output from this should be as follows.

Before starting simulation view the defined objects

```
idebug (cpul) > view

platform
  Events:
    begin: Start of simulation event
    finish: Finish of simulation event
  cpul: processor
    Events:
      modeswitch: Mode switch event
      exception: Exception event
    type = orlk
    id = 0 (0x00000000)
    addressException: Address exception
      Events:
        address: Address exception event trigger
        eear: = 3735928559 (0xdeadbeef)
        epc: = 3735928559 (0xdeadbeef)
        esr: = 3735928559 (0xdeadbeef)
```

The new view object, addressException, is visible, containing an event and views of the eear, epc and esr registers. Note that the processor also contains automatically-created modeswitch and exception events. The modeswitch event is triggered when the processor switches mode using vmirtSetMode. The exception event is triggered when the processor updates its execution address using vmirtSetPCException.

Set an event point for the address exception event

```
idebug (cpul) > event address
Created eventpoint 1 on /platform/cpul/exception/address
```

Run the simulation

```
idebug (cpul) > continue
Info 'cpul', 0x0000000000010000: l.ori    r30,r0,0x0
Info 'cpul', 0x0000000000010004: l.movhi  r1,0x8000
Info 'cpul', 0x0000000000010008: l.movhi  r2,0x1234
Info 'cpul', 0x000000000001000c: l.ori    r2,r2,0x5678
Info 'cpul', 0x0000000000010010: l.ori    r3,r0,0x0
Info 'cpul', 0x0000000000010014: l.sb     0x0(r1),r2
Info 'cpul', 0x0000000000010018: l.sh     0x0(r1),r2
Info 'cpul', 0x000000000001001c: l.sw     0x0(r1),r2
Info 'cpul', 0x0000000000010020: l.addi   r3,r3,0x1
```

```
Info 'cpul', 0x00000000000010024: l.sfeqi  r3,0xa
Info 'cpul', 0x00000000000010028: l.bnf    0x00010014
Info 'cpul', 0x0000000000001002c: l.addi   r1,r1,0x1
Info 'cpul', 0x00000000000010014: l.sb     0x0(r1),r2
Info 'cpul', 0x00000000000010018: l.sh     0x0(r1),r2
```

Simulation stops on the event point i.e. an address exception

```
Eventpoint 1 for /platform/cpul/exception/address triggered
0x00000200 in ?? ()
```

Examining the view provides all the information defined in the model. In this case it is simply the register values associated with an address exception.

```
idebug (cpul) > view

platform
  Events:
    begin: Start of simulation event
    finish: Finish of simulation event
  cpul: processor
    Events:
      modeswitch: Mode switch event
      exception: Exception event
    type = orlk
    id = 0 (0x00000000)
    exception: Address exception
      Events:
        address: Address exception event trigger
        eear:   = 2147483649 (0x80000001)
        epc:    = 65560 (0x00010018)
        esr:    = 32769 (0x00008001)

idebug (cpul) >
```

25 Processor Configuration

A processor model is configured by variables which can be set by the platform or simulation environment. They are called *model parameters*. Model parameters are collected into a structure which is initialized by the simulator then passed to the model's constructor function. The structure exists only when the constructor is executing. To specify the type and constraints of each parameter, the model must provide an iterator function which returns each parameter specification in turn. This allows the simulator to set and check each parameter and allows other tools to discover the configuration interface of the model.

25.1 Example of a Configurable Processor

A model for the OR1K processor with model parameters can be found in:

```
$IMPERAS_HOME/Examples/Models/Processor/22.or1kConfigurable
```

Take a copy of the template model:

```
cp -r $IMPERAS_HOME/Examples/Models/Processor/22.or1kConfigurable .
```

Compile the model using make:

```
cd 22.or1kConfigurable
make
```

25.2 The Parameters Structure

The parameters structure is defined in `or1kParameters.h`:

```
typedef struct or1kParamValuesS {
    VMI_BOOL_PARAM(verbose);
    VMI_UN32_PARAM(extinterrupts);
    VMI_STRING_PARAM(extintlogfile);
} or1kParamValues, *or1kParamValuesP;
```

The example defines boolean, integer and string parameters. This is the complete list:

macro	data type
VMI_BOOL_PARAM	boolean
VMI_INT32_PARAM	32bit signed
VMI_UN32_PARAM	32bit unsigned
VMI_UN54_PARAM	64bit unsigned
VMI_DBL_PARAM	floating point
VMI_STRING_PARAM	0 terminated string
VMI_ENUM_PARAM	0 terminated string
VMI_ENDIAN_PARAM	0 terminated string

25.3 Parameters Specification

The parameters are specified in `orlkParameters.c`.

25.3.1 Structure Size

The simulator allocates this structure so need to know its size. A function must be defined (using the `VMI_PROC_PARAM_TABLE_SIZE_FN`) and set in the model attributes table:

```
//  
// Get the size of the parameter values table  
//  
VMI_PROC_PARAM_TABLE_SIZE_FN(orlkParamValueSize) {  
    return sizeof(paramValues);  
}  
  
//  
// Add function to the model attributes table  
//  
const vmiIASAttr modelAttrs = {  
    ...  
    .paramValueSizeCB = orlkParamValueSize,  
    ...  
};
```

25.3.2 Specification Objects

A function must be defined to supply the attribute specifications to the simulator (using the `VMI_PROC_PARAM_SPECS_FN` macro) and set in the model attributes table. In this model, the parameter specifications are built as a static list, and the function iterates over the list.

Each element of the list is initialized using macros defined in `vmiParameters.h`:

```
//  
// Table of parameter specs  
//  
static vmiParameter formals[] = {  
    VMI_BOOL_PARAM_SPEC (orlkParamValues, verbose, 0,  
        "Extra messages"),  
  
    VMI_UN32_PARAM_SPEC (orlkParamValues, extinterrupts, 1, 1, 8,  
        "Number of interrupts" ),  
  
    VMI_STRING_PARAM_SPEC(orlkParamValues, extintlogfile, 0,  
        "Model writes to its own log file" ),  
  
    VMI_END_PARAM  
};  
  
//  
// Function to iterate over the parameter specs  
//  
VMI_PROC_PARAM_SPECS_FN(orlkGetParamSpec) {  
    if(!prev) {  
        return formals;  
    } else {  
        prev++;  
        if (prev->name)
```

```

        return prev;
    else
        return 0;
    }
}

//
// Add function to the model attributes table
//
const vmiIASAttr modelAttrs = {
    ...
    .paramSpecsCB      = orlkGetParamSpec,
    ...
};

```

This is the complete list of parameter specification macros corresponding to the parameter definition macros:

macro	data type	limits
VMI_BOOL_PARAM_SPEC	boolean	0 or 1
VMI_INT32_PARAM_SPEC	32bit signed	specified min / max
VMI_UN32_PARAM_SPEC	32bit unsigned	specified min / max
VMI_UN54_PARAM_SPEC	64bit unsigned	specified min / max
VMI_DBL_PARAM_SPEC	floating point	specified min / max
VMI_STRING_PARAM_SPEC	0 terminated string	any string (0 if not specified)
VMI_ENUM_PARAM_SPEC	0 terminated string	string must be a member of the specified list
VMI_ENDIAN_PARAM_SPEC	0 terminated string	"big" or "little"

25.3.3 Using the Parameters

The parameter structure is allocated and assigned by the simulator then passed to the model constructor. In this example the constructor is in `orlkMain.c`:

```

//
// OR1K processor constructor
//
VMI_CONSTRUCTOR_FN(orlkConstructor) {

    orlkP          orlk  = (orlkP)processor;
    orlkParamValuesP params = parameterValues;

    ...
    // Copy parameters to the model instance
    orlk->numExtInts = params->extinterrupts;
    orlk->noisy      = params->verbose;

    // Open a log file and report or not
    if (params->extintlogfile) {
        orlk->logFile = fopen(params->extintlogfile, "w");
        if (params->verbose) {
            if (orlk->logFile) {
                ... (etc)
            }
        }
    }
}

```

The values `noisy`, `numExtInts` and `logFile` are on the model instance, so can be used anywhere in the processor model. In `orlkExceptions.c`, `numExtInts` is used to control the number of external interrupt nets:

```

////////////////////////////////////
// NET PORTS
////////////////////////////////////

static vmiNetPort netPorts[] = {
    {"intr0", vmi_NP_INPUT, (void*)1,  externalInterrupt, 0, "External interrupt" },
    {"intr1", vmi_NP_INPUT, (void*)2,  externalInterrupt, 0, "External interrupt" },
    {"intr2", vmi_NP_INPUT, (void*)4,  externalInterrupt, 0, "External interrupt" },
    {"intr3", vmi_NP_INPUT, (void*)8,  externalInterrupt, 0, "External interrupt" },
    {"intr4", vmi_NP_INPUT, (void*)16, externalInterrupt, 0, "External interrupt" },
    {"intr5", vmi_NP_INPUT, (void*)32, externalInterrupt, 0, "External interrupt" },
    {"intr6", vmi_NP_INPUT, (void*)64, externalInterrupt, 0, "External interrupt" },
    {"intr7", vmi_NP_INPUT, (void*)128, externalInterrupt, 0, "External interrupt" },
    VMI_END_NET_PORTS
};

//
// Get the next net port.
// The processor configuration sets the number of ports implemented
//
vmiNetPortP orlkGetNetPortSpec(vmiProcessorP processor, vmiNetPortP current) {
    if (!current) {
        return netPorts;
    }
    orlkP orlk = (orlkP)processor;
    current++;
    Uns32 num = current - netPorts;
    if (current->name && num < orlk->numExtInts) {
        return current;
    }
    return 0;
}

```

The `logFile` variable is used to control logging of interrupt events to a special file:

```

static VMI_NET_CHANGE_FN(externalInterrupt) {

    orlkP orlk      = (orlkP)processor;
    Uns32 deviceId = (Uns32)userData;

    if(newValue) {
        if(orlk->logFile) {
            fprintf(orlk->logFile, "Taking interrupt with ID=0x%x\n", deviceId);
        }
        ...
    }
    ....
}

```

25.4 Using a parameterized model

The example platform (`platform/platform.c`) creates two model instances. `procA` uses default parameter values, `procB` has its parameters overridden.

```

//
// Platform attributes become model parameters

```

```
//
static icmAttrListP attrsForB(void) {
    icmAttrListP userAttrs = icmNewAttrList();
    icmAddUns64Attr (userAttrs, "verbose", 1);
    icmAddUns64Attr (userAttrs, "extinterrupts", 2);
    icmAddStringAttr(userAttrs, "extintlogfile", "test.log");
    return userAttrs;
}

int main(int argc, char ** argv) {

    ...
    icmProcessorP procA = icmNewProcessor(
        "procA",           // CPU name
        TYPE_NAME,         // CPU type
        0,                 // CPU cpuId
        0,                 // CPU model flags
        32,                // address bits
        modelFile,         // model file
        "modelAttrs",      // morpher attributes
        MODEL_FLAGS,       // enable tracing
        0,                 // user-defined attributes
        semihostFile,      // semi-hosting file
        "modelAttrs"       // semi-hosting attributes
    );
    icmProcessorP procB = icmNewProcessor(
        "procB",           // CPU name
        TYPE_NAME,         // CPU type
        0,                 // CPU cpuId
        0,                 // CPU model flags
        32,                // address bits
        modelFile,         // model file
        "modelAttrs",      // morpher attributes
        MODEL_FLAGS,       // enable tracing
        attrsForB(),       // user-defined attributes
        semihostFile,      // semi-hosting file
        "modelAttrs"       // semi-hosting attributes
    );
}
```

Consequently, `procB` can be connected to two nets while `procA` can be connected to only one.

```
icmNetP n1 = icmNewNet("n1");
icmNetP n2 = icmNewNet("n2");
icmNetP n3 = icmNewNet("n3");

icmConnectProcessorNet(procA, n1, "intr0", ICM_INPUT);
icmConnectProcessorNet(procB, n2, "intr0", ICM_INPUT);
icmConnectProcessorNet(procB, n3, "intr1", ICM_INPUT);
```

26 Making High-Performance Processor Models

Previous sections in this document have shown some of the techniques required to make high performance processor models. This section summarizes some key points, not all of which have been mentioned previously.

Do as much work as possible at morph time. If it is possible to simplify an instruction by making a decision at morph time, always do so. As an example, for some OR1K control register accesses it was possible to determine at morph time which control register was being accessed, and special code was generated to do such accesses. As another example, instructions that have different behavior in different modes (e.g. user and supervisor) need only morph the behavior required for the current mode in a modal processor.

Use morph time constructs wherever possible and reasonable in preference to embedded calls. With every instruction, there is a fundamental choice of implementation strategy: should it be implemented by a morpher primitive, or by an embedded call to a C function? The interface in `vmimt.h` contains a rich set of morph time primitives, and these should almost always be used in preference to embedded calls. Embedded calls are slower, and the DFA optimizer in the simulator cannot propagate optimizations across calls, which means the code around the call is also less efficient.

The one exception is when an instruction requires many morph primitives to implement – in this case, a call to a native function may be faster (and the JIT compiled code will be smaller).

In particular, use `vmimtUncondJump` etc in preference to `vmirtSetPC` in an embedded call. Jumps described using the morph-time functions are hugely faster.

In particular, use `vmimtLoadRRO` etc in preference to `vmirtReadNByteDomain` etc in an embedded call. Loads and stores described using the morph-time functions are hugely faster.

Position the most-frequently-used simulated registers in the first 128 bytes of the processor structure. This section of the processor model is accessed in JIT code using a byte offset from the processor object, so the generated code is most compact.

Always use `VMI_REG_TEMP` to describe intermediates that are not true processor registers. The JIT code generator can often assign these entirely to native registers or optimize away references to them altogether.

Do not explicitly assign the program counter on every instruction. The program counter can be found when required using `vmirtGetPC` etc, so do not maintain its value explicitly in the processor structure.

Derive complex register values on demand. Infrequently-accessed but frequently-changing registers like status registers should have their values assembled on demand when required. This is especially true if the registers contain flags.

Master flag register values as bytes in the processor structure. Do not master them as bits in a status register, as this will require complex code to correct the status register on every flag change.

Use `vmirtAliasMemoryVM/vmirtUnaliasMemoryVM` to model virtual memory pages. The simulator will generate JIT code in such regions so that it is *physically* mapped and relocatable, which dramatically improves simulation performance of VM systems.

Use `vmirtAliasMemory/vmirtUnaliasMemory` to model fixed-mapped pages. The simulator will generate JIT code in such regions so that it is *virtually* mapped and not relocatable, which is even faster than relocatable physically mapped sections, provided that the mapping remains constant.

26.1 Processor Model Efficiency Analysis

When a processor model is fully implemented there are a some further checks that can be made to ensure that it is working effectively as possible. The first step is to specify the `ICM_VERBOSE_DICT` attribute when initializing ICM in the test harness:

```
icmInitPlatform(ICM_VERSION, ICM_VERBOSE_DICT, NULL, 0, "platform");
```

When this attribute is used additional summary statistics are printed at the end of simulation for each of the code dictionaries.

```
Info -----
Info DICTIONARY 'KERNEL (ERL=0)' STATISTICS
Info Simulated code size      : 220 bytes
Info Native code size        : 530 bytes
Info      bloat factor        : 2.4
Info Active code blocks       : 8
Info MEMORY STATISTICS
Info Static mapped memory     : 0 bytes
Info Dynamic mapped memory    : 73,728 bytes
Info Master regions (by size)
Info      smaller than 1Kb    : 0
Info      smaller than 4Kb    : 0
Info      at least 4Kb        : 11
Info Memory model efficiency
Info      uncached reads      : 1
Info      - which set region  : 1
Info      uncached writes     : 3
Info      - which set region  : 3
Info -----
Info DICTIONARY 'KERNEL (ERL=1)' STATISTICS
Info Simulated code size      : 48 bytes
Info Native code size        : 138 bytes
Info      bloat factor        : 2.9
Info Active code blocks       : 3
Info Block ejections
Info      set mode block      : 1
```

```
Info MEMORY STATISTICS
Info Static mapped memory      : 0 bytes
Info Dynamic mapped memory    : 73,728 bytes
Info Master regions (by size)
Info   smaller than 1Kb       : 0
Info   smaller than 4Kb       : 0
Info   at least 4Kb           : 11
Info Memory model efficiency
Info   uncached reads         : 1
Info   - which set region     : 1
Info   uncached writes        : 3
Info   - which set region     : 3
Info -----
```

The fields of most interest are the *bloat factor* and the categories listed beneath *Block ejections*.

'**bloat factor**' gives the ratio of *JIT-translated code size* to the *original simulated code size* for all currently-active code blocks (for example, if a simulated instruction sequence of 20 bytes gets translated into a native block of 100 bytes, the bloat factor is 5).

Typically, JIT-translated code is somewhat more verbose than the original source, so this number is usually greater than 1. If the ratio is much higher than 5 or 6, then it could mean that some instruction translations are producing very long native code equivalents, which could be a cause of inefficiency. Create small (probably assembler) test cases to help identify which instructions may be suffering from poor translation.

'**Block ejections**' indicates how many code blocks were ejected during simulation, grouped by the reason for the ejection. If a large number of code blocks are ejected during a simulation run, that could indicate an error in the model (an inappropriate use of block masks, for example). By writing small, probably assembler, test cases and using ICM_VERBOSE_DICT, individual instructions and/or families of instructions can be verified for efficiency.

You may wish to **morph instructions into separate code blocks** to be able to look at individual instructions morphed code generation rather than, as is normally the case, multiple instructions morphed into a code block and the simulator applying inter-instruction optimizations, that results in much more efficient code but may mask an issue. This is achieved in the test harness by running one instruction at a time using the `icmSimulate()` function with an instruction count of 1:

```
icmSimulate(processorHandle, 1);
```

```
##
```